

# General Concepts

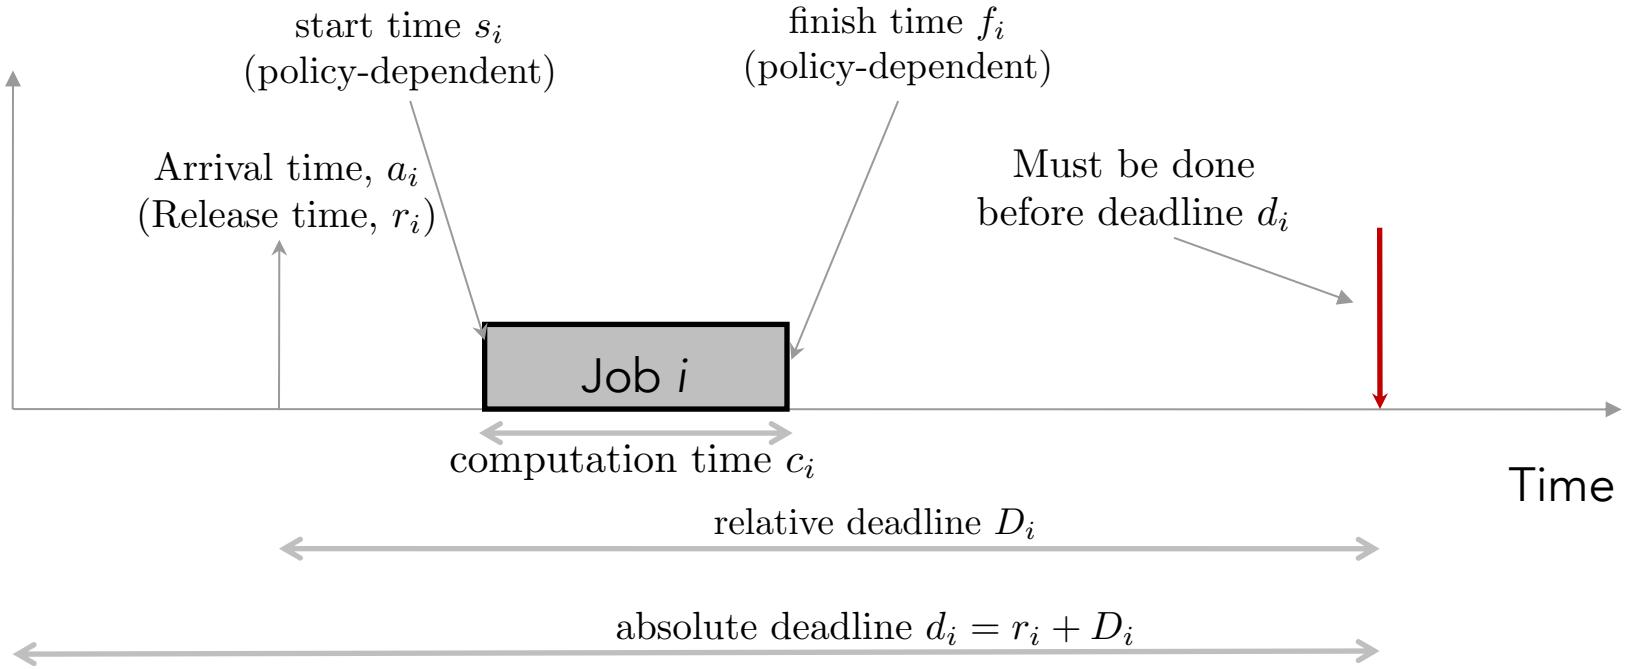
---

Introduction to real-time systems

# Abstraction

- To reason about timeliness we need an abstraction that is:
  1. Expressive enough to capture realistic workload models and parameters
  2. Constrained enough to derive useful results (tractable)

## Job model



# Online vs. Offline scheduling

---

- What do we know about the jobs to be scheduled?
- **Offline:** All job parameters  $r_i, c_i, D_i, \dots$  and number of jobs  $n$  are known a priori
- **Online:** Jobs arrive as a stream
  - Might not know arrival times
  - We do not know number of jobs
  - We might not know job execution time requirements until job arrives or even until job finishes execution!

# Job Scheduling: Policies

---

- A rule that specifies which job occupies which processor at every time instant
- May be
  - Deterministic
  - Randomized
  - Preemptive or non-preemptive
  - *Priority-driven* or *clock-driven* (or a hybrid of both, or none of which!)

# Job Scheduling: Policies

---

- Deterministic scheduling policy:  $n$  jobs

A function  $S: \mathbb{R}_+ \rightarrow \{1, \dots, n\}$  where  $S(t)$  is the index of the job that occupies the processor at time  $t$

- $S(t) = 0$  if the processor is kept idle
- Some authors impose the requirement that a policy be a *step function*
  - $\forall t > 0, \exists t_1, t_2 \geq 0$  with  $t \in [t_1, t_2]$  such that  $S(t') = S(t)$  for every  $t' \in [t_1, t_2]$

# Job Scheduling: Priority-driven Policies

---

- Priority-driven policies
  - Assign priorities to jobs so that the available job with the highest priority occupies the processor
  - Priority can be static or vary through time as job execution advances (dynamic priority)
  - Q: With  $n$  jobs, how many static priority assignments exist? (where each job has a unique priority)
  - Q: Can an optimal static priority assignment be found in time that is polynomial in the problem parameters?
  - What does **optimal** mean?
    - For now, minimizes some scheduling objective → will consider more sophisticated definitions of optimality shortly
  - Priority-driven Policies: also called **work-conserving, list scheduling, greedy** → processor is never left idle on purpose, always schedule a job if one is available

# Job Scheduling: Objectives (cost functions)

We have  $n$  jobs  $J_1, \dots, J_n$  with parameters  $r_i, c_i, D_i$ , etc

- **Objective related to timing properties**
  - minimize *total completion (finish) time*:  $\sum_{i=1}^n f_i$
  - minimize *total weighted completion time*:  $\sum_{i=1}^n w_i f_i$
  - Minimize **makespan**  $\max_i f_i$ 
    - *Makespan*: finish time of last job to leave the system (length of schedule)
    - Is “machine owner oriented”: A minimum makespan implies a good utilization of the machine(s).
    - Not suitable for interactive applications: Jobs released early might be delayed till end of an optimal-makespan schedule → not acceptable by “users”
  - Minimize total flow (response) time:  $\sum_{i=1}^n (f_i - r_i)$ 
    - $(f_i - r_i)$ : total time job  $i$  is in system = *waiting time* + *processing time*
    - Beneficial from “users” perspective
    - Also minimizes total completion time and total waiting time
    - **Drawback:** Can lead to *starvation*: Might still cause some jobs to be delayed unboundedly (no *fairness*)

# Job Scheduling: Objectives (cost functions)

- Minimize max. flow time:  $\max_i(f_i - r_i)$ 
    - Schedule responsive to each job
    - Starvation-free
  - **Lateness:**  $L_i = f_i - d_i \rightarrow$  Minimize maximum lateness  $L_{\max} = \max_i L_i$ 
    - Worst case violation of deadlines. The earlier the better!
  - **Tardiness:**  $T_i = \max\{0, L_i\} \rightarrow$  Minimize maximum tardiness  $T_{\max} = \max_i T_i$ 
    - Doesn't care how early (before the deadline) jobs finish
  - **# of tardy jobs**  $\sum_{i=1}^n \mathbf{1}\{L_i > 0\} = \sum_{i=1}^n \mathbf{1}\{f_i > d_i\}$
- 
- Deadline-related**

# Job Scheduling: Stochastic Models

---

- What if we have only **statistical data** of job parameters? E.g.,
  - Jobs (events) *arrive* randomly
  - Job execution time (**demand**) is random
    - **Inventory** where product **demands** are not known exactly but we know their distribution
    - Machine might *fail* randomly, causing execution times to be “stretched” randomly by (down time + time to service machine)
- **Questions of interest:**
  - What is the *probability* that jobs miss their deadlines under some policy?
  - What is the *expected max lateness*  $\mathbb{E}L_{\max}$  under some fixed policy?
  - What policy *minimizes* the expected max lateness?

# Job Scheduling: Preemption

---

- Can we interrupt a job once it starts execution?
- Preemptible jobs/resources
  - Example: execution of programs on your computer may be interleaved to implement multitasking
- Non-preemptible jobs/resources:
  - You might be on an important phone call so you block all other incoming calls
  - Disable interrupts during the execution of an ISR
  - When using resources with complex or costly state associated with it → prohibitive context switching costs → *blocks on disk*
- No-preemption problems are usually computationally harder (**NP-Hard**)
- Allowing preemption gives more flexibility to designer → sometimes more tractable
  - Example: Minimizing maximum lateness with arbitrary release times is **NP-hard** when no preemptions allowed → becomes **poly-time** solvable when preemptions allowed
  - Downside: preemption (context switching) costs might be prohibitive

## Example: Minimize total weighted completion time

---

- Given have  $n$  jobs  $J_1, \dots, J_n$ 
  - $J_i$  has execution time  $e_i > 0$  and weight (importance)  $w_i \geq 0$
  - All jobs arrive at time 0  $\rightarrow r_i = 0 \ \forall i \in \{1, \dots, n\}$
  - One processor
  - No Preemption
- For instance: Job might be message to be broadcasted, all messages ready for transmission at time 0, message  $i$  has length  $e_i > 0$ , and a fraction  $w_i \geq 0$  of receivers are interested in message  $i$ 
  - Finish time  $f_i$  is the time at which message is fully transmitted

$$\text{minimize } \sum_{i=1}^n w_i f_i$$

## Example: Minimize total weighted completion time

---

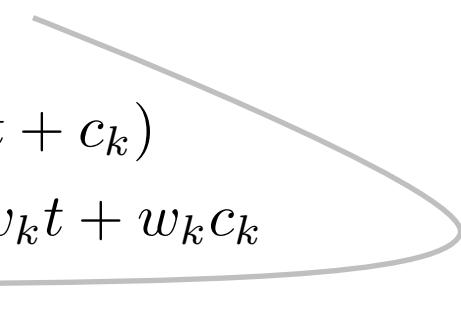
- Is there an optimal priority assignment?
- If so, how does the priority function look like?
- Think ... greedy
- Key: Job interchange argument
- Will do proof of optimality in class
- Optimal policy: **weighted shortest processing time first (WSPT)**
- Order jobs such that  $\frac{w_1}{c_1} \geq \frac{w_2}{c_2} \geq \dots \geq \frac{w_n}{c_n}$  : ratio  $\frac{w_i}{c_i}$  is job  $i$ 's index under WSPT
- At any time instant schedule the available job with the highest index

# Optimality of WSPT for total weighted completion time

- Let  $S$  be an optimal schedule that is not WSPT
  - Then there are two adjacent jobs  $i$  and  $k$  (assuming  $i$  is before  $k$  in  $S$ ) such that  $\frac{w_i}{c_i} < \frac{w_k}{c_k} \rightarrow$  some pair of adjacent jobs violates the WSPT ordering
- Let  $t$  be the start time of  $J_i$  in  $S \rightarrow$  under  $S$ :
  - $f_i = (t + c_i), \quad f_k = (t + c_i + c_k)$
- Interchange the order of  $J_i$  and  $J_k$  to produce a new schedule  $S'$ 
  - Interchanging those jobs doesn't affect  $\sum_j w_j f_j$  of job sequence before and after  $J_i$  and  $J_k \rightarrow$  difference in objectives of  $S$  and  $S'$  is solely due to  $J_i$  and  $J_k$
- Under  $S'$ :
  - $f'_i = (t + c_i + c_k), \quad f'_k = (t + c_k)$

## Optimality of WSPT for total weighted completion time

- But we know that  $\frac{w_i}{c_i} < \frac{w_k}{c_k} \Rightarrow w_i c_k < w_k c_i$

$$\begin{aligned} w_i f'_i + w_k f'_k &= w_i(t + c_i + c_k) + w_k(t + c_k) \\ &= w_i t + w_i c_i + \underbrace{w_i c_k}_{< w_k c_i} + w_k t + w_k c_k \\ &< w_i(t + c_i) + w_k(t + c_i + c_k) \\ &= w_i f_i + w_k f_k \end{aligned}$$


- Objective under  $S' <$  Objective under  $S \rightarrow$  contradicts optimality of  $S$

# Variations

---

Allow preemptions and arbitrary release times

**Problem:**  $1|r_j, \text{prmt}|\sum w_j f_j$

1. Is the Weighted Shortest *Remaining* Processing Time rule optimal for this problem?
2. What if all weights are unity but we allow arbitrary release times and preemptions?  $1|r_j, \text{prmt}|\sum f_j$
3. What if we disallow preemptions in 2.?  $1|r_j|\sum f_j$

## Optimality of SRPT for $1 |r_j, \text{prmt}| \sum f_j$

- SRPT is a **dynamic** priority rule
  - Priority of  $J_i$  at time  $t = \frac{1}{c_i(t)}$  if  $t \geq r_i$  and is undefined otherwise
    - $c_i(t)$  is the *remaining* execution time of  $J_i$  at time  $t$
  - The job currently running can be preempted only by a job with higher priority that arrives during its execution
    - **Why?** A job's priority is a non-decreasing function of time → a job's priority increases as it is given more processor time
    - The job that is in possession of the processor cannot be preempted by a job that has arrived prior to its start time
    - This shows that (1) we need to check only arrival times to determine which job occupies the processor at every time instant, and (2) # of preemptions of SRPT is finite

## Optimality of SRPT for $1 |r_j, \text{prmt}| \sum f_j$

- Let  $S$  be a policy that is optimal but not SRPT
- Then there is  $t > 0$  and available job  $i$  with highest priority (least RPT) that is not scheduled at  $t$  but instead another available job  $k \neq i$  is allocated the processor
- Let  $c'_i$  and  $c'_k$  be the remaining processing times of  $J_i$  and  $J_k$ 
  - Then  $c'_i < c'_k$  and a total of  $c'_i + c'_k$  time is expended on  $J_i$  and  $J_k$  past time  $t$
- Assume wlog that  $f_i < f_k$  under  $S$  (the other case is identical)
- **Pairwise interchange:** more intricate because of preemptions
  1. Use the first  $c'_i$  time units spent on only  $i$  and  $k$  past time  $t$  to execute  $J_i$
  2. Use the remaining  $c'_i$  time units spent on only  $i$  and  $k$  past time  $t$  to execute  $J_k$

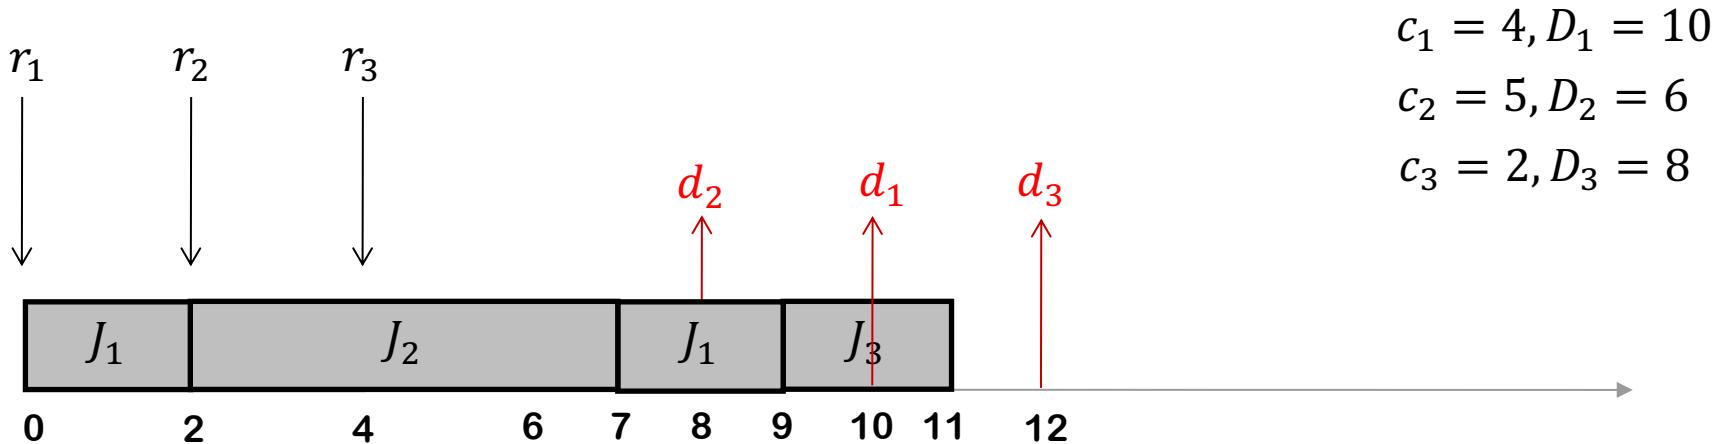
## Optimality of SRPT for $1 |r_j, \text{prmt}| \sum f_j$

- **Observe:**
  - Because of preemptions, both  $J_i$  and  $J_k$  might be preempted after time  $t$ ; however, exactly the total of their remaining processing times  $c'_i + c'_k$  is spent on their combined execution after time  $t$
  - The interchange of “execution time distribution” we did does not affect jobs other than  $J_i$  and  $J_k$
  - In new schedule  $S'$ :  $f'_i < f_i$  and  $f'_{k'} = f_k \rightarrow$  finish time of  $J_k$  doesn’t change
  - Difference in objectives of  $S$  and  $S'$  is due only to  $J_i$  and  $J_k$
  - $\sum_j f'_j - \sum_j f_j = (f'_i + f'_k) - (f_i + f_k) = f'_i + f_k - f_i - f_k = f'_i - f_i < 0$ 
    - Value of objective of  $S'$  is strictly less than value of  $S \rightarrow$  contradicts optimality of  $S$

# Deadline-driven Example: $1|\text{prmt}, r_j|L_{\max}$

**Problem:** Schedule  $n$  jobs with arbitrary arrival times  $r_1, \dots, r_n$  and relative deadlines  $D_1, \dots, D_n$  on a single processor with preemption allowed to minimize **maximum lateness**

- $1|\text{prmt}, r_j|L_{\max}$  is an example of the **Graham notation** for job scheduling problems
- **Earliest Deadline First (EDF)** is optimal → At every time instant, schedule the available job whose *absolute deadline*  $d_j = r_j + D_j$  is nearest
- **Preemptive:** if job  $i$  is running when job  $j$  arrives at time  $r_j$  and  $d_j < d_i$ , then remove job  $i$  from processor and start running job  $j$ , otherwise keep running job  $i$



# Optimality of preemptive EDF for $1|\text{prmt}, r_j|L_{\max}$

---

**Problem:** Schedule  $n$  jobs with arbitrary arrival times  $r_1, \dots, r_n$  and relative deadlines  $D_1, \dots, D_n$  on a single processor with preemption allowed to minimize **maximum lateness**

- Optimality of preemptive EDF follows by the two following claims:
- **Claim1:**  $L_{\max} \geq r(S) + c(S) - d(S)$  for any  $S \subset \{1, \dots, n\}$ , where  
$$r(S) = \min_{j \in S} r_j, \quad c(S) = \sum_{j \in S} c_j, \quad d(S) = \max_{j \in S} d_j$$
 (regardless of scheduling algorithm)
- **Claim2:** preemptive EDF gives a schedule with  $L_{\max}(\text{EDF}) = \max_{S \subset \{1, \dots, n\}} \{ r(S) + c(S) - d(S) \}$   
(achieves the min possible  $L_{\max}$ )
- **Proof:** HW1
- Another optimal policy: Least **Laxity** First (LLF)
- **Laxity** (slack) at time  $t$  is  $\ell_i(t) = d_i - t - c_i(t)$ ;  $c_i(t)$ : remaining execution-time at time  $t$

# Notes on preemptive EDF

- Is an *online* policy
- How many preemptions for  $n$  jobs?
  - equals the number of **distinct** *release times*
- Can be implemented in  $O(n \log n)$

## Stochastic $1||L_{\max}$

- Suppose job execution times are positive random variables on some probability space and deadlines are deterministic
- What rule minimizes the **expected max** lateness  $\mathbb{E}(L_{\max}) = \mathbb{E}(\max(L_1, \dots, L_n))$ ?
- Hint: very similar to deterministic case!
- Does the same rule minimize the **max expected** lateness  $\max(\mathbb{E}(L_1), \dots, \mathbb{E}(L_n))$ ?

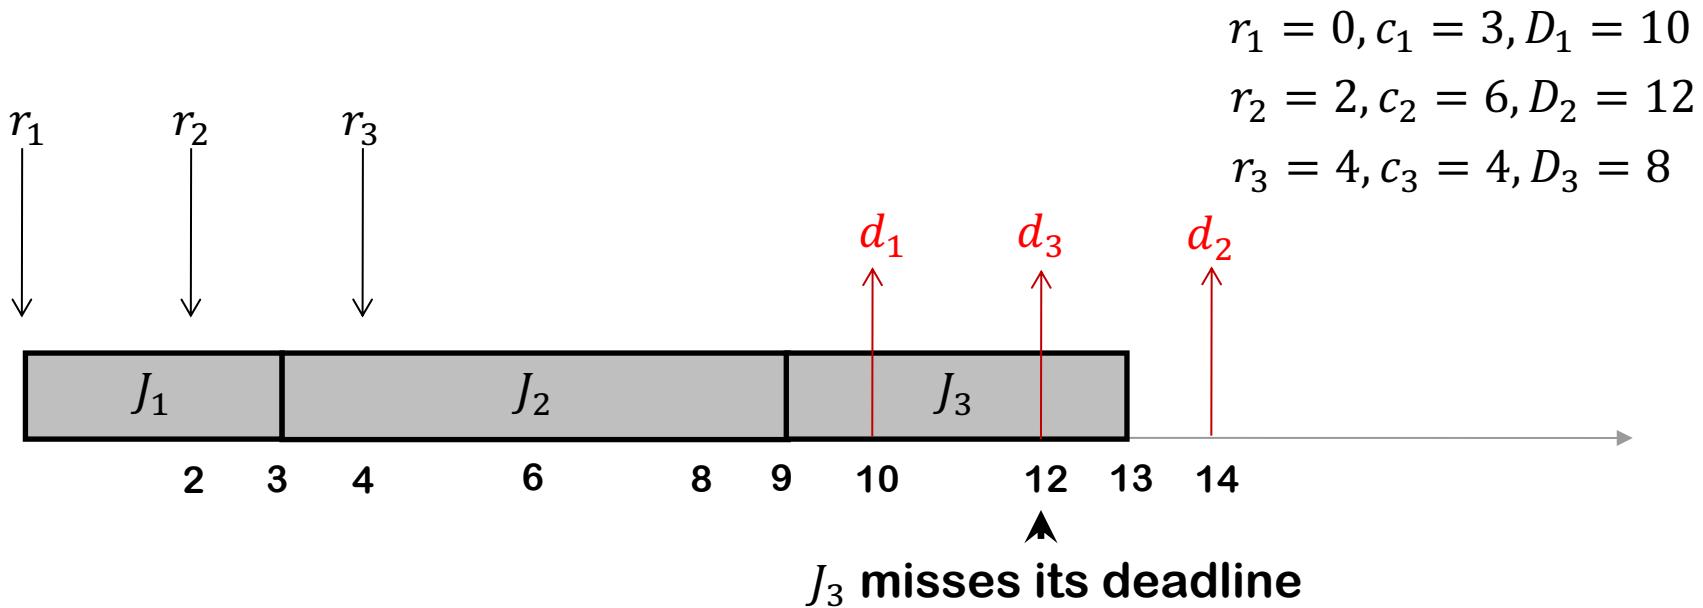
# Schedulability (feasibility) problems

- We have  $n$  jobs  $J_1, \dots, J_n \rightarrow J_i$  has parameters  $(c_i, r_i, D_i)$
- Problem: how to schedule jobs so that they meet their deadlines if this is possible.
- There is no cost function!
  - Such problems are called **feasibility** (or **schedulability**) problems
- How do we define optimality then?
  - A rule  $S$  is optimal in the following sense: *For every instance of the problem, it can always meet the deadlines if some other rule can*
  - Equivalent: *If rule  $S$  cannot meet deadlines, then no other rule can*
  - Formally: Algorithm  $S$  is optimal for problem  $\Pi$  if  $\forall S' \neq S \ \forall I \in \Pi (I \text{ schedulable under } S' \rightarrow I \text{ schedulable under } S)$

# Non-preemptive scheduling with arbitrary release times

We have  $n$  non-preemptible jobs  $J_1, \dots, J_n \rightarrow J_i$  has parameters  $c_i, r_i, D_i$

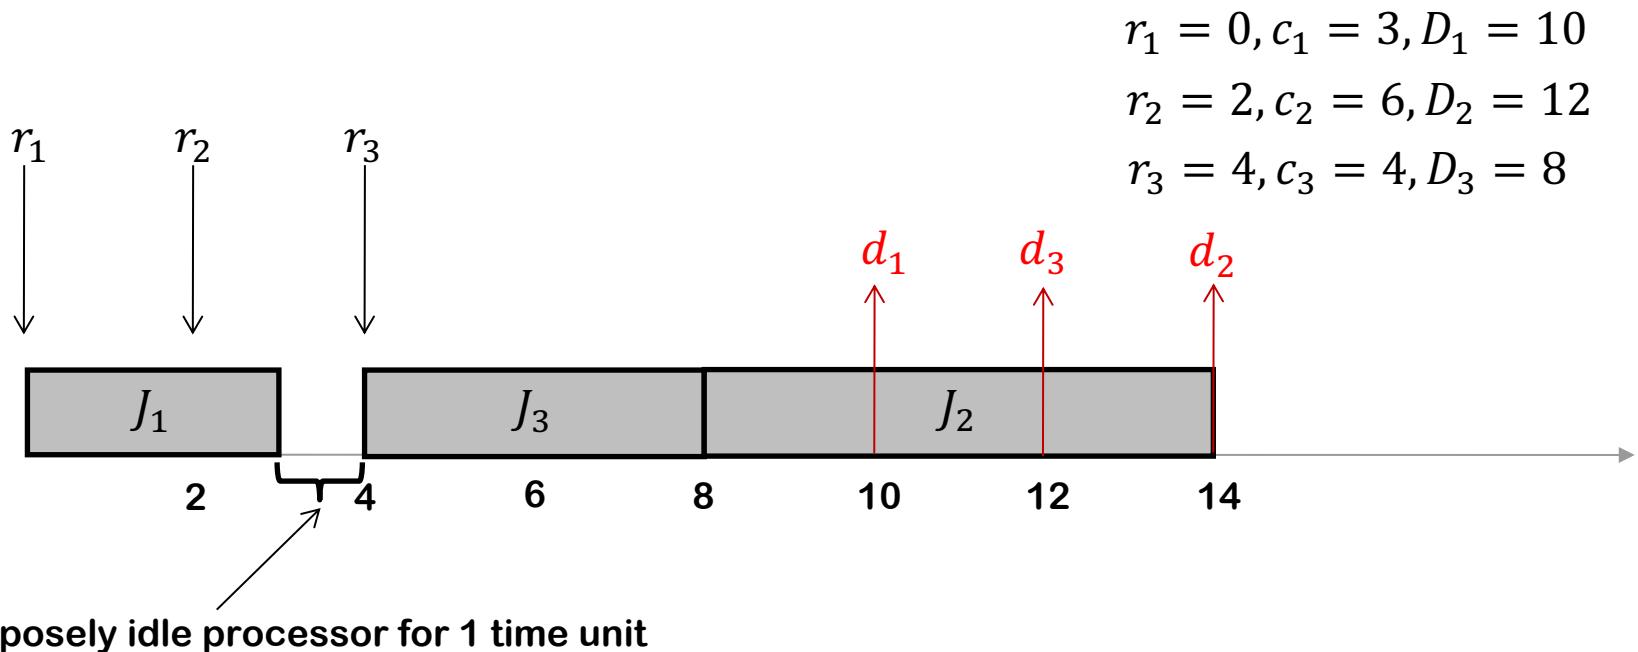
- Problem: how to schedule jobs **non-preemptively** so that all jobs meet their deadlines if this is possible.
- Is non-preemptive EDF optimal?



# Non-preemptive scheduling with arbitrary release times

We have  $n$  non-preemptible jobs  $J_1, \dots, J_n \rightarrow J_i$  has parameters  $c_i, r_i, D_i$

- Problem: how to schedule jobs **non-preemptively** so that all jobs meet their deadlines if this is possible.
- Is there a policy under which this job set can meet its deadlines?



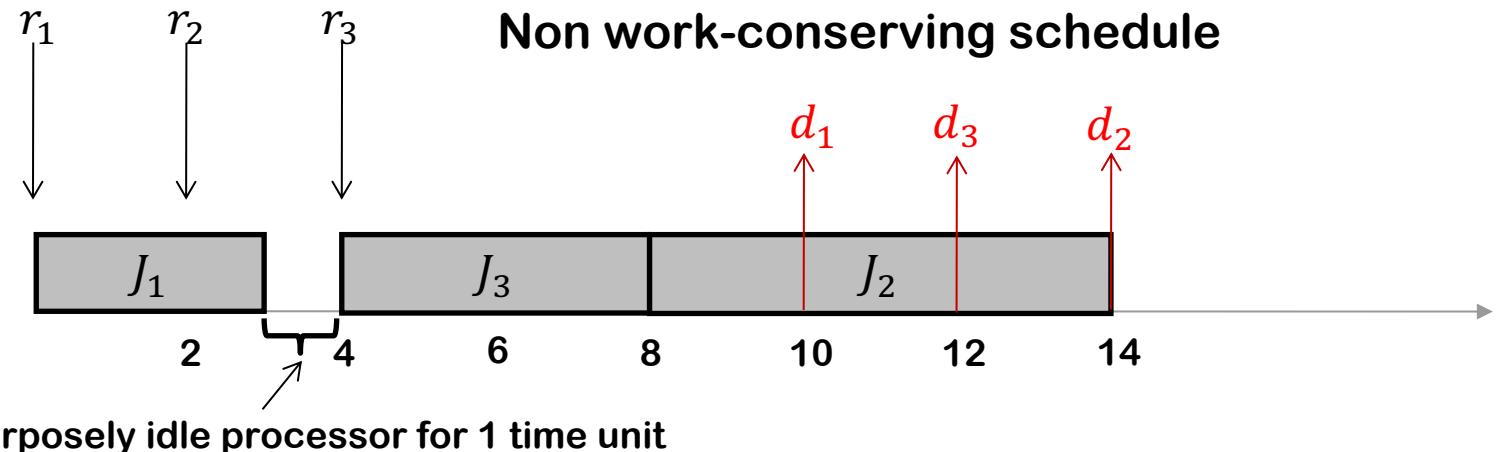
# Non-preemptive scheduling with arbitrary release times

- This schedule introduces idle time on purpose to meet deadlines
  - Schedule is *non-work-conserving* (not greedy, not priority-drive)
- **This example shows more:** no priority-driven algorithm is optimal for the non-preemptive problem with arbitrary deadlines, execution times, and release times

$$r_1 = 0, c_1 = 3, D_1 = 10$$

$$r_2 = 2, c_2 = 6, D_2 = 12$$

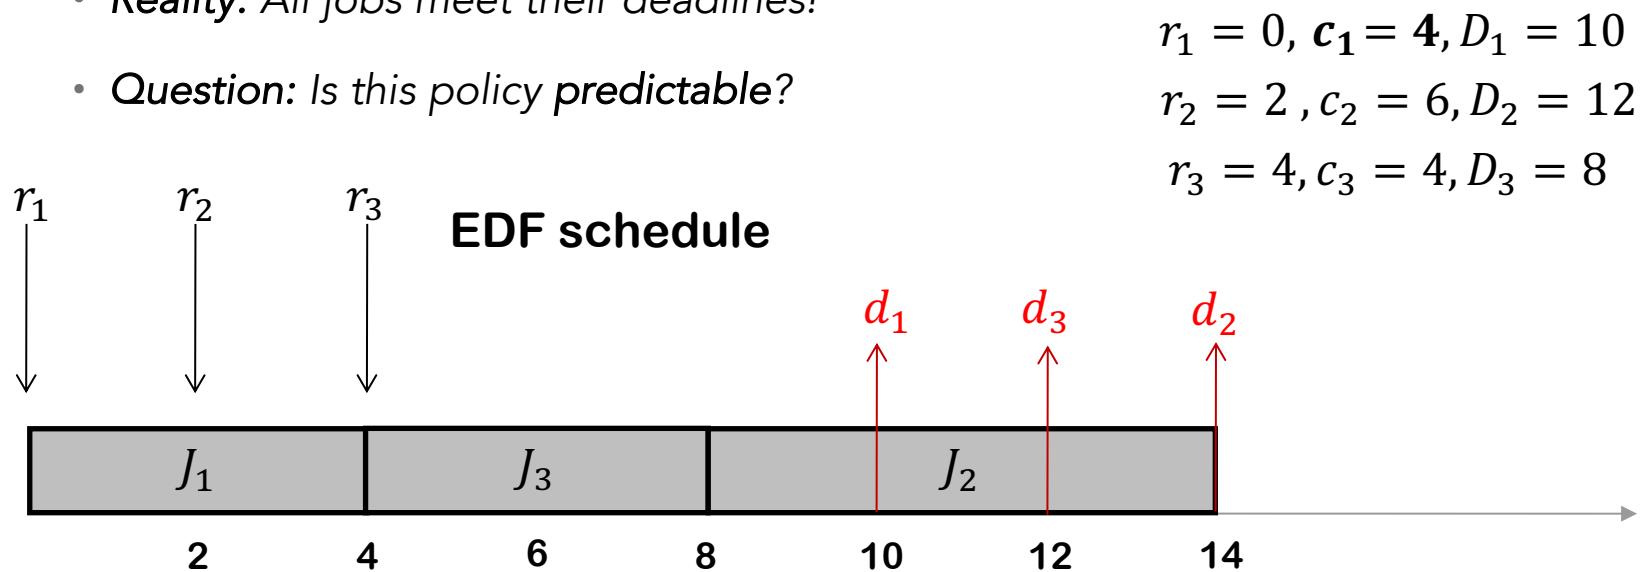
$$r_3 = 4, c_3 = 4, D_3 = 8$$



# A scheduling anomaly

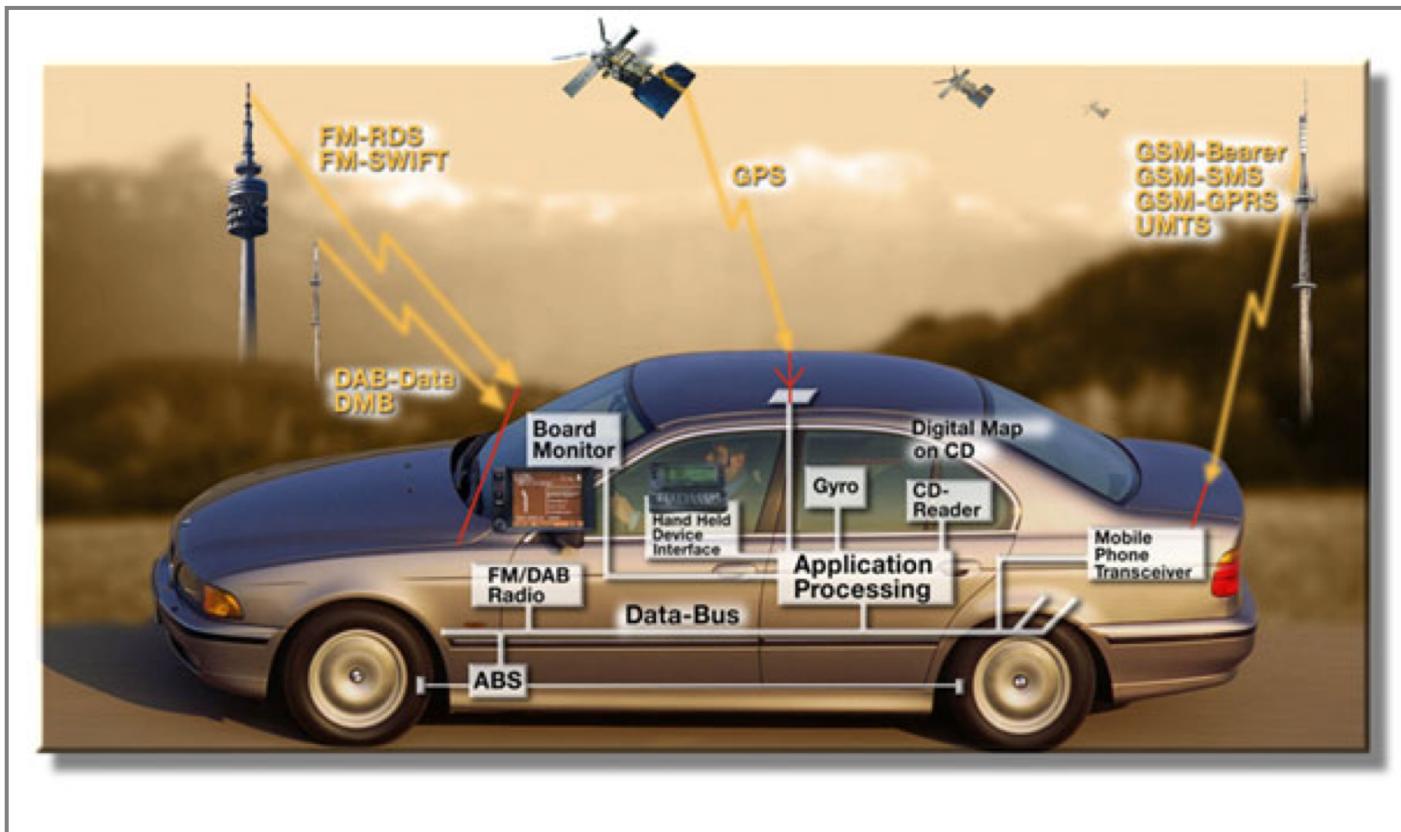
We have  $n$  non-preemptible jobs  $J_1, \dots, J_n \rightarrow J_i$  has parameters  $c_i, r_i, D_i$

- Problem: how to schedule jobs **non-preemptively** so that all jobs meet their deadlines if this is possible.
- What happens under EDF if we increase  $c_1$  from 3 to 4?
  - Intuition: More workload  $\rightarrow$  more deadlines should be missed under same policy
  - Reality: All jobs meet their deadlines!
  - Question: Is this policy predictable?



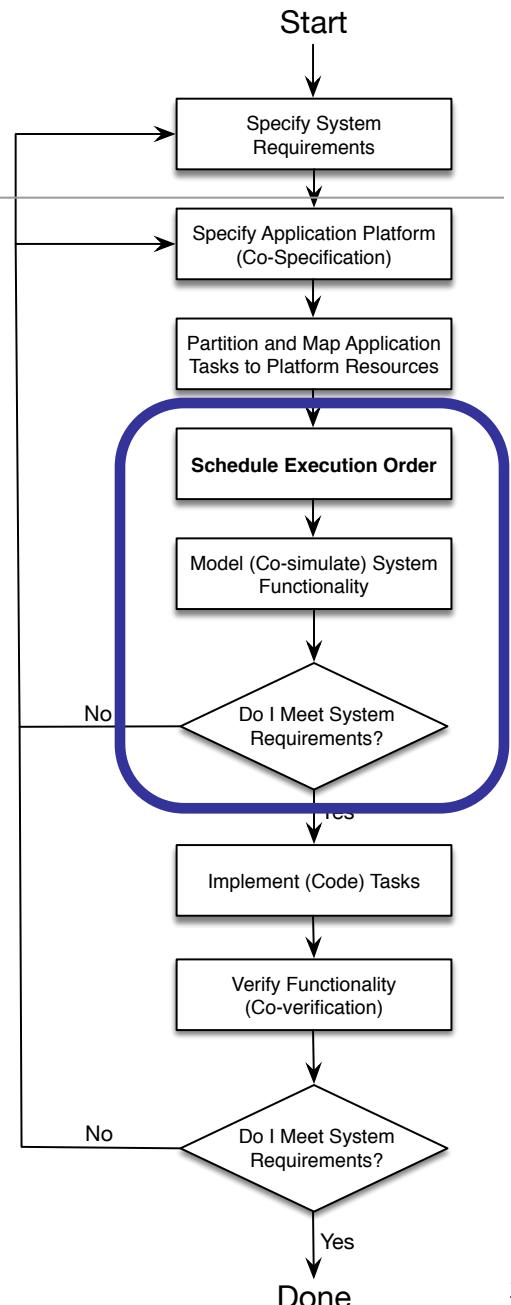
# Review

- What is a real-time system?
- What is an embedded system?
- What characteristic of a real-time system is probably the most important?



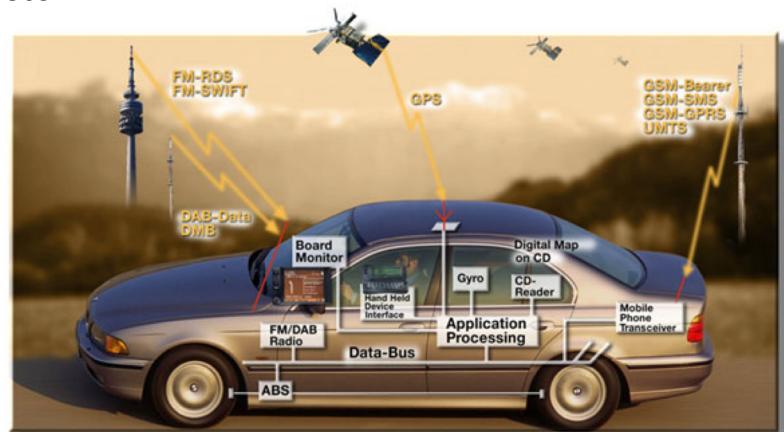
# The system design process

- Designing any computer system involves many steps.
- Some steps are common to many types of systems.
- A few steps are more important in a real-time system.
  - **Scheduling** is one such operation.
  - How do we know if a set of tasks can be scheduled in a predictable manner?
- We will touch upon other parts of the design process later in the course.



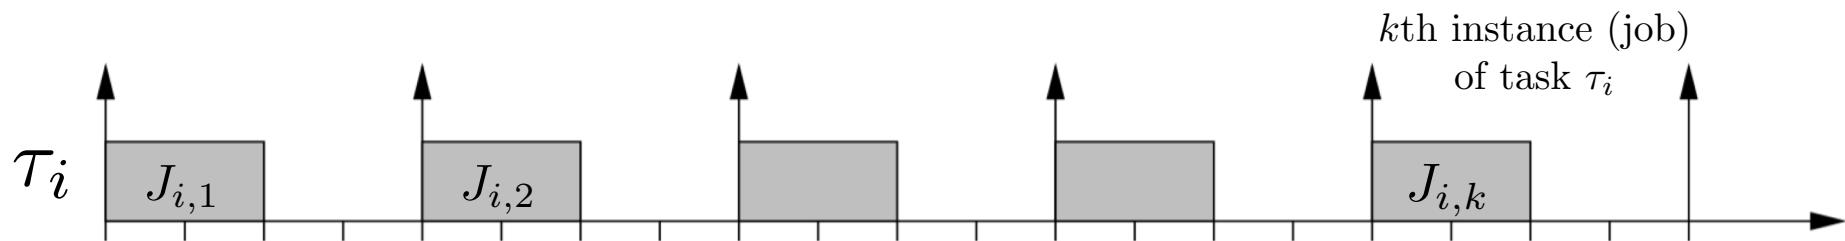
# The schedulability question: Drive-by-Wire Example

- Consider a control system in a future vehicle
  - Steering wheel sampled every **10 ms** – wheel positions adjusted accordingly (computing the adjustment takes **4.5 ms** of CPU time)
  - Brakes sampled every **4 ms** – brake pads adjusted accordingly (computing the adjustment takes **2 ms** of CPU time)
  - Velocity is sampled every **15 ms** – acceleration is adjusted accordingly (computing the adjustment takes **0.45 ms**)
  - For safe operation, adjustments must always be computed before the next sample is taken
- Is it possible to always compute all adjustments in time?
- The underlying computer system is a **uniprocessor** system.



# Tasks

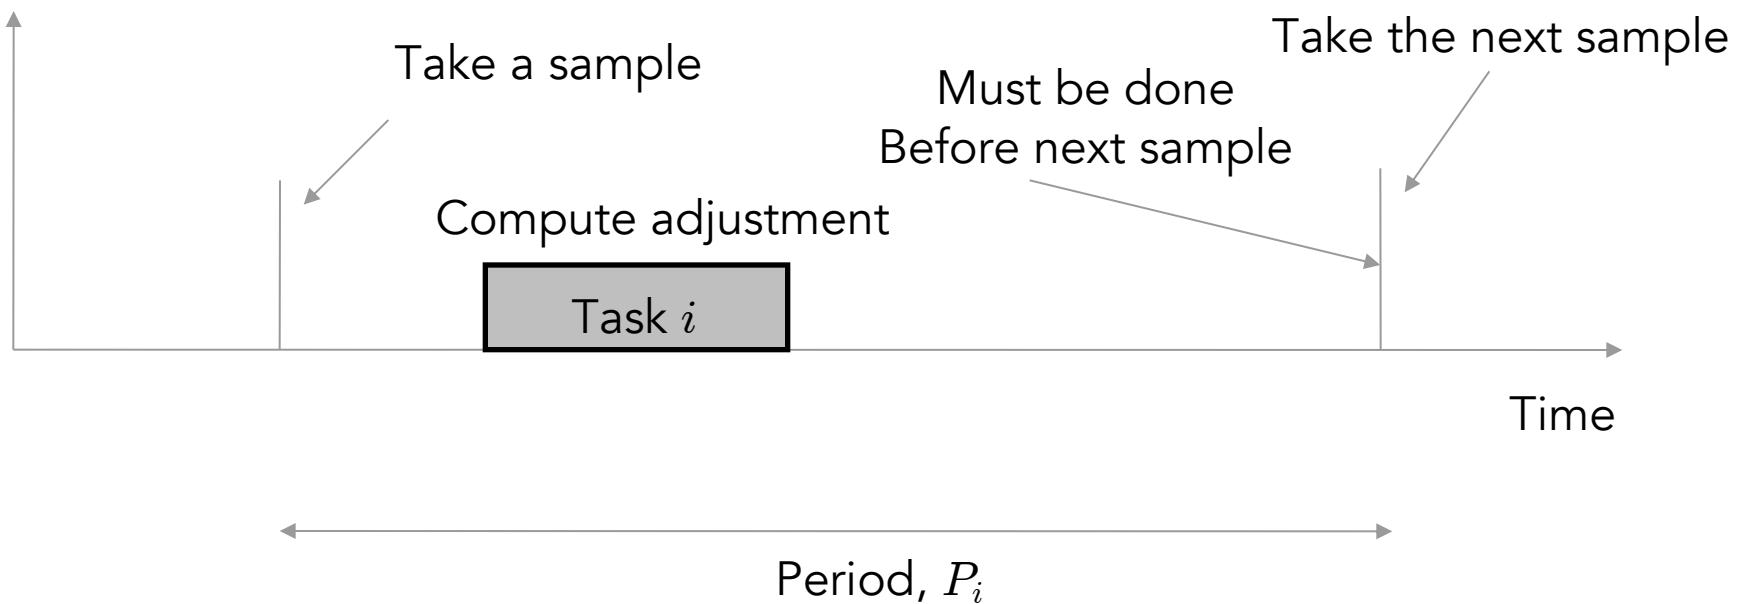
- Basic RT model: Periodic **Task model**
  - A task releases an infinite (indefinitely long) succession of jobs



# Some terminology

---

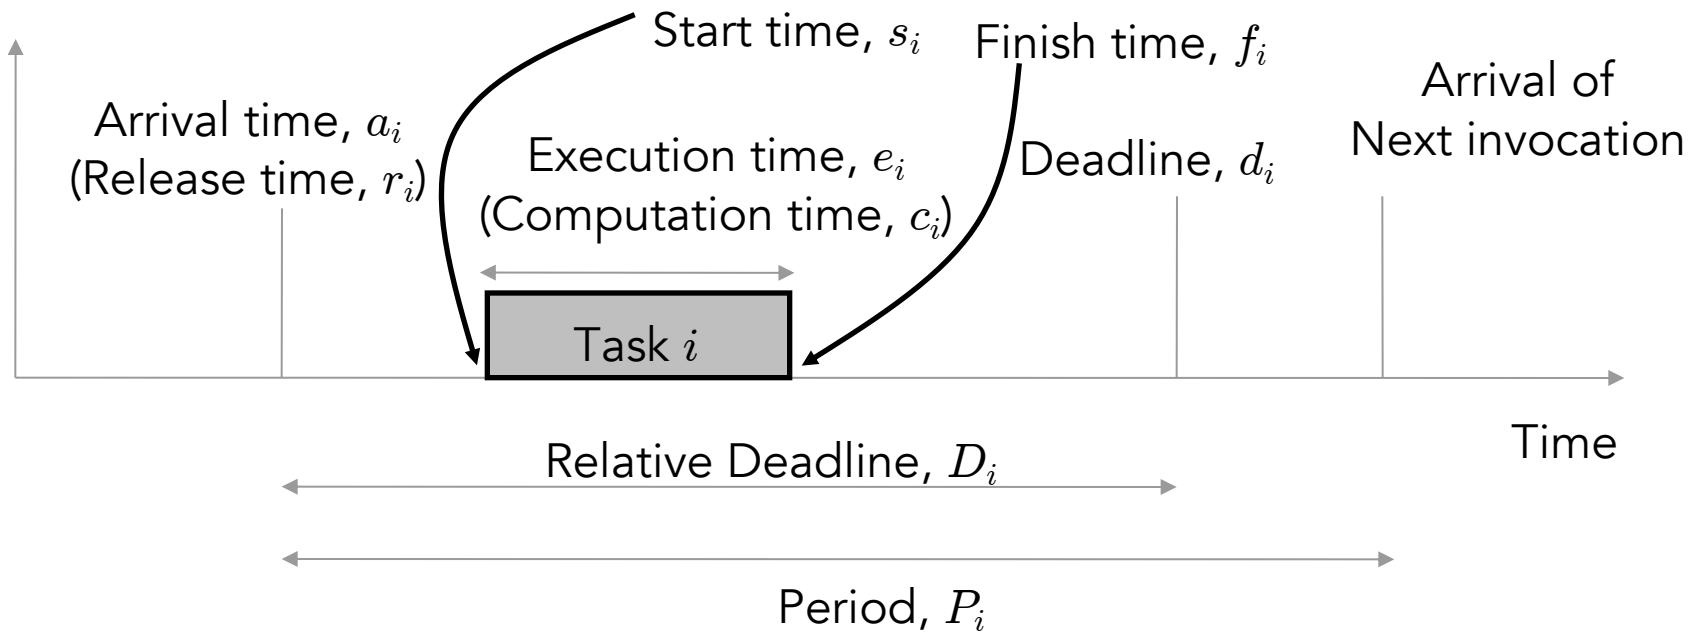
- Tasks, periods, arrival-time, deadline, execution time, etc.



# Some terminology

---

- Tasks, periods, arrival-time, deadline, execution time, etc.

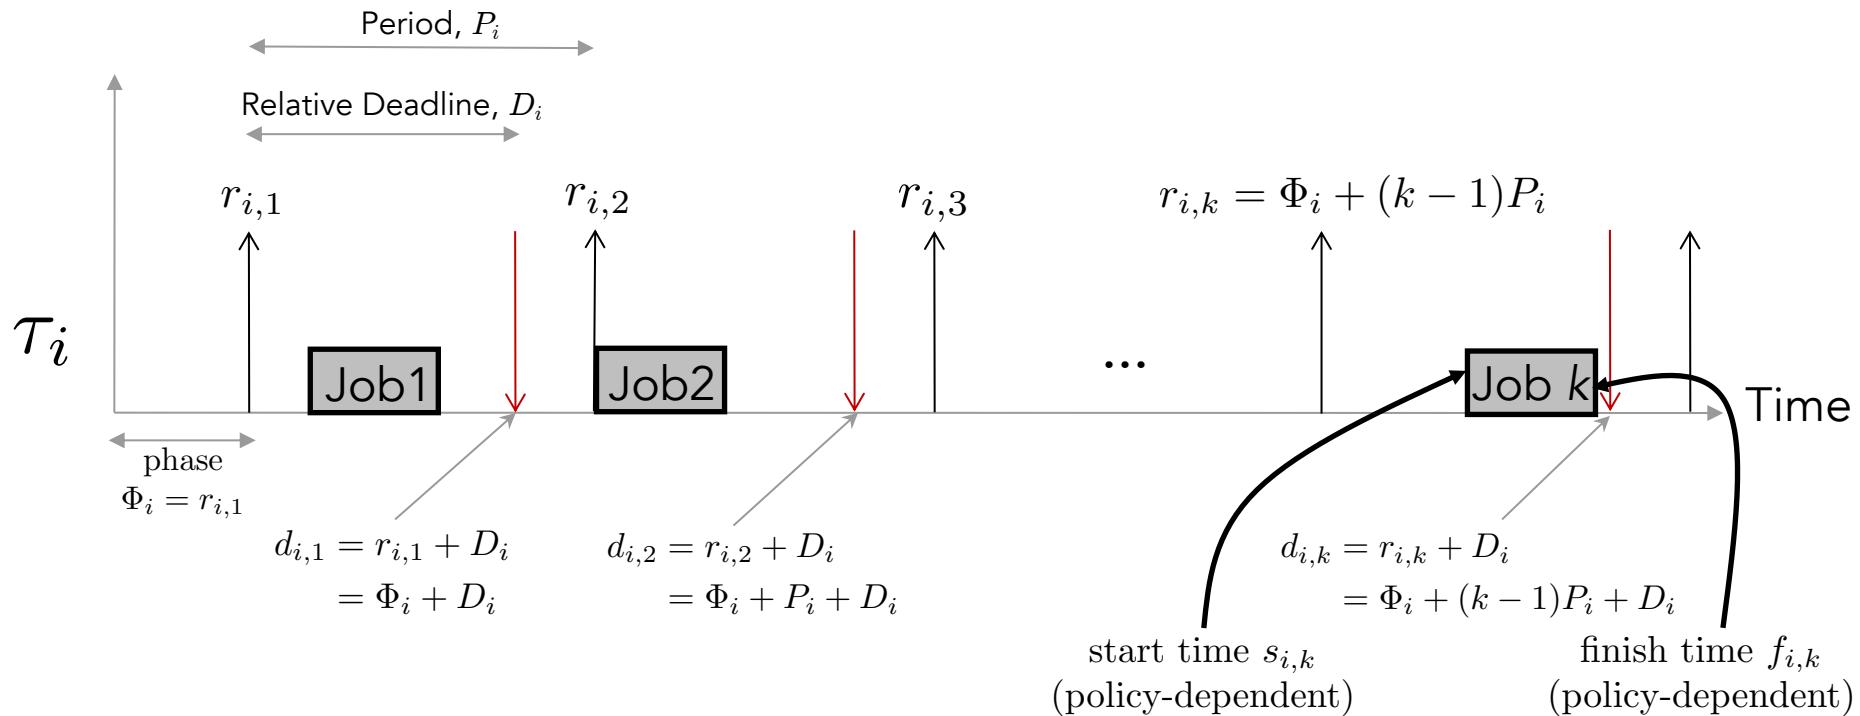


# Some terminology

Tasks, periods, arrival-time, deadline, execution time, etc.

Each invocation of a task is called a "job."

A common assumption is that arrival times for the first job of all tasks is 0.

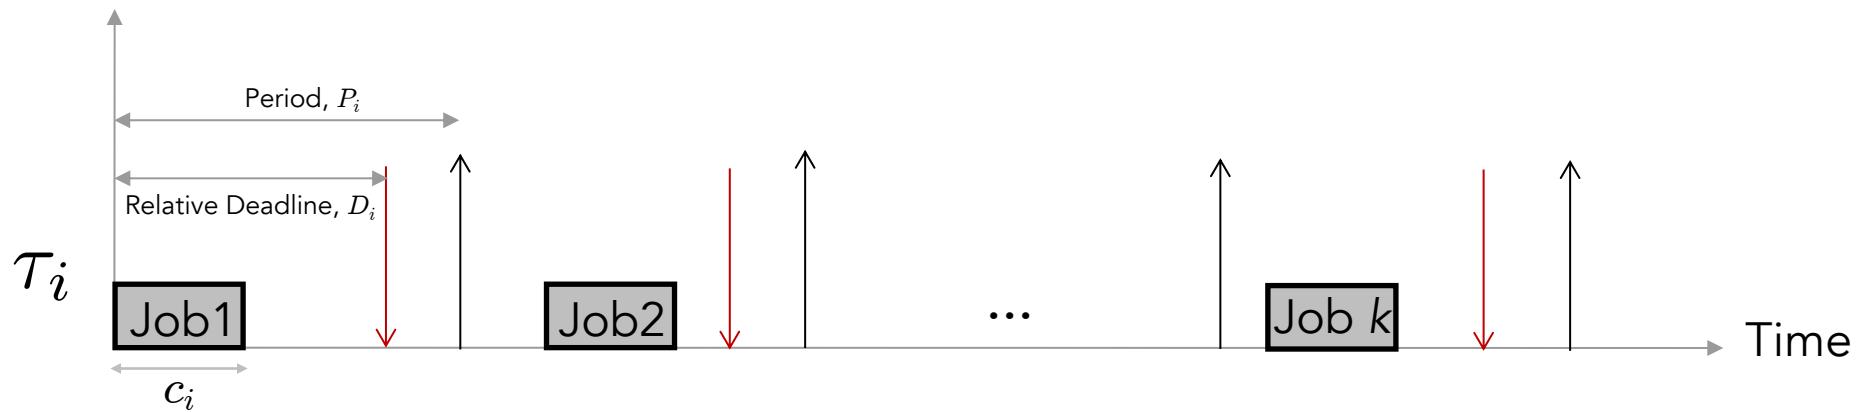


# Some terminology

Tasks, periods, arrival-time, deadline, execution time, etc.

Each invocation of a task is called a “job.”

A common assumption is that arrival times for the first job of all tasks is 0.

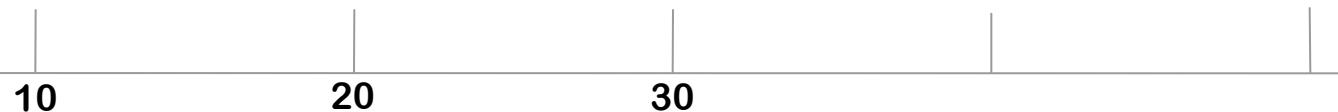


# Back to the Drive-by-Wire example

---

- Find a schedule that makes sure all task invocations meet their deadlines
- Often, **relative deadlines are equal to the period lengths (implicit-deadline)**
- **Here each task's phase is 0**

Steering wheel task (4.5 ms every 10 ms)



Brakes task (2 ms every 4 ms)



Velocity control task (0.45 ms every 15 ms)



# Back to the Drive-by-Wire example

---

- Sanity check #1: Is the processor over-utilized? (e.g., if you have 5 assignments due this time tomorrow and each takes 6 hours, then  $5 \times 6 = 30 > 24 \rightarrow$  you are overutilized)
  - Hint: Check if processor utilization > 100%

Steering wheel task (4.5 ms every 10 ms)



Brakes task (2 ms every 4 ms)



Velocity control task (0.45 ms every 15 ms)



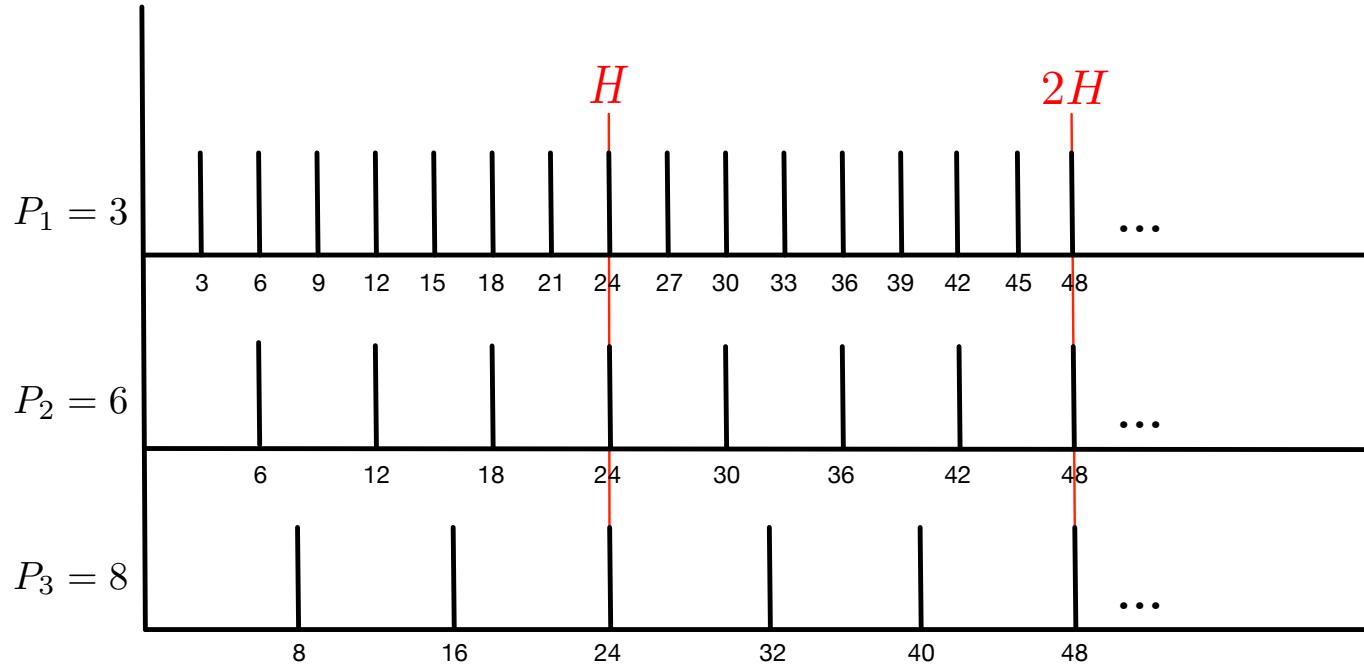
# Utilization of a task set

---

- For a set of tasks  $\{T_i\}$  with execution times  $\{e_i\}$  and periods  $\{P_i\}$ , the utilization,  $U$ , is the fraction of time, in the long run, for which the task set will use the system.

$$U = \sum_i \frac{e_i}{P_i}$$

# Why care about task utilization factor $U = \sum_i \frac{e_i}{P_i}$ ?



When does the **collective (system-wide)** pattern of arrivals repeat itself?

When a task arrives it *requests*  $c_i$ :  $\text{request}_i([0, t]) = \left\lceil \frac{t}{P_i} \right\rceil c_i$

[# of requests increases by 1 at the *right limit* of  $kP_i$ ]

Requests become *demands* that must be completed at absolute deadline boundaries  
[actually at the *left limit* of  $kP_i$ , just before the next invocation arrives]

$$\text{demand}_i([0, t]) = \left\lfloor \frac{t}{P_i} \right\rfloor c_i$$

At hyperperiod boundaries:  $\text{demand}_i([0, t]) = \text{request}_i([0, t]) = H/P_i$

# Task scheduling: Clock-driven scheduling

- In what order should tasks be executed?
  - Hand-crafted schedule (fill timeline by hand)
  - Cyclic executive scheduling

Steering wheel task (4.5 ms every 10 ms)



Brakes task (2 ms every 4 ms)



Velocity control task (0.45 ms every 15 ms)



# Task scheduling: Clock-driven scheduling

- Cyclic executive scheduling
  - Why is it called a “cyclic” executive?
  - What are the problems with cyclic executive scheduling?
    - Hard to adjust the schedule if tasks change
    - Difficult to specify

Steering wheel task (4.5 ms every 10 ms)



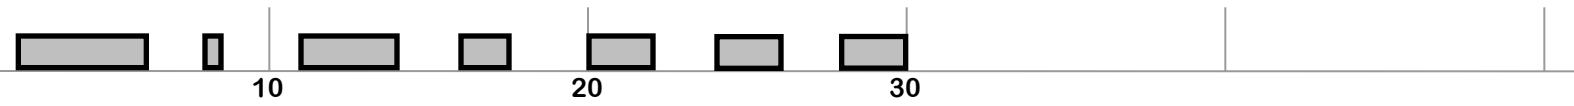
Brakes task (2 ms every 4 ms)



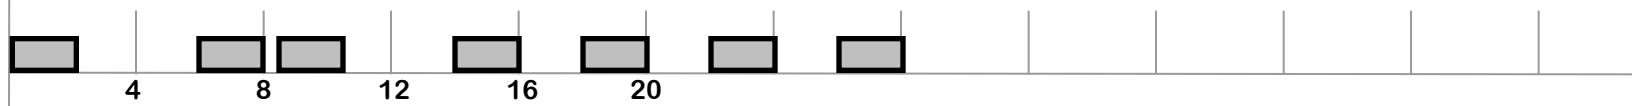
Velocity control task (0.45 ms every 15 ms)



Steering wheel task (4.5 ms every  $P_1 = 10$  ms)

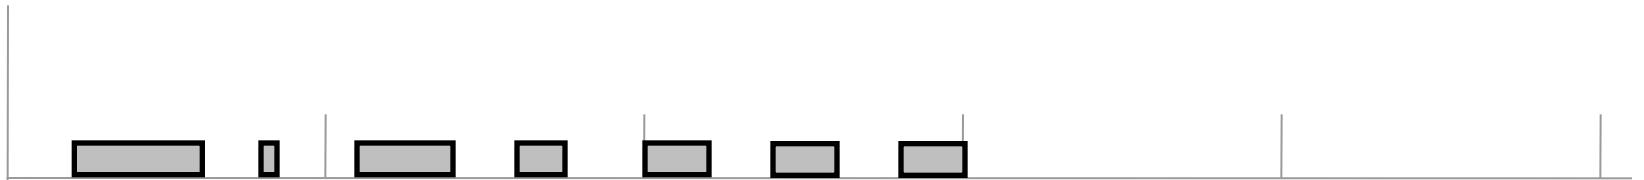


Brakes task (2 ms every  $P_2 = 4$  ms)

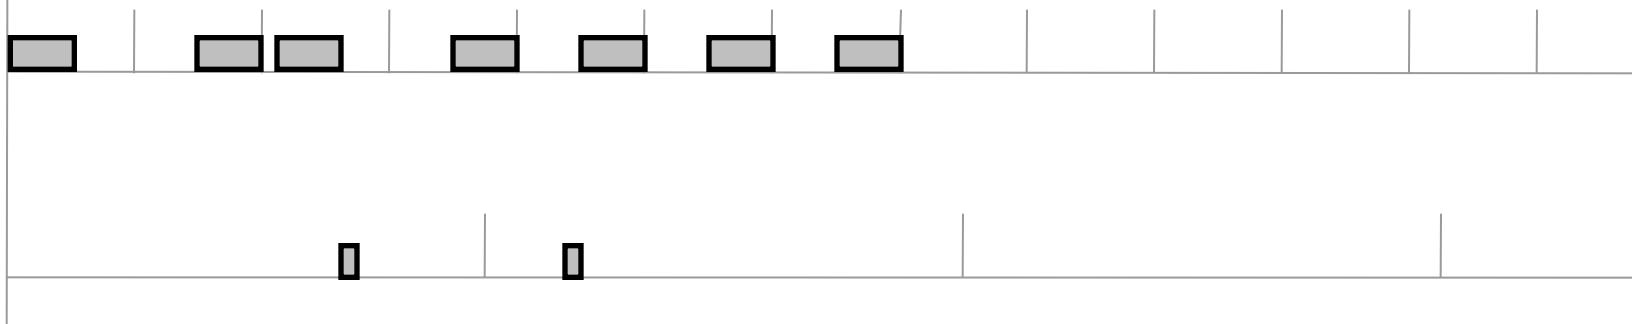


Velocity control task (0.45 ms every  $P_3 = 15$  ms)



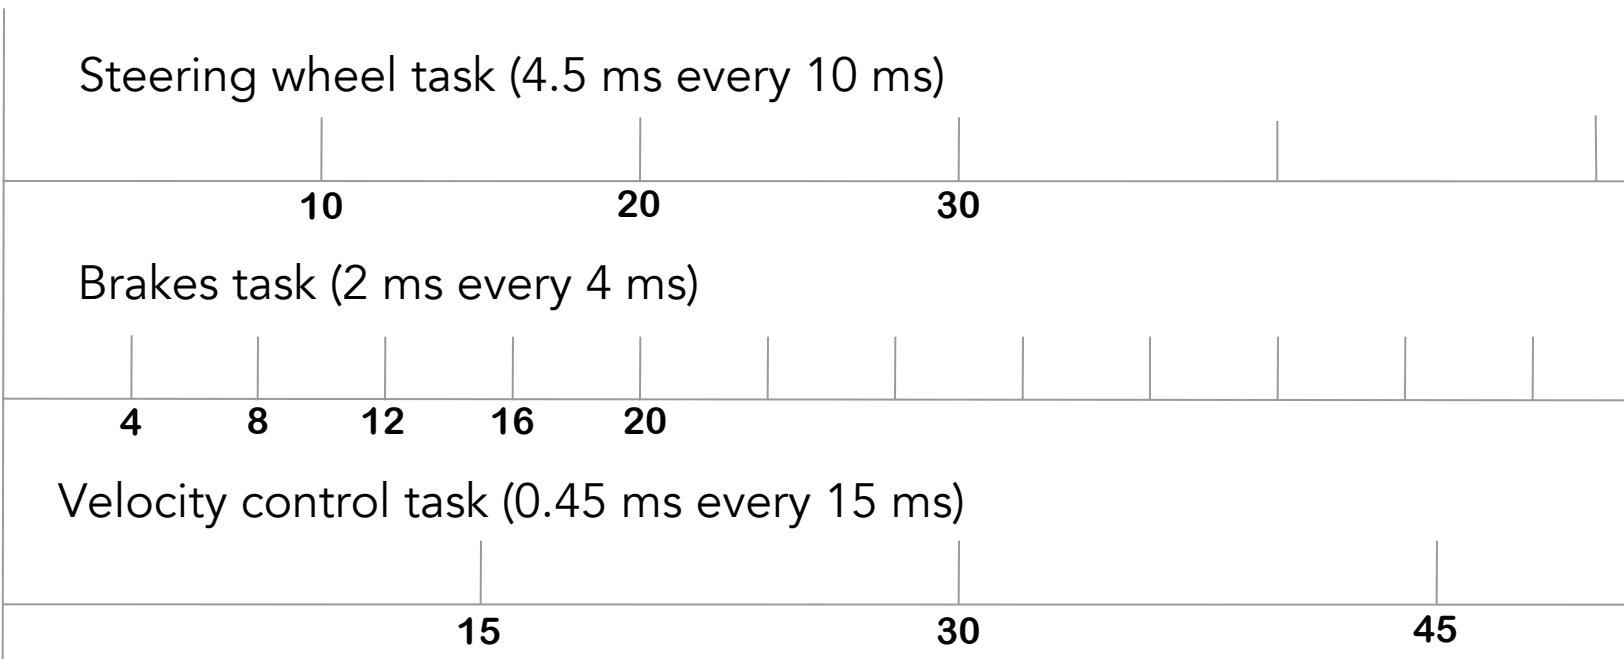


Brakes task (2 ms every 4 ms)



# Task scheduling

- In what order should tasks be executed?
  - Cyclic executive scheduling or
  - Priority based schedule (assign priorities; schedule is implied)

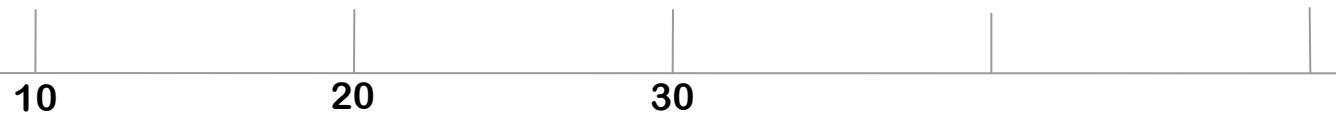


**Intuition: Urgent tasks should be higher in priority**

## Task scheduling: Preemptive versus non-preemptive

- Preemptive: Higher-priority tasks can interrupt lower-priority ones
- Non-preemptive: Can the tasks meet deadlines?

Steering wheel task (4.5 ms every 10 ms)



Brakes task (2 ms every 4 ms)



Velocity control task (0.45 ms every 15 ms)



In this example, will non-preemptive scheduling work?

Hint: Compare relative deadlines of tasks to execution times of others

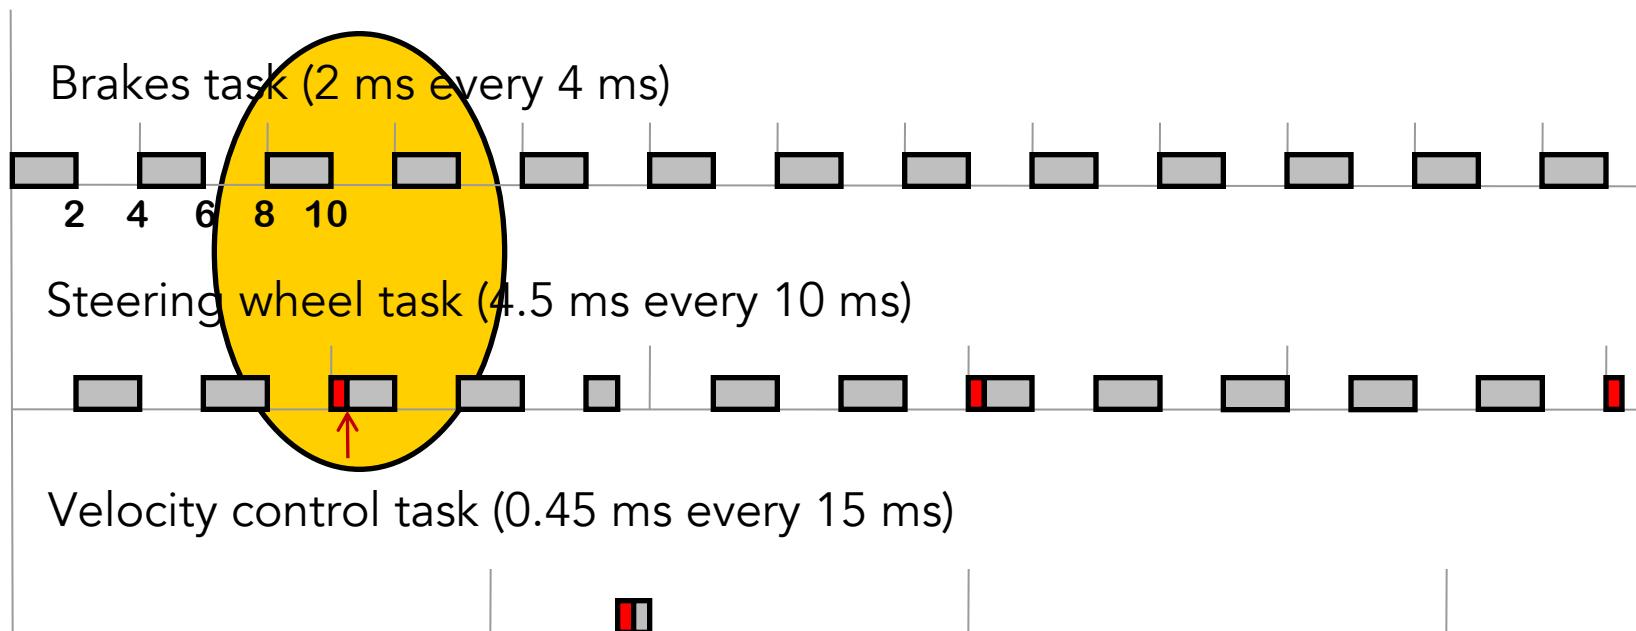
# Timeline

- Even with preemption, deadlines are missed!
- Average utilization < 100%



# Timeline

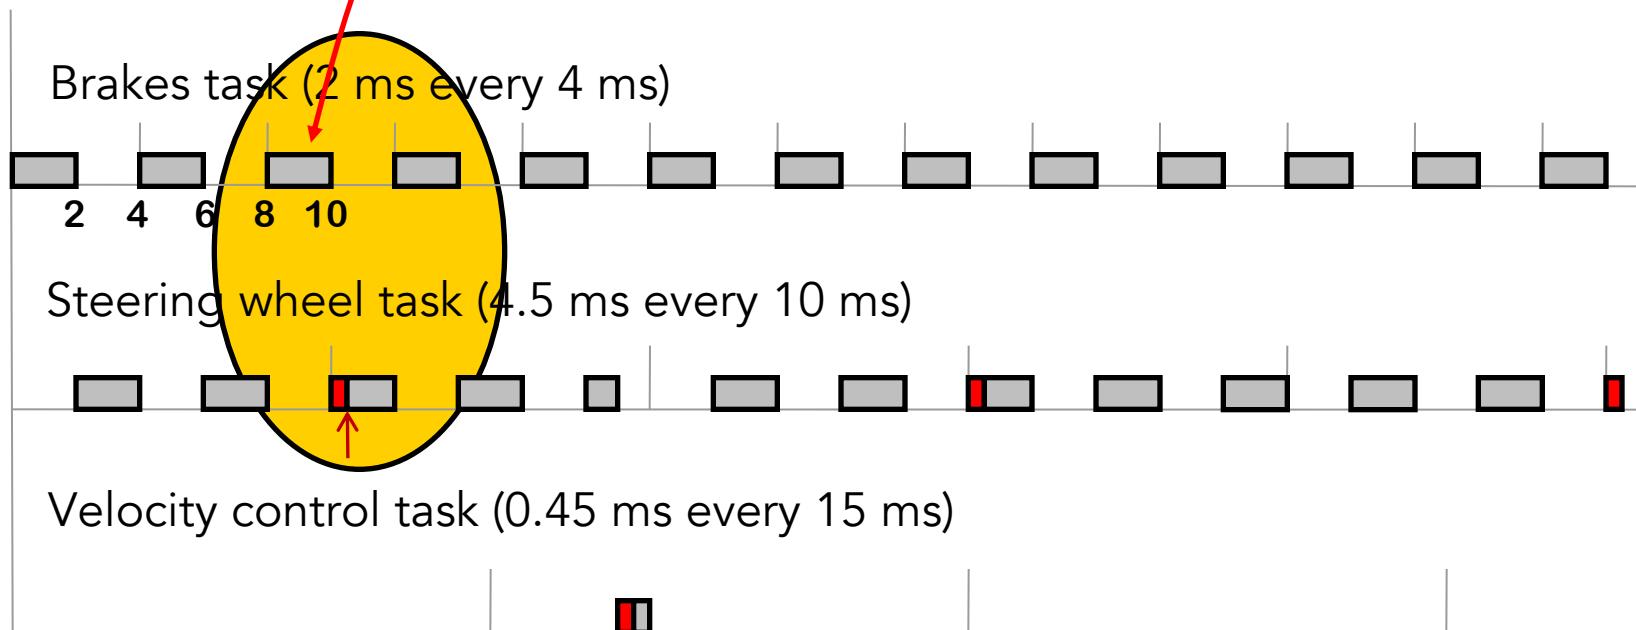
- Deadlines are missed!
- Average utilization < 100%



# Timeline

- Deadlines are missed!
- Average utilization < 100%

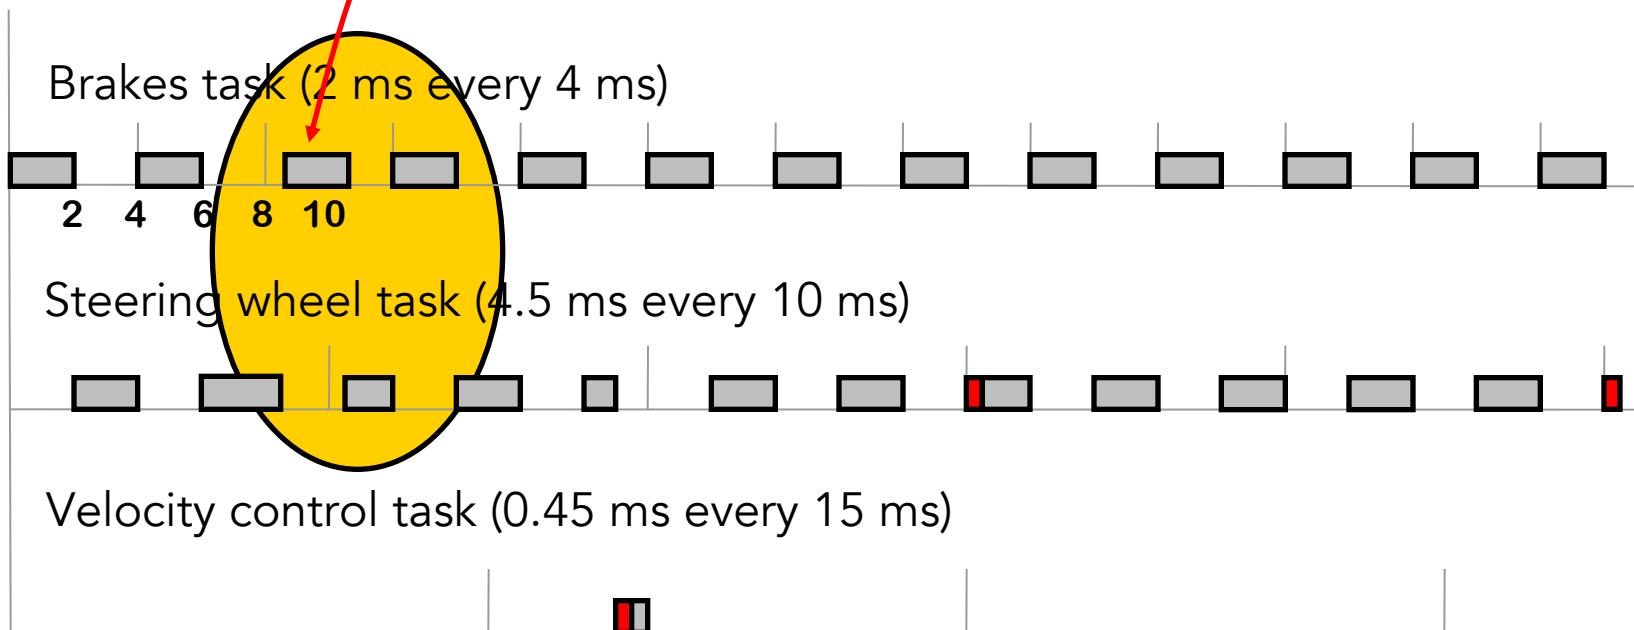
Fix:  
**Give this task invocation  
a lower priority**



# Timeline

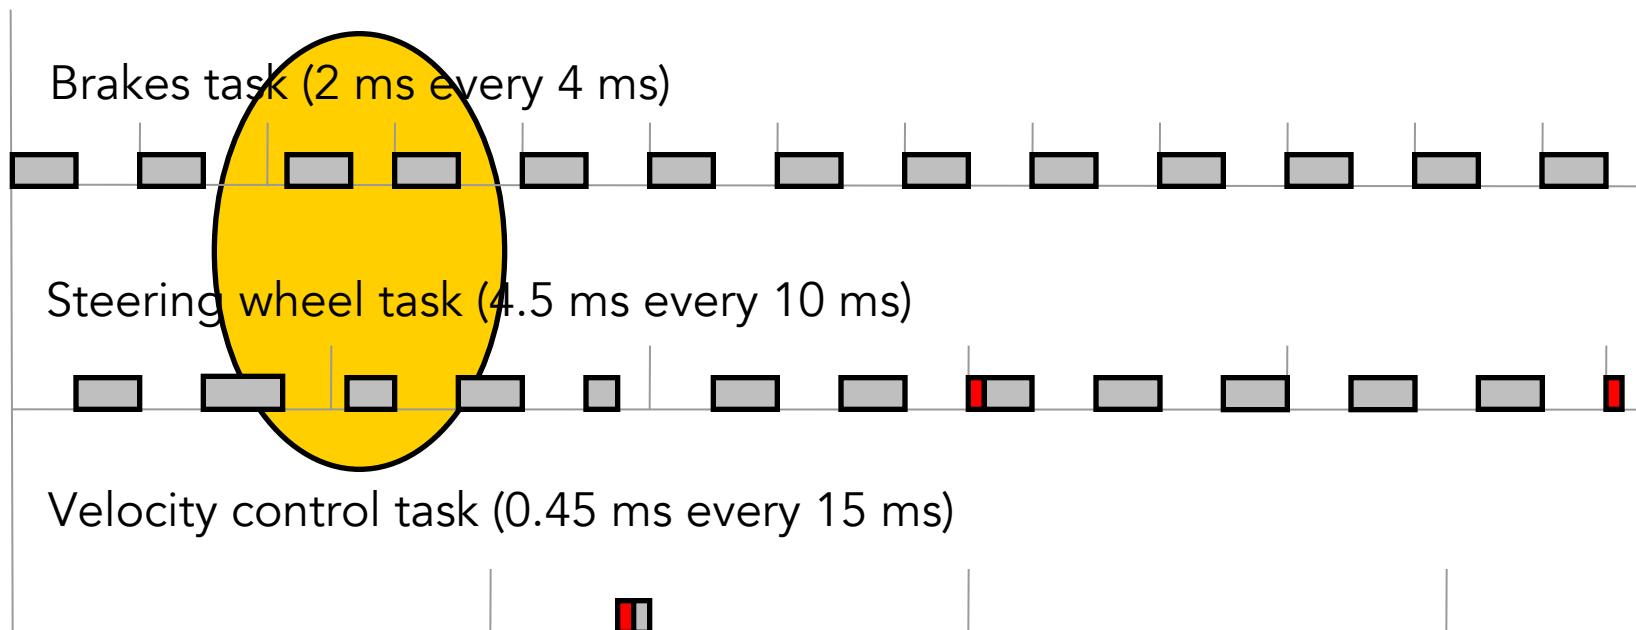
- Deadlines are missed!
- Average utilization < 100%

Fix:  
**Give this task invocation  
a lower priority**



# Task scheduling

- Static versus Dynamic priorities?
  - Static: All jobs (instances) of the same task have the same priority
  - Dynamic: Jobs (instances) of same task may have different priorities



**Intuition:** Dynamic priorities offer the designer more flexibility and hence are more capable of meeting deadlines

# Examples of policies

---

- Static-task priority policies
  - **Rate monotonic priority:** tasks with shorter periods (higher frequency or **rate**) get higher priority
  - **Deadline monotonic priority:** tasks with shorter relative deadlines get higher priority
  - Rate monotonic priorities and deadline monotonic priorities are identical if relative deadlines are equal to the periods
- Dynamic-**task** (static-**job**) priority policies
  - **Earliest deadline first:** jobs with the earliest absolute deadline take highest priority
  - **First In, First Out:** jobs with earliest arrival time take highest priority
- Dynamic-**job** priority policies: priority of **job** might change through time
  - **Shortest Remaining Processing Time (SRPT), Least Laxity First (LLF)**

# Interesting questions

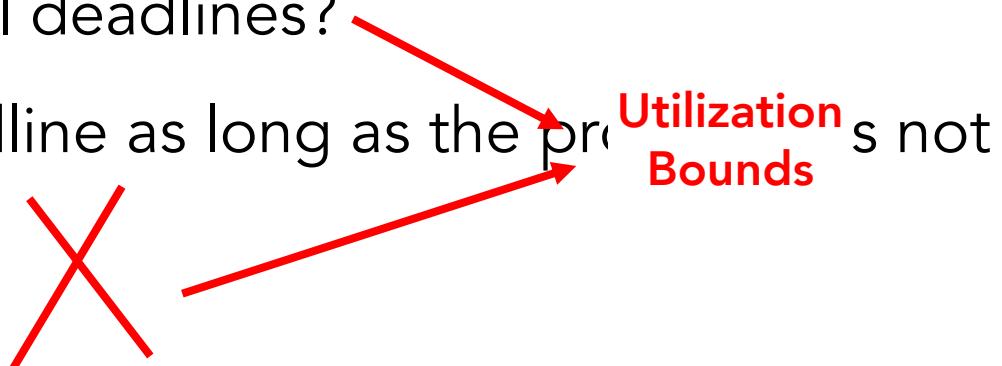
---

- What is the optimal dynamic priority scheduling policy?  
(Optimal: meets all deadlines as long as any other policy in its class can)
  - Can it meet all deadlines as long as the processor is not over-utilized?  $[U \leq 1]$
- What is the optimal static priority scheduling policy?
  - When can it meet all deadlines?
  - Can it meet all deadline as long as the processor is not over-utilized?

# Interesting questions

---

- What is the optimal dynamic priority scheduling policy?  
(Optimal: meets all deadlines as long as any other policy in its class can)
  - Can it meet all deadlines as long as the processor is not over-utilized?  $[U \leq 1]$
- What is the optimal static priority scheduling policy?
  - When can it meet all deadlines?
  - Can it meet all deadline as long as the processor is not over-utilized?



# Utilization bounds for schedulability

---

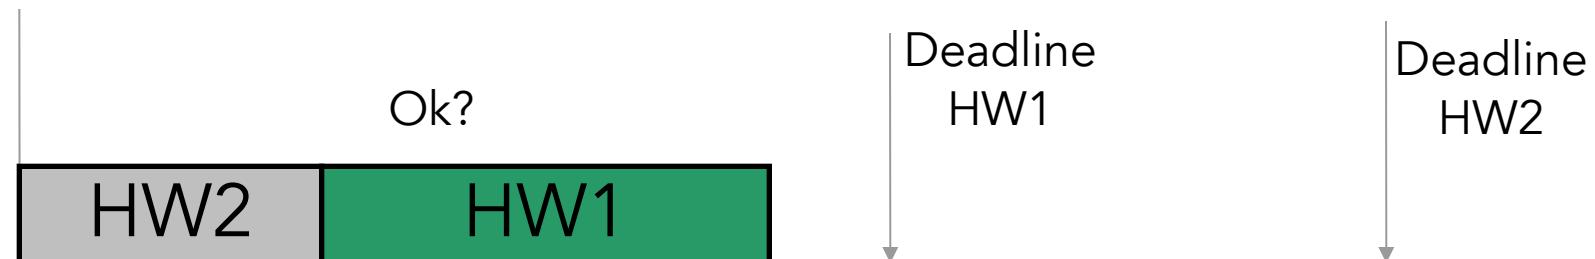
- $U_S$  is called a **utilization bound** for a given scheduling policy  $S$  if
  - All task sets with utilization factor  $\leq U_S$  can be scheduled using policy  $S$
- $U_S$  is **tight** if, in addition, for a given scheduling policy  $S$  the following holds:

*For every  $\epsilon > 0$ , there exists at least one task set with utilization  $(U_S + \epsilon)$  that cannot be scheduled using policy  $S$*
- A tight bound is the best (largest) possible utilization bound: If  $U_S$  is tight, then no other  $U > U_S$  can be a utilization bound for scheduling policy  $S$
- Of course, the maximum value that  $U_S$  can attain for any  $S$  is 1. Why? In class
- $U_S$  is also called the **schedulable utilization** of algorithm  $S$

# Optimality result for EDF scheduling

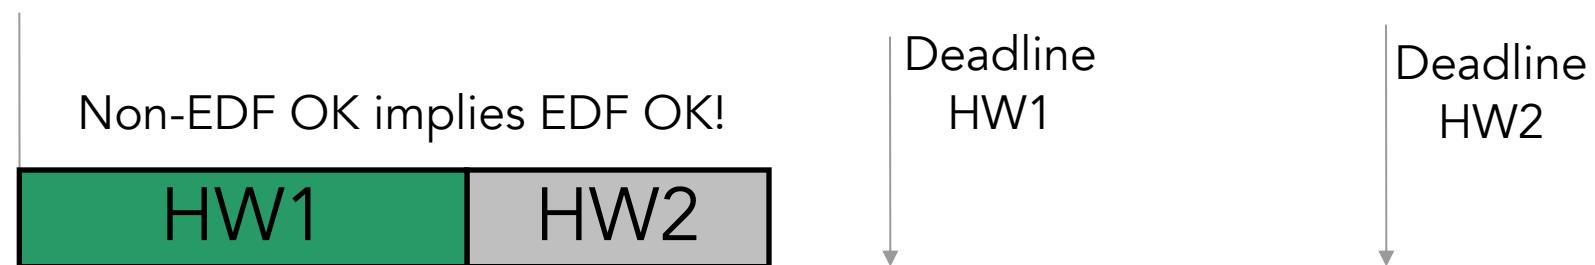
---

- EDF is the optimal dynamic priority scheduling policy
  - Priorities correspond to absolute deadlines
  - It can meet all deadlines whenever the processor utilization is less than 100%
  - Intuition:
    - You have HW1 due tomorrow and HW2 due the day after, which one do you do first?
    - If you started with HW2 and met both deadlines you could have started with HW1 (in EDF order) and still met both deadlines
    - EDF can meet deadlines whenever anyone else can



# Optimality result for EDF scheduling

- EDF is the optimal dynamic priority scheduling policy
  - It can meet all deadlines whenever the processor utilization is less than 100%
  - Intuition:
    - You have HW1 due tomorrow and HW2 due the day after, which one do you do first?
    - If you started with HW2 and met both deadlines you could have started with HW1 (in EDF order) and still met both deadlines
    - EDF can meet deadlines whenever anyone else can



# Why study static-priority policies?

---

- EDF is the optimal dynamic scheduling policy and has a utilization bound of 1.
  - The utilization bound is 1 (or 100%) when tasks have periods equal to their relative deadlines.
- EDF, however, is hard to implement in most systems.
  - Complexity is high.
  - Job queues need to be reordered often (high overhead!).
  - Most hardware subsystems allow only static priorities.

# What you should know

---

- Definitions
  - Tasks
  - Task invocations
  - Release/arrival time,
  - Absolute deadline and relative deadline
  - Period
  - Start time and finish time
- Preemptive versus non-preemptive scheduling
- Priority-based scheduling
- Static versus dynamic priorities
- Utilization and Schedulability
- Optimality of EDF

## Another example

---

- You are the system administrator of Ameritrade.com (online stock trading site)
- You offer the following guarantee to your premium customers:
  - Stock trades of less than a \$100,000 value go through in 8 seconds or you charge no commission.
  - Stock trades of more than a \$100,000 value go through in 3 seconds or you charge no commission.
- Non-premium customers do not enjoy these guarantees
- Your job is to ensure that the premium customers are always served within their agreed-upon maximum latencies. What needs to be done?

## For later: Aperiodic tasks

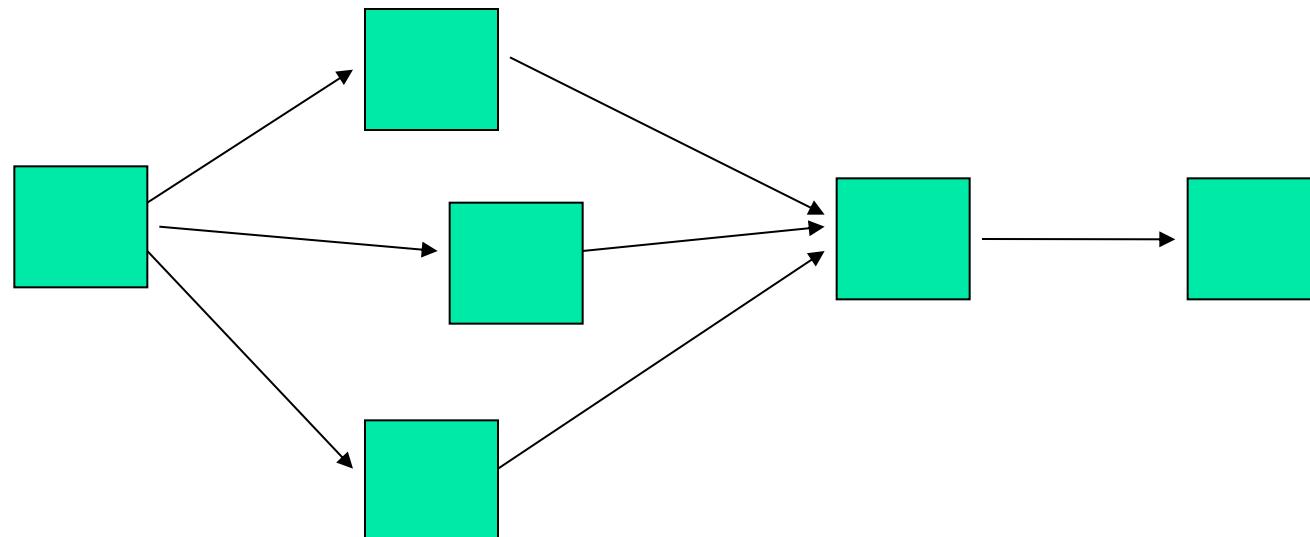
---

- Periodic tasks vs. aperiodic tasks
  - Aperiodic tasks may arrive at any time
  - Periodic tasks arrive at regular intervals [strictly  $P_i$ ]
- Sporadic tasks
  - Successive arrivals have a *minimum separation distance* [greater than or equal to  $P_i$ ].
- How does the lack of periodicity affect scheduling, and schedulability analysis?

## For later: Precedence constraints

---

- In the discussion thus far, we focused on tasks that have no dependencies.
- What if tasks have precedence constraints?
  - Tasks can execute only if their predecessors have finished execution.



# For later: Resource constraints

---

- In addition to the CPU, tasks may need resources
  - Memory
  - Disk
  - Shared data structures
- Types of resources
  - Space multiplexed: An example is the memory system. Different tasks may use different portions of the resource.
  - Time multiplexed: Only one task can access the resource at a time. An example is a data structure protected by a lock.
- How do resource constraints affect scheduling?