

# Periodic task scheduling

---

Optimality of rate monotonic scheduling (among static priority policies)

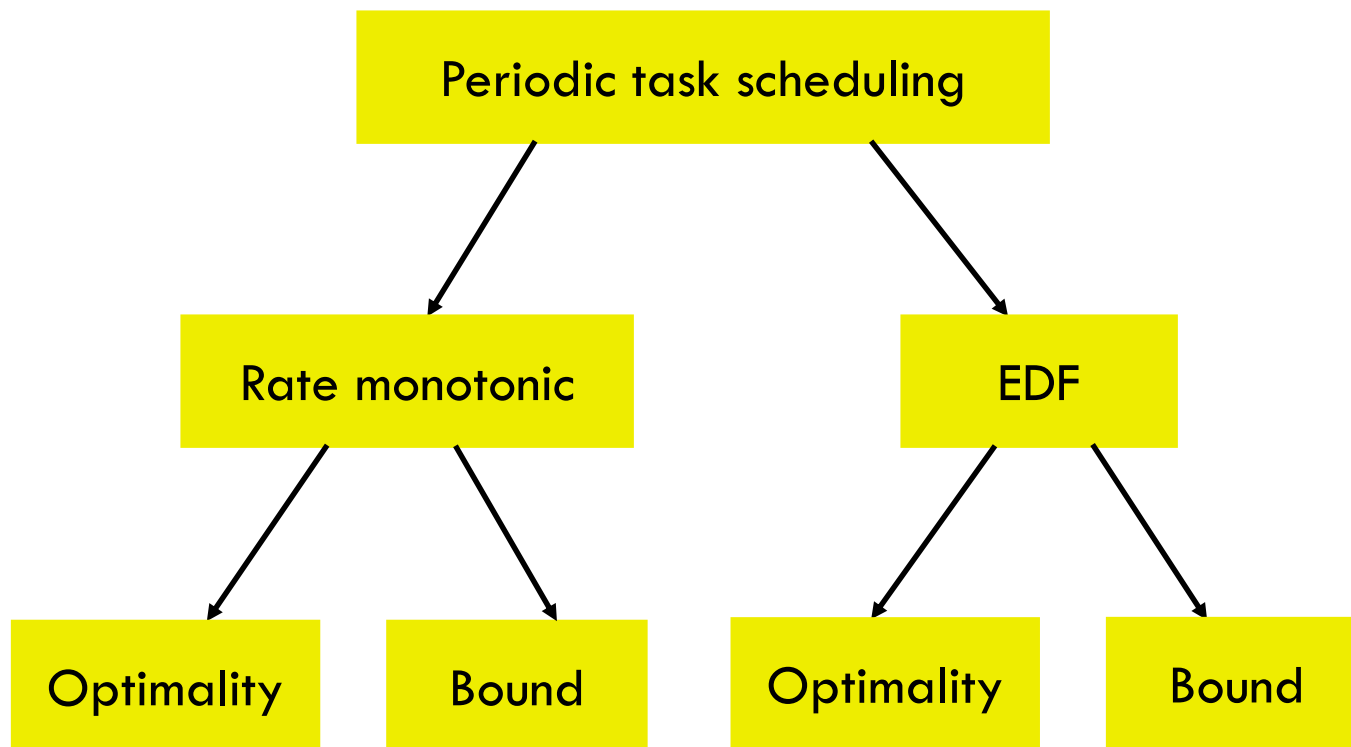
Utilization bound for EDF

Optimality of EDF (among dynamic priority policies)

Tick-driven scheduling (OS issues)

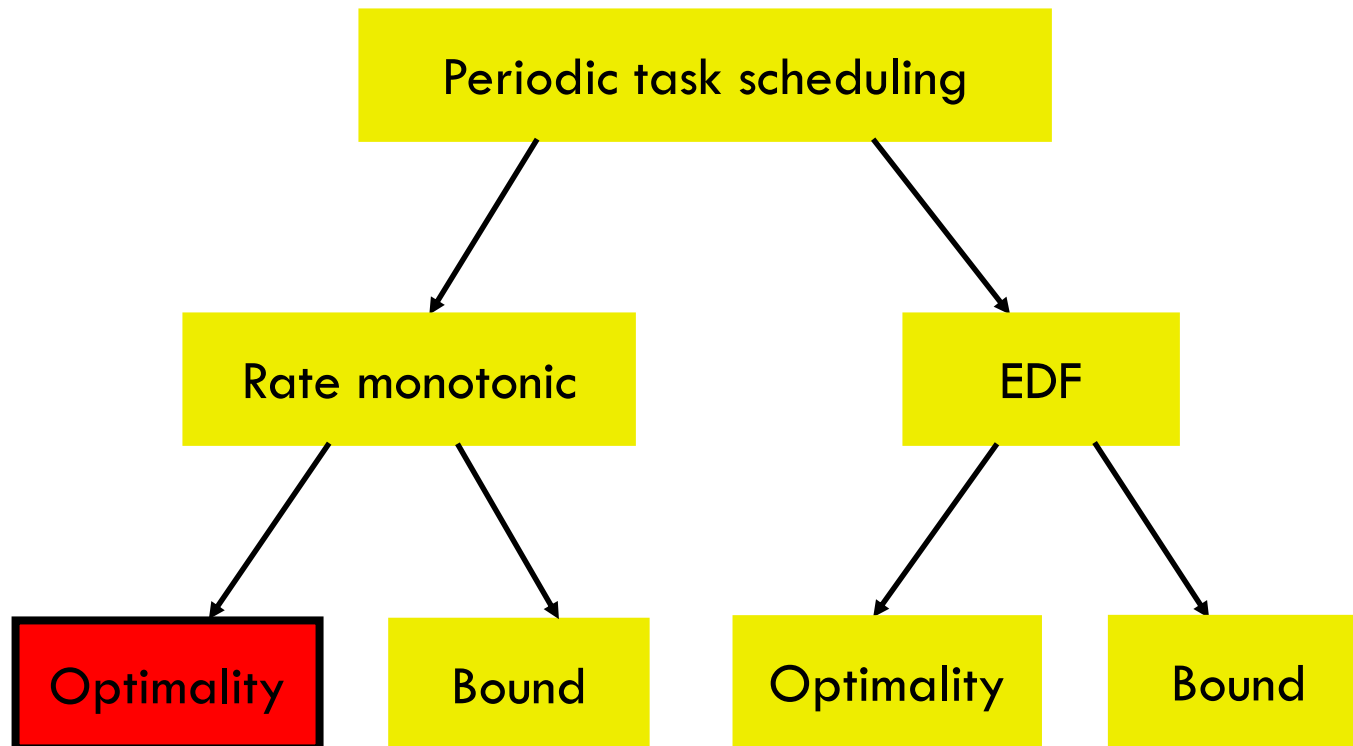
# Lecture outline

---



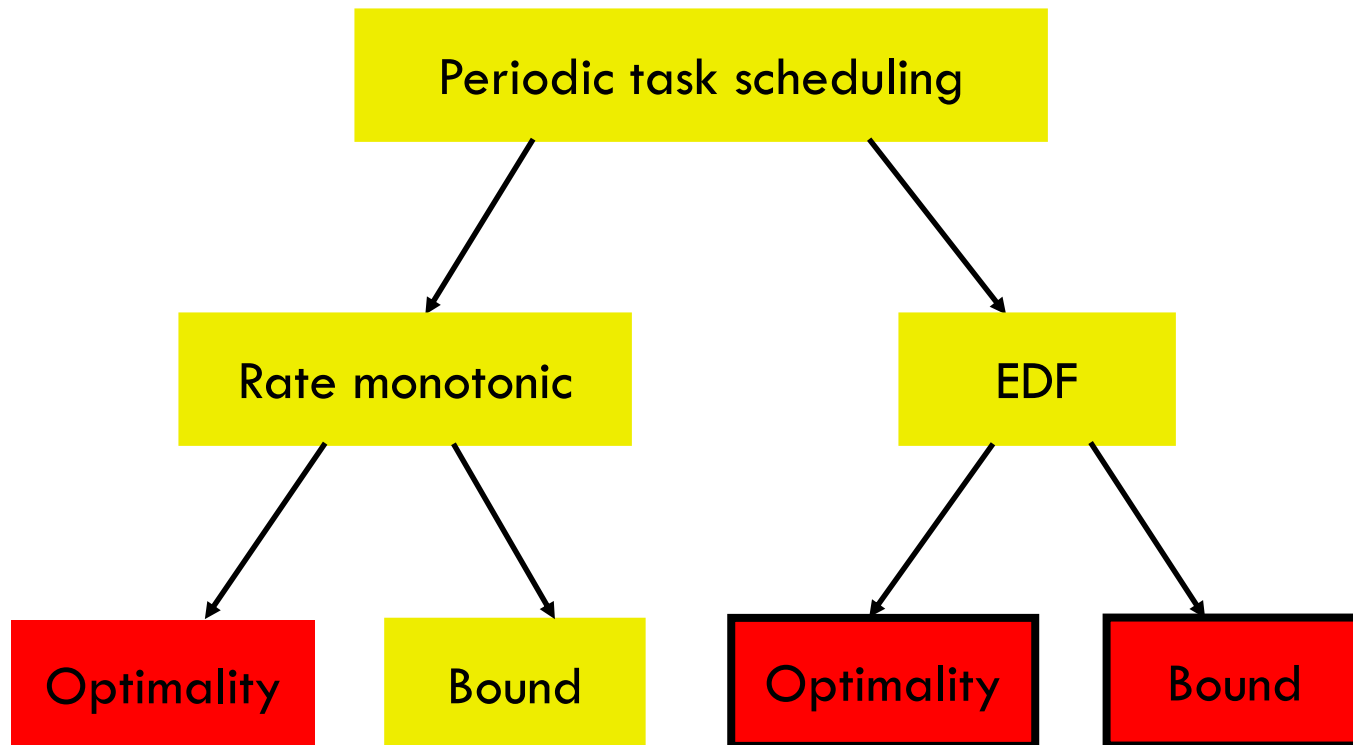
# Lecture outline

---



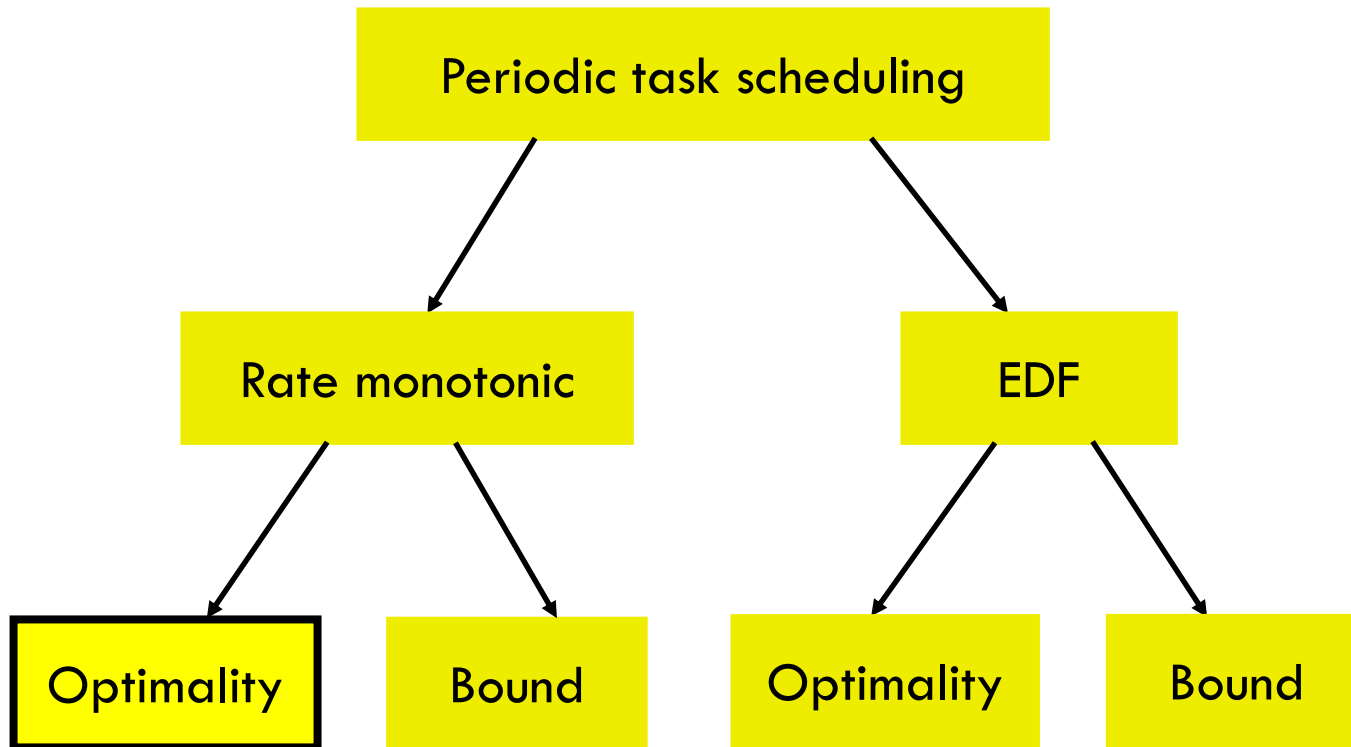
# Lecture outline

---



# Next

---



# Rate monotonic scheduling

---

- Rate monotonic scheduling is the optimal fixed-priority (or static-priority) scheduling policy for periodic tasks.
  - Optimality (Trial #1):

# Rate monotonic scheduling

---


- Rate monotonic scheduling is the optimal fixed-priority (or static-priority) scheduling policy for periodic tasks.
  - Optimality (Trial #1): If any other fixed-priority scheduling policy can meet deadlines, so can RM.

# Rate monotonic scheduling

---

- Rate monotonic scheduling is the optimal fixed-priority (or static-priority) scheduling policy for periodic tasks.
  - Optimality (Trial #1): If any other fixed-priority scheduling policy can meet deadlines, so can RM

$$P_1 < P_2$$






# Rate monotonic scheduling

---

- Rate monotonic scheduling is the optimal fixed-priority (or static-priority) scheduling policy for periodic tasks.
  - Optimality (Trial #1): If any other fixed-priority scheduling policy can meet deadlines, so can RM

$$P_1 < P_2$$




# Rate monotonic scheduling

---


- Rate monotonic scheduling is the optimal fixed-priority (or static-priority) scheduling policy for periodic tasks.
  - Optimality (Trial #1): If any other fixed-priority scheduling policy can meet deadlines, so can RM

$$P_1 < P_2$$



# Rate monotonic scheduling

- Rate monotonic scheduling is the optimal fixed-priority (or static-priority) scheduling policy for periodic tasks.
  - Optimality (Trial #1): If any other fixed-priority scheduling policy can meet deadlines, so can RM



# Rate monotonic scheduling

---

- Rate monotonic scheduling is the optimal fixed-priority (or static-priority) scheduling policy for periodic tasks.
  - Optimality (Trial #2): If any other fixed-priority scheduling policy can meet deadlines **in the worst case scenario**, so can RM.
- How do we prove it?

# Rate monotonic scheduling

---

- Rate monotonic scheduling is the optimal fixed-priority (or static-priority) scheduling policy for periodic tasks.
  - Optimality (Trial #2): If any other fixed-priority scheduling policy can meet deadlines **in the worst case scenario**, so can RM.
- How do we prove it?
  - Consider the worst case scenario
  - Show that if someone else can schedule then RM can

# The worst-case scenario


---

- **Q:** When does a periodic task,  $T$ , experience the maximum delay?
  - Which arrival time produces the largest *response time* for  $T$ ?
- **A:** When it arrives together with all the higher-priority tasks (critical instant)
  - Liu and Layland
- Idea for the proof
  - If some higher-priority task does not arrive together with  $T$ , aligning the arrival times can only increase the completion time of  $T$ .

Critical instant theorem

## Critical Instant: Proof (Case 1)


---



Case 1: Higher priority task 1 is running when task 2 arrives.

## Critical Instant: Proof


---




Case 1: higher priority task 1 is running when task 2 arrives  
→ shifting task 1 right will increase completion time of 2



## Critical Instant: Proof




Case 1: higher priority task 1 is running when task 2 arrives  
→ shifting task 1 right will increase completion time of 2



## Critical Instant: Proof (Case 2)


---



Case 2: processor is idle when task 2 arrives


## Critical Instant: Proof (Case 2)

---




Case 2: processor is idle when task 2 arrives  
→ shifting task 1 left cannot decrease completion time of 2

## Critical Instant: Proof (Case 2)



Case 2: processor is idle when task 2 arrives  
→ shifting task 1 left cannot decrease completion time of task 2



# Critical Instant: Remarks

- *All analyses hereafter will assume the critical instant theorem in effect*
- *Why is it important to identify the critical instant?*
  - *Characterizes the worst case scenario when a task experiences the max delay (remember the pitfall in trial #1 in proving RM optimality earlier?)*
  - *For task schedulability, need to reason only about the feasibility of the job arriving at the critical instant*

## Critical Instant: Remarks

- If task **phases** are **not all** 0, does there always exist a point of simultaneous release?  
Can you come up with a counterexample?
  - It does not necessarily exist in this case and a simple counterexample of 3 tasks exists
- If not, then how easy it is to determine whether a point of simultaneous release exists for non-zero phase task sets?
  - An algorithm exists! (naïve approach takes exponential time, however)

## Critical Instant: Remarks

- Is RM optimal if a simultaneous release does not exist?
- What about EDF?
  - **Hint:** study the proof of optimality of EDF later and see whether the critical instant theorem was used.
- Does the critical instant theorem still hold in **non-preemptive** scheduling? (Assuming a simultaneous release exists)

# Assumptions


- All scheduling is preemptive
- A simultaneous release exists even if tasks have non-zero phases
  - And thus critical instant theorem assumed
- Implicit deadlines (deadline = period)
- A task does not suspend itself (on I/O, for instance)
- All tasks in a task set are *independent* (there are no precedence relations and no resource constraints.)
- All overheads in the kernel are assumed to be zero (context switching and others)



# Optimality of the RM policy

---

- If any other fixed-priority policy can meet deadlines so can RM




Policy X meets deadlines?

# Optimality of the RM policy

---


- If any other policy can meet deadlines so can RM



Policy X meets deadlines?


**YES**

→ RM meets deadlines



# Optimality of the RM policy: Proof

---



**Two tasks scheduled not according to RM**

For feasibility in a non-RM policy, we need  $C_1 + C_2 \leq P_1$  to hold at critical instant

Why?

# Optimality of the RM policy: Proof

---


For feasibility in a non-RM policy, we need  $C_1 + C_2 \leq P_1$  to hold at critical instant

- Now **exchange priorities** of tasks to make it into an RM assignment
- Plan:
  - identify all possible cases
  - In each case derive feasibility condition
  - Show that if  $C_1 + C_2 \leq P_1$  then derived feasibility condition in RM holds

# Optimality of the RM policy: Case 1

For feasibility in a non-RM policy, we need  $C_1 + C_2 \leq P_1$  to hold at critical instant


- **Case 1:** computation time of  $\tau_1$  is short enough that all its *requests* are completed before the second request of  $\tau_2$ 
  - Number of periods  $P_1$  **entirely contained** in  $P_2$  is  $\left\lfloor \frac{P_2}{P_1} \right\rfloor \rightarrow$  Let  $F = \left\lfloor \frac{P_2}{P_1} \right\rfloor$
  - **Case 1** translates to  $C_1 + FP_1 \leq P_2$
  - **Feasibility:** All computation requested by  $\tau_1$  during  $P_2$ , in addition to  $C_2$ , should be completed by  $P_2$
  - $(F + 1)C_1 + C_2 \leq P_2 \quad (*)$
  - Need to show that  $C_1 + C_2 \leq P_1$  implies  $(*)$



# Optimality of the RM policy: Case 2

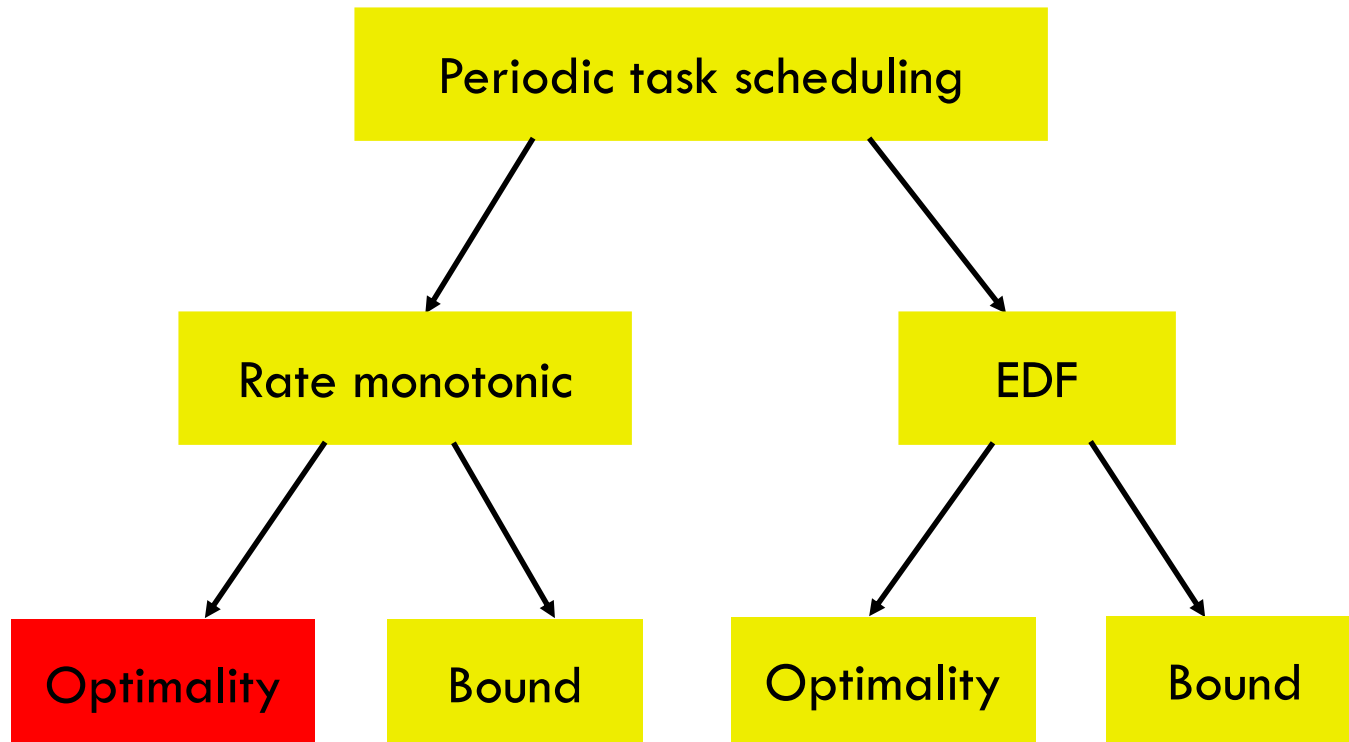
For feasibility in a non-RM policy, we need  $C_1 + C_2 \leq P_1$  to hold at critical instant

- **Case 2:** computation time of  $\tau_1$  is long enough to overlap with the second request of  $\tau_2$ 
  - **Case 2 translates to**  $C_1 + FP_1 \geq P_2$
  - **Feasibility condition:**  $FC_1 + C_2 \leq FP_1$  (\*\*) Why?
    - For feasibility  $\tau_2$  must finish before the start of the  $(FP_1)$ -th request of  $\tau_1$  because  $\tau_1$  has higher priority than  $\tau_2$  so  $\tau_1$  will occupy the processor until  $P_2$  by the condition in **case 2** and thus  $P_2$  cannot execute in  $[FP_1, P_2]$
- Need to show that  $C_1 + C_2 \leq P_1$  implies (\*\*)



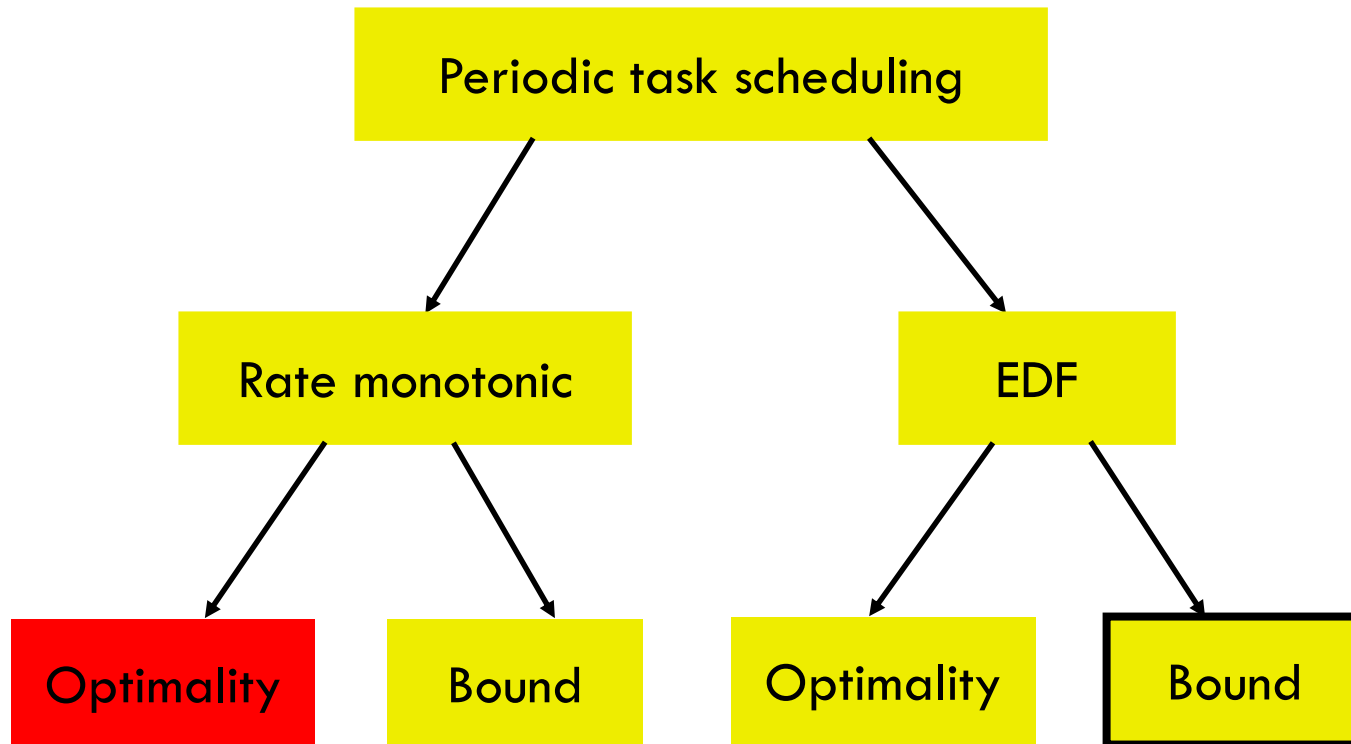
# What have we achieved?

---



# Next

---





## Recall: Utilization bounds for schedulability

---

- $U_S$  is called a **utilization bound** for a given scheduling policy  $S$  if

All task sets with utilization factor  $\leq U_S$  can be scheduled using policy  $S$

- $U_S$  is **tight** if, in addition, for a given scheduling policy  $S$  the following holds:

*For every  $\epsilon > 0$ , there exists at least one task set with utilization  $(U_S + \epsilon)$  that **cannot** be scheduled using policy  $S$*

- A tight bound is the best (largest) possible utilization bound: If  $U_S$  is tight, then no other  $U > U_S$  can be a utilization bound for scheduling policy  $S$
- Of course, the maximum value that  $U_S$  can attain for any  $S$  is 1. Why? In class
- $U_S$  is also called the **schedulable utilization** of algorithm  $S$


# Utilization bound for EDF

---

- Why is it 100%?
- Consider a task set where:


$$\sum_i \frac{C_i}{P_i} = 1$$

- Imagine a policy that reserves for each task  $i$  a fraction  $u_i$  of each clock tick, where  $u_i = C_i / P_i$



# Utilization bound for EDF

- Imagine a policy that reserves for each  $T_i$  a fraction  $u_i$  of each time unit, where  $u_i = C_i/P_i$
- Divide time into, say,  $L = \text{GCD}(P_1, \dots, P_n)$ -length ticks after proper scaling of periods to integers




- This policy meets all deadlines, because:
  - Time given to  $T_i$  in its period =  $u_i \times (\# \text{ ticks/period}) \times \text{tick length (time/tick)}$   
 $= u_i(P_i/L)L = (C_i/P_i) P_i = C_i \text{ time/period}$   
(i.e., enough to finish)

## Utilization bound for EDF


---

- Pick any two execution chunks that are not in EDF order and swap them



## Utilization bound for EDF


- Pick any two execution chunks that are not in EDF order and swap them



- Still meets deadlines! Why?

# Utilization bound for EDF

- Pick any two execution chunks that are not in EDF order and swap them



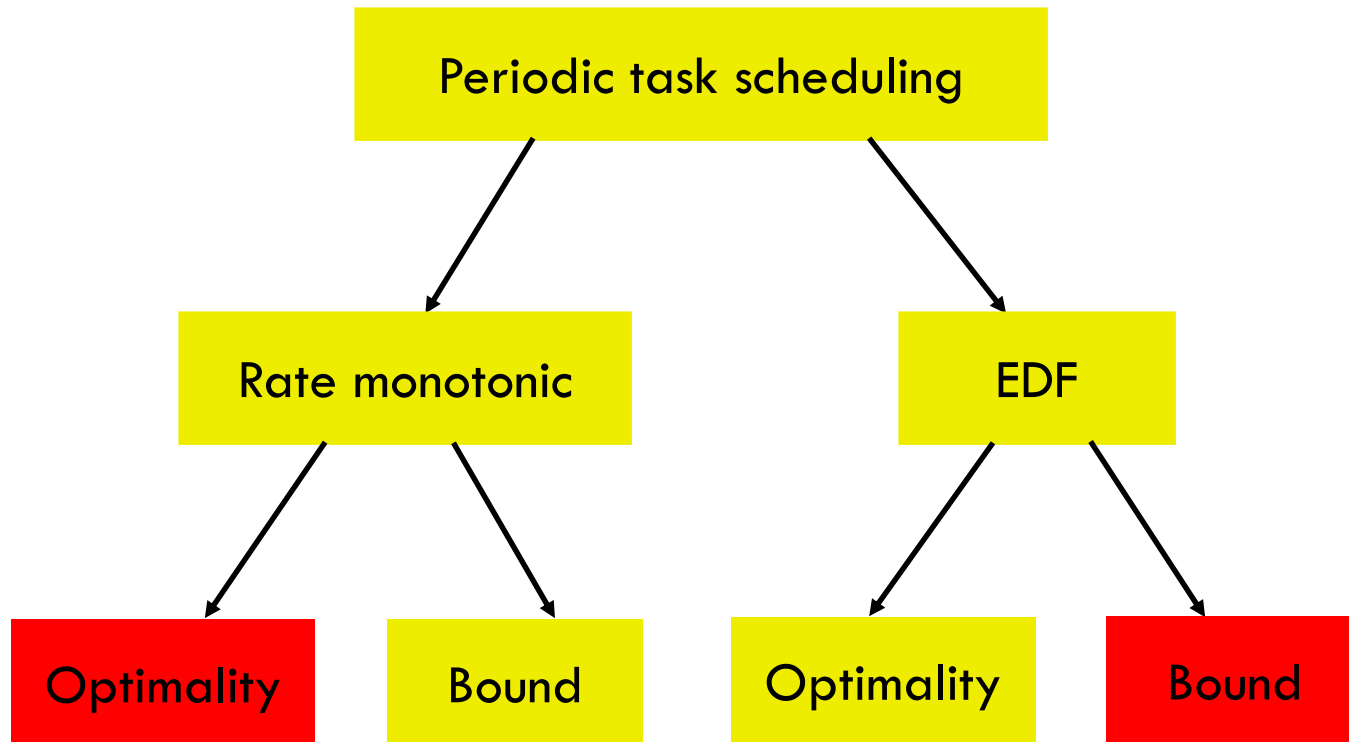
- Still meets deadlines!
- Repeat swap until all in EDF order
  - → EDF meets deadlines

# Utilization bound for EDF

- Why does this prove that the utilization bound of EDF is 1?
  - We showed that every taskset with  $U = 1$  is feasible under EDF
  - Must also show that every taskset with  $U < 1$  is also feasible under EDF
  - This is not needed! Previous argument follows for any  $U \leq 1$
- Consequences:
  - EDF is optimal!
  - EDF is able to schedule every task set who utilization is 1 or less

# Next

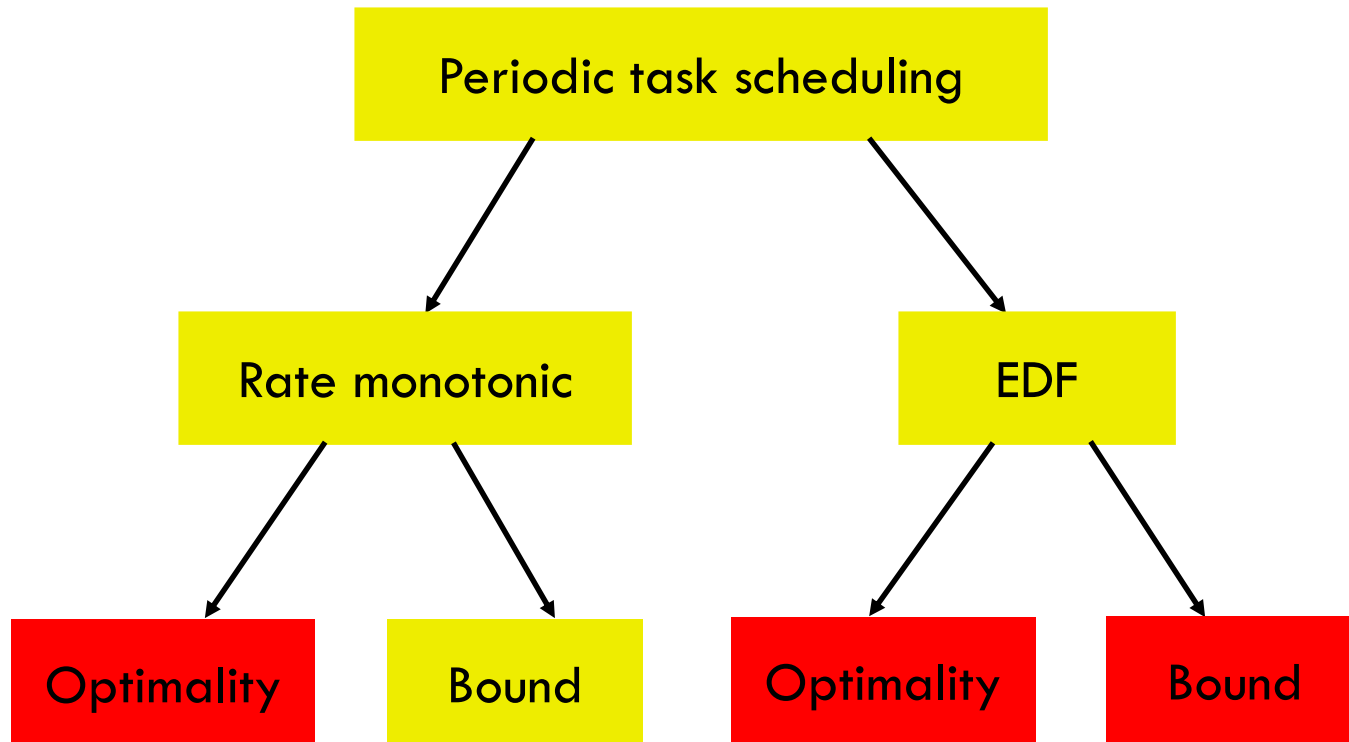
---





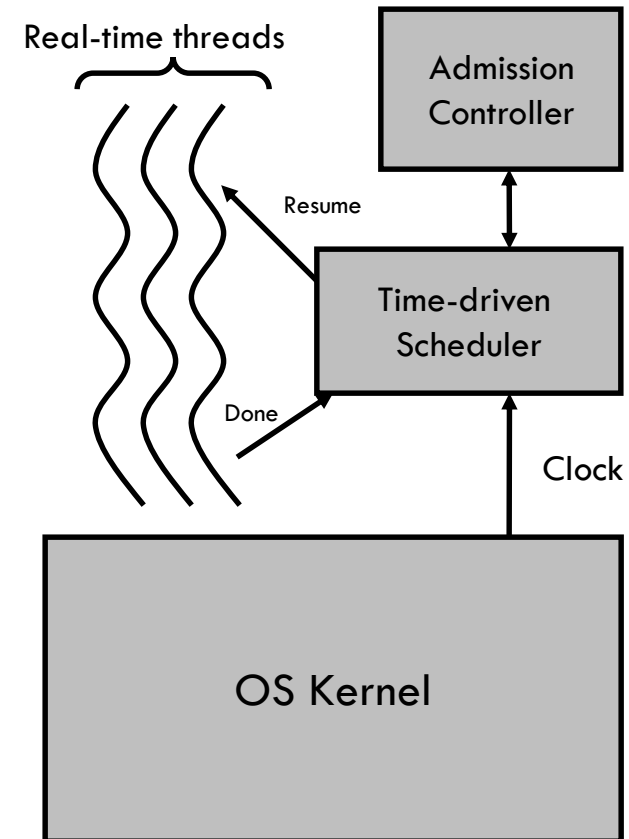
# Next

---



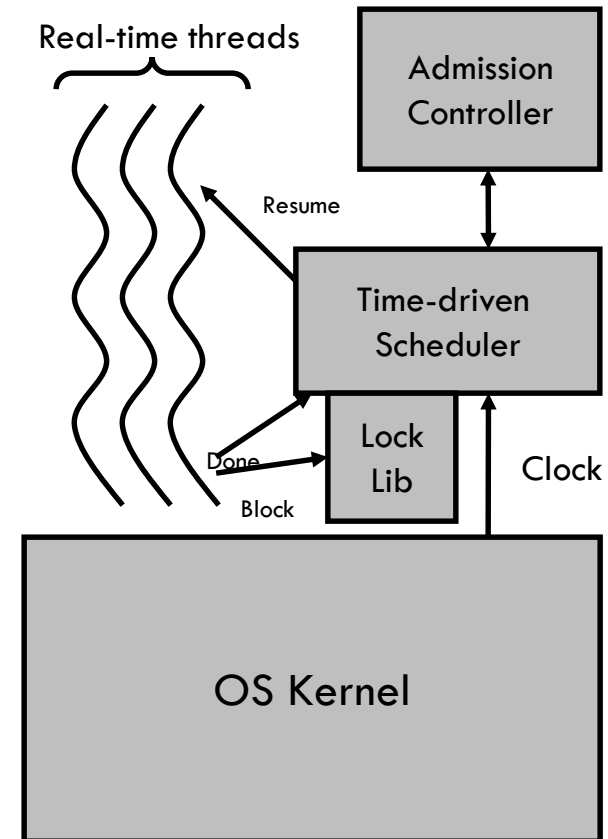
# Tick-based scheduling within an OS

- A real-time library for periodic tasks on Linux or Windows
  - There is need to provide approximate real-time guarantees on common operating systems (as opposed to specialized real-time OSes)
  - A high-priority “real-time” thread pool is created and maintained
  - A higher-priority scheduler is invoked periodically by timer-ticks to check for periodic invocation times of real-time threads. The scheduler resumes threads whose arrival times have come.
  - Resumed threads execute one invocation then block.
  - Scheduling is preemptive
  - The scheduler can implement arbitrary scheduling policies including EDF, RM, etc.
  - An admission controller is responsible for spawning new periodic threads if the new task set can meet its deadlines.



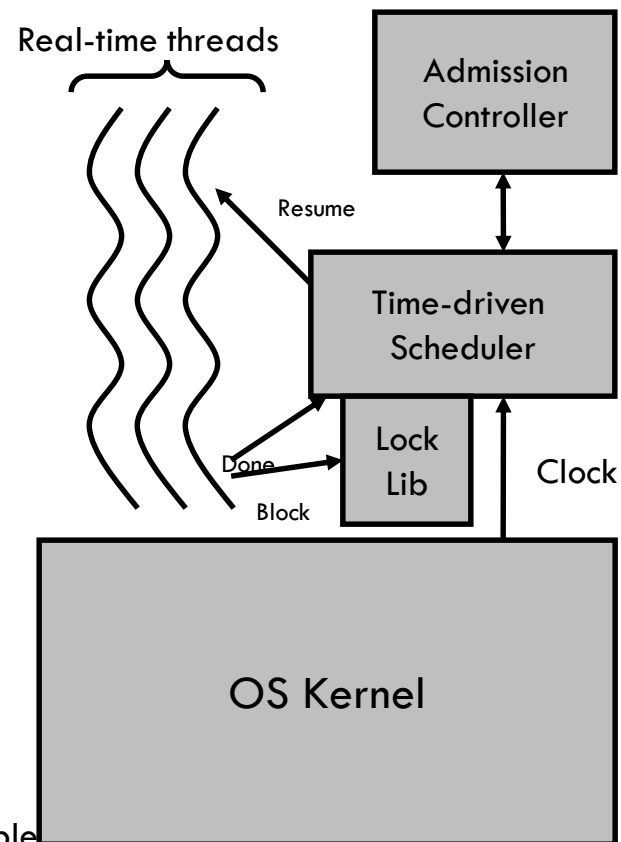
# Tick-based scheduling within an OS

- A real-time library for periodic tasks on Linux or Windows
  - There is need to provide approximate real-time guarantees on common operating systems (as opposed to specialized real-time OSes)
  - A high-priority “real-time” thread pool is created and maintained
  - A higher-priority scheduler is invoked periodically by timer-ticks to check for periodic invocation times of real-time threads. The scheduler resumes threads whose arrival times have come.
  - Resumed threads execute one invocation then block.
  - Scheduling is preemptive
  - The scheduler can implement arbitrary scheduling policies including EDF, RM, etc.
  - An admission controller is responsible for spawning new periodic threads if the new task set can meet its deadlines.
- Scheduler implements wrappers for blocking primitives




# The time-driven scheduler

- `/* N is the number of periodic tasks */`
- For  $i=1$  to  $N$
- if (`current_time = next_arrival_time` of task  $i$ )
- put task  $i$  in `ready_queue`
- `/* ready_queue is a priority queue that implements`
- the desired scheduling policy. `*/`
- Inspect top task from ready queue, call it  $j$
- If (a task is running and its priority is higher than priority of  $j$ ) return
- Else resume task  $j$  (and put the running task into the ready queue if applicable),  
return




# An example schedule


---




# An example schedule



# An example schedule




# An example schedule







# An example schedule



# An example schedule

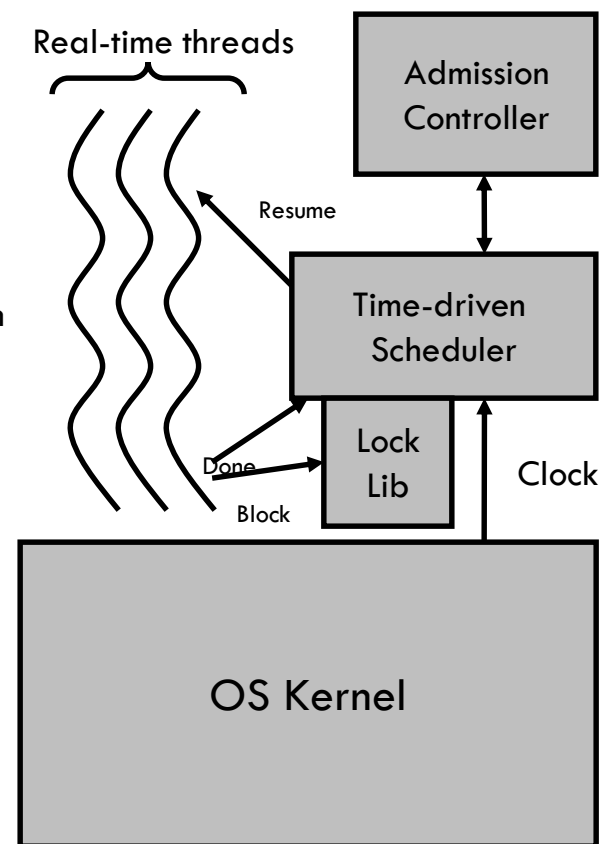


# An example schedule



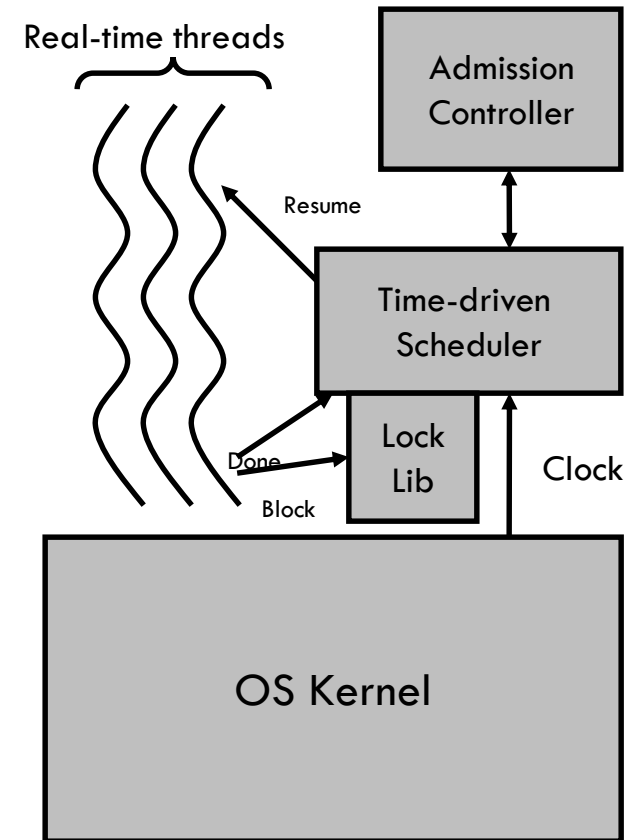
# Admission controller

- Implements schedulability analysis
  - If  $U + C_{new}/P_{new} < U_{bound}$  admit task
  - Must account for various practical overheads. How?
  - Examples of overhead:
    - How to account for the overhead of running the time-driven scheduler on every time-tick?
    - How to account for the overhead of running the scheduler after task termination?
- If new task admitted
  - $U = U + C_{new}/P_{new}$
  - Create a new thread
  - Register it with the scheduler



# Library with lock primitives

- Lock (S) {
  - Check if semaphore S = locked
  - If locked
    - enqueue running tasks in semaphore queue
  - Else
    - let semaphore = locked
- }
- Unlock (S) {
  - If semaphore queue empty then
    - semaphore = unlocked
  - Else
    - Resume highest-priority waiting task
- }



Problem: some threads may execute blocking OS calls (e.g., disk or network read/write and block without calling your lock/unlock!)