Periodic task scheduling

Static priority scheduling
Rate monotonic priority assignment
Derivation of the RM utilization bound

Impact of GRMS

- GRMS: Generalized Rate Monotonic Scheduling
- Cited in the Selected Accomplishment section of the National Research Council's report on A Broader Agenda for Computer Science and Engineering in 1992.
- "Through the development of Rate Monotonic Scheduling [theory], we now have a system that will allow [Space Station] Freedom's computers to budget their time, to choose between a variety of tasks, and decide not only which one to do first but how much time to spend in the process." [Deputy Administrator of NASA, Aaron Cohen]
- "The navigation payload software for the next block of Global Positioning System upgrade recently completed testing. ... This design would have been difficult or impossible prior to the development of rate monotonic theory." [L. Doyle, and J. Elzey, "Successful Use of Rate Monotonic Theory on A Formidable Real-Time System"]

Review

- Terminology
  - Definitions of tasks, task invocations, release/arrival time, absolute deadline, relative deadline, period, start time, finish time, …
  - Preemptive versus non-preemptive scheduling
  - Priority-based scheduling
  - Static versus dynamic priorities
- Utilization (*U*) and schedulability
  - Main problem: Find *Bound* for scheduling policy such that
    - $U < Bound \rightarrow$ All deadlines met!
- Optimality of EDF scheduling
  - $Bound_{EDF}$ = 100%

A quick refresher

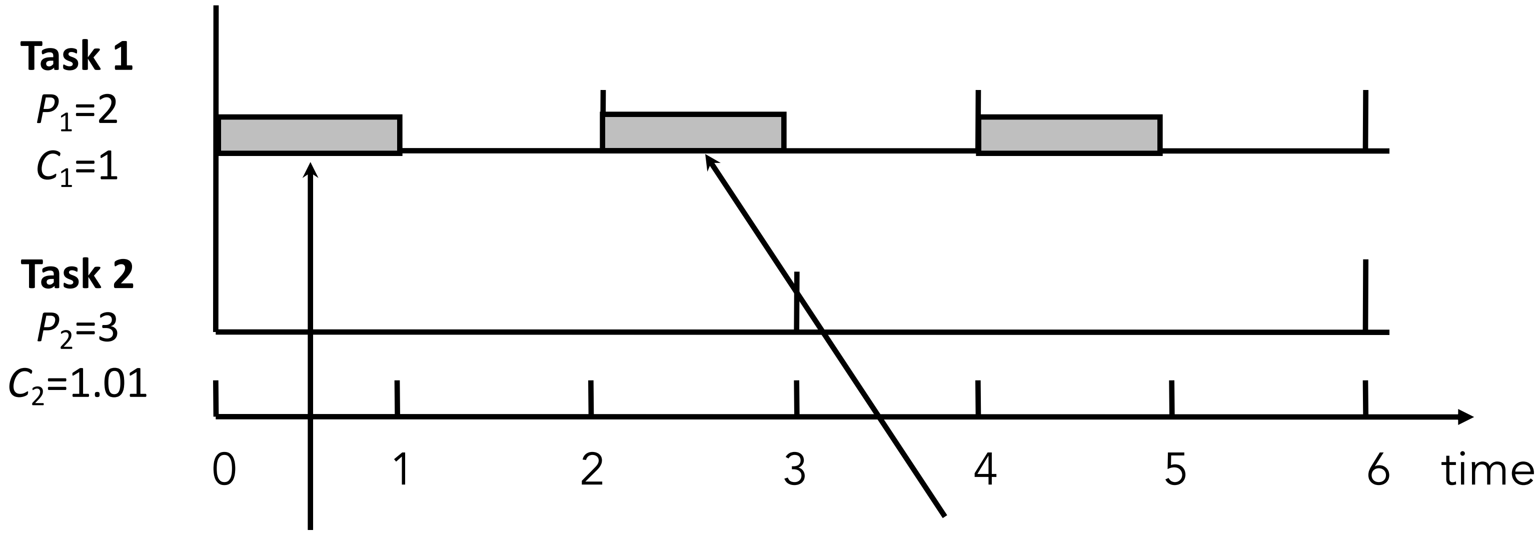

Job 1 of Task 1
Release time = 0
Relative deadline = 2
Absolute deadline = 0+2 = 2

Job 2 of Task 1
Release time = 2
Relative deadline = 2
Absolute deadline = 2+2 = 4

The release time of the first job of a task is also known as the **phase** of the task. The phase of Task 1 is 0.

A quick refresher

Job 2 of Task 1 preempts Job 1 of Task 2 at time=2 because we are following the rate monotonic priority assignment: a static priority policy.

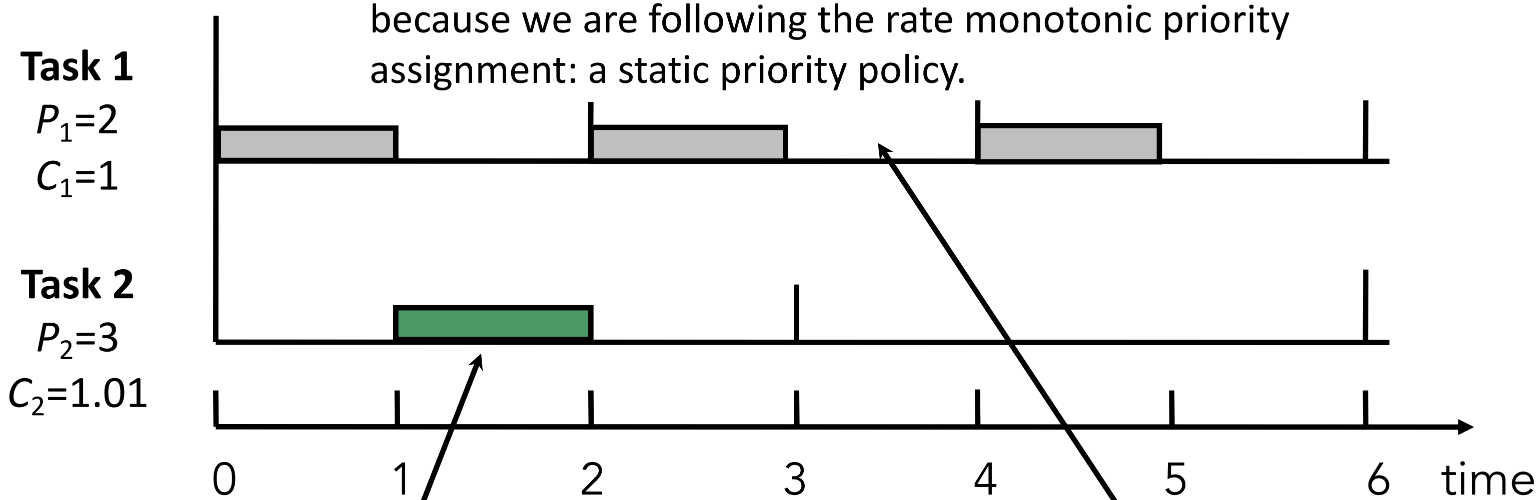

Job 1 of Task 2
Release time = 0
Start time = 1
Relative deadline = 3
Absolute deadline = 0+3 = 3

Job 2 of Task 1
Release time = 2
Relative deadline = 2
Absolute deadline = 2+2 = 4

Schedulability analysis of periodic tasks

- Main problem
  - Given a set of periodic tasks, can they meet their deadlines?
  - Depends on scheduling policy
- Solution approaches
  - Utilization bounds (simplest)
  - Exact analysis (NP-Hard)
  - Approximation schemes with provable bounds on approximation error
  - (Meta)Heuristics
- Two most important scheduling policies
  - Earliest deadline first (dynamic)
  - Rate monotonic (static)

Schedulability analysis of periodic tasks

- Main problem
  - Given a set of periodic tasks, can they meet their deadlines?
  - Depends on scheduling policy
- Solution approaches
  - Utilization bounds (simplest)
  - Exact analysis (NP-Hard)
  - Approximation schemes with provable bounds on approximation error
  - (Meta)Heuristics
- Two most important scheduling policies
  - Earliest deadline first (Dynamic)
  - Rate monotonic (static)

Utilization bounds

- Intuitively,
  - The lower the processor utilization, U, the easier it is to meet deadlines.
  - The higher the processor utilization, U, the more difficult it is to meet deadlines.
- **Question:** Is there a threshold U_{bound} such that
  - When $U < U_{bound}$ deadlines are met
  - When $U > U_{bound}$ deadlines are missed

Example (Rate monotonic scheduling)

Task 1
P_1=2
C_1=1

Task 2
P_2=3
C_2=1.01

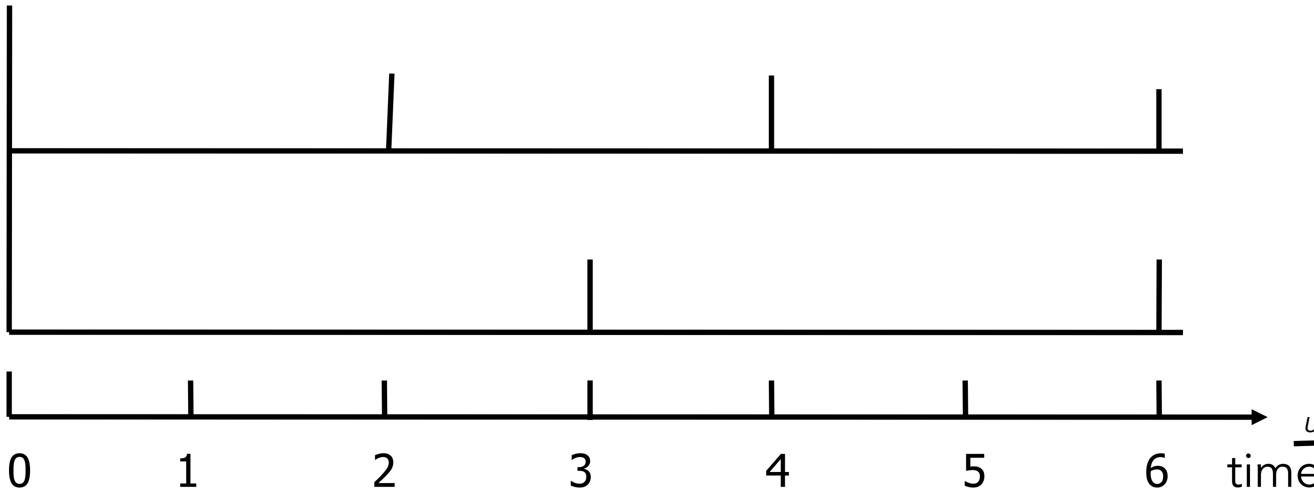

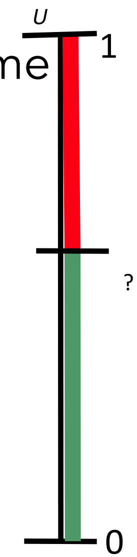

$$U = \frac{C_1}{P_1} + \frac{C_2}{P_2} = \frac{1}{2} + \frac{1.01}{3} \approx 0.833$$

- Question: Is there a threshold U_{bound} such that
  - When $U < U_{bound}$ deadlines are met
  - When $U > U_{bound}$ deadlines are missed

Example (Rate monotonic scheduling)

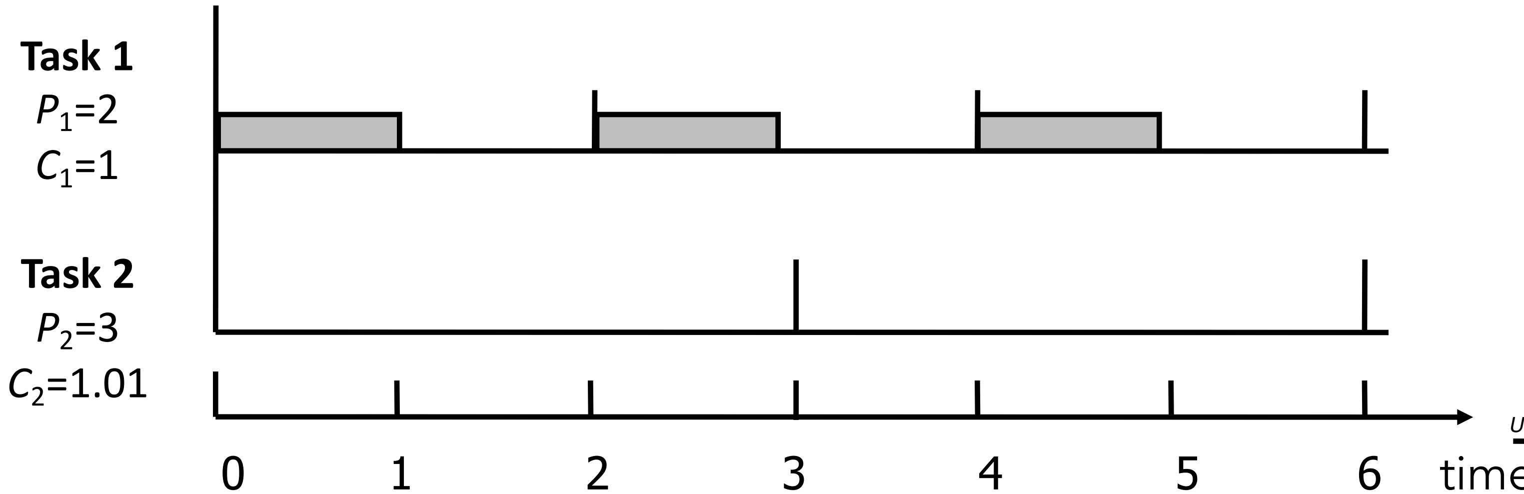

$$U = \frac{C_1}{P_1} + \frac{C_2}{P_2} = \frac{1}{2} + \frac{1.01}{3} \approx 0.833$$

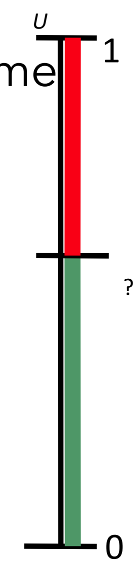

- Question: Is there a threshold U_{bound} such that
- When $U < U_{bound}$ deadlines are met
- When $U > U_{bound}$ deadlines are missed

Example (Rate monotonic scheduling)

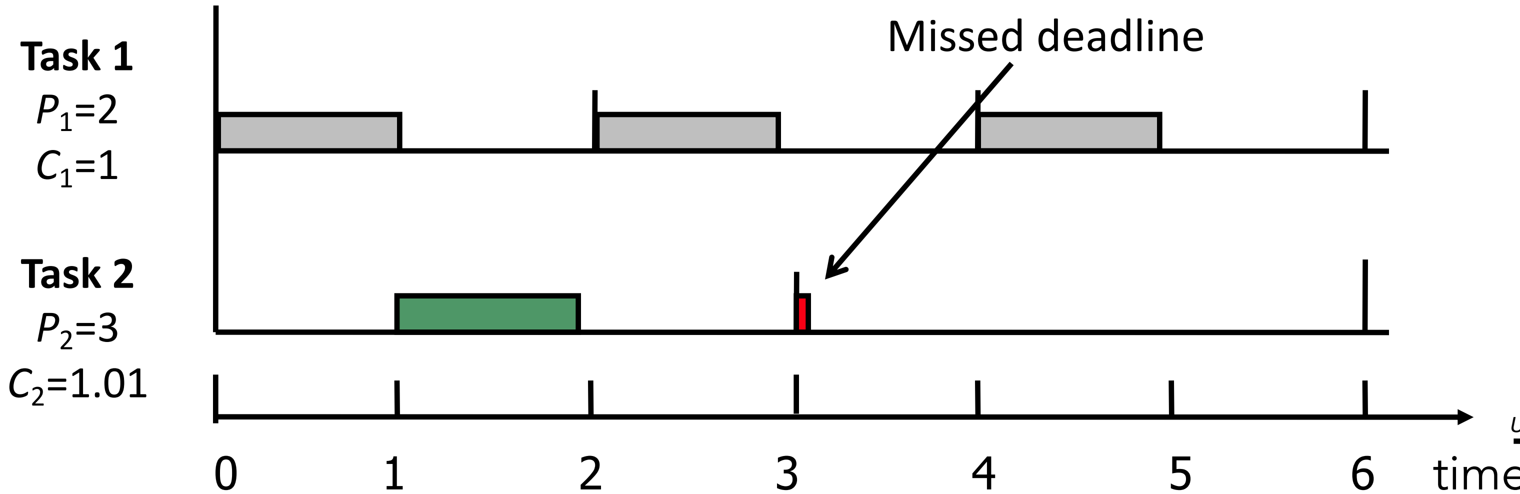

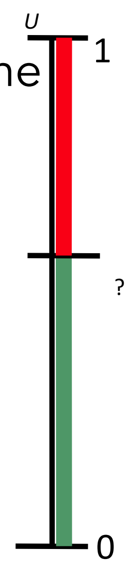

$$U = \frac{C_1}{P_1} + \frac{C_2}{P_2} = \frac{1}{2} + \frac{1.01}{3} \approx 0.833$$

Unschedulable

- Question: Is there a threshold U_{bound} such that
- When $U < U_{bound}$ deadlines are met
- When $U > U_{bound}$ deadlines are missed

Another example (Rate monotonic scheduling)

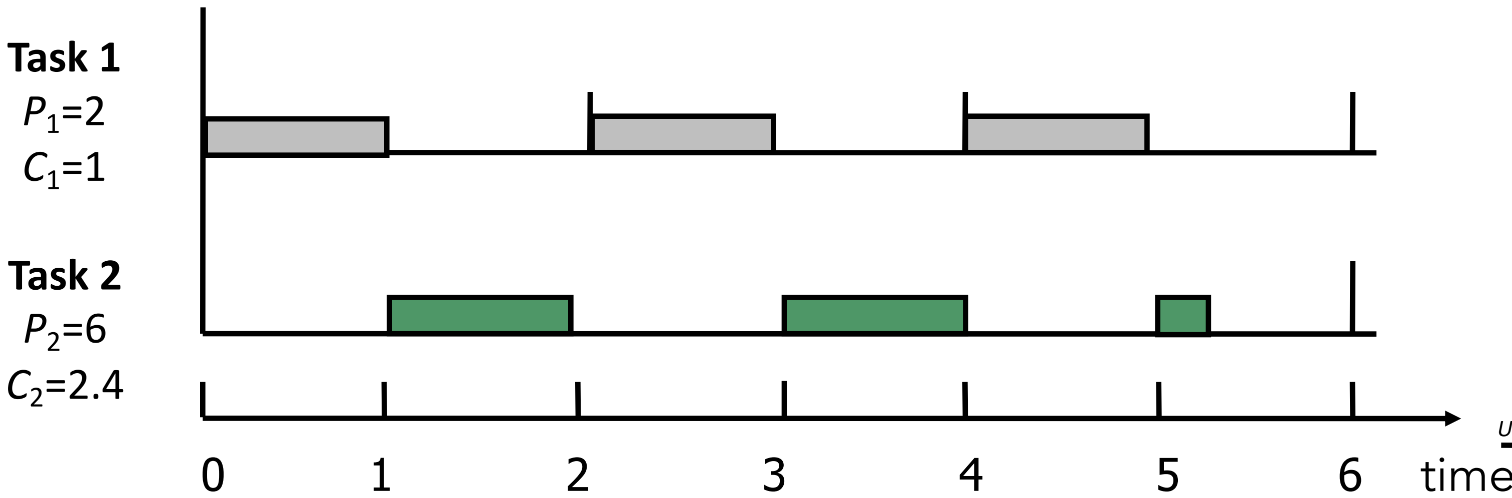

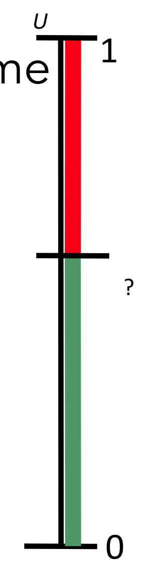

$$U = \frac{C_1}{P_1} + \frac{C_2}{P_2} = \frac{1}{2} + \frac{2.4}{6} = 0.9$$

Schedulable

- Question: Is there a threshold U_{bound} such that
  - When $U < U_{bound}$ deadlines are met
  - When $U > U_{bound}$ deadlines are missed

Another example (Rate monotonic scheduling)

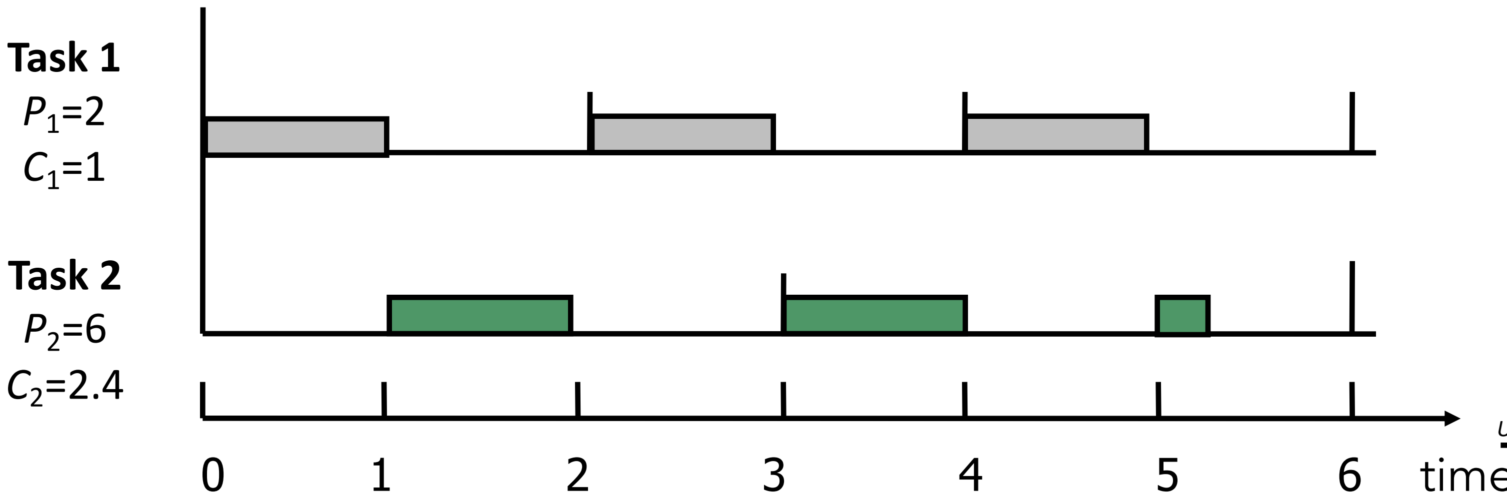

$$U = \frac{C_1}{P_1} + \frac{C_2}{P_2} = \frac{1}{\ } + \frac{2.4}{\ }$$

Schedulability depends on the task set.
No clean relationship between utilization and schedulability.

- Question: Is there … that
  - When $U < U_{bo}$… deadlines are met
  - When $U > U_{bound}$ deadlines are missed

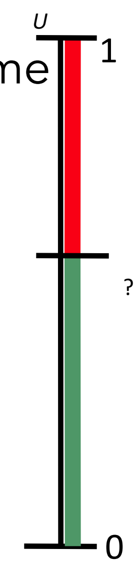

Visualizing schedulability

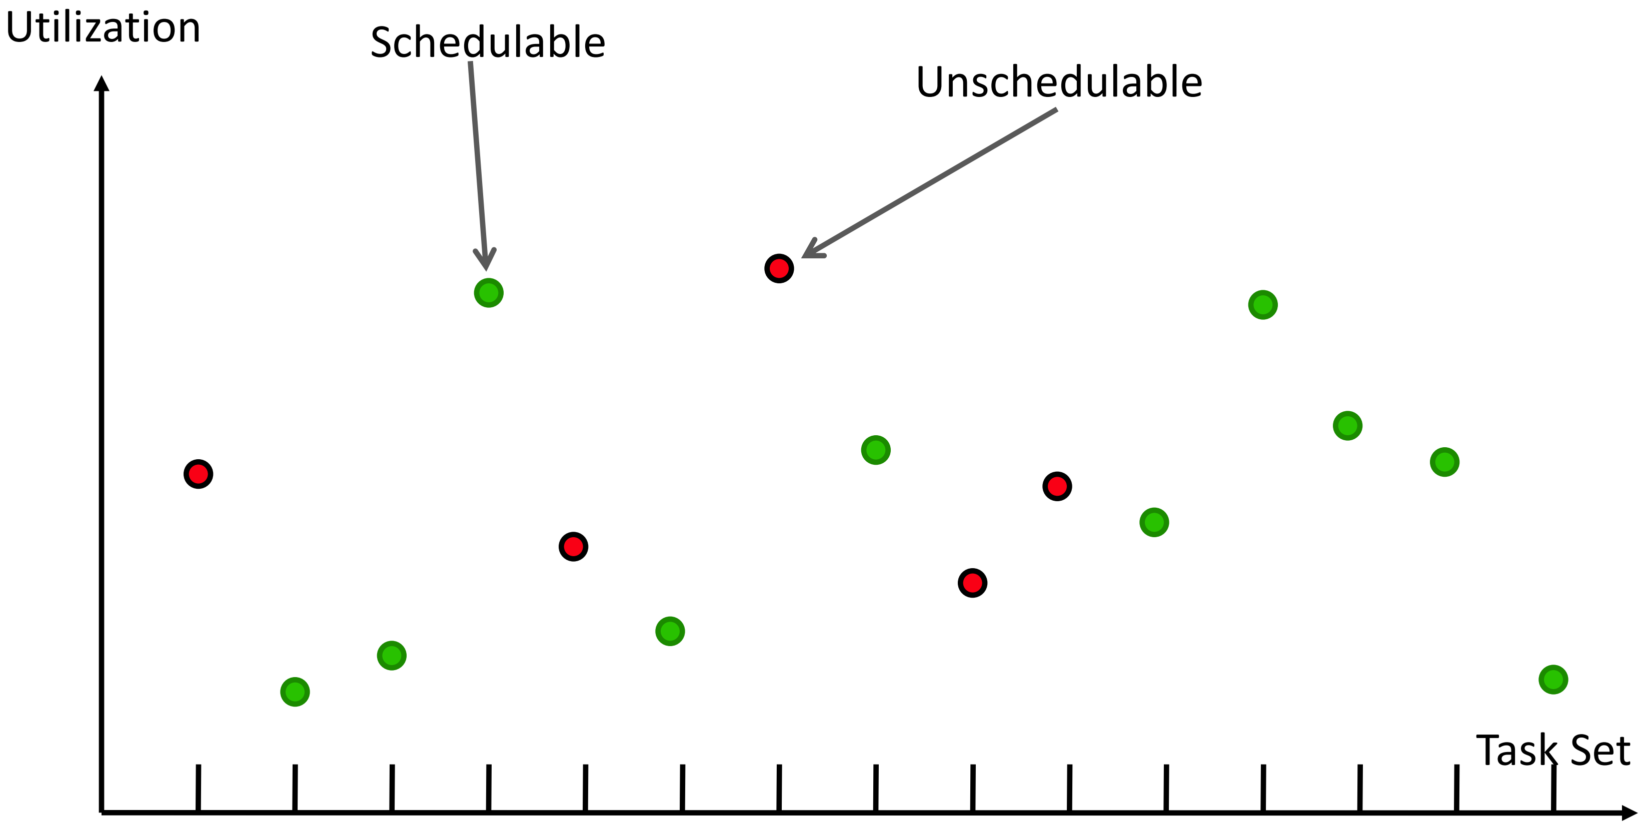

- Question: is there a threshold U_{bound} such that
  - When $U < U_{bound}$ deadlines are met
  - When $U > U_{bound}$ deadlines are missed

Visualizing schedulability

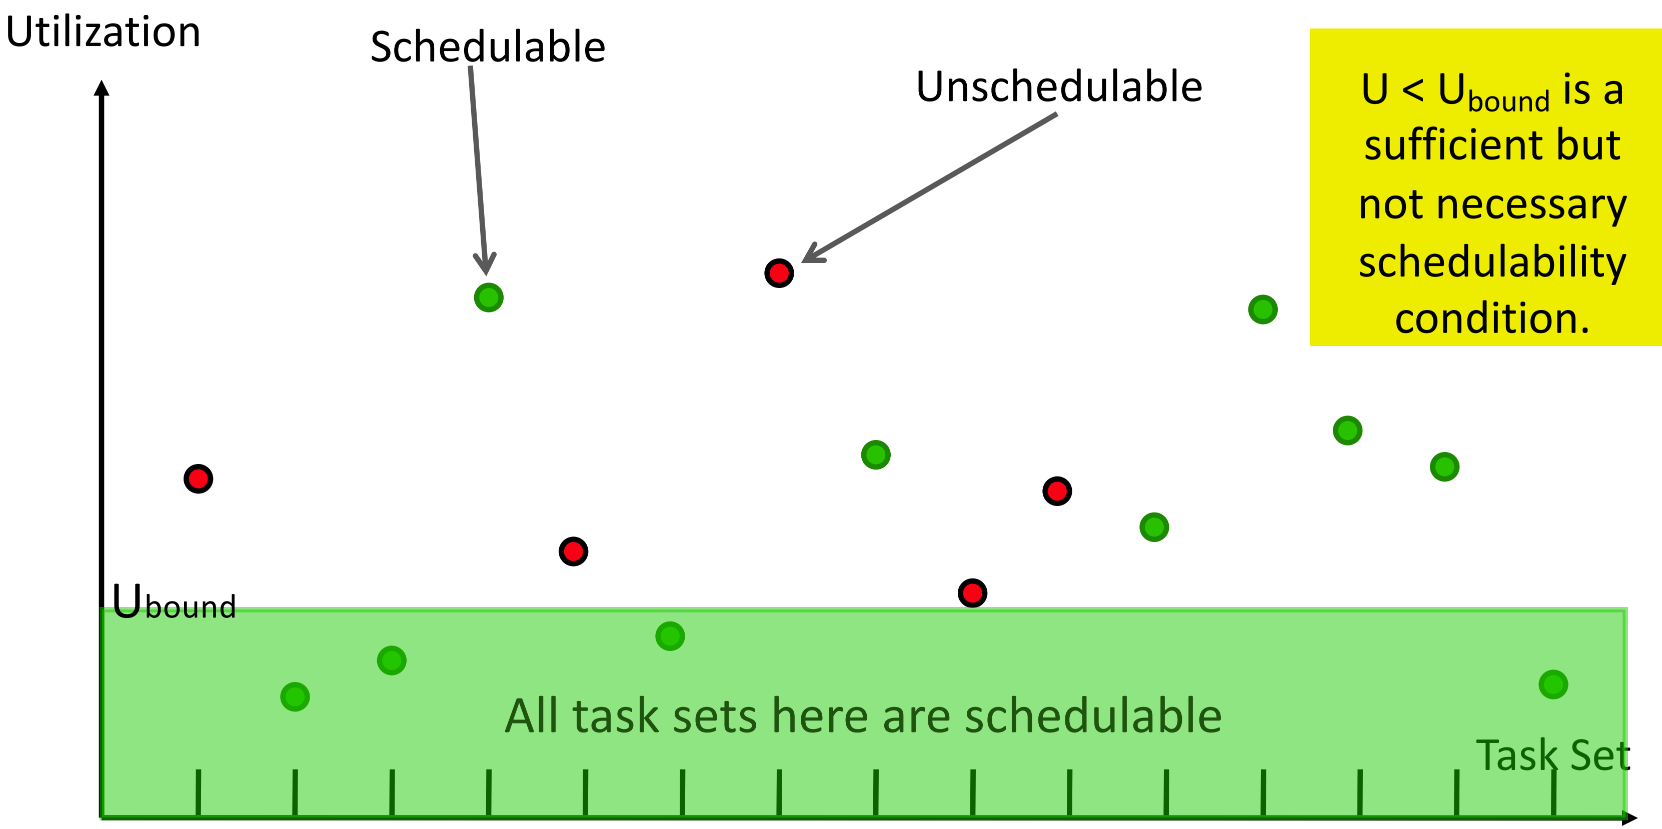

- Modified question: Is there a threshold U_{bound} such that
- When $U < U_{bound}$ deadlines are met (false negatives are OK, but false positives are NOT OK)
- When $U > U_{bound}$ deadlines might or might not be missed

Visualizing schedulability

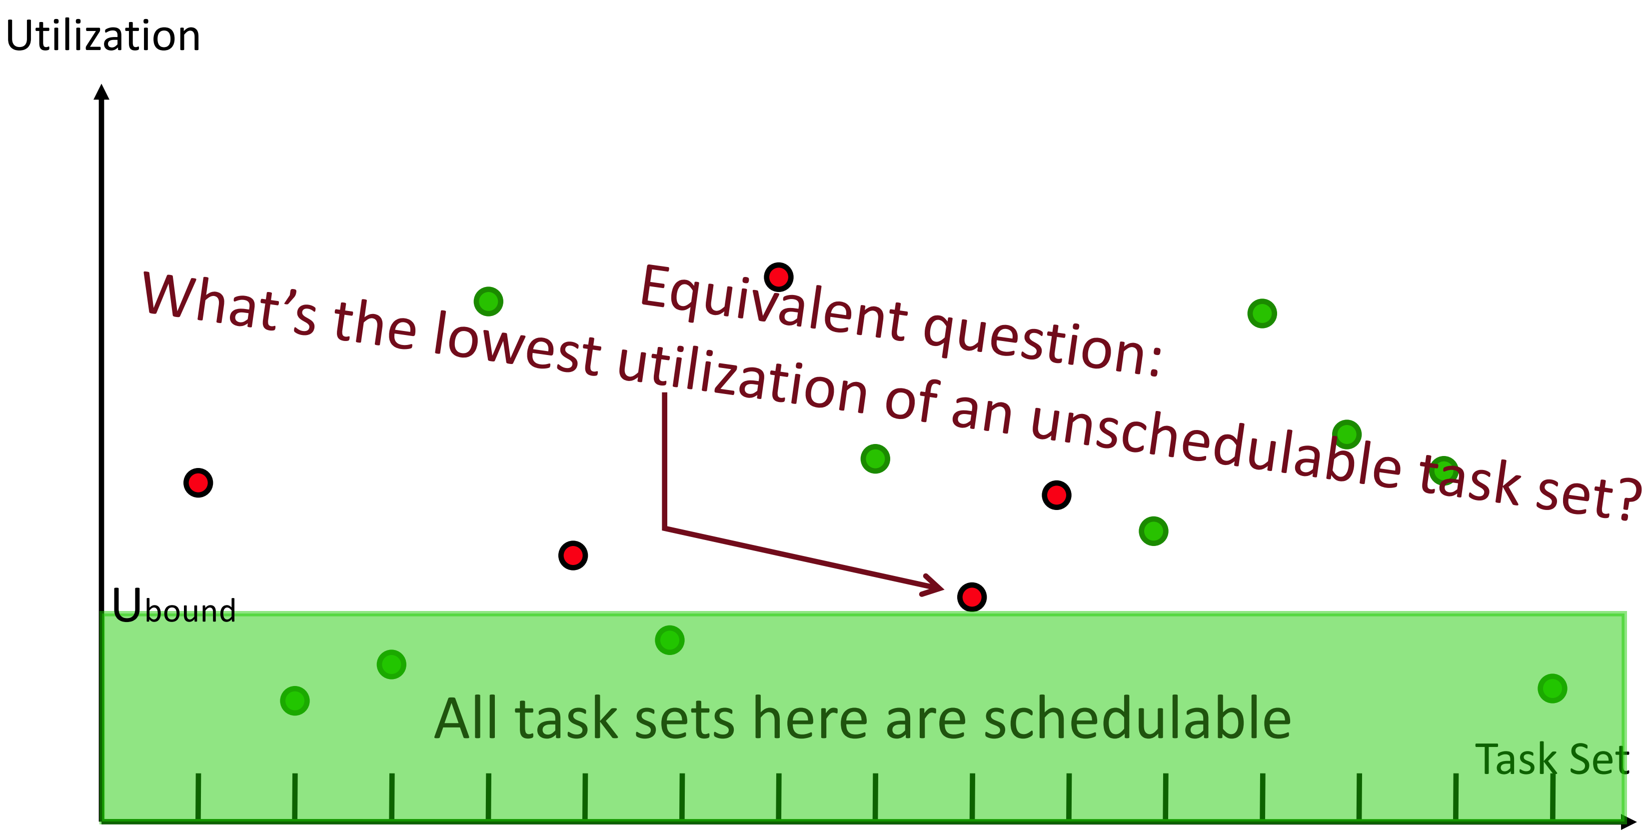

- Modified question: Is there a threshold U_{bound} such that
- When $U < U_{bound}$ deadlines are met
- When $U > U_{bound}$ deadlines might or might not be missed

Visualizing schedulability

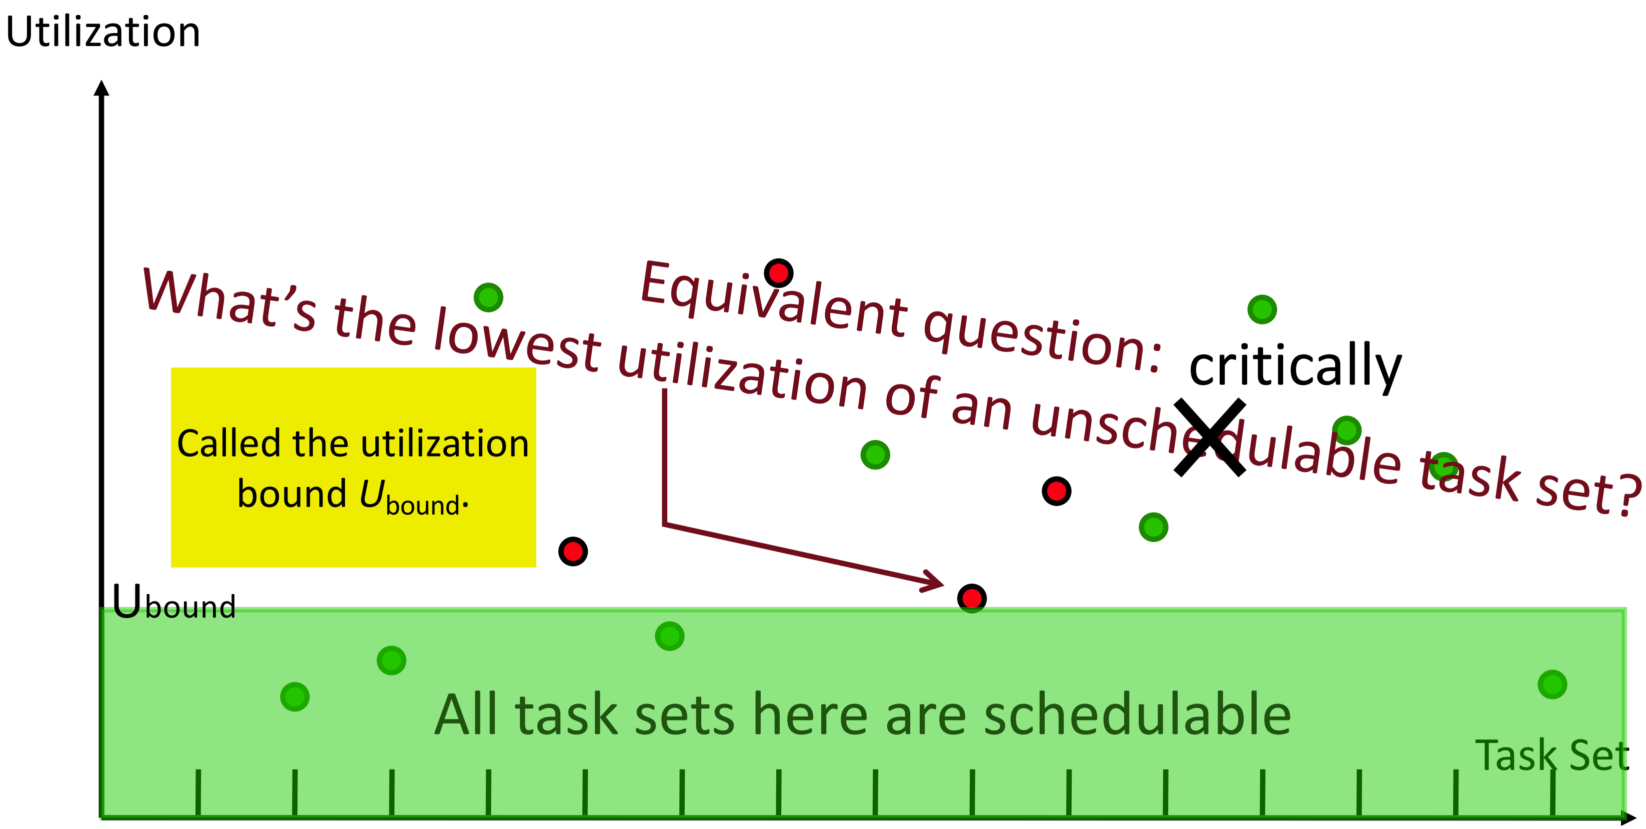

- Modified question: Is there a threshold U_{bound} such that
- When $U < U_{bound}$ deadlines are met
- When $U > U_{bound}$ deadlines might or might not be missed

Utilization bound: The Plan

1. Consider a fixed priority assignment S

2. Consider a task set Γ with utilization $U = \sum_{i=1}^{n} \frac{c_i}{P_i}$ that is **feasible** under S

3. **Inflate** execution times of the tasks until Γ **fully utilizes** the processor under S

   - **Meaning:** If execution times of inflated Γ are further increased by any $\epsilon > 0$, then the inflated Γ will become infeasible under S

   - Inflated Γ is said to ***fully utilize*** the processor under S
     (Also called ***critically schedulable*** under S)

   - We are **not** touching the **periods** of tasks in Γ

Utilization bound: The Plan

4. Now given the set

$$\mathcal{U} = \left\{ U(\Gamma) : \Gamma \text{ critically schedulable under } S \right\}$$

Q: What utilization value U is such that all task sets whose utilization is $\leq U$ are feasible under S?

$$U_{\text{lub}}(S) = \inf \mathcal{U}$$

least **u**pper **b**ound

***Intuition**: Any taskset that is generated from $\mathcal{U}$ by reducing execution times and holding the periods fixed is feasible (less workload → less interference)*

How we will carry out the plan

For an arbitrary taskset that is ***critically* schedulable**:
Find some taskset parameter x for which $U(x)$, the utilization of the **critically schedulable** task set as a function of parameter x, is such that:
(a) for $x < x_0$, $U(x)$ decreases as x increases
(b) for $x > x_0$, $U(x)$ increases as x increases

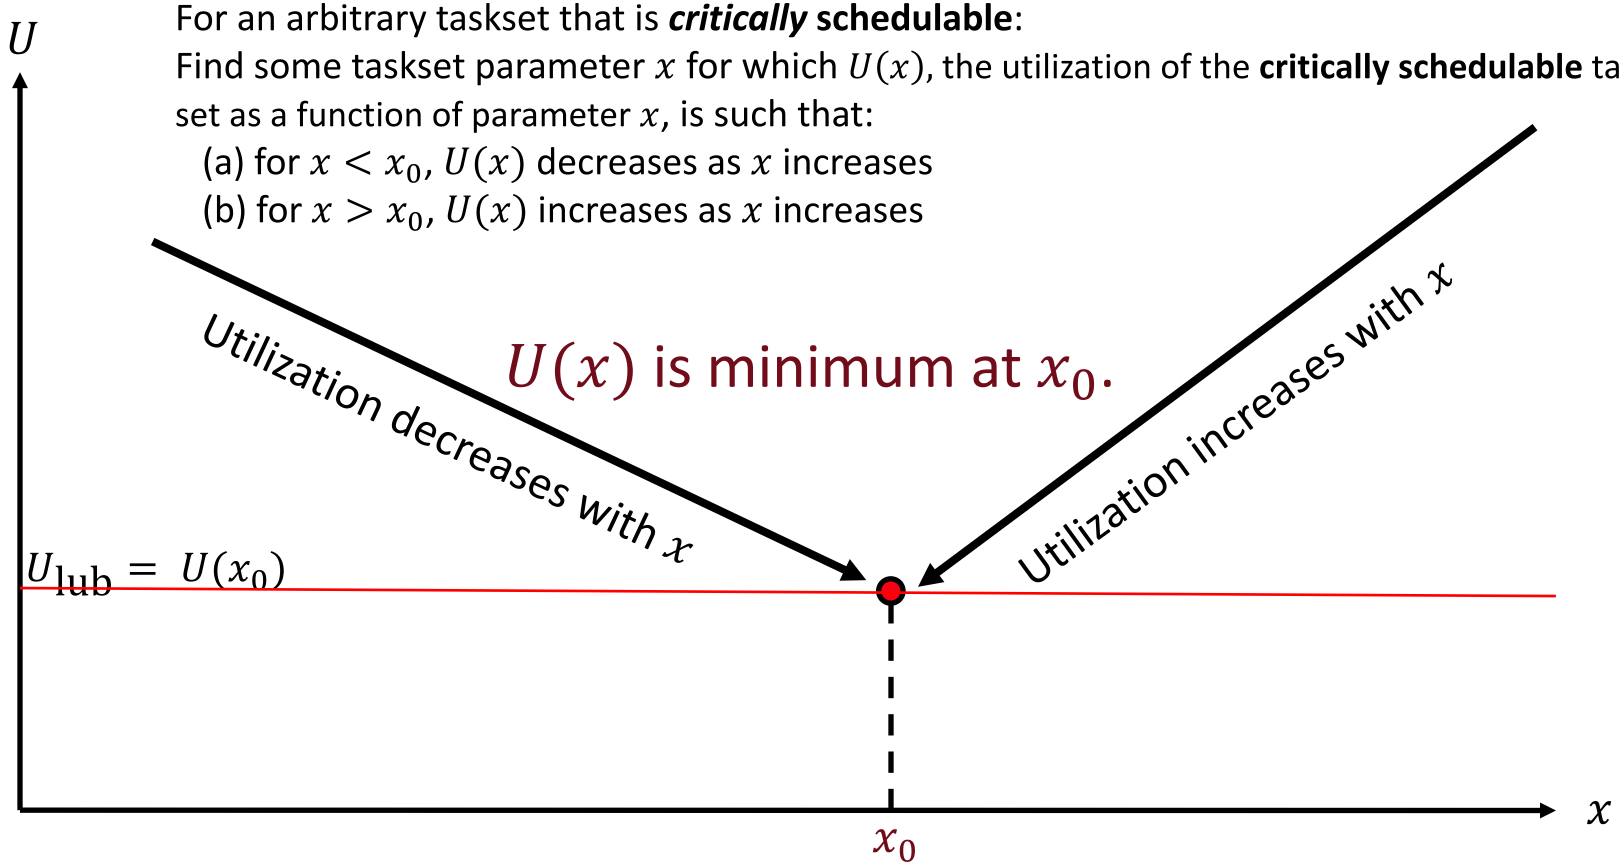

Deriving $U_{\text{lub}}(\text{RM})$

- Two tasks case first
- Consider the same two cases as in proof of optimality in previous lecture

Deriving $U_{\text{lub}}(\text{RM})$

- Two tasks case first
- Consider the same two cases as in proof of optimality in previous lecture

Finding the utilization bound for RM scheduling

Find some task set parameter x such that

Case (a): $x<x_o$ such that $U(x)$ decreases as x increases

Case (b): $x>x_o$ such that $U(x)$ increases as x increases

Thus $U(x)$ is minimum when $x=x_o$ → Find $U(x_o)$

$$C_1 \leqslant P_2 - \left\lfloor \frac{P_2}{P_1} \right\rfloor P_1$$

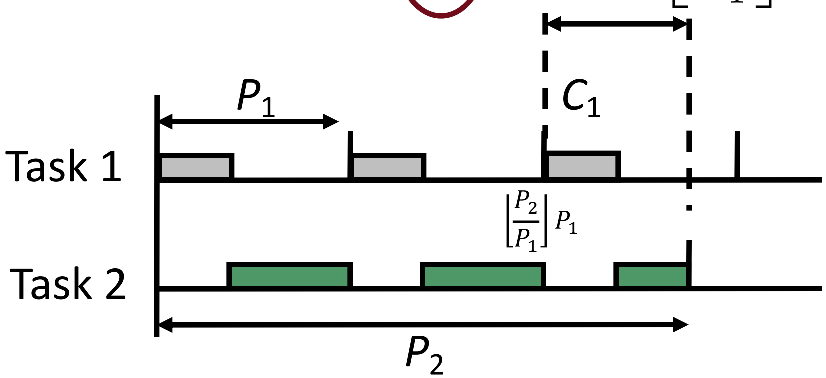

For fixed such C_1, inflate:

$$C_2 = P_2 - C_1 \left\lceil \frac{P_2}{P_1} \right\rceil = P_2 - C_1 \left(\left\lfloor \frac{P_2}{P_1} \right\rfloor + 1 \right)$$

$$U = \frac{C_1}{P_1} + \frac{C_2}{P_2} = 1 + \frac{C_1}{P_2} \left(\frac{P_2}{P_1} - \left\lfloor \frac{P_2}{P_1} \right\rfloor - 1 \right)$$

$$C_1 > P_2 - \left\lfloor \frac{P_2}{P_1} \right\rfloor P_1$$

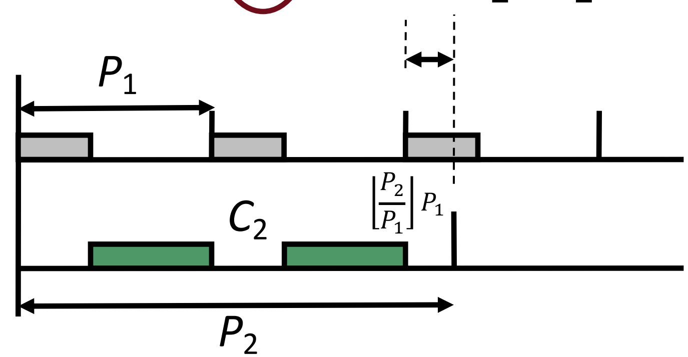

$$C_2 = (P_1 - C_1) \left\lfloor \frac{P_2}{P_1} \right\rfloor$$

$$U = \frac{P_1}{P_2} \left\lfloor \frac{P_2}{P_1} \right\rfloor + \frac{C_1}{P_2} \left(\frac{P_2}{P_1} - \left\lfloor \frac{P_2}{P_1} \right\rfloor \right)$$

Finding the utilization bound for RM scheduling

Case 1:

$$C_1 \leqslant P_2 - \left\lfloor \frac{P_2}{P_1} \right\rfloor P_1$$

$$C_2 = P_2 - C_1 \left\lceil \frac{P_2}{P_1} \right\rceil = P_2 - C_1 \left(\left\lfloor \frac{P_2}{P_1} \right\rfloor + 1 \right)$$

$$U = \frac{C_1}{P_1} + \frac{C_2}{P_2} = 1 + \frac{C_1}{P_2} \left(\frac{P_2}{P_1} - \left\lfloor \frac{P_2}{P_1} \right\rfloor - 1 \right)$$

$$U = 1 + \frac{P_1}{P_2} \left(\frac{P_2}{P_1} - \left\lfloor \frac{P_2}{P_1} \right\rfloor \right) \left(\frac{P_2}{P_1} - \left\lfloor \frac{P_2}{P_1} \right\rfloor - 1 \right)$$

What value of C_1 minimizes U?
- Term in ( ) is < 0
- $C_1 \mapsto U$ is a decreasing map ↓

$$C_1 = P_2 - \left\lfloor \frac{P_2}{P_1} \right\rfloor P_1 = P_1 \left(\frac{P_2}{P_1} - \left\lfloor \frac{P_2}{P_1} \right\rfloor \right)$$

The same C_1 minimizes U of **case 2** and the resulting $U(C_1)$ for case 2 is the same as that of case 1

Finding the utilization bound for RM scheduling

Case 1:

$$U = 1 + \frac{P_1}{P_2}\left(\frac{P_2}{P_1} - \left\lfloor\frac{P_2}{P_1}\right\rfloor\right)\left(\frac{P_2}{P_1} - \left\lfloor\frac{P_2}{P_1}\right\rfloor - 1\right)$$

To minimize *U*, we must have

$$\left\lfloor\frac{P_2}{P_1}\right\rfloor = 1$$

Why?

$0 \leq G < 1$ and so is $1 - G \Rightarrow G(1-G) \geq 0$

Then $\left\lfloor\frac{P_2}{P_1}\right\rfloor \mapsto U$ is increasing $\uparrow$

U minimized at minimum possible $\left\lfloor\frac{P_2}{P_1}\right\rfloor$

But $P_2 \geq P_1$ (RM) $\Rightarrow \left\lfloor\frac{P_2}{P_1}\right\rfloor \geq 1$

Then min. possible $\left\lfloor\frac{P_2}{P_1}\right\rfloor$ is 1

$$U = 1 - \frac{\left(\frac{P_2}{P_1} - \lfloor\frac{P_2}{P_1}\rfloor\right)\left(1 - (\frac{P_2}{P_1} - \lfloor\frac{P_2}{P_1}\rfloor)\right)}{P_2/P_1}$$

G

$$U = 1 - \frac{\left(\frac{P_2}{P_1} - \lfloor\frac{P_2}{P_1}\rfloor\right)\left(1 - (\frac{P_2}{P_1} - \lfloor\frac{P_2}{P_1}\rfloor)\right)}{(\frac{P_2}{P_1} - \lfloor\frac{P_2}{P_1}\rfloor) + \lfloor\frac{P_2}{P_1}\rfloor}$$

$$\Rightarrow \quad U = 1 - \frac{G(1-G)}{G + \lfloor\frac{P_2}{P_1}\rfloor}$$

Finding the utilization bound for RM scheduling

The minimum utilization case

$$U = 1 - \frac{\left(\frac{P_2}{P_1} - \lfloor\frac{P_2}{P_1}\rfloor\right)\left(1 - (\frac{P_2}{P_1} - \lfloor\frac{P_2}{P_1}\rfloor)\right)}{P_2/P_1}$$

To minimize U, we must have $\left\lfloor \frac{P_2}{P_1} \right\rfloor = 1$

$$U = 1 + \frac{(P_2/P_1 - 1)(P_2/P_1 - 2)}{P_2/P_1}$$

Real-valued function in one variable P_2/P_1

Differentiable!

Then $\frac{dU}{d(P_2/P_1)} = 0 \Rightarrow \frac{P_2}{P_1} = \sqrt{2}$

Finally, $U = 0.83$

Finding the utilization bound for RM scheduling

$$U = 1 + \frac{P_1}{P_2}\left(\frac{P_2}{P_1} - \left\lfloor \frac{P_2}{P_1} \right\rfloor\right)\left(\frac{P_2}{P_1} - \left\lfloor \frac{P_2}{P_1} \right\rfloor - 1\right)$$

Observe:

If P_2 is an integer multiple of $P_1 \Rightarrow \frac{P_2}{P_1} = \left\lfloor \frac{P_2}{P_1} \right\rfloor \Rightarrow U = 1$

Generalization for *n* tasks

What were the worst case conditions for the 2-task case?

$$\begin{cases} \left\lfloor \frac{P_2}{P_1} \right\rfloor = 1 \\ C_1 = P_2 - \left\lfloor \frac{P_2}{P_1} \right\rfloor P_1 = P_2 - P_1 \\ C_2 = (P_1 - C_1) \left\lfloor \frac{P_2}{P_1} \right\rfloor = P_1 - C_1 \end{cases} \Longleftrightarrow \begin{cases} P_1 < P_2 < 2P_1 \\ C_1 = P_2 - P_1 \\ C_2 = P_1 - C_1 = 2P_1 - P_2 \end{cases}$$

Generalizing to n tasks

$$\begin{cases} P_1 < P_n < 2P_1 \\ C_1 = P_2 - P_1 \\ C_2 = P_3 - P_2 \\ C_3 = P_4 - P_3 \\ \qquad \vdots \\ C_{n-1} = P_n - P_{n-1} \\ C_n = P_1 - (C_1 + \cdots + C_{n-1}) = 2P_1 - P_n \end{cases}$$

Generalization for n tasks

$$\begin{cases} P_1 < P_n < 2P_1 \\ C_1 = P_2 - P_1 \\ C_2 = P_3 - P_2 \\ C_3 = P_4 - P_3 \\ \qquad \vdots \\ C_{n-1} = P_n - P_{n-1} \\ C_n = P_1 - (C_1 + \cdots + C_{n-1}) = 2P_1 - P_n \end{cases}$$

$$\begin{aligned} U &= \frac{C_1}{P_1} + \frac{C_2}{P_2} + \cdots \frac{C_n}{P_n} \\ &= \frac{P_2 - P_1}{P_1} + \frac{P_3 - P_2}{P_2} + \cdots + \frac{2P_1 - P_n}{P_n} \\ &= \frac{P_2}{P_1} + \frac{P_3}{P_2} + \cdots + \frac{2P_1}{P_n} - n \end{aligned}$$

Generalization for n tasks

$$U = \frac{P_2}{P_1} + \frac{P_3}{P_2} + \cdots + \frac{2P_1}{P_n} - n$$

$$\frac{\partial U}{\partial (P_2/P_1)} = 0; \frac{\partial U}{\partial (P_3/P_2)} = 0; \ldots ; \frac{\partial U}{\partial (P_n/P_1)} = 0$$

We can then obtain

Liu & Layland, 1973

$$\frac{P_{i+1}}{P_i} = 2^{1/n} \Rightarrow U = n(2^{1/n} - 1)$$

For large *n*: $\lim_{n \to \infty} U = \lim_{n \to \infty} n(2^{1/n} - 1) = \ln 2 \approx 0.69$

Sanity check

Let $b_n = \inf\{U(\Gamma) : \Gamma \text{ is a critically schedulable set of } n \text{ tasks}\}$

Let U_n be utilization factor of n tasks

Question: If $U_n \leq b_n$, is the task set schedulable?

Lecture summary

- Understanding utilization bounds
- The utilization bound for rate-monotonic scheduling
- For RM scheduling the bound decreases with the number of tasks, approaching an asymptotic limit of 0.69
- Coming up: Why is RM priority assignment the optimal static priority policy? Are there better schedulability tests?