

Periodic task scheduling

Static priorities

- ★ Better utilization bounds
- ★ Deadlines less than periods
- ★ Exact test for schedulability

Quick review

- Why is rate monotonic scheduling optimal (among static priority policies)?
 - **Critical instant theorem:** The worst-case execution time of a job when tasks are scheduled with fixed priorities occurs when jobs belonging to all tasks release at the same instant
 - It is sufficient, then, to verify that the job that is released at the **critical instant** meets its deadline
 - In this worst case, rate monotonic scheduling is optimal (easy to see; if tasks are feasibly scheduled in any other order, swap based on deadlines)
- Utilization bound and optimality of EDF
 - The utilization bound is 1 (or 100%)
 - EDF is optimal because no policy can do better (may do as well but not better)

Exercise

Know Your Worst Case Scenario

- Consider a periodic system of two tasks
- Let $U_i = C_i/P_i$ (for $i = 1, 2$)
- What is the maximum value of $\prod_i (1 + U_i)$ for a schedulable system?
- **Motivation:** There may be other functions of a task set rather than just utilization that also indicate schedulability.

Hyperbolic bound for RM

worst case conditions
for schedulability
of 2 tasks under RM

Critically schedulable

$$C_1 = P_2 - P_1$$

$$C_2 = P_1 - C_1 = 2P_1 - P_2$$

$$U_1 + 1 = \frac{C_1}{P_1} + 1 = \frac{C_1 + P_1}{P_1} = \frac{P_2}{P_1}$$

$$U_2 + 1 = \frac{C_2}{P_2} + 1 = \frac{C_2 + P_2}{P_2} = 2\frac{P_1}{P_2}$$

$$\prod_i (U_i + 1) = 2$$

Schedulable

$$\prod_i (U_i + 1) \leq 2$$

Hyperbolic bound

Solutions

Critically schedulable

$$\left\{ \begin{array}{l} C_1 = P_2 - P_1 \\ C_2 = P_1 - C_1 = 2P_1 - P_2 \\ U_1 + 1 = \frac{C_1}{P_1} + 1 = \frac{C_1 + P_1}{P_1} = \frac{P_2}{P_1} \\ U_2 + 1 = \frac{C_2}{P_2} + 1 = \frac{C_2 + P_2}{P_2} = 2 \\ \prod_i (U_i + 1) = 2 \end{array} \right.$$

Generalizes to
task sets with n
tasks

Schedulable

$$\prod_i (U_i + 1) \leq 2$$

Hyperbolic bound

Hyperbolic bound for **rate monotonic** scheduling

- A set of periodic tasks is schedulable if

$$\prod_i (U_i + 1) \leq 2$$


Hyperbolic bound for rate monotonic scheduling

A set of periodic tasks is schedulable if

$$\prod_i (U_i + 1) \leq 2$$

- It is a better bound than the Liu and Layland bound $U \leq n(2^{1/n} - 1)$
- **Example:** consider a system with two tasks such that $U_1 = 0.8$ and $U_2 = 0.1$
 - $U = 0.9 > 0.83$ (**unschedulable** according to the Liu and Layland bound)
 - $(1 + U_1)(1 + U_2) = (1.8)(1.1) = 1.98 < 2$ (**schedulable** according to the hyperbolic bound)
- **Question:** What happens to the hyperbolic bound if task utilizations are equal?

Feasibility regions in U-space




Hyperbolic bound is *tight*

- It is the best possible bound with only knowledge of utilization factors and the number of jobs
- A utilization-based condition $C(u_1, \dots, u_n)$ for a scheduling algorithm is ***tight*** if for every utilization set (u_1, \dots, u_n) with $0 \leq u_i \leq 1$ for which $C(u_1, \dots, u_n)$ does ***not*** hold, there exists a task set T_1, \dots, T_n with utilizations u_1, \dots, u_n that is ***not*** schedulable by the scheduling algorithm
 - We can construct a task set with the prescribed utilizations (which violate the schedulability condition) that is ***infeasible*** under the given algorithm
- Tightness was proved for H-bound \rightarrow With the algorithm being RM and $C(u_1, \dots, u_n) \equiv \prod_{i=1}^n (1 + u_i) \leq 2$
- **Q:** Is the LL bound tight?

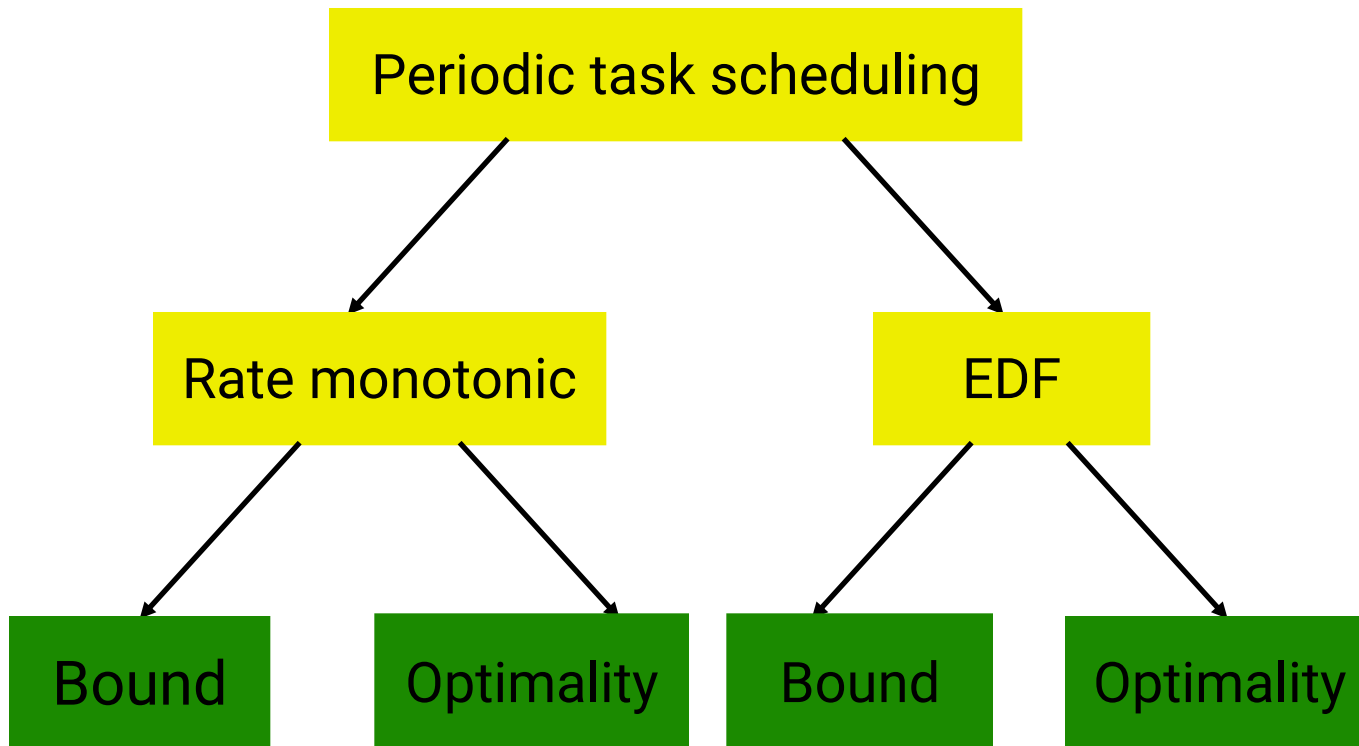
How much “better” is Hyperbolic bound relative LL?

- How do we measure “better” or the “gain” of H-bound over LL-bound?
- Consider the utilization space
 - **U-space:** Subset of n -dimensional Euclidean space consisting of vectors $(u_1, \dots, u_n) \in [0,1]^n$
- Fix number of jobs n
- Volume $\text{vol}^n(A)$: n -dimensional Lebesgue measure of (measurable) set $A \subset \mathbb{R}^n$
- Take volume of H-bound region
 - Here need to find $\text{vol}^n(H)$, $H = \{u \in \mathbb{R}^n: u_i \in [0,1], \prod_{i=1}^n (1 + u_i) \leq 2\}$
- Take volume of LL-bound region
 - need to find $\text{vol}^n(LL)$, $LL = \{u \in \mathbb{R}^n : u_i \in [0,1], \sum_{i=1}^n u_i \leq n(2^{1/n} - 1)\}$

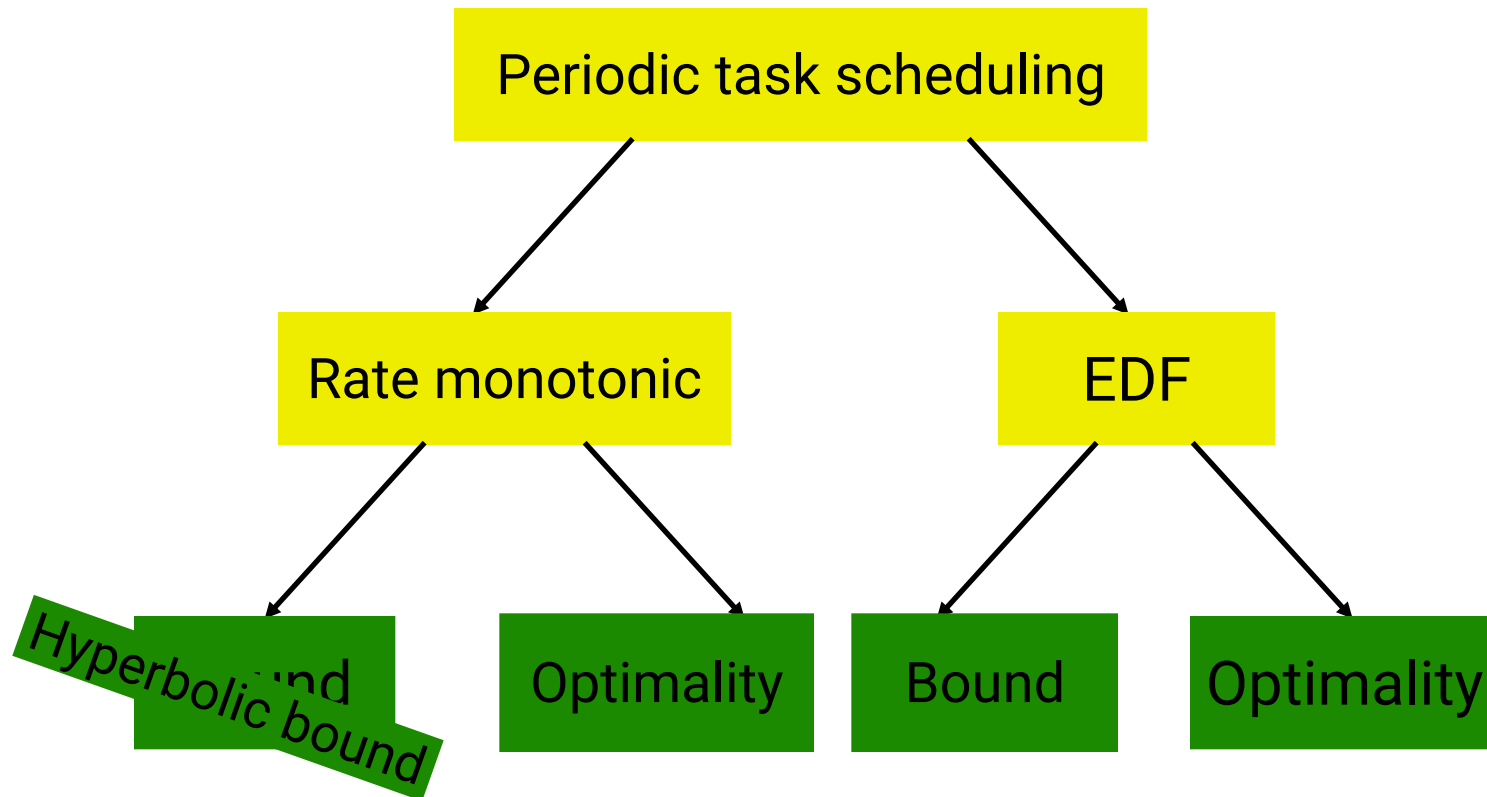


$$\text{Asymptotic Gain: } \rho_n = \frac{\text{vol}^n(H)}{\text{vol}^n(LL)} = \sqrt{2} + O(n^{-1})$$

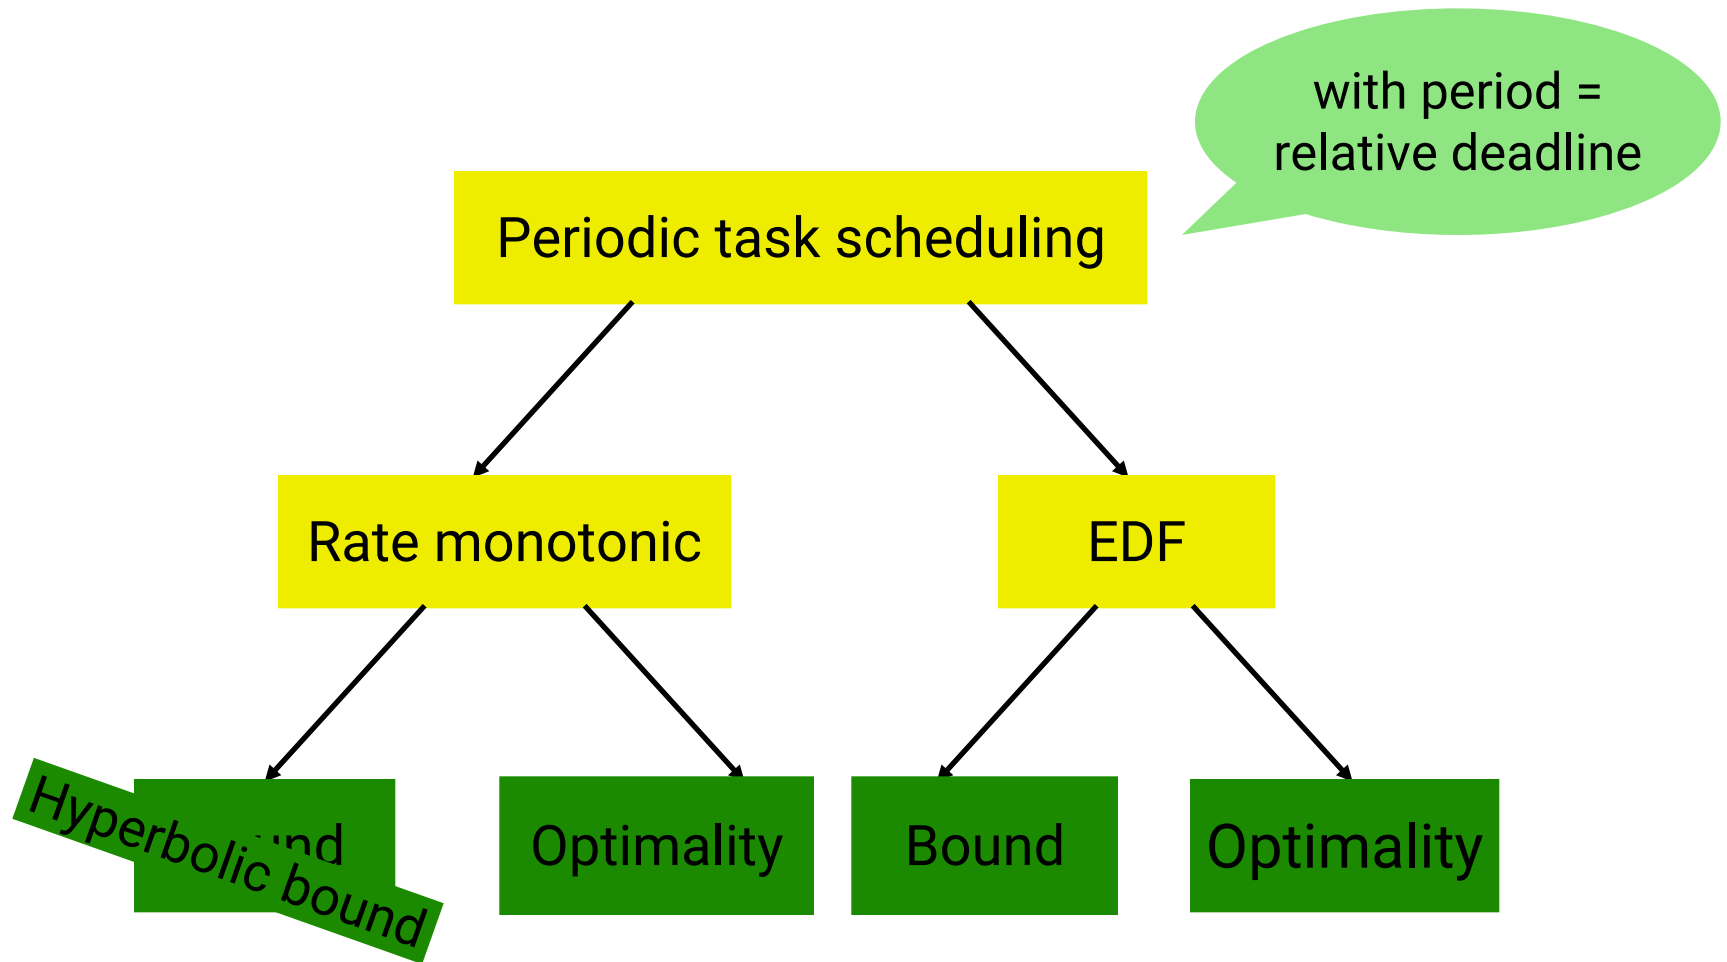
Scheduling taxonomy



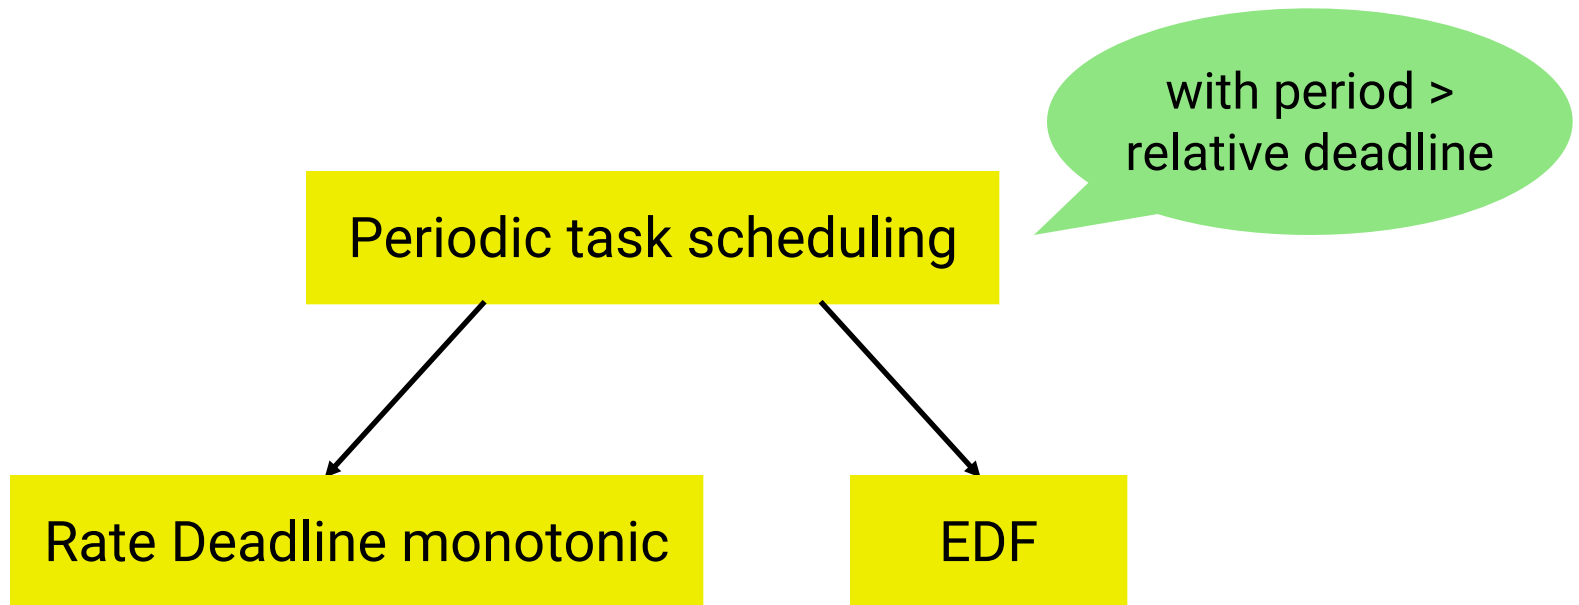
Scheduling taxonomy



Scheduling taxonomy




Scheduling taxonomy




Deadline monotonic scheduling

- Consider a set of periodic tasks where each task, i , has a computation time, C_i , a period, P_i , and a relative deadline $D_i < P_i$.



Deadline monotonic scheduling


- Consider a set of periodic tasks where each task, i , has a computation time, C_i , a period, P_i , and a relative deadline $D_i < P_i$.



- What is the schedulability condition?
- Can not be worse than when the period of each task is reduced to D_i .

Deadline monotonic scheduling

- Consider a set of periodic tasks where each task, i , has a computation time, C_i , a period, P_i , and a relative deadline $D_i < P_i$.




- What is the schedulability condition?
- Can not be worse than when the period of each task is reduced to D_i .

$$\sum_i \frac{C_i}{D_i} \leq n(2^{1/n} - 1)$$

Deadline monotonic scheduling

- Consider a set of periodic tasks where each task, i , has a computation time, C_i , a period, P_i , and a relative deadline $D_i < P_i$.



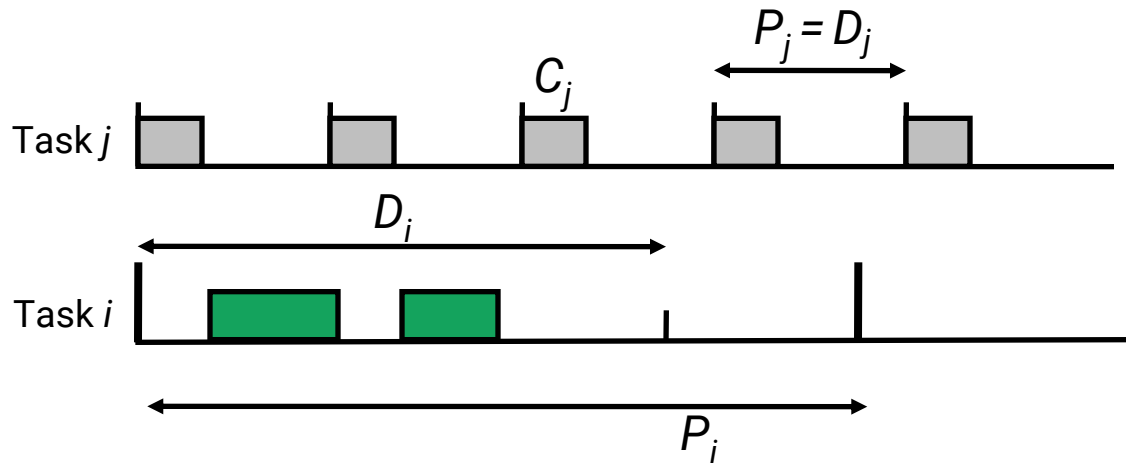
- What is the schedulability condition?
- Can not be worse than when the period of each task is reduced to D_i .

$$\sum_i \frac{C_i}{D_i} \leq n(2^{1/n} - 1)$$

What is the problem?

A better condition for schedulability


- Worst case interference from a higher priority task, j ?



A better condition for schedulability

- Worst case interference from a higher priority task, j ?

Time required by a higher priority task in an interval of length that corresponds to the relative deadline of task i .




Worst case Interference task j exercises on task i
 → upper bound on total workload requested
 by task j during D_i at the critical instant

$$I_i(j) = \underbrace{\left\lceil \frac{D_i}{P_j} \right\rceil}_{\text{Number of execution requests of task } j \text{ in duration of length } D_i \text{ assuming critical instant}} C_j$$

Number of execution requests of task j in duration of length D_i assuming critical instant ²⁰

A better condition for schedulability

- Worst case interference from a higher priority task, j ?




$$I_i(j) = \left\lceil \frac{D_i}{P_j} \right\rceil C_j$$

$$\sum_{j \in \text{hp}(i)} \left\lceil \frac{D_i}{P_j} \right\rceil C_j + C_i \leq D_i$$

A better condition for schedulability

- Worst case interference from a higher priority task, j ?



$$I_i(j) = \left\lceil \frac{D_i}{P_j} \right\rceil C_j$$

Interference from higher priority tasks \longrightarrow

$$\sum_{j \in \text{hp}(i)} \left\lceil \frac{D_i}{P_j} \right\rceil C_j + \boxed{C_i} \leq D_i$$

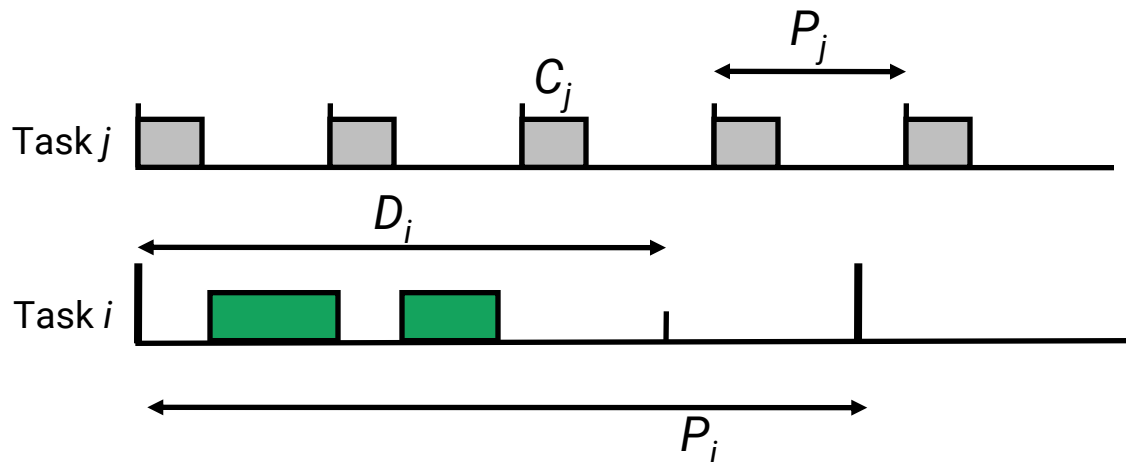
↑

Execution time of T_i 22

A better condition for schedulability

- Worst case interference from a higher priority task, j ?

There still is a problem!



$$I_i(j) = \left\lceil \frac{D_i}{P_j} \right\rceil C_j$$


Interference from higher priority tasks →

$$\sum_{j \in \text{hp}(i)} \left\lceil \frac{D_i}{P_j} \right\rceil C_j + \underbrace{C_i}_{\substack{\uparrow \\ \text{Execution time of } T_i}} \leq D_i$$

Execution time of T_i 23

An exact condition for schedulability

- Interference exists only till a job completes execution, i.e., up to the response time R_i
- Not necessarily up to the relative deadline D_i



$$I_i(j) = \left\lceil \frac{\cancel{D_i} + R_i}{P_j} \right\rceil C_j$$

An exact condition for schedulability

(1) $\left\lceil \frac{R_i}{P_j} \right\rceil C_j$: the *exact* workload interfering with task i
in an interval of length R_i starting at the latest critical instant


But $\left\lceil \frac{R_i}{P_j} \right\rceil C_j$ is the workload *requested* by higher priority task τ_j
in an interval of length R_i starting at the latest critical instant

And (1) is saying that this workload indeed *completes* by R_i . Why?

Because task j has higher priority than task i so all instances of task j
that arrive in interval of length R_i finish before task i

An exact condition for schedulability

- Interference exists only till a job completes execution, i.e., up to the response time R_i
- Not necessarily up to the relative deadline D_i




$$I_i = \sum_{j \in \text{hp}(i)} I_i(j) = \sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i}{P_j} \right\rceil C_j$$

Interference on task i from all higher priority tasks

An exact condition for schedulability

- Interference exists only till a job completes execution, i.e., up to the response time R_i
- Not necessarily up to the relative deadline D_i




We also have the following relation:

$$R_i = I_i + C_i$$

$$I_i = \sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i}{P_j} \right\rceil C_j$$

An exact condition for schedulability

- Interference exists only till a job completes execution, i.e., up to the response time R_i
- Not necessarily up to the relative deadline D_i



We also have the following relation:

$$R_i = I_i + C_i$$

Solve iteratively for the smallest R_i to satisfy both relations

$$I_i = \sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i}{P_j} \right\rceil C_j$$

Computing Response Time

$$R_i = \underbrace{\sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i}{P_j} \right\rceil C_j + C_i}_{f(R_i)}$$

variable

- What is a solution to this recurrence?
- Need to find a **fixed point**

Computing Response Time

variable

$$R_i = \underbrace{\sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i}{P_j} \right\rceil C_j + C_i}_{f(R_i)}$$

- **What is a solution to this recurrence?**

- The smallest $t > 0$ such that $t = \sum_{j \in \text{hp}(i)} \left\lceil \frac{t}{P_j} \right\rceil C_j + C_i \rightarrow$ called *fixed-point*

- Does a solution always exist?

- If so, how can a solution be computed in a finite number of steps?
(convergence)

**Finding Response Times in a Real-Time System M. Joseph P. Pandya.
The Computer Journal, Volume 29, Issue 5, 1 January 1986, Pages 390–395.**

Computing Response Time


$$\text{Recurrence: } R_i^{(n+1)} = \underbrace{\sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i^{(n)}}{P_j} \right\rceil C_j}_{I_i^{(n)}} + C_i$$

- What is a proper initial value $R_i^{(0)}$?
 - Any *lower* bound on the response time \rightarrow take $R_i^{(0)} = C_i$ or $R_i^{(0)} = \sum_{j \in \text{hp}(i)} C_j$
 - Affects the rate of convergence

Exact Response Time: Solution Existence

- It was shown that recurrence converges if $\sum_{j \in \text{hp}(i)} u_j \leq 1$
- Easy to see that $R_i^{(n+1)} \geq R_i^{(n)} \rightarrow$
 - Induction on n + reason about $(R_i^{n+1} - R_i^n)$ + use fact that $x \mapsto \lceil x \rceil$ is increasing
- Stop at first n for which $R_i^{(n+1)} = R_i^{(n)}$
- Recurrence might not converge if $\sum_{j \in \text{hp}(i)} u_j > 1$
 - If only want to know whether or not taskset is schedulable \rightarrow Terminate as soon as $R_i^{(n)} > D_i$ or $R_i^{(n)} > P_i$

Example




$$R_i^{(n+1)} = \underbrace{\sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i^{(n)}}{P_j} \right\rceil C_j}_{I_i^{(n)}} + C_i$$

Consider a system of two tasks:

Task 1: $P_1 = 1.7$, $D_1 = 0.5$, $C_1 = 0.5$

Task 2: $P_2 = 8$, $D_2 = 3.2$, $C_2 = 2$

Example



$$R_i^{(n+1)} = \underbrace{\sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i^{(n)}}{P_j} \right\rceil C_j}_{I_i^{(n)}} + C_i$$

$$R_2^{(0)} = C_2 = 2$$


$$I_i^{(0)} = \lceil 2/1.7 \rceil (0.5) = 1$$

Consider a system of two tasks:

Task 1: $P_1 = 1.7$, $D_1 = 0.5$, $C_1 = 0.5$

Task 2: $P_2 = 8$, $D_2 = 3.2$, $C_2 = 2$

Example



$$R_i^{(n+1)} = \underbrace{\sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i^{(n)}}{P_j} \right\rceil C_j}_{I_i^{(n)}} + C_i$$

$$R_2^{(0)} = C_2 = 2$$

$$I_2^{(0)} = \lceil 2/1.7 \rceil (0.5) = 1$$

$$R_2^{(1)} = I_2^{(0)} + C_2 = 3$$


$$I_2^{(1)} = \lceil 3/1.7 \rceil (0.5) = 1$$

Consider a system of two tasks:

Task 1: $P_1 = 1.7$, $D_1 = 0.5$, $C_1 = 0.5$

Task 2: $P_2 = 8$, $D_2 = 3.2$, $C_2 = 2$

Example



Consider a system of two tasks:

Task 1: $P_1 = 1.7, D_1 = 0.5, C_1 = 0.5$

Task 2: $P_2 = 8, D_2 = 3.2, C_2 = 2$

$$R_i^{(n+1)} = \underbrace{\sum_{j \in \text{hp}(i)} \left\lceil \frac{R_i^{(n)}}{P_j} \right\rceil C_j}_{I_i^{(n)}} + C_i$$

$$R_2^{(0)} = C_2 = 2$$

$$I_2^{(0)} = \lceil 2/1.7 \rceil (0.5) = 1$$

$$R_2^{(1)} = I_2^{(0)} + C_2 = 3$$

$$I_2^{(1)} = \lceil 3/1.7 \rceil (0.5) = 1$$

$$R_2^{(2)} = I_2^{(1)} + C_2 = 3$$

$$R_2^{(2)} = R_2^{(1)}$$

$3 < 3.2$; Task 2 is schedulable. 36

RTA Algorithm

```
DM_guarantee ( $\Gamma$ ) {  
  for (each  $\tau_i \in \Gamma$ ) {  
     $I_i = \sum_{k=1}^{i-1} C_k$ ;  
    do {  
       $R_i = I_i + C_i$ ;  
      if ( $R_i > D_i$ ) return(UNSCHEDULABLE);  
       $I_i = \sum_{k=1}^{i-1} \left\lceil \frac{R_i}{T_k} \right\rceil C_k$ ;  
    } while ( $I_i + C_i > R_i$ );  
  }  
  return(SCHEDULABLE);  
}
```

Response Time Analysis: Time Complexity

- Is this test efficient?
- Assuming all instance parameters are integers:
 - Inner loop adds at most 1 to interference until deadline is reached
- Runs in time $O(nP_{\max})$, where P_{\max} is the largest period
 - test runs in **pseudo-polynomial** time, not efficient as periods become larger
 - Not suitable for online admission control

Lecture summary

- There are better utilization bounds than the Liu & Layland utilization bound: the hyperbolic bound
- When the relative deadline of a task is less than its period, we can apply utilization bounds
 - But such tests are even more pessimistic than normal
- We can apply **exact tests for schedulability** when deadlines are less than or equal to periods
 - Such tests require more computation
 - Iterative process