

# Periodic task scheduling

---

An exact test for EDF

Relative deadlines less than periods

Processor demand criterion

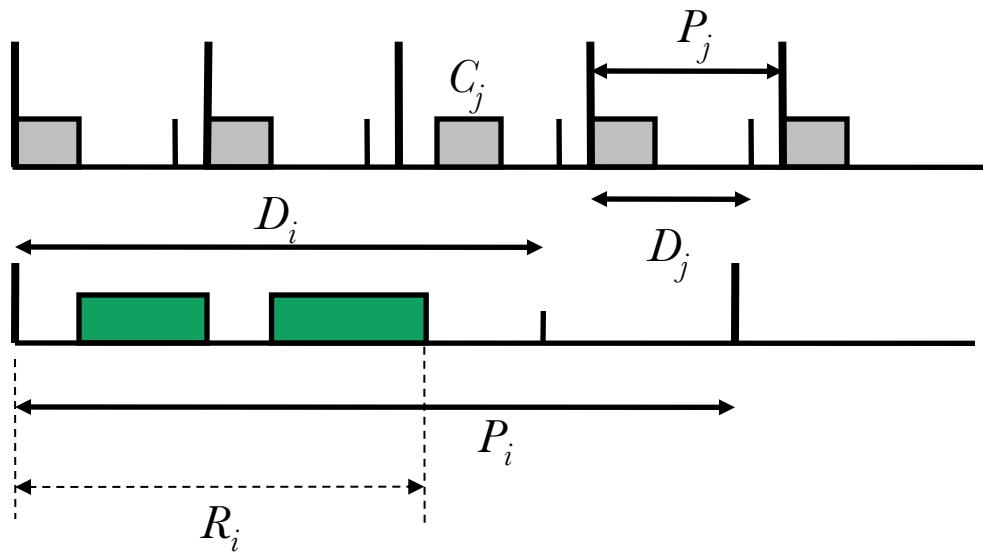
# The earliest deadline first policy

---

- Optimal scheduling policy when relative deadlines are equal to task periods
- Is a task-dynamic priority policy (but not *job*-dynamic)
- How does it behave ***when relative deadlines are less than periods?***
- Exact analysis for EDF

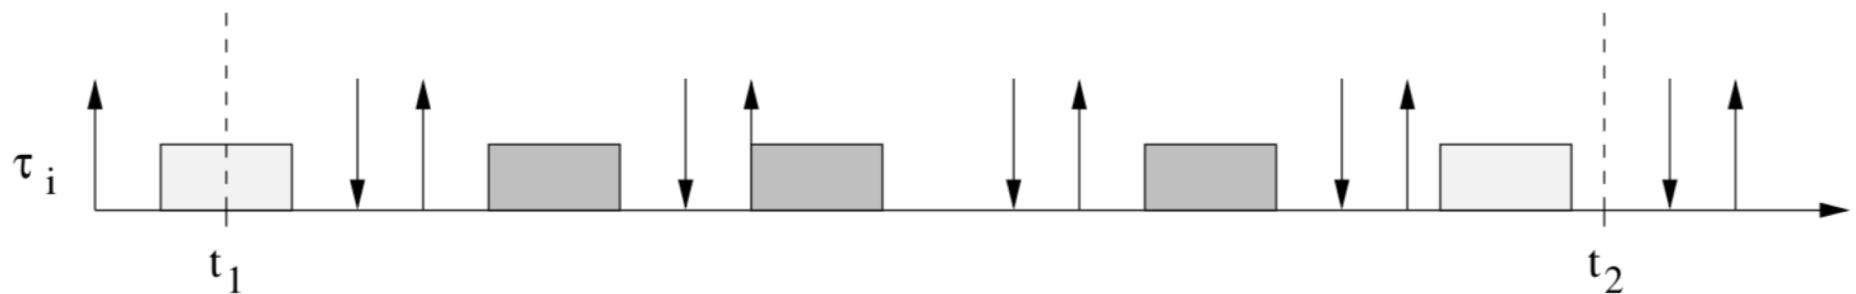
# Why not analyze task interference as we did with DM?

- Interference is due to only those jobs with **earlier absolute deadlines**
- Task  $j$  exercises interference on task  $i$ , and task  $i$  exercises interference on task  $j$
- Interference relations are complex due to dynamic priorities  $\rightarrow$  *will abandon this view*



# EDF and processor demand

- Task  $i$  **requests** service (work)  $C_i$  at every arrival  $r_{i,k}$
- A request becomes **demand** only at the deadline  $d_{i,k}$ 
  - Before the deadline, we are not obliged to complete the request from a feasibility perspective

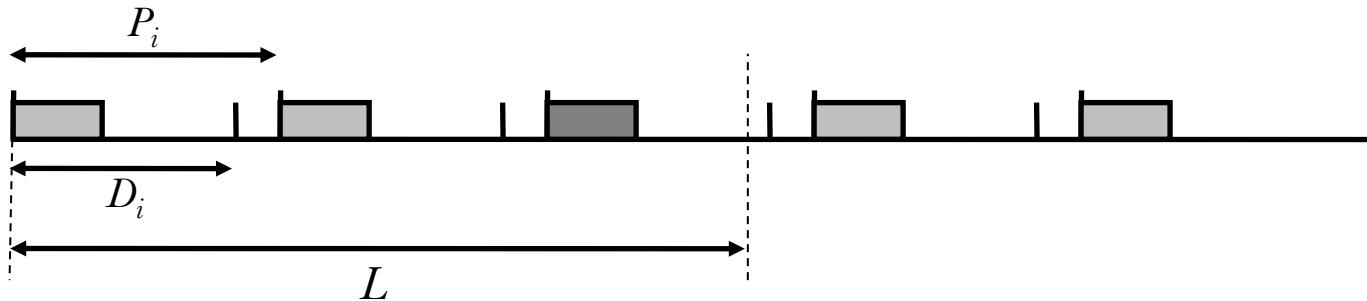


$$\text{demand}_i(t_1, t_2) = \sum_{\substack{k \geq 0 : r_{i,k} \geq t_1, \\ d_{i,k} \leq t_2}} C_i$$

# EDF and processor demand

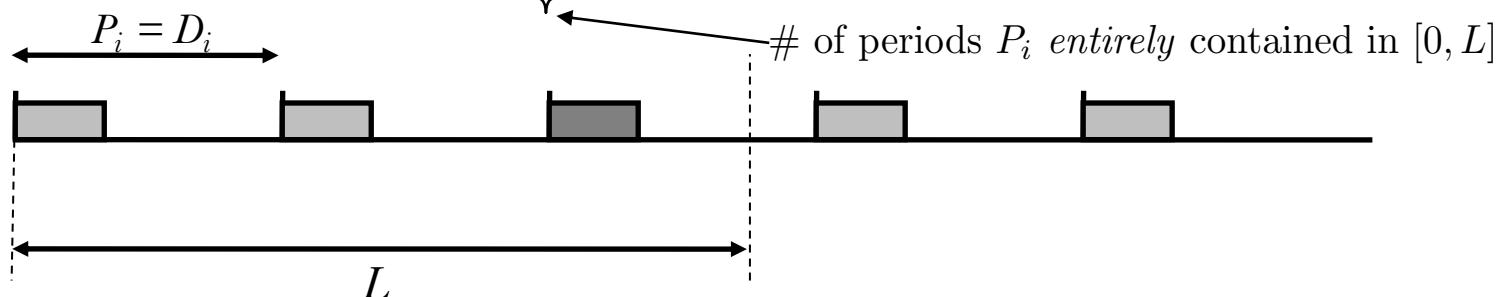
---

- Task  $i$  **requests** service (work)  $C_i$  at every arrival  $r_{i,k}$
- What is the demand in  $[0, L]$  for  $L > 0$ ? ( $t_1 = 0$  and  $t_2 = L$ )



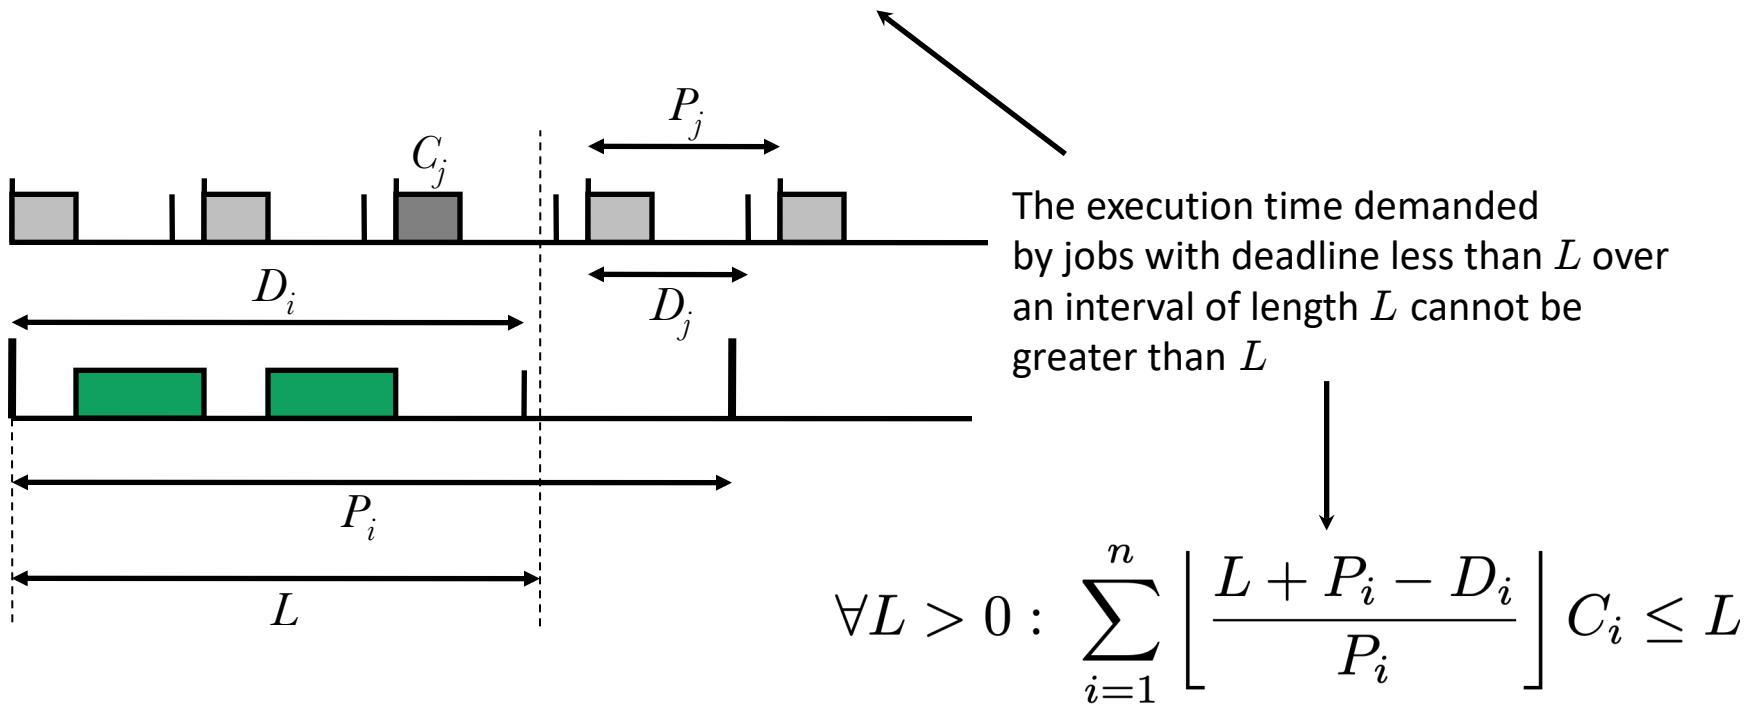
$$\text{demand}_i(0, L) = \sum_{k \geq 0: d_{i,k} \leq L} C_i = \left\lfloor \frac{L + P_i - D_i}{P_i} \right\rfloor C_i$$

If  $P_i = D_i \Rightarrow \text{demand}_i(0, L) = \underbrace{\left\lfloor \frac{L}{P_i} \right\rfloor}_{\# \text{ of periods } P_i \text{ entirely contained in } [0, L]} C_i$



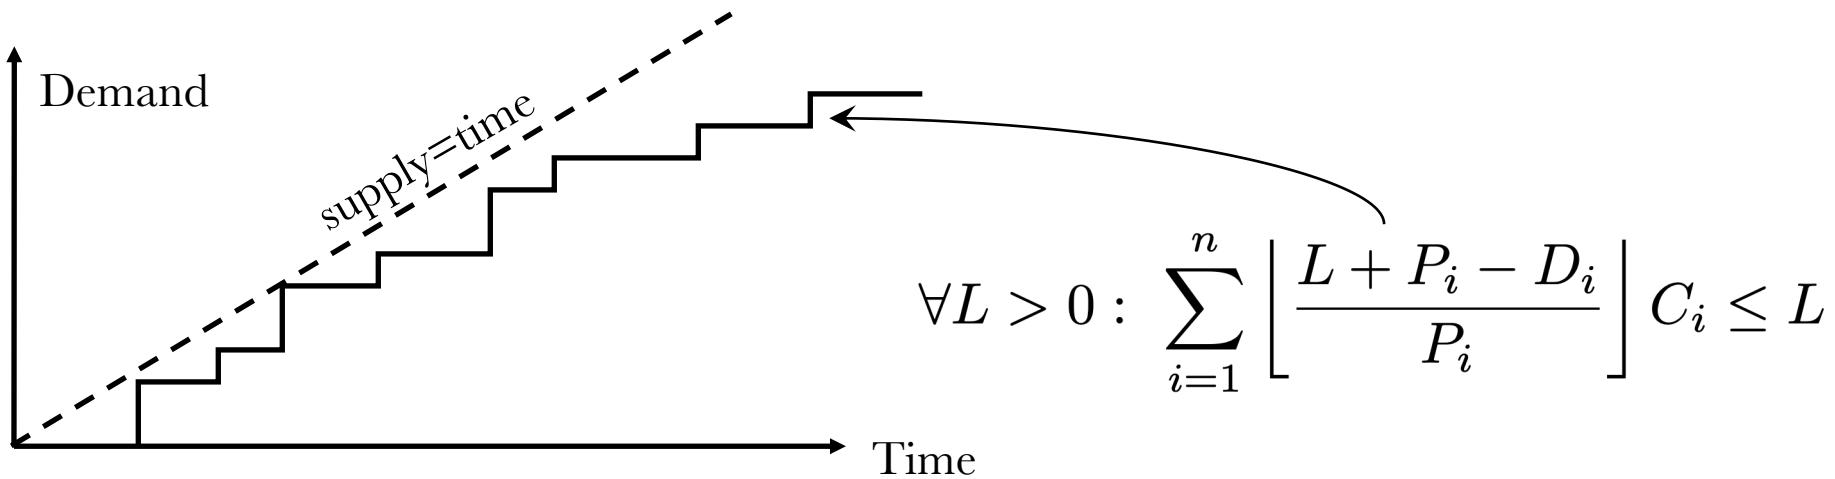
# EDF and processor demand

- Consider demand on the processor due to tasks whose deadlines have passed
- Within any time interval,  $L$ , the demand must be less than  $L$ .



# Processor demand

- Checking the schedulability for all time instants is not possible
  - Overwhelming complexity
- **Observation 1:** Sufficient to check up to the hyperperiod (schedule repeats itself)
- **Observation 2:** Check only at absolute deadlines. Why?

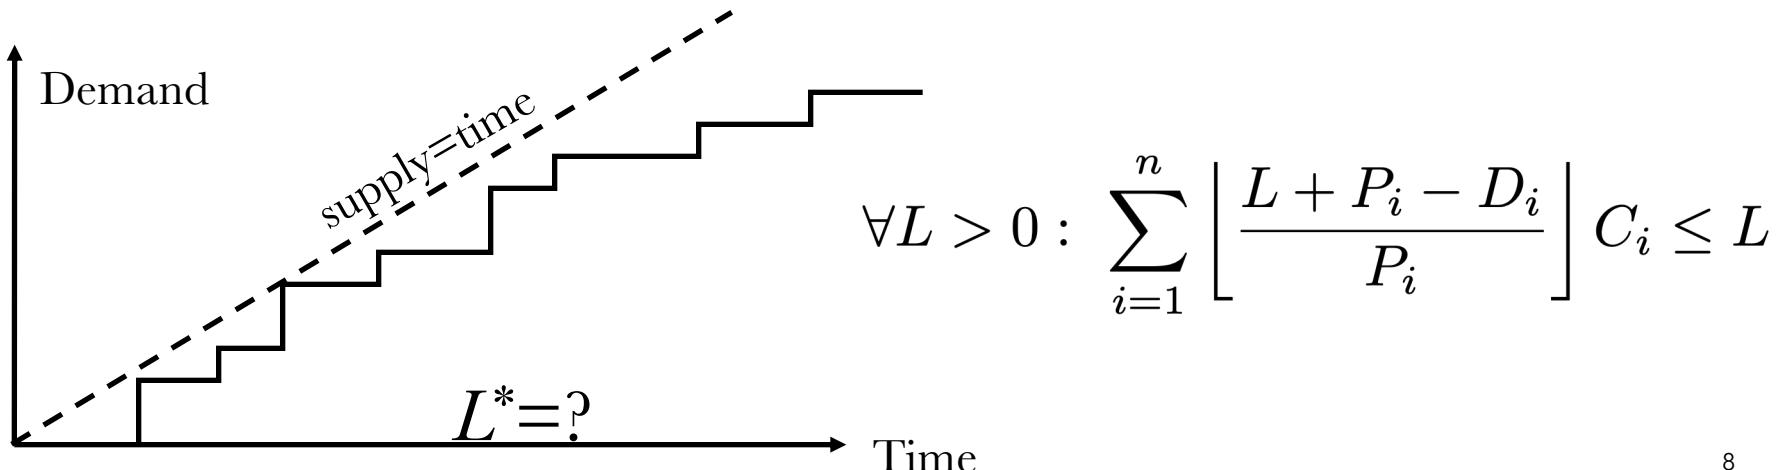


The map  $t \mapsto \left\lfloor \frac{t+P_i-D_i}{P_i} \right\rfloor$  is a *staircase* function with jumps (discontinuities) at absolute deadlines  $d_{i,k} = (k-1)P_i + D_i$

# Processor demand

- Checking the schedulability for at all time instants is not possible
  - Overwhelming complexity
- **Observation 1:** Sufficient to check up to the hyperperiod (schedule repeats itself)
- **Observation 2:** Check only at absolute deadlines
- Can we do better?

Find a point  $L^*$  such that if  $\text{demand}(0, L^*) \leq L$  for all  $L \leq L^*$ , then  $\text{demand}(0, L) \leq L$  for all  $L \geq L^*$   $\rightarrow$  sufficient to test only up to  $L^*$

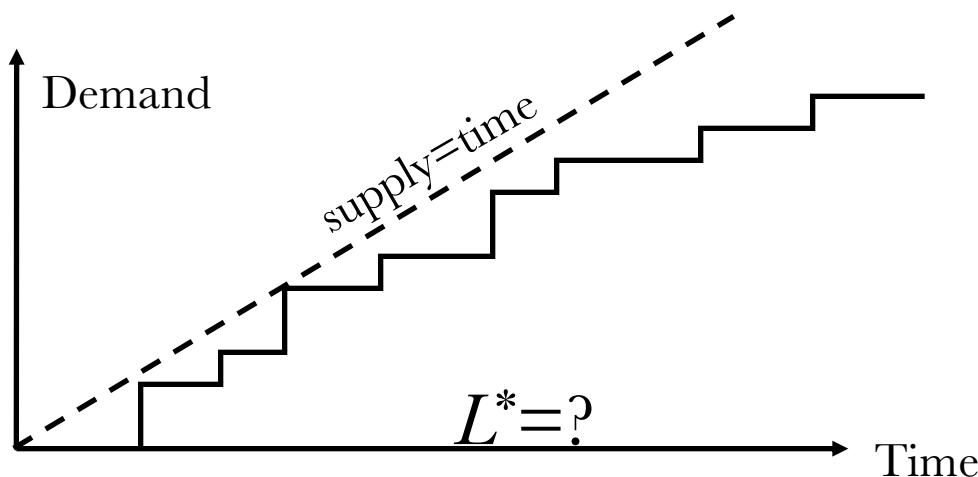


# Processor demand

- Deriving  $L^*$

$$\sum_{i=1}^n \left\lfloor \frac{L + P_i - D_i}{P_i} \right\rfloor C_i \leq L$$

$$\sum_{i=1}^n \left\lfloor \frac{t + P_i - D_i}{P_i} \right\rfloor C_i \leq \sum_{i=1}^n \frac{t - D_i + P_i}{P_i} C_i = tU + \sum_{i=1}^n (P_i - D_i)U_i$$

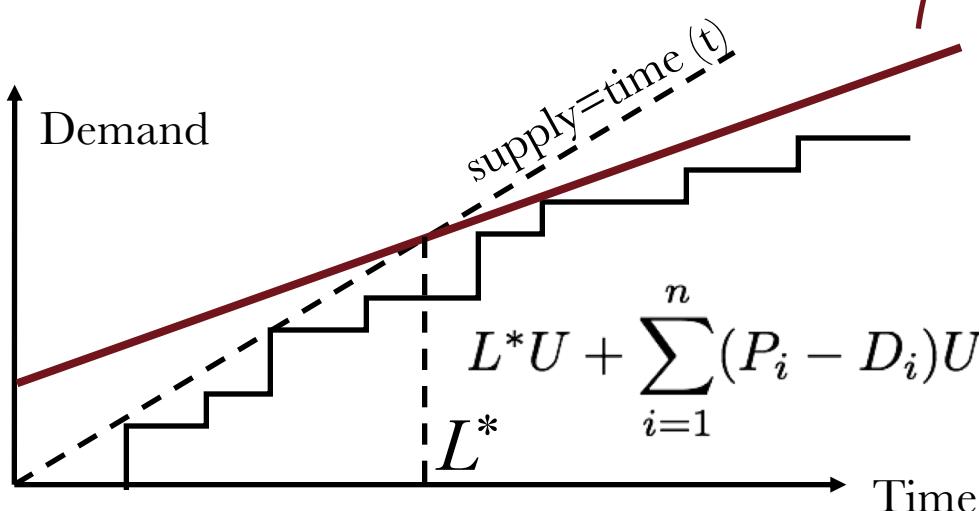


# Processor demand

- Deriving  $L^*$

$$\sum_{i=1}^n \left\lfloor \frac{L + P_i - D_i}{P_i} \right\rfloor C_i \leq L$$

$$\sum_{i=1}^n \left\lfloor \frac{t + P_i - D_i}{P_i} \right\rfloor C_i \leq \sum_{i=1}^n \frac{t - D_i + P_i}{P_i} C_i = tU + \sum_{i=1}^n (P_i - D_i)U_i$$



$$L^*U + \sum_{i=1}^n (P_i - D_i)U_i = L^* \Rightarrow L^* = \frac{\sum_{i=1}^n (P_i - D_i)U_i}{1 - U}$$

# Processor demand criterion

---

- Check if  $\sum_{i=1}^n \left\lfloor \frac{L + P_i - D_i}{P_i} \right\rfloor C_i \leq L$

- for all  $L$  that are absolute deadlines in the interval  $[0, L^*]$
- where

$$L^* = \frac{\sum_{i=1}^n (P_i - D_i) U_i}{1 - U}$$

- This is an **exact** test for EDF when *relative deadlines are less than task periods*
- *What is the complexity of this test? Is it polynomial-time?*
- **Complexity:** Testing constrained-deadline EDF-schedulability is **coNP-hard**.  
(Eisenbrand & Rothvoß - SODA'10)

# Example using processor demand

---

- Consider the following task set
  - $T_1: (C_1=1, P_1=3, D_1=2)$
  - $T_2: (C_2=2, P_2=7, D_2=5.5)$
  - $T_3: (C_3=2, P_3=10, D_3=6)$
- The task set has a hyperperiod of 210
- However, we only need to test deadlines up to  $L^* = \frac{\sum_{i=1}^n (P_i - D_i)U_i}{1 - U}$
- $U = 86/105 = 0.8190$ ;  $L^* = 8.63$
- At  $L = 2$ :

$$\lfloor \frac{L + P_1 - D_1}{P_1} \rfloor C_1 = \lfloor \frac{2 + 3 - 2}{3} \rfloor 1 = 1$$

# Example using processor demand

---

- Consider the following task set
  - $T_1: (C_1=1, P_1=3, D_1=2)$
  - $T_2: (C_2=2, P_2=7, D_2=5.5)$
  - $T_3: (C_3=2, P_3=10, D_3=6)$
- The task set has a hyperperiod of 210
- However, we only need to test deadlines up to  $L^* = \frac{\sum_{i=1}^n (P_i - D_i)U_i}{1 - U}$
- $U = 86/105 = 0.8190$ ;  $L^* = 8.63$
- At  $L=5.5$ : (also need to check at  $L=5$ ; not shown here)

$$\lfloor \frac{L + P_1 - D_1}{P_1} \rfloor C_1 + \lfloor \frac{L + P_2 - D_2}{P_2} \rfloor C_2 = \lfloor \frac{5.5 + 3 - 2}{3} \rfloor 1 + \lfloor \frac{5.5 + 7 - 5.5}{7} \rfloor 2 = 2 + 2 = 4$$

# Example using processor demand

---

- Consider the following task set
  - $T_1: (C_1=1, P_1=3, D_1=2)$
  - $T_2: (C_2=2, P_2=7, D_2=5.5)$
  - $T_3: (C_3=2, P_3=10, D_3=6)$
- The task set has a hyperperiod of 210
- However, we only need to test deadlines up to  $L^* = \frac{\sum_{i=1}^n (P_i - D_i)U_i}{1 - U}$
- $U = 86/105 = 0.8190$ ;  $L^* = 8.63$
- At  $L=6$ :

$$\lfloor \frac{6+3-2}{3} \rfloor 1 + \lfloor \frac{6+7-5.5}{7} \rfloor 2 + \lfloor \frac{6+10-6}{10} \rfloor 2 = 6$$

# Example using processor demand

---

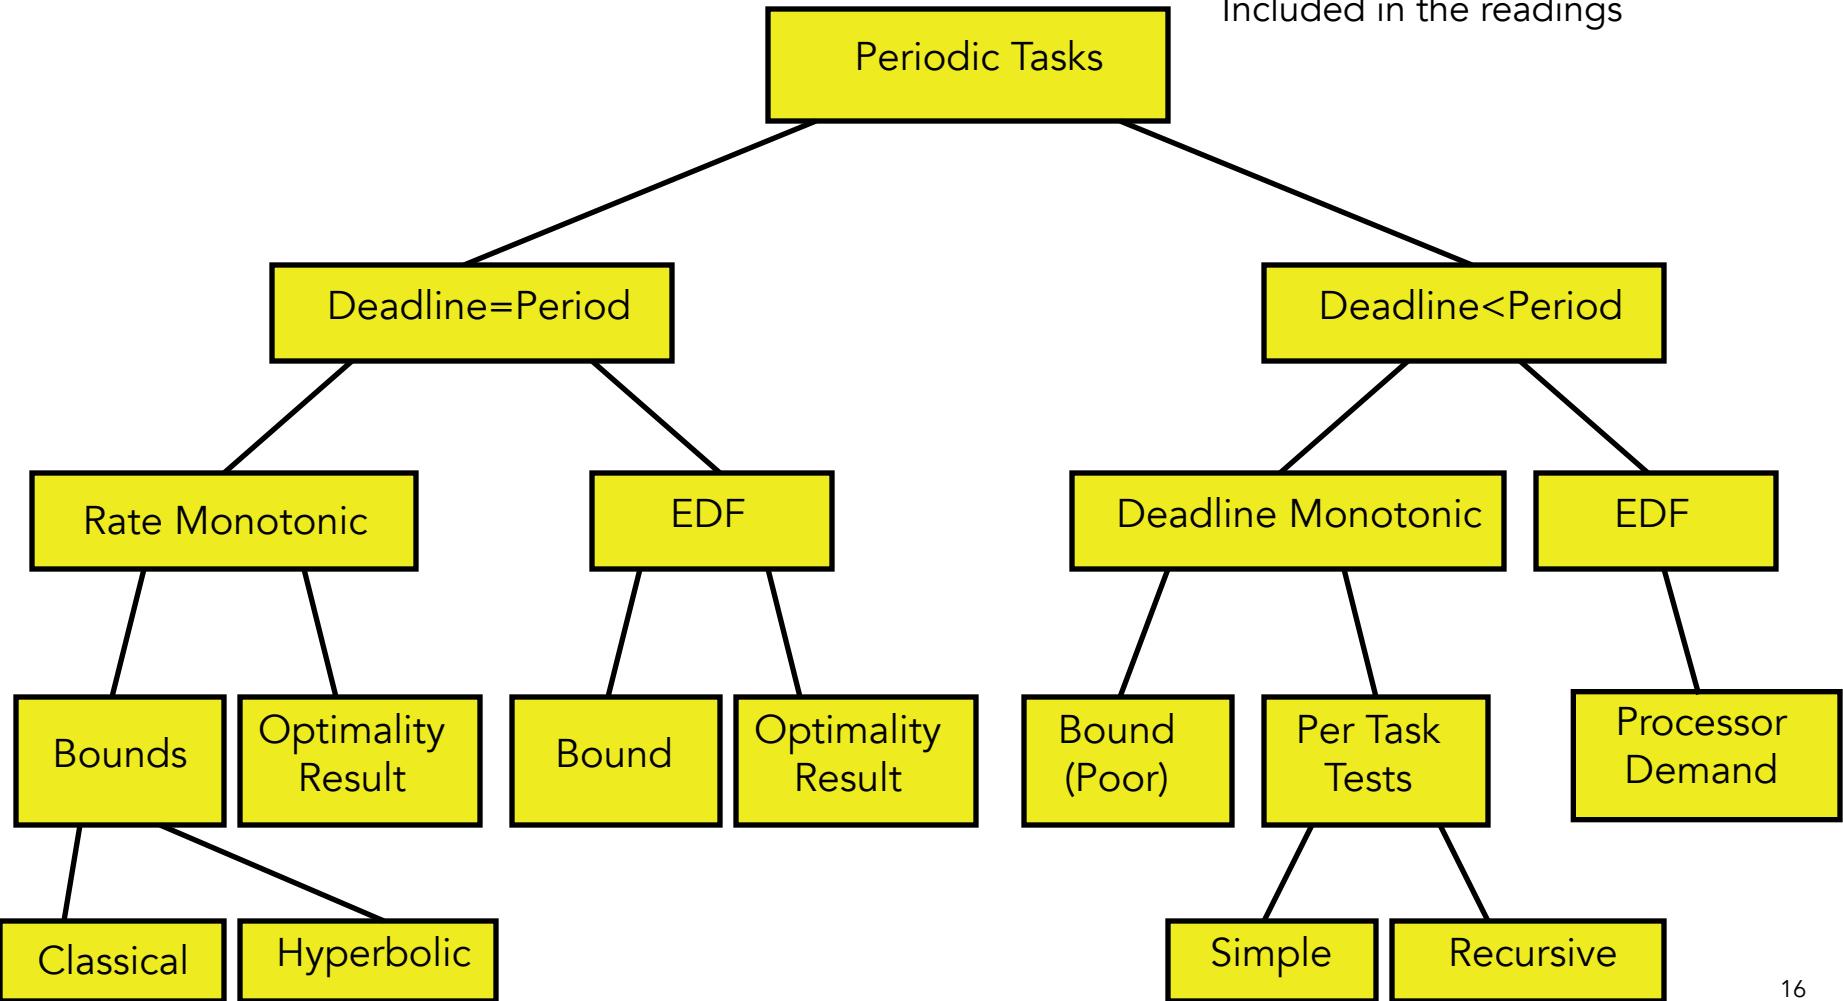
- Consider the following task set
  - $T_1: (C_1=1, P_1=3, D_1=2)$
  - $T_2: (C_2=2, P_2=7, D_2=5.5)$
  - $T_3: (C_3=2, P_3=10, D_3=6)$
- The task set has a hyperperiod of 210
- However, we only need to test deadlines up to  $L^* = \frac{\sum_{i=1}^n (P_i - D_i)U_i}{1 - U}$
- $U = 86/105 = 0.8190$ ;  $L^* = 8.63$
- At  $L=8$ :
$$\lfloor \frac{8+3-2}{3} \rfloor 1 + \lfloor \frac{8+7-5.5}{7} \rfloor 2 + \lfloor \frac{8+10-6}{10} \rfloor 2 = 3 + 2 + 2 = 7$$

No more absolute deadlines < 8.63. We are done!

# Topics covered so far

---

Chapter 4 of Buttazzo's text  
Included in the readings



# RM vs. EDF: Judgement Day

---

- Which policy incurs more preemptions?
- Which has the edge in implementation simplicity and running time complexity?
- How does each behave under *transient* and *permanent* overload conditions?
  - Is one more predictable than the other when overload happens? In what sense?
- Release jitter

# Common misconceptions

---

- EDF incurs more preemptions than RM
- RM is easier to analyze than EDF
- RM introduces less runtime overhead
- EDF is not predictable in transient and permanent overloads
- RM is more predictable in overload conditions, and causes less jitter in task execution

All these statements are either wrong or not precise!

We need to look more carefully into the situation

# Debunking the Myths

# Implementation Cost

---

- For a precise comparison, must distinguish two situations
  - **Case 1:** Scheduler is developed on top of a generic priority-based operating system
  - **Case 2:** Scheduler is developed from scratch

# Implementation Cost

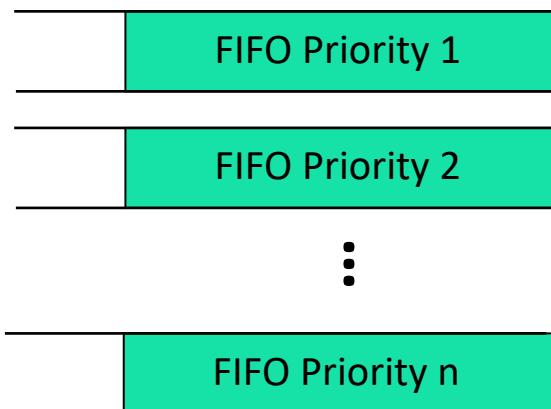
---

- **Case 1:** Scheduler is developed on top of a generic priority-based operating system (limited number of priority levels available)
- **RM wins!**
  - Most commercial kernels contain a *fixed* number of priority levels
  - Consider the case in which two deadlines  $d_a$  and  $d_b$  are mapped into two adjacent priority levels
  - A new periodic instance is released with absolute deadline  $d_c$  s.t.  $d_a < d_c < d_b$ .
  - There is no priority level that can be selected to map  $d_c$ , even when the number of active tasks is less than the number of priority levels in the kernel.
- **Solution:** remap  $d_a$  and  $d_b$  into two new priority levels that are not consecutive → in worst case, all current deadlines may need to be remapped (high cost)

# Implementation Cost

---

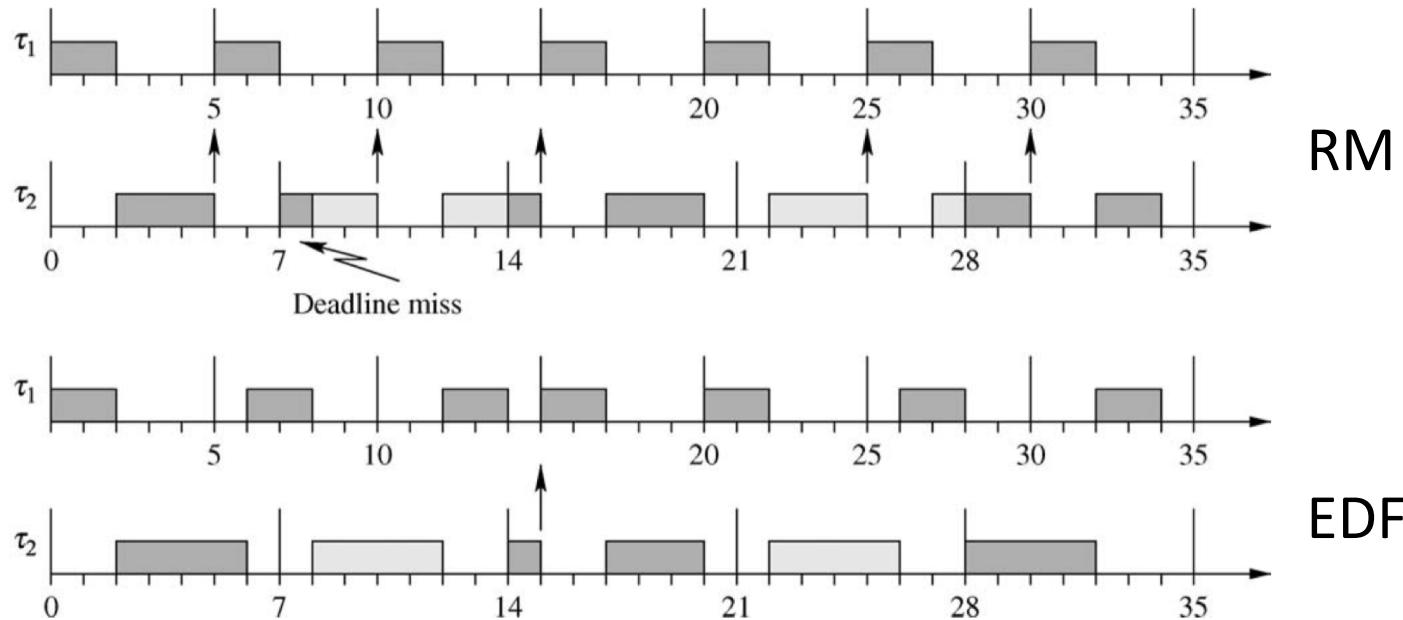
- **Case 2:** Scheduler is developed from scratch
- **A tie!**
- **RM has a tiny edge:** If # of priority levels is low, can split ready queue into separate FIFO queues, one for each priority level → Insertion takes  $O(1)$



- This is not possible for EDF → number of queues would be too large (e.g., equal to  $2^{32}$  if system **time** is represented by four byte variables).

# Runtime Cost

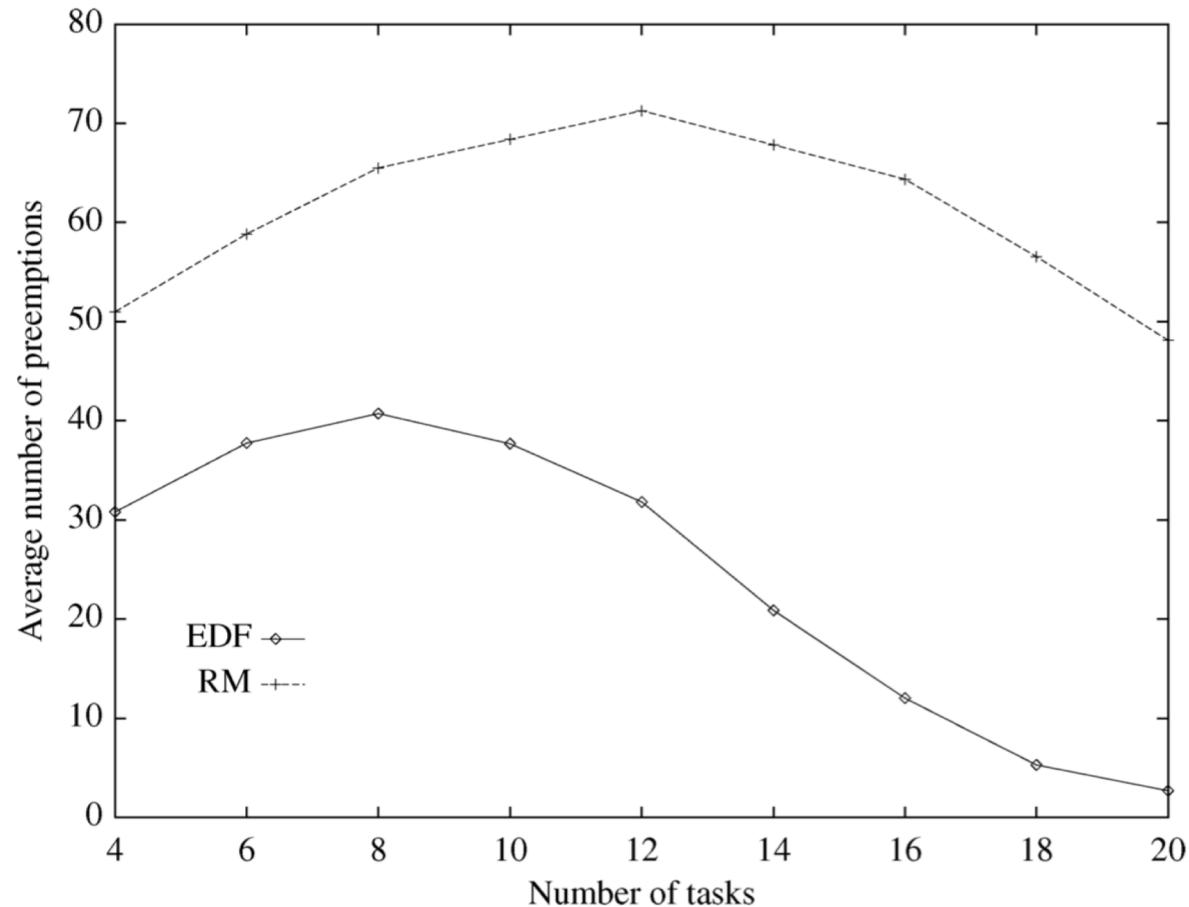
- EDF incurs *less* context switches than RM!



# Context Switching Cost: Interesting Phenomenon

---

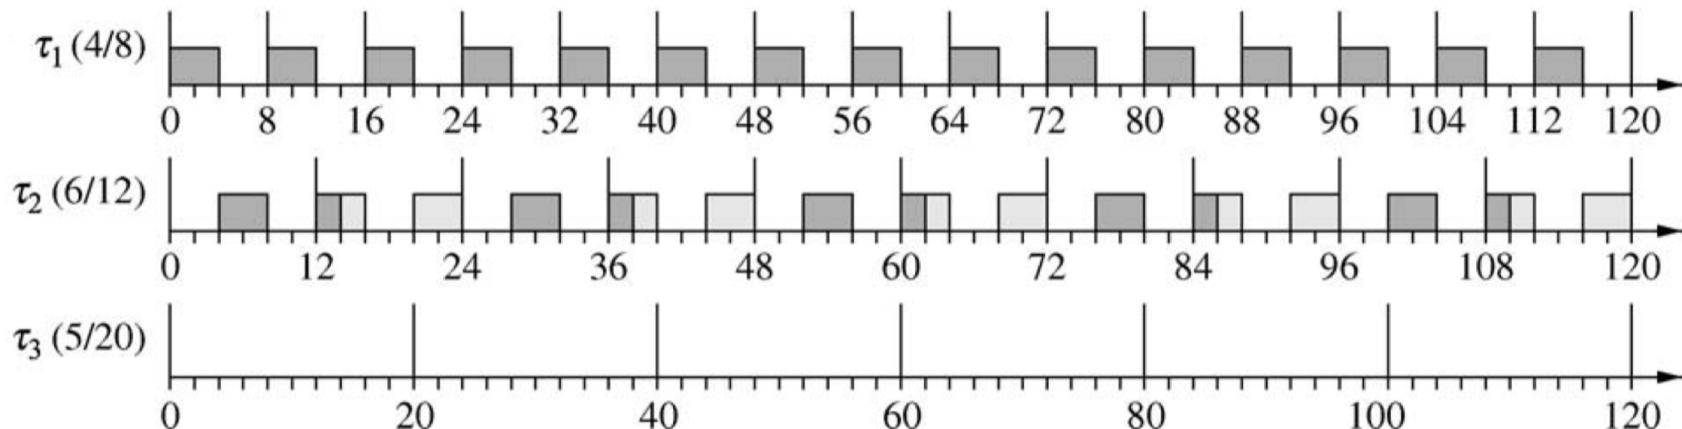
- Utilization kept at 0.9 as number of tasks increases



# Predictability and Stability: Permanent Overload

---

- Permanent Overload  $\rightarrow U > 1$
- Behavior of both RM and EDF *is predictable* in some sense
- RM: Might cause complete blocking of *lower* priority tasks



$$U = 1.25$$

RM

# Predictability and Stability: Permanent Overload

---

- Permanent Overload  $\rightarrow U > 1$
- EDF: Interesting phenomenon happens; *Period rescaling*
  - In the long run, EDF will behave such that it appears that tasks were released at a lower rate  $\rightarrow \bar{P}_i = UP_i$

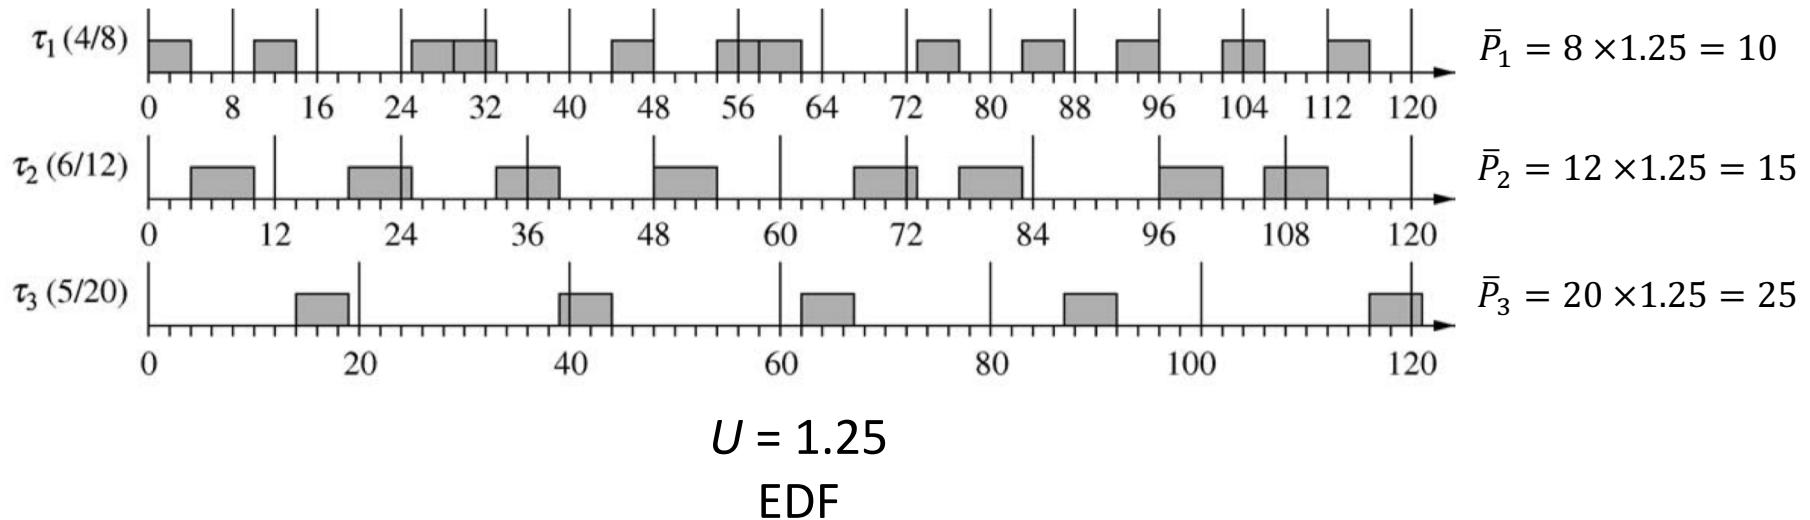
**Theorem 1 [Cervin].** Assume a set of  $n$  periodic tasks, where each task is described by a fixed period  $T_i$ , a fixed execution time  $C_i$ , a relative deadline  $D_i$ , and a release offset  $\Phi_i$ . If  $U > 1$  and tasks are scheduled by EDF, then, in stationarity, the average period  $\bar{T}_i$  of each task  $\tau_i$  is given by  $\bar{T}_i = T_i U$ .

$$\bar{T}_i := \lim_{t \rightarrow \infty} \frac{t}{k_i(t)} = T_i U$$

$k_i(t)$ : # of instances of task  $\tau_i$  that completed up to time  $t > 0$

# Predictability and Stability: Permanent Overload

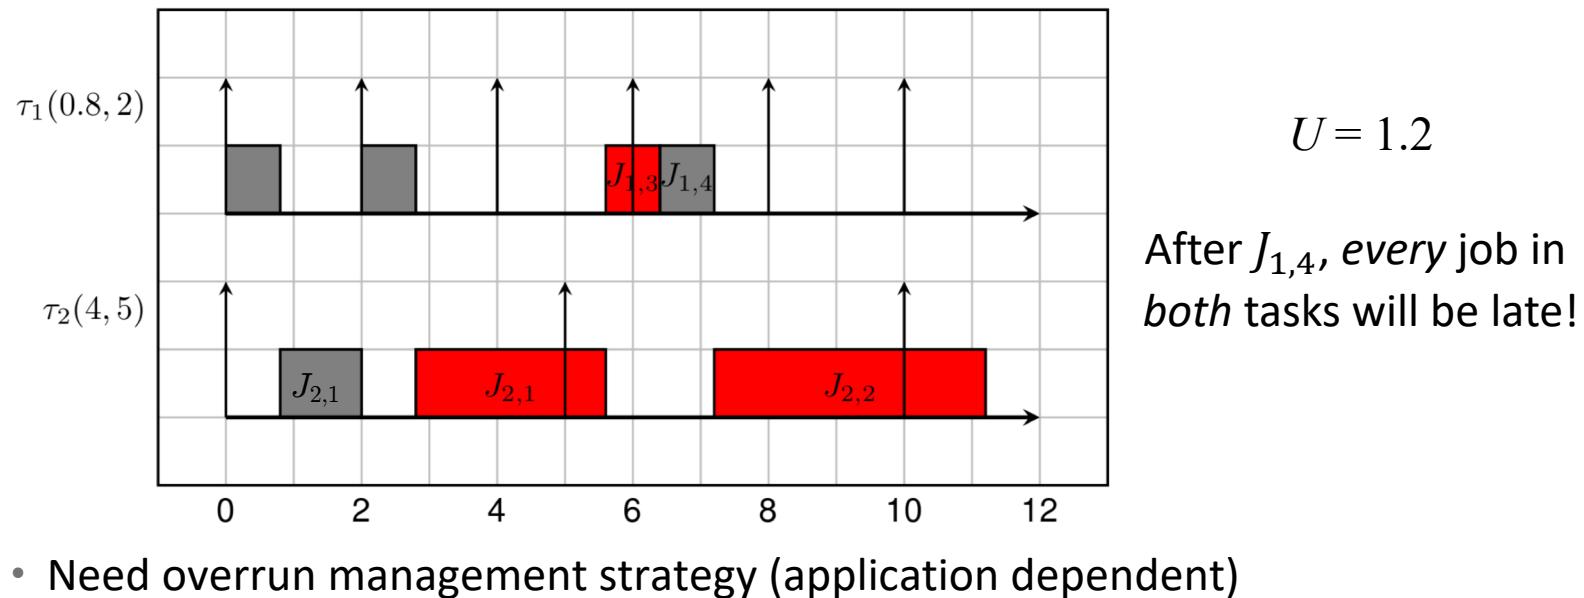
- Permanent Overload  $\rightarrow U > 1$
- EDF Period rescaling



$$\overline{U} = \frac{c_1}{\bar{P}_1} + \frac{c_2}{\bar{P}_2} + \frac{c_3}{\bar{P}_3} = \frac{4}{10} + \frac{6}{15} + \frac{5}{25} = 1$$

# A deadline miss in EDF can be dangerous

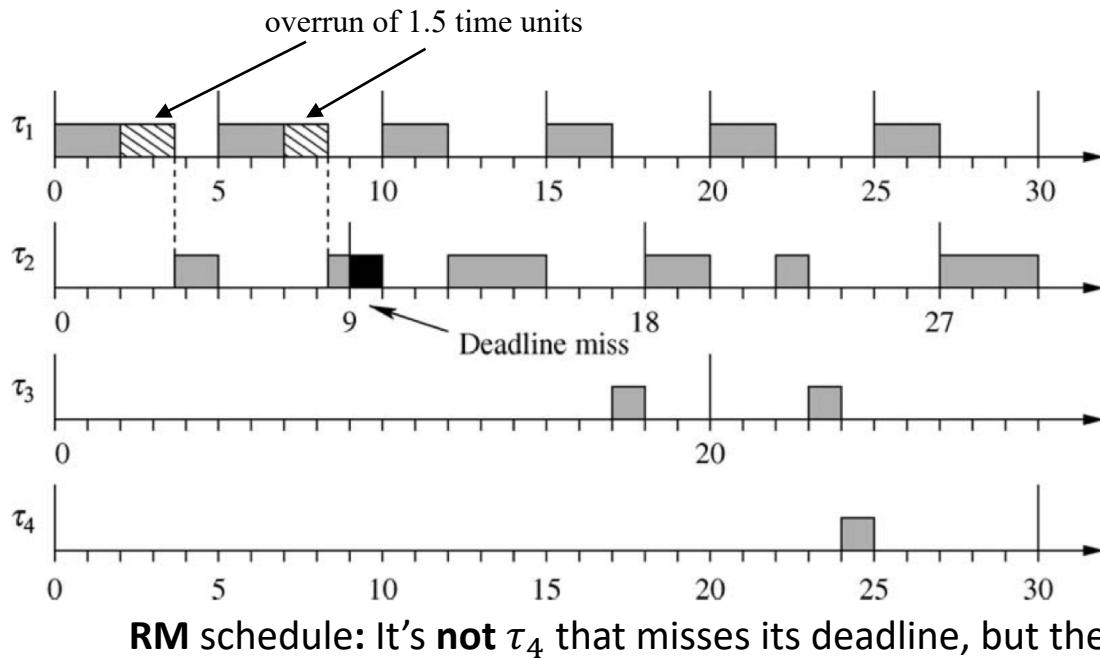
- If job misses deadline it executes at a very high priority level (higher than all tasks that have not missed their deadlines)
- If late job allowed to continue, will cause other jobs to miss their deadlines!



- Need overrun management strategy (application dependent)
  - E.g., schedule each late job at a lower priority than jobs that are not late
  - Drop late job altogether

# Predictability and Stability: *Transient* Overload

- **Misconception:** In RM, in the presence of transient overload conditions, deadlines are missed predictably, that is, the *first tasks that fail are those with the longest period* (lowest priority tasks fail first)



If a task  $\tau_i$  overruns, only the tasks with *higher priority than  $\tau_i$*  are protected under **RM**, but nothing can be ensured for the other tasks.

- **EDF:** situation is *not* better than RM!

# Handling aperiodic work

---

- EDF has the edge since it has higher utilization
- **Slack-stealing:** If  $U_p$  is utilization of periodic work, then the remaining  $1 - U_p$  utilization can be used for aperiodic tasks
  - Increased utilization given to aperiodic servers (discussed later)

# Lecture summary

---

- Exact test for EDF scheduling (relative deadlines < periods)
- Processor demand criterion
- RM vs. EDF

Giorgio C. Buttazzo. **Rate monotonic vs. EDF: Judgment day.**  
*Real-Time Systems*, 29:5–26. 2005