

Schedulability with resource sharing

Priority inheritance protocol

Priority ceiling protocol

Stack resource policy

Lecture overview

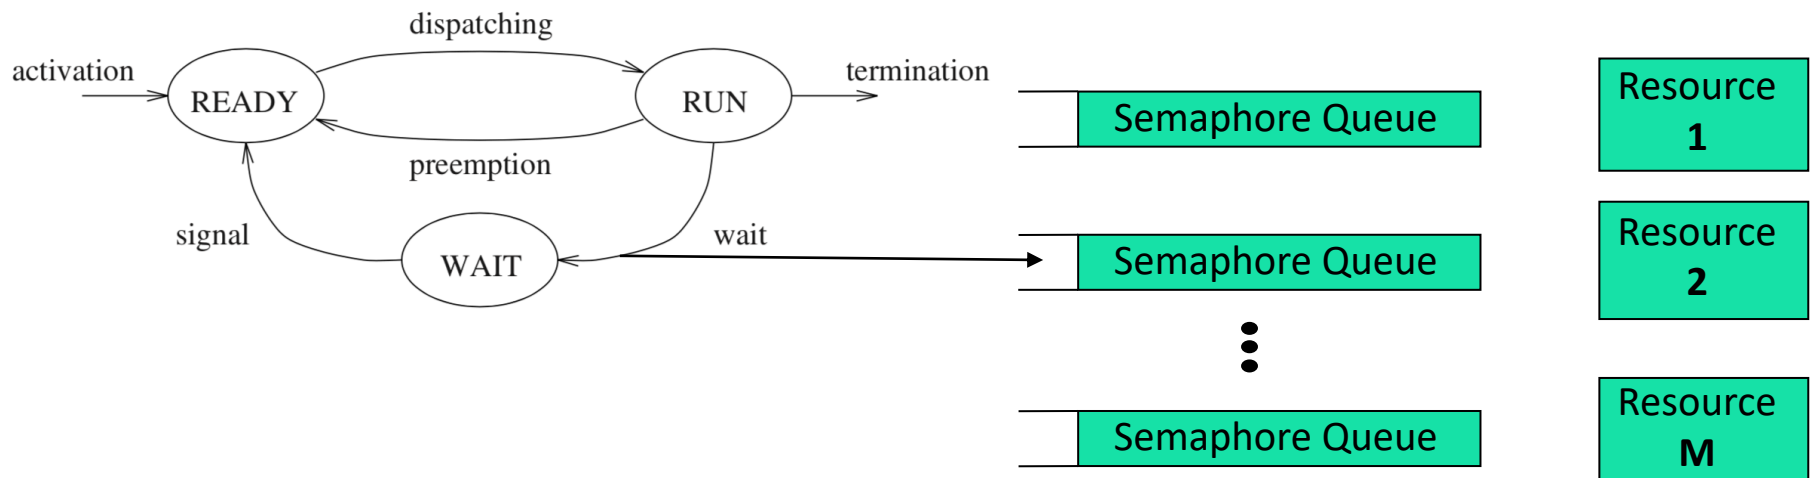
- We have discussed the occurrence of **unbounded priority inversion**
- We know about **blocking** and **blocking times**
- Now: Evaluating schedulability in combination with protocols for avoiding unbounded priority inversion
- **Priority ceiling protocol** to prevent deadlocks
- **Stack-based resource policy**
 - Improves on other policies
 - Extends to EDF

Blocking

- Tasks have synchronization constraints
 - Use semaphores to protect critical sections
- Blocking can cause a ***higher priority task to wait for a lower priority task to unlock a resource***
 - We always assumed that higher priority tasks can preempt lower priority tasks
 - To make rules consistent, we discussed the priority inheritance approach

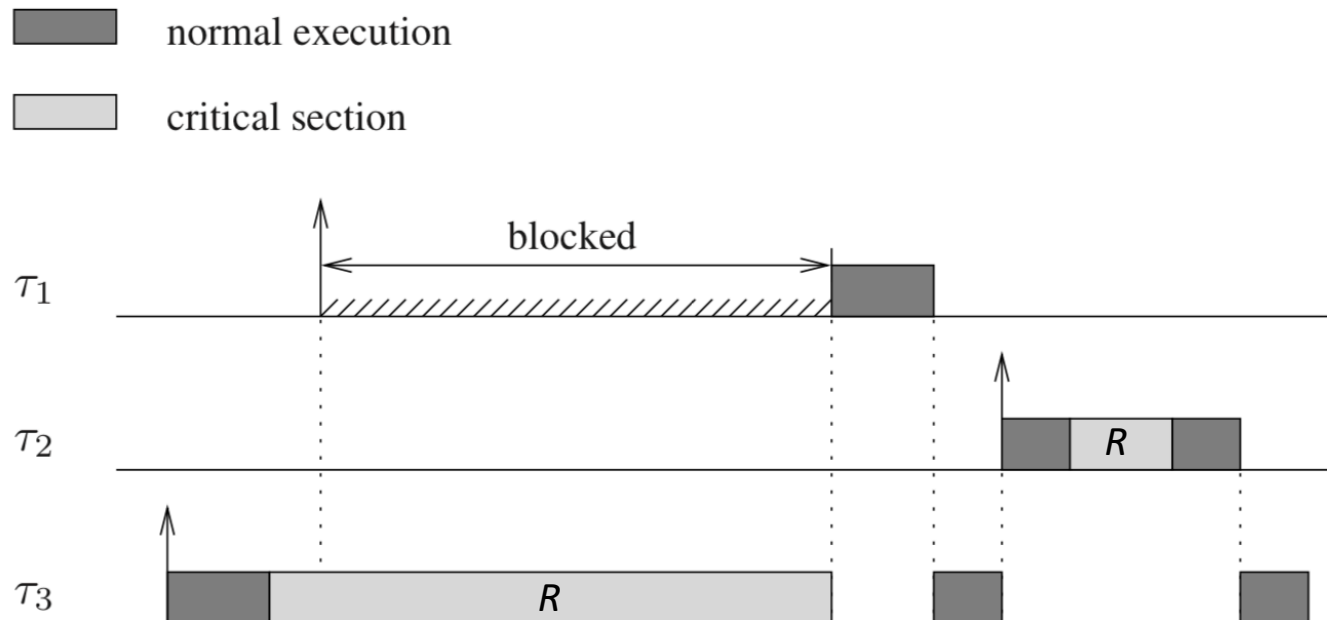
General Model and Assumptions

- **Assumption:** Each resource has one instance only (binary semaphores)
- **Assumption:** Resource requests are *properly nested*
- **Assumption:** *We have perfect knowledge of all task resource requirements*
- *Except for SRP, all protocols are designed for **static-priority** scheduling*
- Each resource has a semaphore queue



Approach #1: Non-Preemptive Protocol (NPP)

- Whenever a task requests a resource, make it the highest priority task *for the duration of its critical section*
- **The good:** Easy to Implement
- **The bad:** unnecessarily blocks higher priority tasks that do not request resources



Non-Preemptive Protocol: Blocking time computation

- A task T_i can be blocked only by a **lower priority** task that has requested a resource before T_i 's arrival
- **Key:** *Whenever a task is in a critical section, it cannot be preempted*
- The lower priority task resumes its static priority as soon as it releases all resources in the critical section
- As soon as the lower priority task releases the resource, the highest priority task available will acquire the processor and it will not be blocked again
- **Conclusion:** Worst case blocking time is the duration of the longest critical section of any lower priority task


$\delta_{j,k}$: duration of *longest* critical section of task j using resource R_k

$$B_i = \max\{\delta_{j,k} : \pi_j < \pi_i, J_j \text{ uses resource } R_k\}$$

Note: If no lower priority task uses any resources, then $B_i = 0$; i.e., $\max \emptyset = 0$

Approach #2: Highest Locker Priority (HLP)

- When a task requests resource R_k , elevate its priority to the priority of the highest priority task that ever shares (uses) resource R_k
- Avoids the unnecessary blocking of higher priority tasks that do not need resources (present in NPP)



Highest Locker Priority: Blocking time computation

- When a task requests resource R_k , elevate its priority to the priority of the highest priority task that ever shares resource R_k
- **Observation:** Task T_i can be blocked only by lower priority tasks that use a resource that is used by a task with priority **greater than or equal** to T_i
- **Claim:** A task T_i can be blocked for at most the duration of a **single** critical section of at most one lower priority task that uses a resource that is used by a task with priority $\geq \pi_i$
- **Ceiling of resource** R_k is the priority of the highest priority task that uses R_k
 - $C(R_k) = \max_{i \in [n]} \{\pi_i : T_i \text{ uses } R_k\}$
- **Claim recast:** A task T_i can be blocked for at most the duration of a **single** critical section of at most one lower priority task that ever uses a resource R_k with $C(R_k) \geq \pi_i$

Highest Locker Priority: Blocking time computation

- **Claim:** A task T_i can be blocked for at most the duration of a **single** critical section of **at most one** lower priority task that ever uses a resource R_k with $C(R_k) \geq \pi_i$

$$B_i = \max\{\delta_{j,k} : \pi_j < \pi_i, T_j \text{ uses } R_k, C(R_k) \geq \pi_i\}$$

Proof of Claim

- Suppose task T_i is blocked by **two** critical sections
- Then both critical sections must belong to two **different** lower priority tasks (why?)

resource R_a used by task T_1

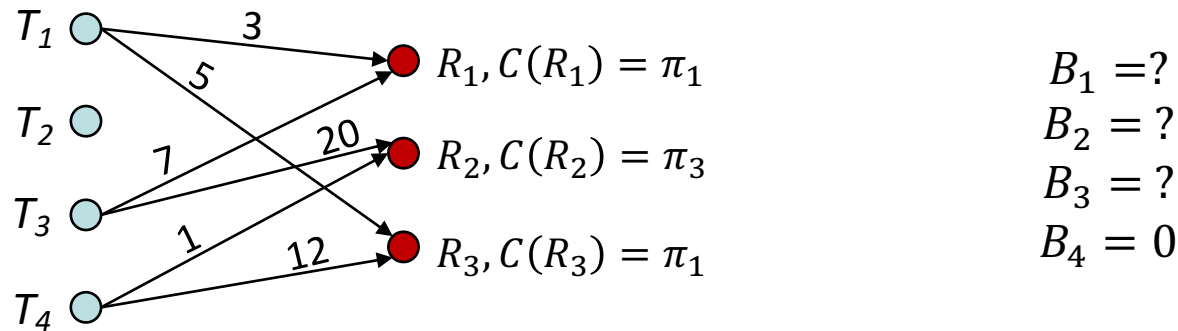
$$\pi_1 < \pi_i \leq C(R_a) \quad (*)$$

resource R_b used by task T_2

$$\pi_2 < \pi_i \leq C(R_b)$$

- Since T_i is blocked on both resources, it must be that tasks T_1 and T_2 were in their critical sections when task T_i arrived
- Then one of T_1 or T_2 must have preempted the other inside its critical section \rightarrow say T_1 preempted T_2 while T_2 is in its critical section using R_b
- This means that $\pi_1 > C(R_b)$
- But $\pi_i > \pi_1 \rightarrow \pi_i > C(R_b)$ *contradicts* (*)

A useful tool: The *resource graph*



$$B_i = \max\{\delta_{j,k} : \pi_j < \pi_i, T_j \text{ uses } R_k, C(R_k) \geq \pi_i\}$$

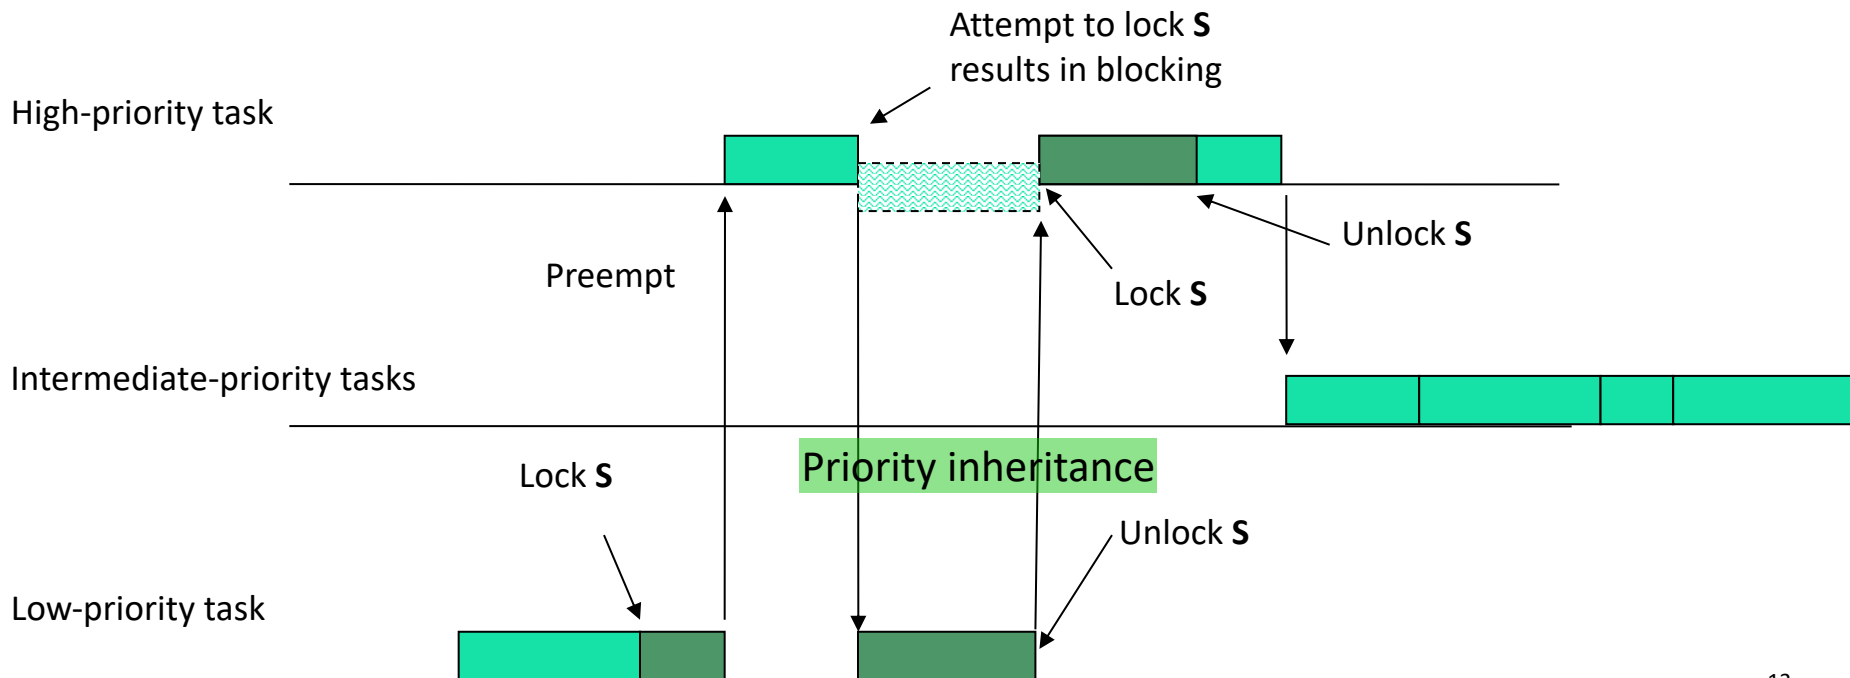
- List jobs in priority order and resources in any order, creating a node for each
- Create an edge between task T_j and resource R_k if T_j uses R_k
- Label arc (T_j, R_k) with the length of the longest critical section of T_j that uses R_k , $\delta_{j,k}$ (even if critical sections are nested)
- Label each resource node R_k by its ceiling $C(R_k)$

Highest Locker Priority: Problems


- HLP causes *unnecessary blocking*
 - A higher priority task is blocked on its *arrival*, not when it attempts to request the resource
 - **Solution:** delay blocking until the task *attempts to request* shared resource → **PIP!**

The priority inheritance protocol

- Allow a task to **inherit the priority** of the highest priority task that it is blocking
- When a hp task is blocked as it attempts to acquire a resource that is held by a lower priority task, it **transfers** its priority to that lower priority task



The priority inheritance protocol

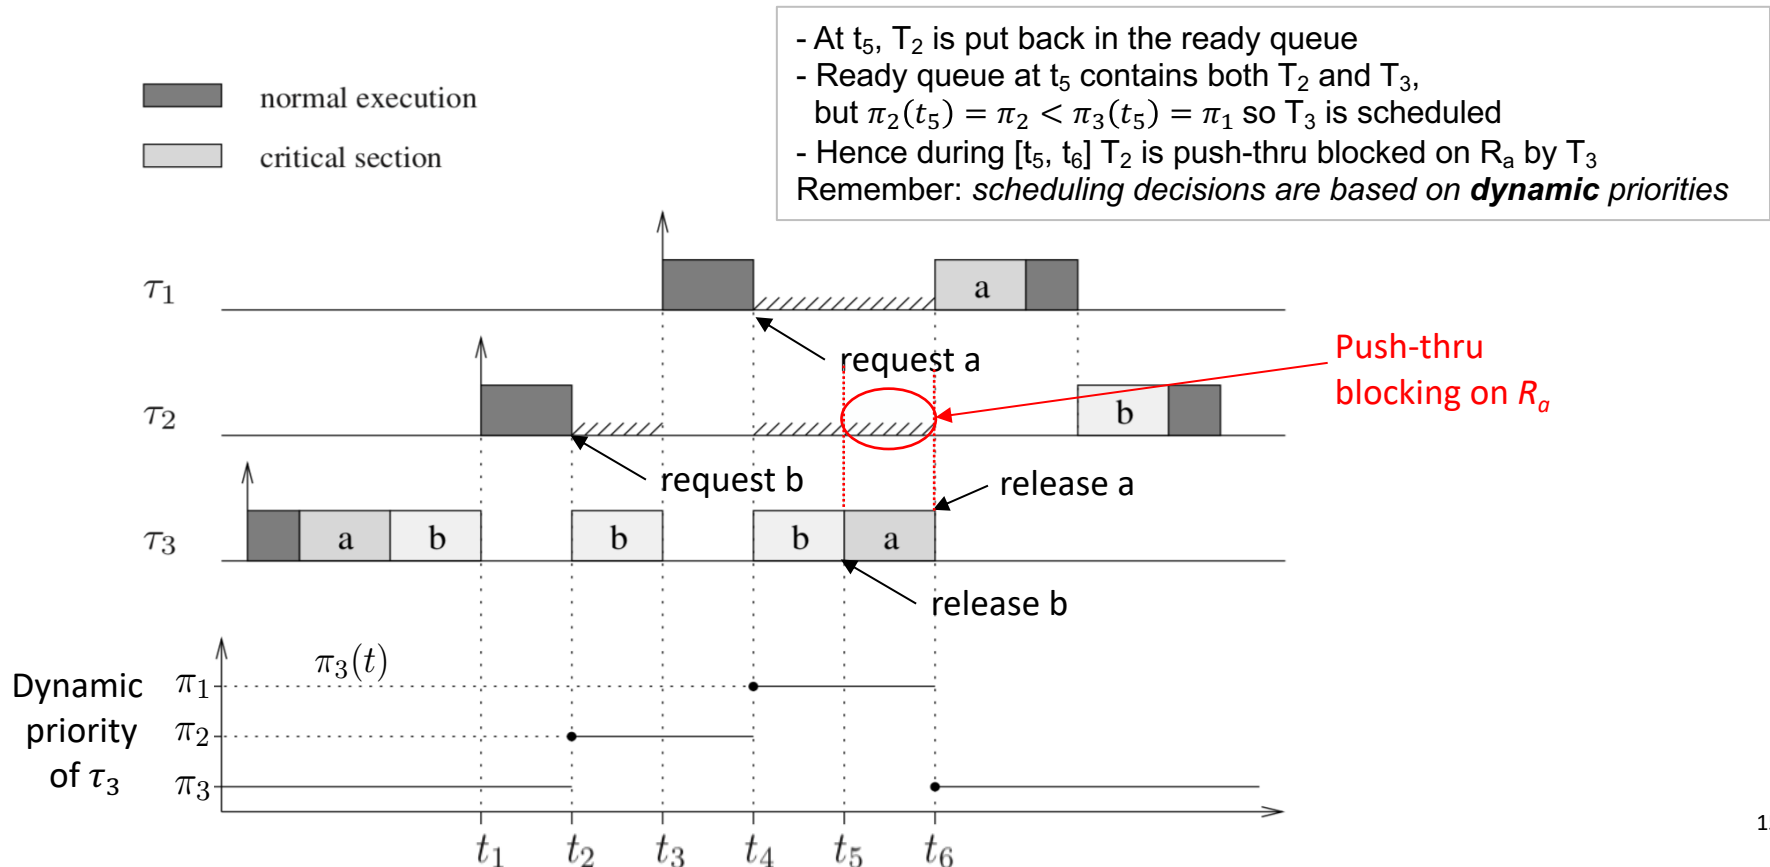


If I am a task, priority inversion occurs when

- (a) Lower priority task holds a resource I need (**direct blocking**)
- (b) Lower priority task inherits a higher priority than me because it holds a resource the higher-priority task needs (**push-through blocking**)

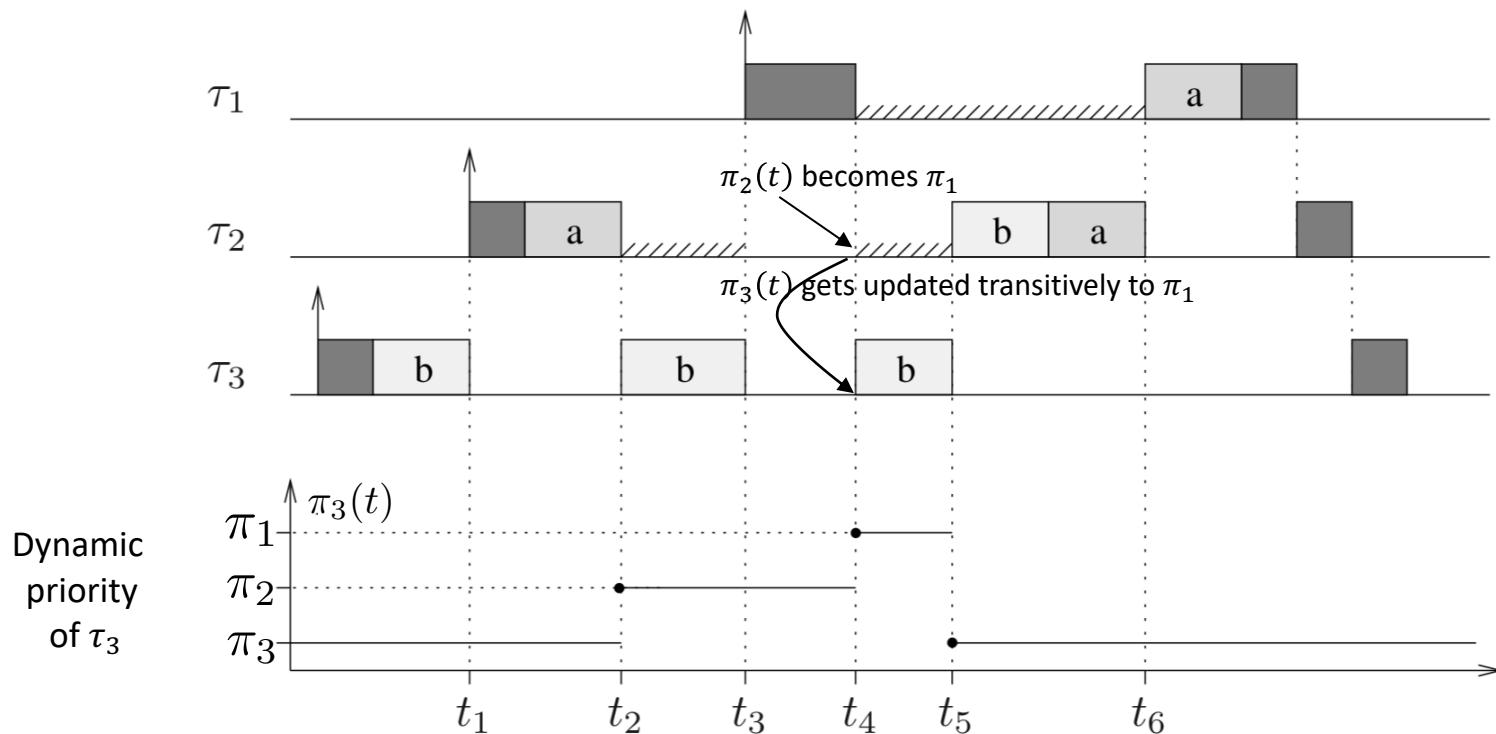
The priority inheritance protocol

- Resource **Release** Rule: When a task **releases** a resource, its dynamic priority $\pi(t)$ is set to the highest priority of the tasks **currently blocked by it**
- Q:** When a task exits a critical section, does it always resume the priority it had when it entered?



The priority inheritance protocol

- Priority inheritance is *transitive*
 - However, transitive priority inheritance can occur only in the presence of **nested critical sections** (proof in book Lemma 7.2)



Maximum blocking time

- **Claim1:** *If there are ℓ_i **lower-priority** tasks that can block task τ_i , then τ_i can be blocked for at most the duration of ℓ_i **critical sections** (one for each of the ℓ_i lower-priority tasks), regardless of the number of semaphores used by τ_i*
 - A critical section $z_{j,k}$ of a lower priority task T_j can block T_i if it causes either direct or push-thru blocking to T_i
- **Claim2:** *If there are s_i **distinct semaphores** that can block task τ_i , then τ_i can be blocked for at most the duration of s_i **critical sections**, one for each of the s_i semaphores, regardless of the number of critical sections used by τ_i*
- Then, if all critical sections are of equal length, b_i
 - Blocking time $B_i = b_i \times \min(\ell_i, s_i)$
- What if the critical sections are of differing lengths?

General approach to computing blocking times

- What if the critical sections are of differing lengths?
- Will consider a safe approximation to blocking time.
- **Assumption:** *no nested critical sections*
- For a high-priority task
 - Examine all tasks with lower priority
 - Determine the worst-case blocking that it may offer (consider the highest priority that it can inherit)
 - Examine all semaphores/resources
 - Determine the worst-case blocking due to that resource
 - Consider lower-priority tasks that may inherit a higher priority when they hold the semaphore

Maximum blocking time

This is just a safe approximation
(upper bound on exact blocking time)
Exact blocking time computation is intractable

- What if the critical sections are of differing lengths?
- $\delta_{j,k}$: length of *longest* critical section among all those of T_j guarded by semaphore S_k
- Let $z_{j,k}$ denote the critical section (CS) whose length is $\delta_{j,k}$
- **(1) Blocking due to lower priority tasks that *can block* T_i (claim1)**
 - A task T_j **can block** T_i if it has lower priority than T_i and uses some resource R_k that is also used by a task with priority greater than or equal to T_i

$$B_i^\ell = \sum_{j=i+1}^n \max_{k \in \{1, \dots, m\}} \{ \delta_{j,k} : z_{j,k} \text{ is max. length CS that can block } T_i \}$$

- **(2) Blocking due to semaphores that can block T_i (claim 2)**
 - A resource **can block** T_i if it is used by a lower priority task **and** a task with priority $\geq \pi_i$

$$B_i^s = \sum_{k=1}^m \max_{j > i} \{ \delta_{j,k} : z_{j,k} \text{ is max. length CS that can block } T_i \}$$

$$B_i = \min(B_i^\ell, B_i^s)$$

Simplifying matters

- Use resource ceilings (very useful device)
- Recall: $C(R_k) = \max_{i \in [n]} \{\pi_i : T_i \text{ uses } R_k\}$
- **Claim:** *In the absence of nested critical sections, a critical section $z_{j,k}$ of τ_j using resource R_k can block τ_i only if $\pi_j < \pi_i \leq C(R_k)$*
 - Proof in text; Lemma 7.5

$$B_i^\ell = \sum_{j=i+1}^n \max_k \{ \delta_{j,k} : z_{j,k} \text{ is max. length CS that can block } T_i \}$$

$$B_i^s = \sum_{k=1}^m \max_{j>i} \{ \delta_{j,k} : z_{j,k} \text{ is max. length CS that can block } T_i \}$$



$$B_i^\ell = \sum_{j=i+1}^n \max_k \{ \delta_{j,k} : C(R_k) \geq \pi_i \}$$

$$B_i^s = \sum_{k=1}^m \max_{j>i} \{ \delta_{j,k} : C(R_k) \geq \pi_i \}$$

Schedulability tests

- For the **fixed-priority** scheduling case
 - We can use the Liu & Layland bound with some modifications
- For task T_k : we need to consider the blocking by lower priority tasks

$$\frac{e_k + B_k}{P_k} + \sum_{i=1}^{k-1} \frac{e_i}{P_i} \leq k(2^{1/k} - 1)$$

Each instance of a task might experience blocking (worst case); equivalent to increasing the execution time of the task by the blocking time.

For task T_k , we need to consider:
(a) preemption by higher priority tasks
(b) blocking from lower priority tasks
bound for T_k involves only k tasks

Why do we test each task separately?
Why can we not have one utilization bound test like we did earlier?

Example: blocking and schedulability

- Consider the following set of tasks, which share resources R_1 , R_2 and R_3
 - Relative deadline are equal to periods; tasks scheduled using RM policy
 - T_1 : $P_1=20$, $e_1=3$, uses R_1 and R_2 separately for 1 time unit each
 - T_2 : $P_2=30$, $e_2=6$, uses R_2 and R_3 **simultaneously** for 2 time units
 - T_3 : $P_3=50$, $e_2=10$, uses R_1 and R_3 **separately** for 3 and 4 time units respectively
 - T_4 : $P_4=80$, $e_2=8$, uses R_2 for 5 time units

Is there a difference?

Without resource constraints

$$U = \frac{3}{20} + \frac{6}{30} + \frac{10}{50} + \frac{8}{80} = 0.65 < 0.69$$

The task set satisfies the Liu and Layland bound; easily schedulable by RM

Example: blocking and schedulability

- Consider the following set of tasks, which uses resources R_1 , R_2 and R_3
 - Relative deadline are equal to periods; tasks scheduled using RM policy
 - T_1 : $P_1=20$, $e_1=3$, uses R_1 and R_2 separately for 1 time unit each
 - T_2 : $P_2=30$, $e_2=6$, uses R_2 and R_3 simultaneously for 2 time units
 - T_3 : $P_3=50$, $e_2=10$, uses R_1 and R_3 separately for 3 and 4 time units respectively
 - T_4 : $P_4=80$, $e_2=8$, uses R_2 for 5 time units

With resource constraints

T_1 can potentially be blocked by T_2 , T_3 and T_4

It can be blocked by T_2 on resource R_2 for up to 6 time units (because it might wait for T_3)

It can be blocked by T_3 on resource R_1 for up to 3 time units

It can be blocked by T_4 on resource R_2 for up to 5 time units

Then maximum wait on lower priority tasks is $B_1^{\ell} = 6 + 3 + 5 = 14$

The worst-case wait for R_1 is 3 units (only T_3 can block T_1)

The worst-case wait for R_2 is 6 units (T_2 can block T_1 for 6 units or T_4 can block T_1 for 5 units)

Then maximum wait for resources is $B_1^s = 3 + 6 = 9$

Then $B_1 = \min(14, 9) = 9$

$$\frac{9}{20} + \frac{3}{20} < 1$$

T_1 is schedulable

$$\frac{B_k}{P_k} + \sum_{i=1}^k \frac{e_i}{P_i} \leq k(2^{1/k} - 1)$$

Example: blocking and schedulability

- Consider the following set of tasks, which uses resources R_1 , R_2 and R_3
 - Relative deadline are equal to periods; tasks scheduled using RM policy
 - T_1 : $P_1=20$, $e_1=3$, uses R_1 and R_2 separately for 1 time unit each
 - T_2 : $P_2=30$, $e_2=6$, uses R_2 and R_3 simultaneously for 2 time units
 - T_3 : $P_3=50$, $e_3=10$, uses R_1 and R_3 separately for 3 and 4 time units respectively
 - T_4 : $P_4=80$, $e_4=8$, uses R_2 for 5 time units

With resource constraints

T_2 can be blocked by T_3 and T_4

T_3 can block T_2 in two ways:

directly on R_3 (upto 4 units)

by obtaining priority of T_1 when using R_1 (upto 3 units) (push-through)

T_4 can block T_2 in two ways:

directly when using R_2 (upto 5 units)

by obtaining priority of T_1 when using R_2 (upto 5 units) (push-through)

The worst-case blocking by T_3 is 4 time units

The worst-case blocking by T_4 is 5 time units

Maximum wait for resources is $B_2 = 5 + 4 = 9 = B_2^{\ell}$ (check for yourself that $B_2^s = 12$)

$$\frac{B_k}{P_k} + \sum_{i=1}^k \frac{e_i}{P_i} \leq k(2^{1/k} - 1)$$

A low priority task can block a high priority task at most once. With priority inheritance, it will get a higher priority and continue till it releases the lock. Therefore, it can block a high priority task at most once.

$$\frac{9}{30} + \left(\frac{3}{20} + \frac{6}{30} \right) = 0.65 < 0.82$$

T_2 is schedulable

Example: blocking and schedulability

- Consider the following set of tasks, which uses resources R_1 , R_2 and R_3
 - Relative deadline are equal to periods; tasks scheduled using RM policy
 - T_1 : $P_1=20$, $e_1=3$, uses R_1 and R_2 separately for 1 time unit each
 - T_2 : $P_2=30$, $e_2=6$, uses R_2 and R_3 simultaneously for 2 time units
 - T_3 : $P_3=50$, $e_3=10$, uses R_1 and R_3 separately for 3 and 4 time units respectively
 - T_4 : $P_4=80$, $e_4=8$, uses R_2 for 5 time units

With resource constraints

T_3 can be blocked by T_4

even when it shares no resource with T_4 (lower priority task)

Notice that T_4 might execute with priority of T_1 (priority inheritance)

T_4 might execute with the priority of T_1 for at most 5 time units

Classic case of push-through blocking

Maximum blocking due to T_4 is 5 time units; $B_3 = 5$

$$\frac{B_k}{P_k} + \sum_{i=1}^k \frac{e_i}{P_i} \leq k(2^{1/k} - 1)$$

$$\frac{5}{50} + \left(\frac{3}{20} + \frac{6}{30} + \frac{10}{50} \right) = 0.65$$

T_3 is schedulable

Example: blocking and schedulability

- Consider the following set of tasks, which uses resources R_1 , R_2 and R_3
 - Relative deadline are equal to periods; tasks scheduled using RM policy
 - T_1 : $P_1=20$, $e_1=3$, uses R_1 and R_2 separately for 1 time unit each
 - T_2 : $P_2=30$, $e_2=6$, uses R_2 and R_3 simultaneously for 2 time units
 - T_3 : $P_3=50$, $e_2=10$, uses R_1 and R_3 separately for 3 and 4 time units respectively
 - T_4 : $P_4=80$, $e_2=8$, uses R_2 for 5 time units

$$\frac{B_k}{P_k} + \sum_{i=1}^k \frac{e_i}{P_i} \leq k(2^{1/k} - 1)$$

With resource constraints

T_4 can never be blocked

because it is the lowest priority task

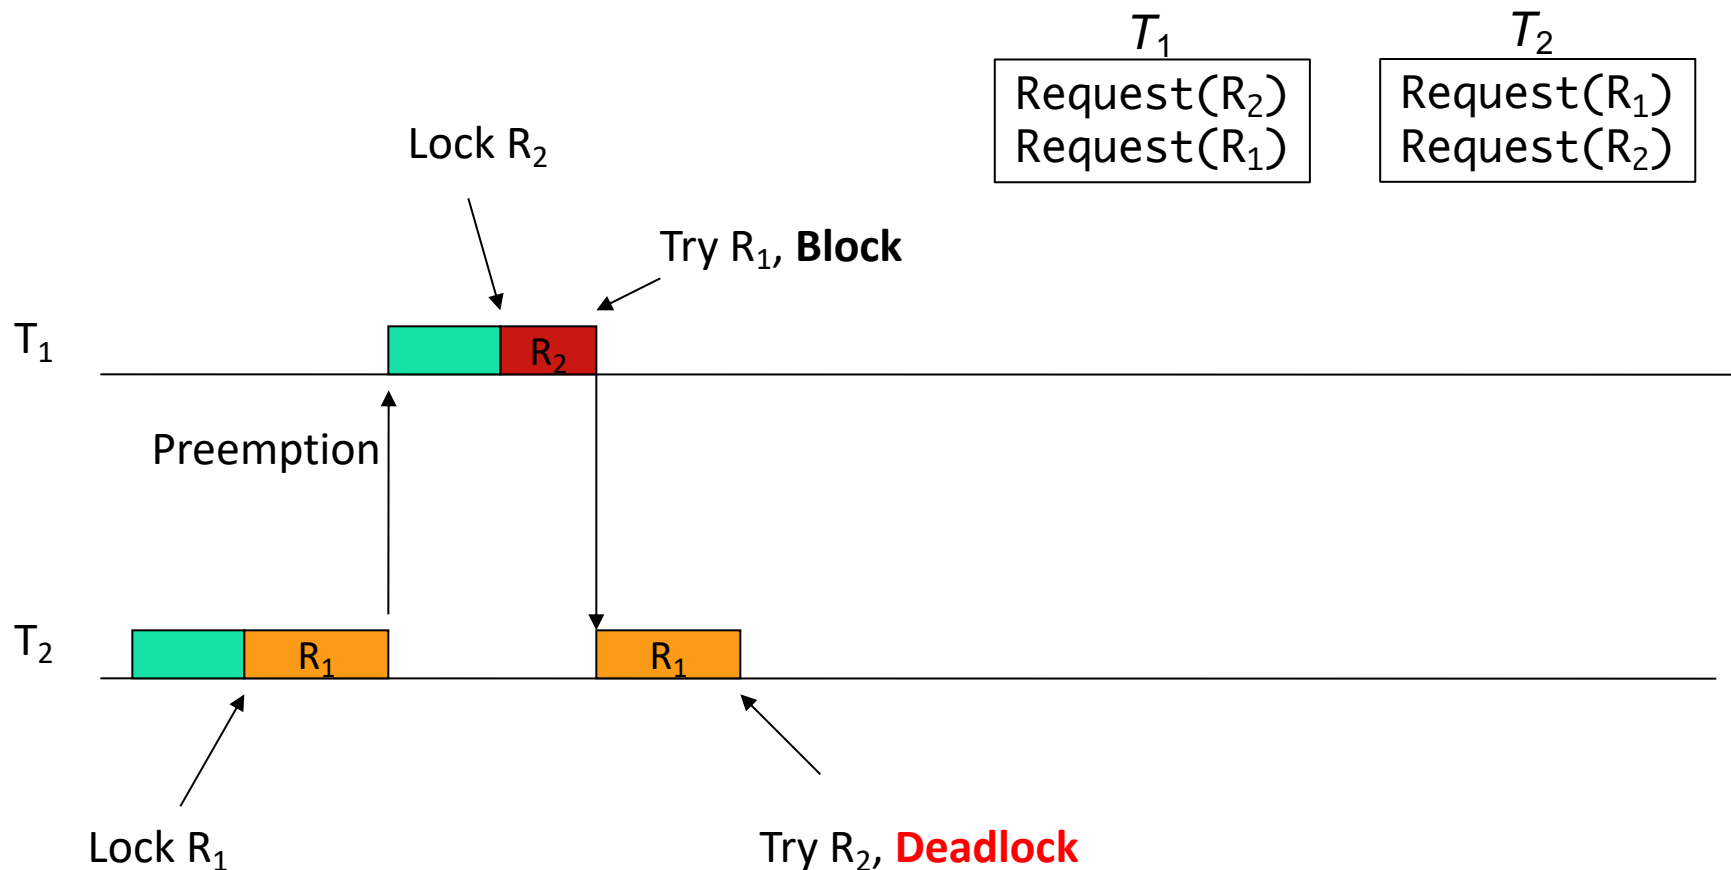
Maximum wait for resources is $B_4 = 0$

$$\left(\frac{3}{20} + \frac{6}{30} + \frac{10}{50} + \frac{8}{80} \right) = 0.65$$

T_4 is schedulable

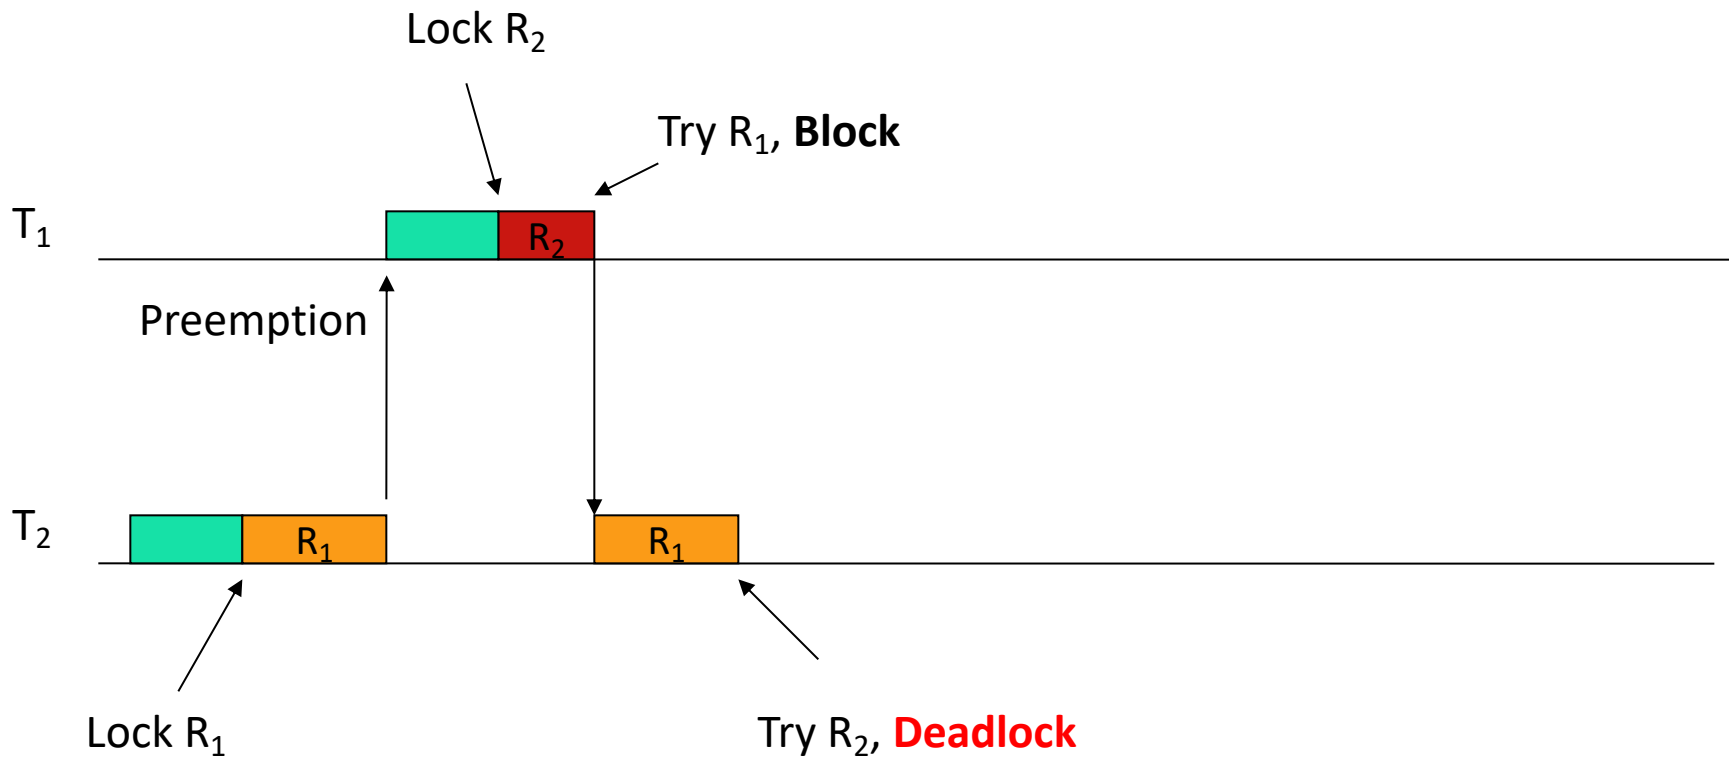
Does priority inheritance solve all problems?

- Actually, not all problems
- We can still have a deadlock if resources are locked in opposing orders



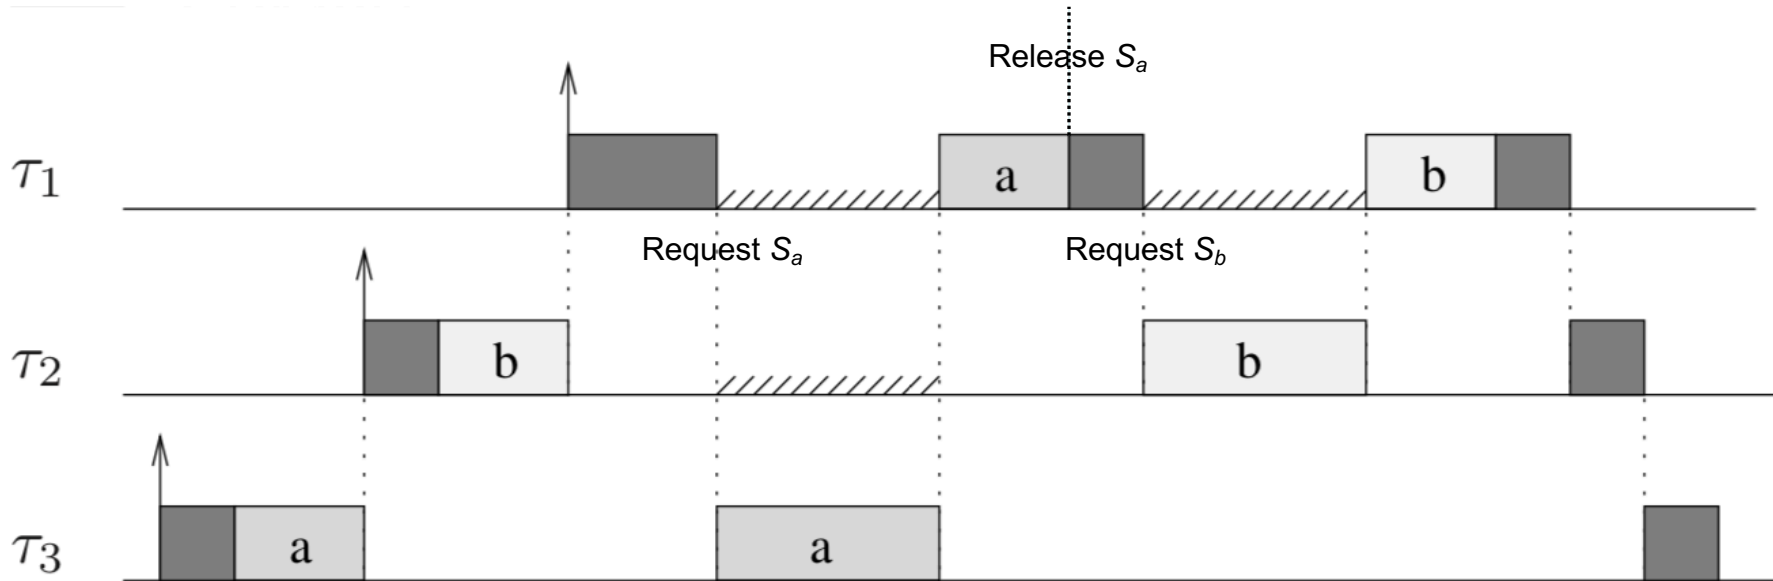
Deadlocks

- Can attribute it to sloppy programming
- But can we solve the problem in a different way
- Avoid deadlocks by designing a suitable protocol



Another problem with PIP: *Chained blocking*

- When τ_1 attempts to use its resources, it is blocked for the duration of **2** critical sections:
 - once to wait for τ_3 to release S_a
 - and then to wait for τ_2 to release S_b
- In the worst case, if τ_1 accesses n distinct semaphores that have been locked by n lower-priority tasks, τ_1 will be blocked for the duration of n critical sections.



Avoiding Multiple Blocking

- When a task enters a critical section, make sure that there are sufficient resources to satisfy its maximum resource requirements
- **Consequence:** When a task enters a critical section, it cannot be blocked on resources
- Do not allow a task to enter a critical section if there are locked resources that can block it
- **Meaning:** do not allow task T_i to enter a critical section at time t if there is a locked resource R_k with $C(R_k) \geq \pi_i$
- Iff allow task T_i to enter a critical section at time t if $\pi_i > C(R_k)$ for **every** locked resource R_k
- Iff allow task T_i to enter a critical section if
$$\pi_i > \max\{C(R_k) : R_k \text{ locked at time } t\} \equiv C(t)$$

Priority ceiling protocol

- **Definition:** the **priority ceiling** of a semaphore is the highest priority among all tasks that can lock the semaphore
- A task that requests lock R_k is denied if its priority is not strictly higher than the highest priority ceiling of all **currently** locked semaphores (let us say this belongs to semaphore R_h ; *Can there be more than one?*)
 - The task is said to be blocked by the task holding semaphore R_h
- A task inherits the priority of the top higher-priority task it is blocking

Priority ceiling protocol

- **Recall: Priority Ceiling** of resource R_k : $C(R_k) = \max_{i \in [n]} \{\pi_i : T_i \text{ uses } R_k\}$
- Suppose task T_i requests a resource R_k at time t
- Let $R_h = \operatorname{argmax}_j \{C(R_j) : \text{resource } R_j \text{ is locked at time } t\}$
 - Can there be more than one such R_h ?
- Define **System Ceiling** as the highest ceiling of currently locked semaphores \rightarrow
 $C(t) = C(R_h) = \max\{C(R_j) : R_j \text{ locked at time } t\}$
 - System ceiling updated whenever a resource is acquired/released
- If $\pi_i \leq C(t)$, then T_i is denied access to the resource
 - **Exception:** If $\pi_i \leq C(t)$ but T_i is the task locking R_h then grant T_i access to R_k (o/w T_i will block itself!)
 - T_i is said to be blocked by the task holding semaphore R_h
 - T_i then transfers its priority to task holding R_h

Priority ceiling protocol

- To avoid multiple blocking, this rule does not allow a task to enter a critical section if there are locked semaphores that could block it.
- This means that *once a task enters its first critical section, it can **never** be blocked by lower-priority tasks until its completion*

Similarity to PIP

Priority Inheritance rule

Fundamental difference from PIP

PIP is *greedy*, PCP is not!

In what sense?

A task can be blocked on a *free* resource in PCP

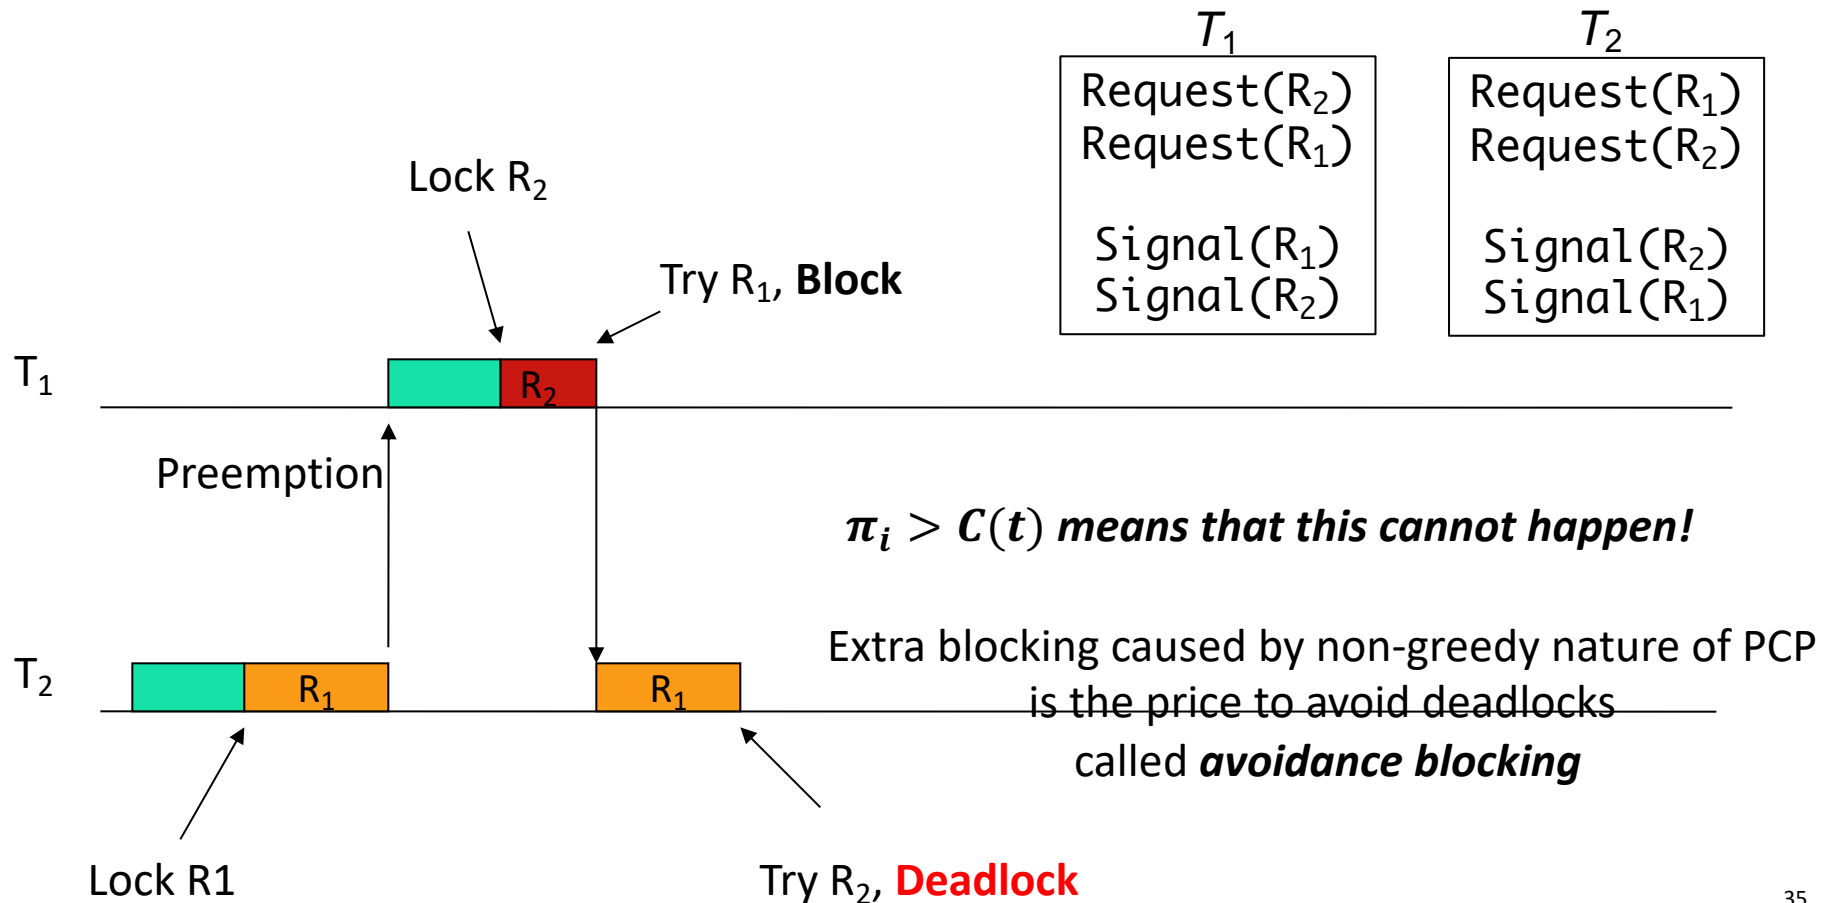
Impossible in PIP



Extra blocking caused by non-greediness of PCP is the price to avoid deadlocks & chained blocking called ***avoidance blocking*** or ***ceiling blocking***

Deadlocks?

- A deadlock can occur if two tasks locked semaphores in opposite order. Can it occur with the priority ceiling protocol?



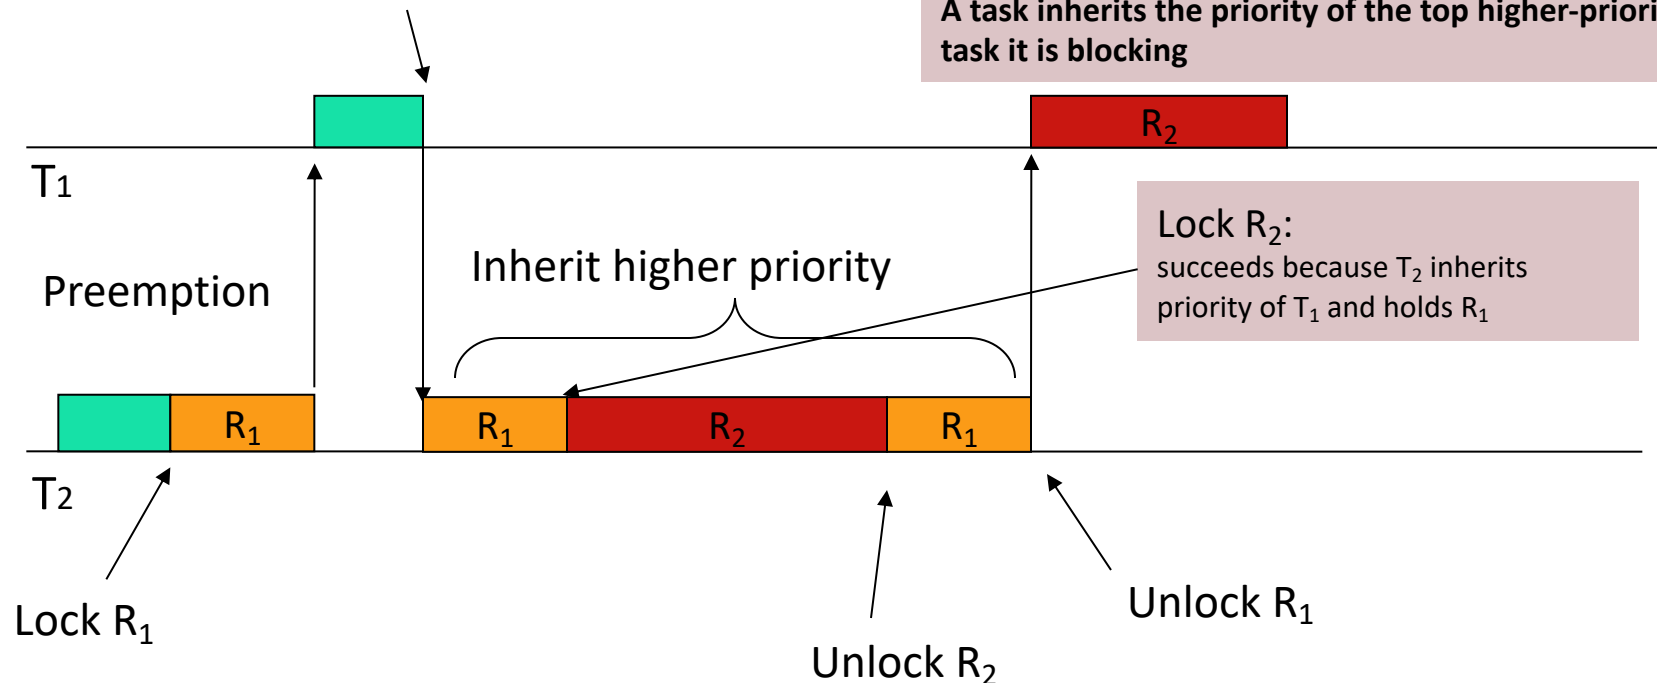
Priority ceilings

- T_1 and T_2 use R_1 and R_2 : the priority ceiling of a resource is the priority of the highest priority task that uses it, therefore the priority ceilings of R_1 and R_2 are the same: the priority of T_1

Lock R_2 : Denied because its priority is not higher than ceiling of R_1

A task that requests lock R_k is denied if its priority is not higher than the highest priority ceiling of all currently locked semaphores

A task inherits the priority of the top higher-priority task it is blocking



PCP blocking time computation

- A task can be blocked by the duration of at most one critical section of at most one lower priority task
- Much simpler to compute than PIP
- Should consider the three types of blocking and take the max of them
- Resource graph to our rescue!

Schedulability test for priority ceiling protocol

- The test is the same as with the priority inheritance protocol
 - Worst-case blocking time may change when compared to PIP

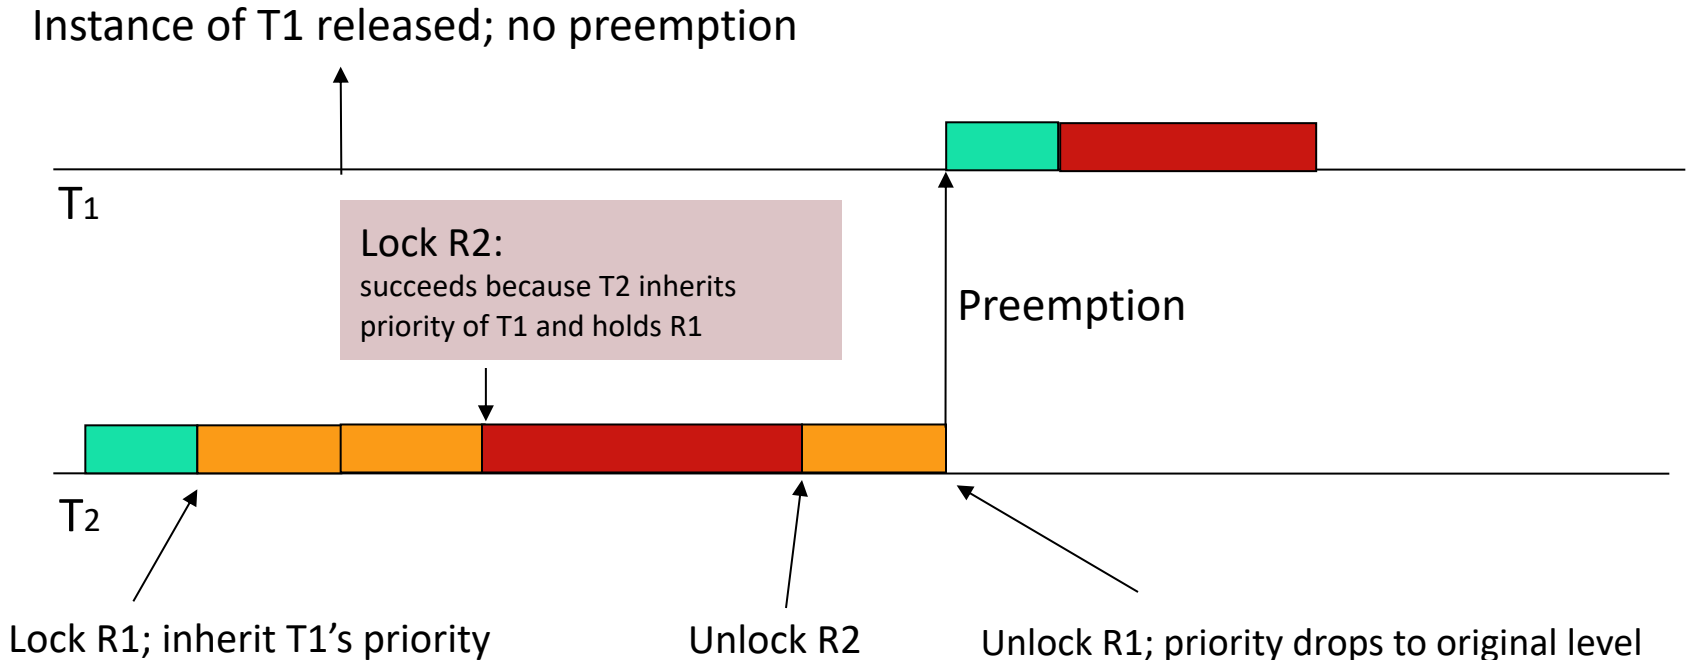
$$\frac{B_k}{P_k} + \sum_{i=1}^k \frac{e_i}{P_i} \leq k(2^{1/k} - 1)$$

For task T_k

Recall: Highest Locking Protocol (HLP)

= PCP with Immediate inheritance

- Priority ceiling protocol with slight difference: when a semaphore is locked, the locking task raises its priority to the ceiling of the semaphore (**immediate inheritance**). When the semaphore is unlocked the task's priority is restored.




Stack-based resource policy

- Let us attempt to support *dynamic-priority* systems
- Does PCP extend directly?
- Task priorities in dynamic-task (equivalently fixed-job) priority systems might change at every invocation
 - Resource ceilings are no longer static: Must be updated potentially at every invocation. High runtime overhead!
- **Observation:** That a job J_h has a higher priority than another job J_l and they both require some resource does not imply that J_l can directly block J_h
 - This blocking can occur **only when it is possible for J_h to preempt J_l**
- When determining whether a free resource can be granted to a job, it is **not necessary** to be concerned with the resource requirements of all higher-priority jobs; *only those that can preempt the job*

Stack-based resource policy

- Since for resource contention purposes we only care about the jobs that a job can possibly preempt, let us *identify the event that causes a job to be preempted in any task-dynamic priority scheduling scheme*
- In a dynamic-task policy, when can a job preempt another job?



Stack-based resource policy


- A quantity that encodes a job's ability to preempt other jobs
- (*) Formally, we want to associate job J_k with quantity ψ_k such that if $\psi_k \leq \psi_i$, then it is **not** possible for J_k to preempt J_i
- J_k cannot preempt $J_i \Leftrightarrow$ either $r_k \leq r_i$ or $\pi_k \leq \pi_i$
- Then (*) translates to:

(**) if $r_k > r_i$ and $\pi_k > \pi_i$, then $\psi_k > \psi_i$ (it's possible for J_k to preempt J_i)
--

- A ψ_k satisfying (**) is called the **preemption level** of job J_k
- **Q:** How does ψ_k look like for EDF?

Stack-based resource policy with EDF

- Priority is inversely proportional to the absolute deadline
- Preemption level is inversely proportional to the relative deadline
- Observe that:
 - If A arrives after B and $\text{Priority}(A) > \text{Priority}(B)$ then
 $\text{PreemptionLevel}(A) > \text{PreemptionLevel}(B)$



Stack-based resource policy

- The preemption level ψ_i of J_i is any quantity satisfying the statement:
if $r_k > r_i$ and $\pi_k > \pi_i$, then $\psi_k > \psi_i$
- **Q:** How does ψ_i look like for EDF?
- EDF:
 - $\pi_k > \pi_i$ iff $r_k + D_k < r_i + D_i$
 - So $r_i < r_k$ implies $r_i + D_k < r_i + D_i \Rightarrow D_k < D_i$
 - $\psi_k > \psi_i \Leftrightarrow D_k < D_i$
 - For EDF, this quantity is for the entire **task**, not only a job!
- *The possibility that a task preempts other tasks remains constant throughout all its invocations*
 - *Task's preemption level is static; can be computed offline once and for all*
- *EDF is one such **fixed preemption-level** system*
 - In such systems, the potentials of resource contentions do not change with time, just as in fixed-priority systems, and hence can be analyzed *statically*

Stack-based resource policy

- In fixed-**preemption** level systems, the set of critical sections that can block T_i are
$$\{z_{j,k}: \psi_i > \psi_j, C(R_k) \geq \psi_i\}$$
- **Stack-based resource policy [SRP]**
 - **Preemption level**: Any fixed value that satisfies the statement “If A arrives after B and $\text{Priority}(A) > \text{Priority}(B)$ then $\text{PreemptionLevel}(A) > \text{PreemptionLevel}(B)$ ”
 - **Resource ceiling** for resource R : Highest **preemption level** of all tasks that may access the resource R
 - **System ceiling**: Highest resource ceiling among all **currently** locked resources
 - A task can preempt another task if
 - it has the highest priority and
 - its preemption level is higher than the system ceiling

Stack-based resource policy

- Resource ceiling $C(R_k) = \max\{\psi_i: T_i \text{ uses } R_k\}$
- System ceiling $C(t) = \max\{C(R_k): \text{resource } R_k \text{ is being used at time } t\}$

SRP Preemption Test

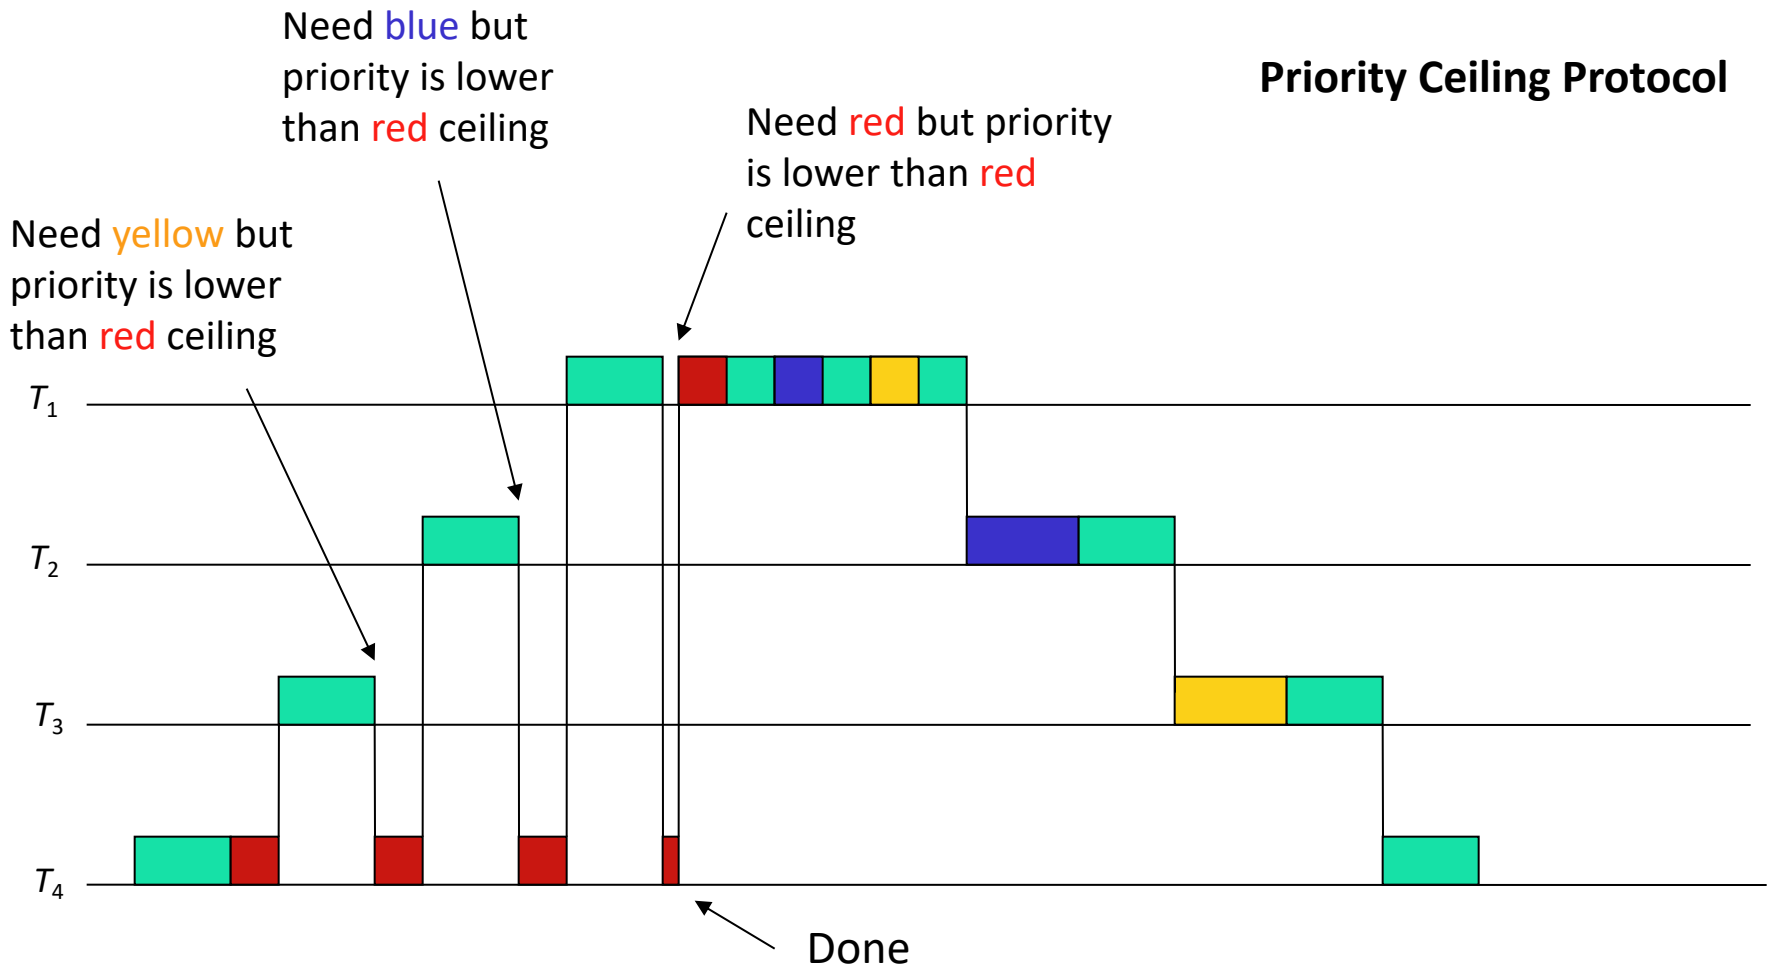
A task can preempt another task if

- it has the highest priority and
- its preemption level is higher than the system ceiling

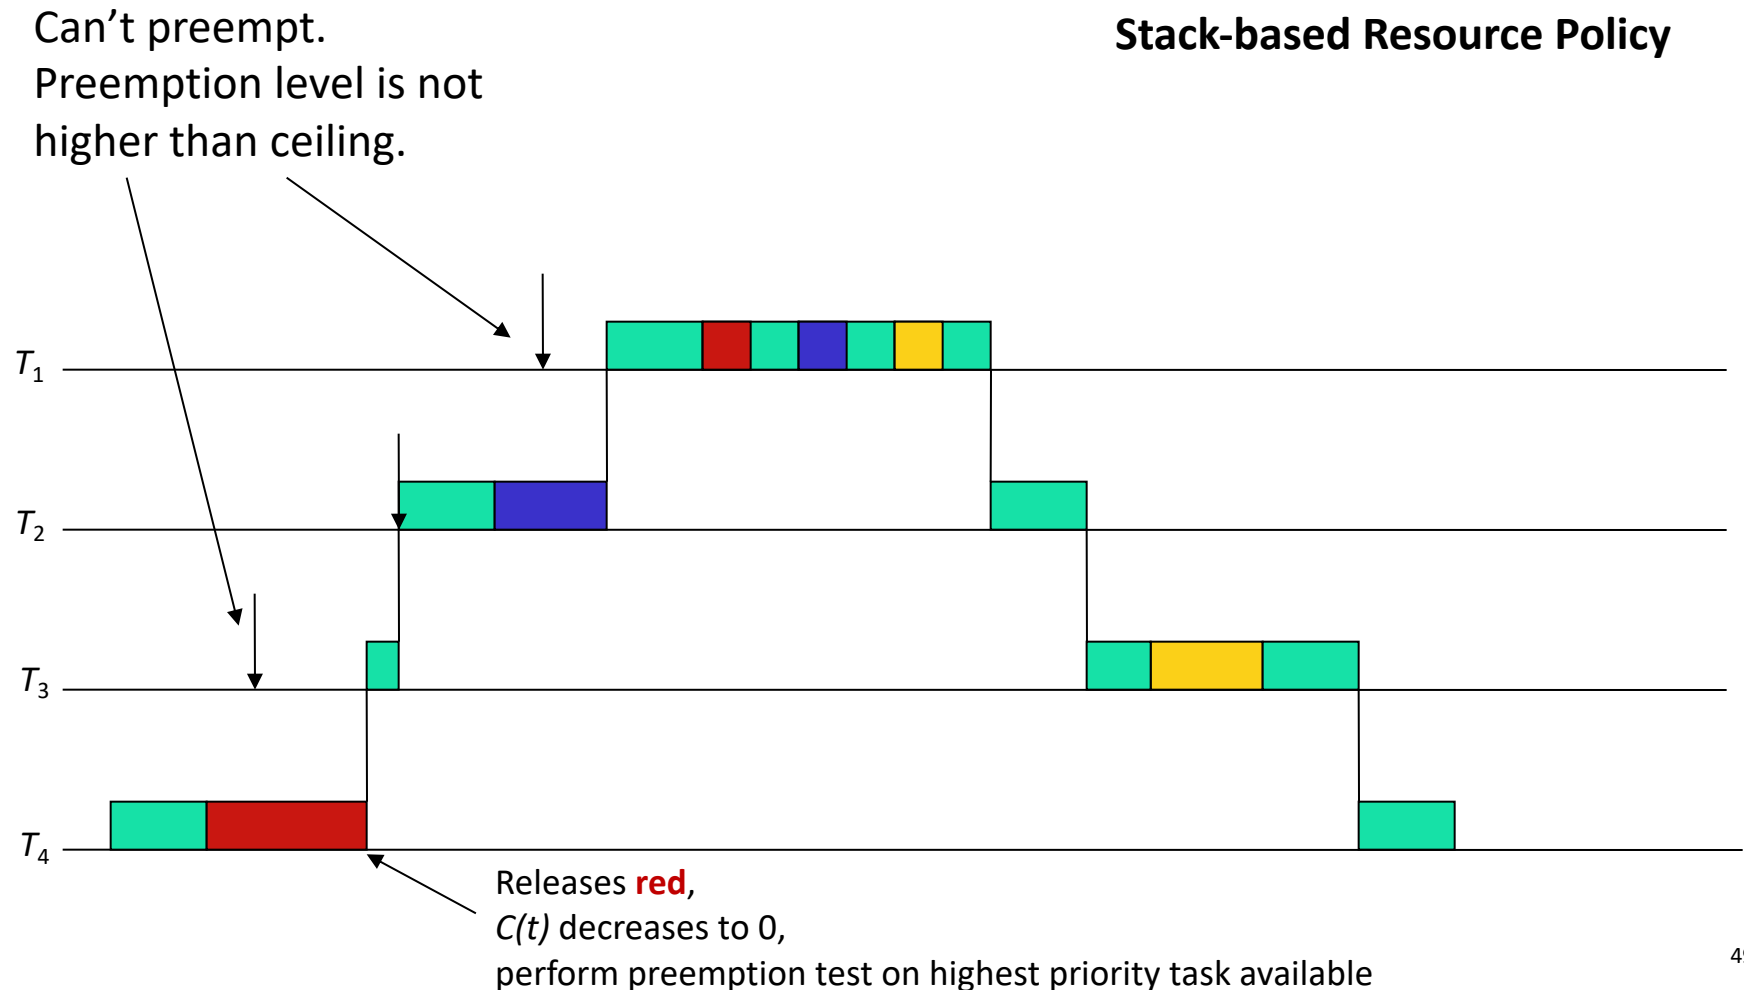
If T_i is the highest priority task at time t and $\psi_i > C(t)$ then allow T_i to preempt, otherwise block it

- Perform preemption test when a task arrives (on the arriving task) and on highest priority task when $C(t)$ decreases (a resource is released)

Priority ceiling vs. stack-based resource policy



Priority ceiling vs. stack-based resource policy

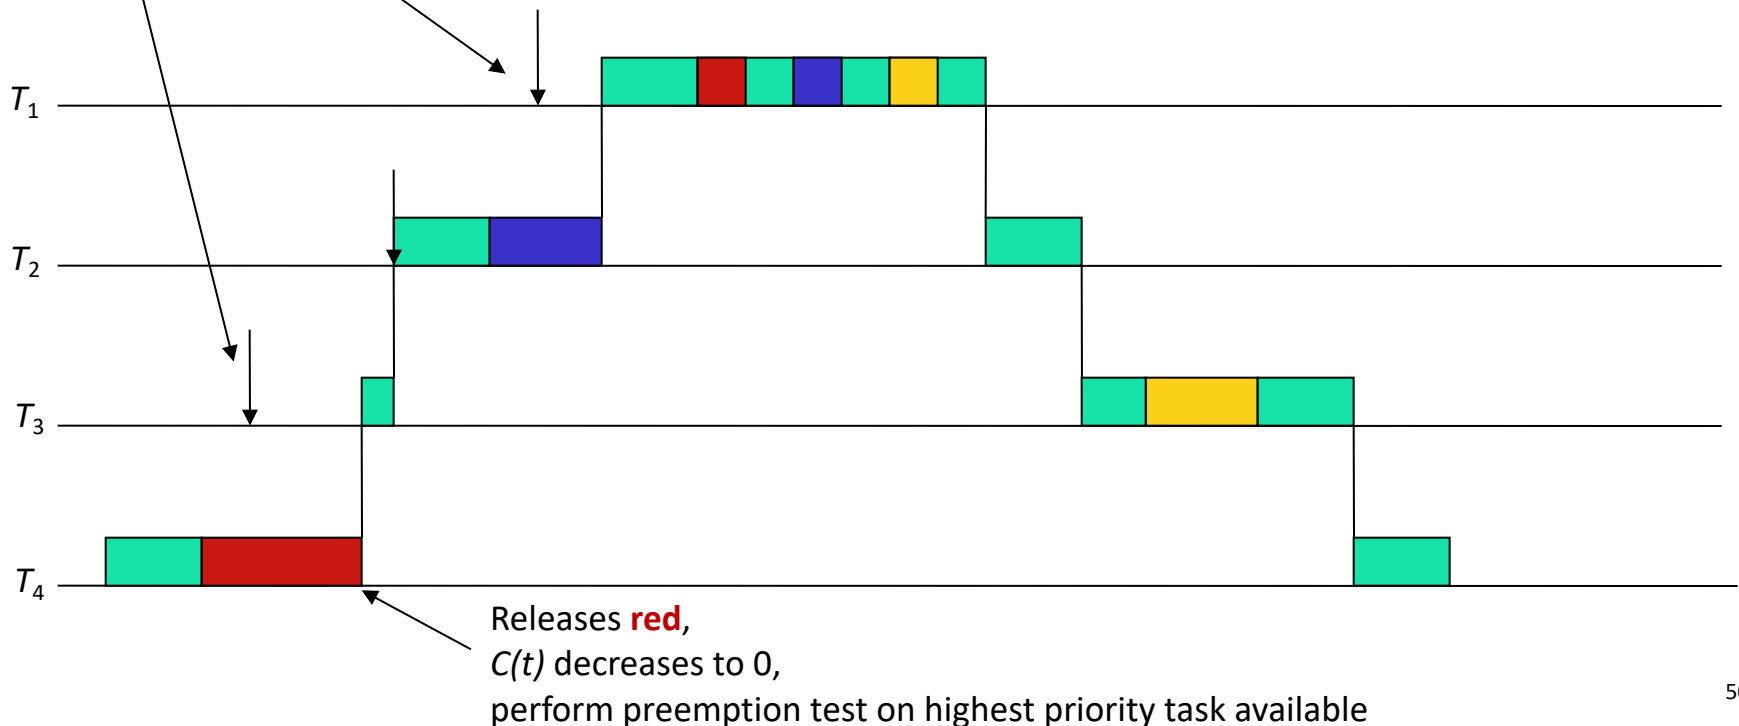


Priority ceiling vs. stack-based resource policy

Stack-based Resource Policy

Can't preempt.
Preemption level is not
higher than ceiling.

Notice that SRP is similar to immediate inheritance in PCP.
However, with no static priority levels, it needs a preemption level.

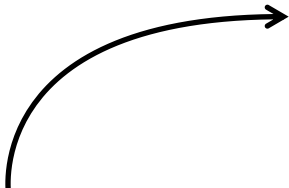


Stack-based resource policy

- **Q:** What does it mean when a task passes the preemption test?
- **A:** the resources that are currently available are sufficient to satisfy the maximum requirement of task T_h and the maximum requirement of every task that could preempt T_h .
 - This means that once T_h starts executing, it will never be blocked for resource contention.
- **Remarks**
 - SRP avoids deadlocks. Why?
 - Resources are only allocated when a task *requests* them, not when it preempts
 - A higher-priority job may preempt and use the resources between these critical sections
 - A task can be blocked by the preemption test even though it does not require any resource. This is needed to avoid unbounded priority inversion.
 - The preemption test has the effect of imposing priority inheritance
 - an executing task that holds a resource modifies the system ceiling and resists preemption as though it inherits the priority of any tasks that might need that resource

Analysis with EDF and SRP

- As simple as other protocols


$$\frac{B_k}{P_k} + \sum_{i=1}^k \frac{e_k}{P_k} \leq 1$$

For task T_k

Maximum blocking due to task with lower preemption level; in the case of EDF: with period P_j such that $P_k < P_j$.

Tasks are sorted such that the task with shortest period is T_1 and so on.

What is the “stack” in Stack-based Resource Sharing Protocol?

- **Two things:**

1. Can be implemented using a stack. How?
2. Allows tasks to share the run-time stack

In-class activity

Determine if the following task set can be scheduled using the *rate monotonic scheduling policy* with the *priority ceiling protocol* to control resource access.

Task	e_i	P_i	Resources used
T_1	4	10	R_1, R_2
T_2	5	20	R_2, R_3
T_3	10	35	R_3
T_4	2	40	R_1

The duration for which each resource is used by the tasks is specified in the following table. You may assume that a task locks only one resource at a time.

Resource	Duration
R_1	2
R_2	1
R_3	2

Highlights

- Schedulability analysis needs to account for blocking due to low priority tasks
- **Priority inheritance protocol (PIP)** may not prevent deadlocks
- **Deadlocks** can be prevented with the **priority ceiling protocol (PCP)**
- To deal with dynamic priority policies (such as EDF), we need a different policy: the **stack-based resource policy (SRP)**
- SRP (and the immediate inheritance version of the PCP) have efficient implementations
 - Reduce the number of context switches
 - SRP also prevents deadlocks (note the similarities between PCP and SRP)