

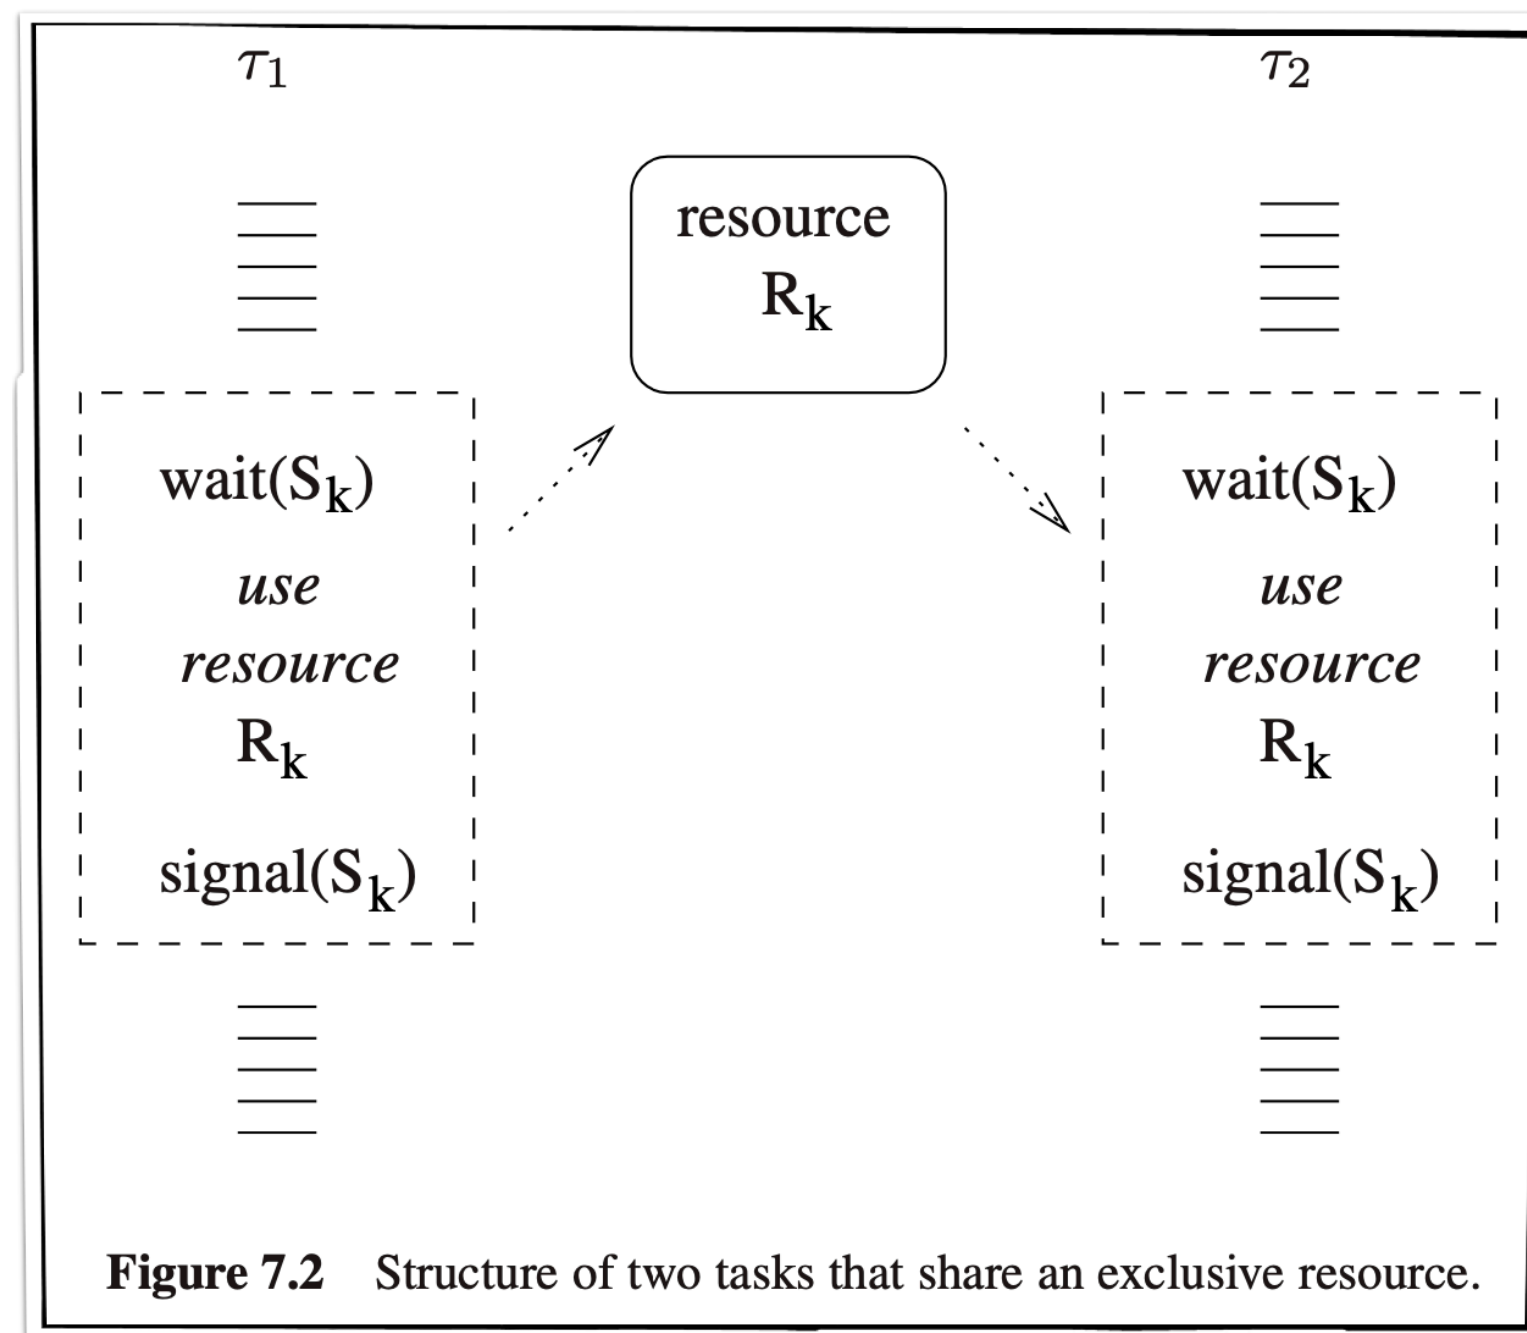
Resource Sharing

CPEN 432 Real-Time System Design

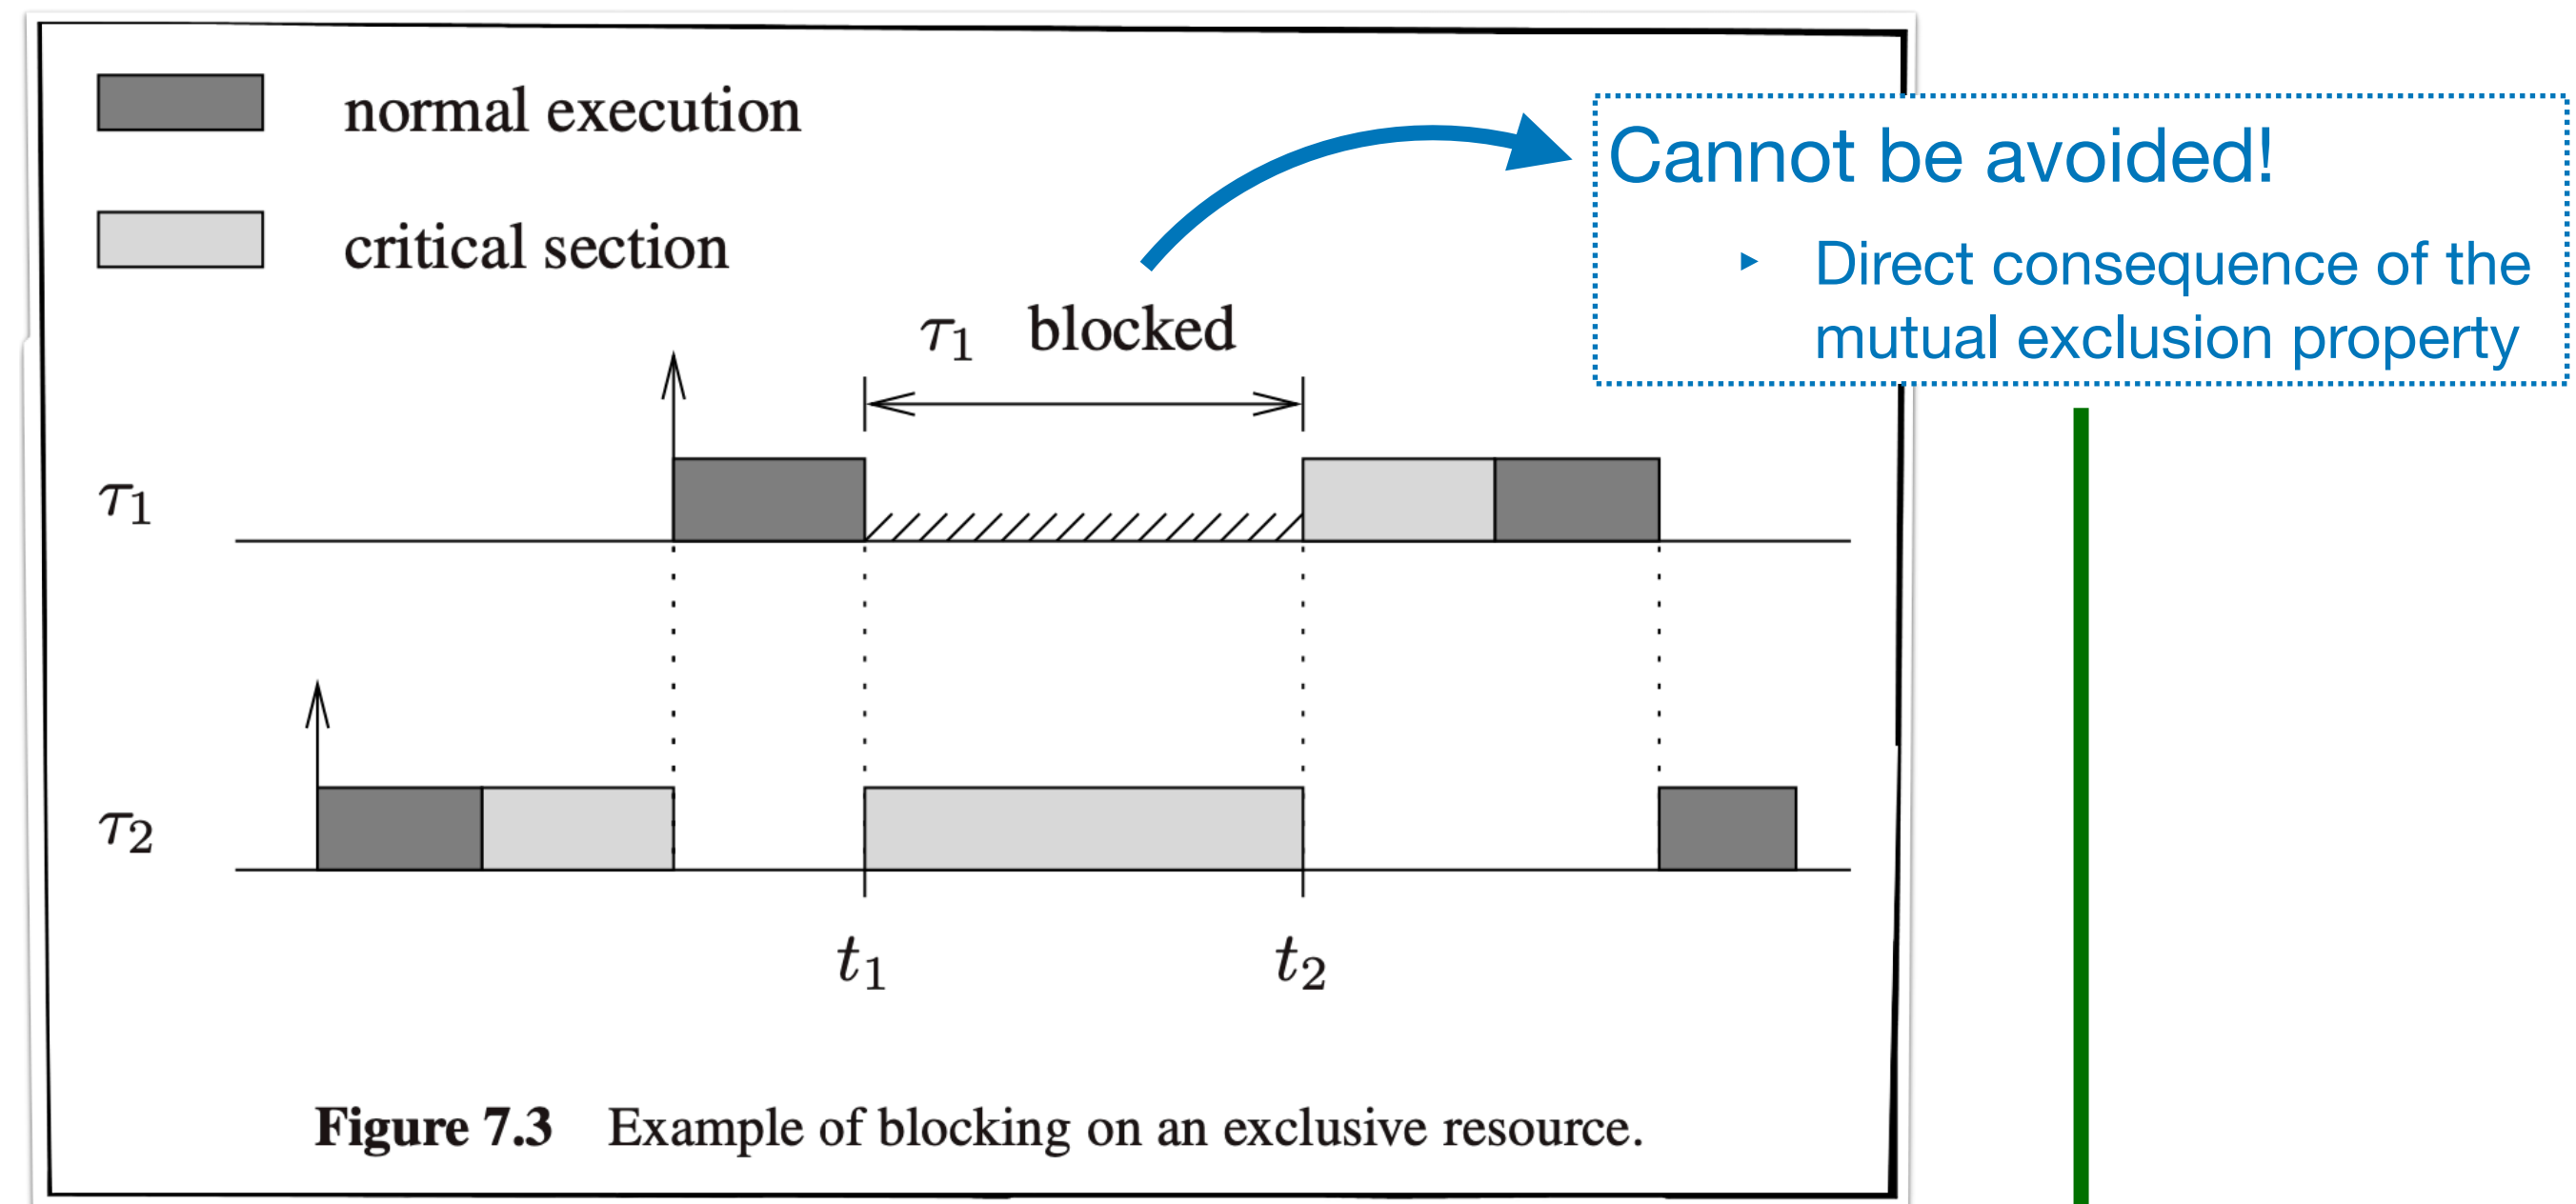
Arpan Gujarati
University of British Columbia

Unavoidable Blocking on an Exclusive Resource

Exclusive resource R_k accessed by waiting and signalling a binary semaphore S_k



Typically, the exclusive resource R_k should not be shared while a critical section using R_k is in progress



Also, easy to bound using the critical section duration

- Like the worst-case completion time C_2 , we could also characterize the worst-case critical section duration of τ_2 while it uses R_k

Unbounded Blocking Due to Priority Inversion

■ normal execution
■ critical section

The duration of priority inversion is unbounded

- ▶ Any intermediate priority task can preempt τ_3 and indirectly block τ_1
- ▶ If we add the completion time $C_{\text{intermediate}}$ of each intermediate-priority task $\tau_{\text{intermediate}}$ as a blocking factor in τ_1 's timing analysis, the resulting analysis will be sound but extremely pessimistic




Figure 7.4 An example of priority inversion.

How to Prevent Unbounded Priority Inversions?

- Key idea ...

Terminology

- Task set $\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$ consists of n periodic tasks
- Each task is characterized by a period T_i and worst-case completion time C_i
- The tasks cooperate through m shared resources R_1, R_2, \dots, R_m
- Each resource R_k is guarded by a distinct **binary semaphore** S_k
 - All critical sections using R_k start and end with operations $wait(S_k)$ and $signal(S_k)$
- Each task is assigned a fixed **base priority** P_i (e.g., using RM)
 - Assumption: priorities are unique and $P_1 > P_2 > \dots > P_n$
- Each task also has an **effective priority** p_i ($\geq P_i$)
 - It is initially set to P_i and can be **dynamically updated**
- B_i denotes the maximum blocking time task τ_i can experience
 - B_i goes into the fixed-priority response-time analysis (recall from previous lectures)
- $z_{i,k}$ denotes any arbitrary critical section of τ_i guarded by semaphore S_k
 - $Z_{i,k}$ denotes the longest among all these critical sections
 - $\delta_{i,k}$ denotes the length of this longest critical section $Z_{i,k}$

Non-Preemptive Protocol (NPP)

Observation

Blocking caused by the **preemption** of a running, **resource-holding** job

- ▶ E.g., τ_3 preempted by τ_2 at time t_4 while holding the shared resource

Key idea: **Disable preemption** before acquiring a shared resource; reenale upon exit of critical section

When a task τ_i acquires a resource R_k , its dynamic priority is raised to the level of the highest priority, i.e., $p_i(R_k) = \max_{\forall h} \{P_h\}$




Figure 7.4 An example of priority inversion.

Example

Priority inversion **bounded by critical section length**

- How can we **formally** define τ_i 's blocking time bound B_i ?

■ normal execution

■ critical section




Figure 7.5 Example of NPP preventing priority inversion.

NPP Benefits & Limitations

- Most **simple** way to prevent unbounded priority inversions
- Can be realized by **disabling/reenabling interrupts**
 - Raising task priorities is a useful abstraction but needn't be implemented in this case
- Limitations
 - Turning off interrupts risks **large interrupt latency**
 - All tasks effected
 - Even **independent tasks blocked** due to priority inversion

What if high-frequency tasks cannot tolerate blocking even due to a single, **long** non-preemptive section?

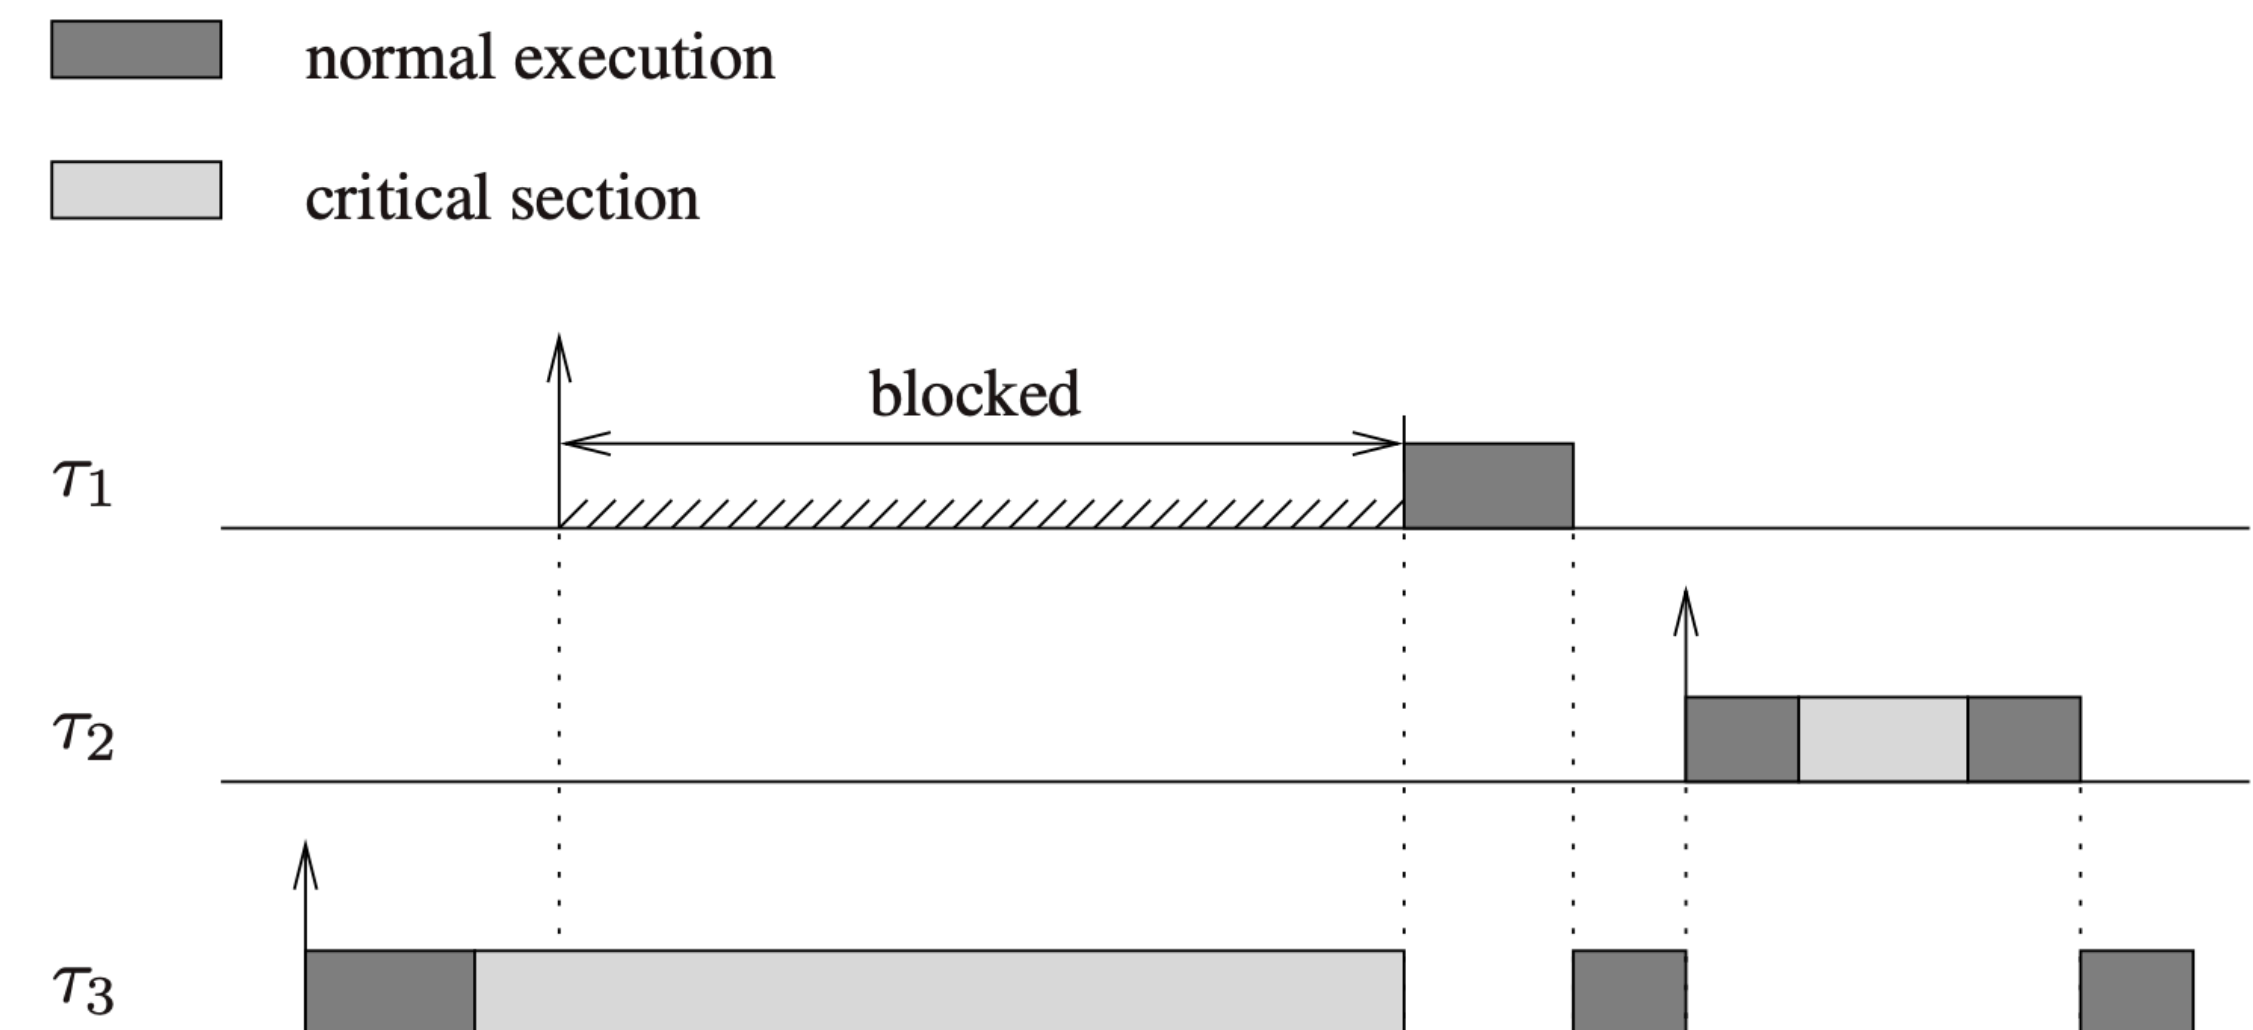





Figure 7.6 Example in which NPP causes unnecessary blocking on τ_1 .






What's next?

The Priority Inheritance Protocol (PIP)


Protocol Definition

- Unlike NPP, resource holding jobs remain **fully preemptive**
- Tasks are scheduled based on their effective priorities
 - For scheduling purposes, τ_i 's priority is considered to be p_i and not P_i
- Suppose task τ_i tries to enter a critical section by acquiring resource R_k
 - Case 1: R_k is already held by a lower-priority task $\tau_j \implies \tau_i$ is **blocked** by τ_j
 - Case 2: R_k is already held by a higher-priority task $\tau_j \implies \tau_i$ is **interfered** by τ_k
 - Case 3: R_k is not held by any task $\implies \tau_i$ **enters** the critical section
- For Case 1, τ_j **inherits** τ_i 's effective priority
 - τ_j 's dynamic priority is updated as $p_j = p_i$
- In general, τ_j inherits the **highest priority of among all tasks that it blocks**
 - At any point of time, $p_j(R_k) = \max \{P_j, \max_{\forall h} \{p_h \mid \tau_h \text{ is blocked on } R_k\}\}$


Example 1




Example 2



Example 3: Nested Blocking




Example 4: Transitive Blocking



PIP Benefits & Limitations

- **No latency penalty** for high-priority independent tasks
- Widely used in practice: POSIX's PTHREAD_PRIO_INHERIT
- Limitations
 - Chained blocking
 - Deadlock



The Priority Ceiling Protocol (PCP)

PCP vs PIP

- The PIP is a **reactive** locking protocol
 - It only kicks in when resource contention already exists
- **Key PCP insight**
 - Better to **prevent** problematic scenarios rather **than resolve** them
- The PCP is an **anticipatory** locking protocol
 - Exploits the knowledge of resource needs at **design time** to avoids excessive blocking at runtime

Key Concepts

- **Priority ceilings**

- Each semaphore S_k is **statically** assigned a priority ceiling $C_{static}(S_k)$
 - $C_{static}(S_k)$ = priority of the highest-priority task that **ever** accesses S_k

- **Current system ceiling**

- At any time t , a global system ceiling $C_{global}(t)$ is dynamically computed
 - $C_{global}(t)$ = highest priority ceiling among all semaphores locked at time t OR
(if no semaphores are locked) sentinel value P_0 that is **smaller** than all task priorities

- **Protocol**

- Task τ_i can acquire semaphore S_k at time t only if
 - Its effective priority $p_i > C_{global}(t)$ OR $p_i = C_{global}(t)$ and τ_i “owns” the ceiling resource
 - OTHERWISE, it transmits its priority to the task τ_j that holds semaphore S_k

Example