




# Resource Sharing

CPEN 432 Real-Time System Design

Arpan Gujarati  
University of British Columbia

# Terminology

- Task set  $\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$  consists of  $n$  periodic tasks
- Each task is characterized by a period  $T_i$  and worst-case completion time  $C_i$
- The tasks cooperate through  $m$  shared resources  $R_1, R_2, \dots, R_m$
- Each resource  $R_k$  is guarded by a distinct **binary semaphore**  $S_k$ 
  - All critical sections using  $R_k$  start and end with operations  $wait(S_k)$  and  $signal(S_k)$
- Each task is assigned a fixed **base priority**  $P_i$  (e.g., using RM)
  - Assumption: priorities are unique and  $P_1 > P_2 > \dots > P_n$
- Each task also has an **effective priority**  $p_i$  ( $\geq P_i$ )
  - It is initially set to  $P_i$  and can be **dynamically updated**
- $B_i$  denotes the maximum blocking time task  $\tau_i$  can experience
  - $B_i$  goes into the fixed-priority response-time analysis (recall from previous lectures)
- $z_{i,k}$  denotes any arbitrary critical section of  $\tau_i$  guarded by semaphore  $S_k$ 
  - $Z_{i,k}$  denotes the longest among all these critical sections
  - $\delta_{i,k}$  denotes the length of this longest critical section  $Z_{i,k}$






What's next?

# Protocol Definition


- Unlike NPP, resource holding jobs remain **fully preemptive**
- Tasks are scheduled based on their effective priorities
  - For scheduling purposes,  $\tau_i$ 's priority is considered to be  $p_i$  and not  $P_i$
- Suppose task  $\tau_i$  tries to enter a critical section by acquiring resource  $R_k$ 
  - Case 1:  $R_k$  is already held by a lower-priority task  $\tau_j \implies \tau_i$  is **blocked** by  $\tau_j$
  - Case 2:  $R_k$  is already held by a higher-priority task  $\tau_j \implies \tau_i$  is **interfered** by  $\tau_k$
  - Case 3:  $R_k$  is not held by any task  $\implies \tau_i$  **enters** the critical section
- For Case 1,  $\tau_j$  **inherits**  $\tau_i$ 's effective priority
  - $\tau_j$ 's dynamic priority is updated as  $p_j = p_i$
- In general,  $\tau_j$  inherits the **highest priority of among all tasks that it blocks**
  - At any point of time,  $p_j(R_k) = \max \{P_j, \max_{\forall h} \{p_h \mid \tau_h \text{ is blocked on } R_k\}\}$

# Example




# Properties of PIP [1/5]

*A semaphore  $S_k$  can cause push-through blocking to task  $T_i$ , only if  $S_k$  is accessed both by a task with priority lower than  $P_i$  and by a task with priority higher than  $P_i$ .*



# Properties of PIP [2/5]

*Transitive priority inheritance can occur only in the presence of nested critical sections.*



# Properties of PIP [3/5]

*If there are  $l_i$  lower-priority tasks that can block a task  $\tau_i$ , then  $\tau_i$  can be blocked for at most the duration of  $l_i$  critical sections, one for each of the  $l_i$  lower-priority tasks, regardless of the number of semaphores used by  $\tau_i$ .*



# Properties of PIP [4/5]

*If there are  $s_i$  distinct semaphores that can block a task  $\tau_i$ , then  $\tau_i$  can be blocked for at most the duration of  $s_i$  critical sections, one for each of the  $s_i$  semaphores, regardless of the number of critical sections used by  $\tau_i$ .*

# Properties of PIP [5/5]

*Under the Priority Inheritance Protocol, a task  $\tau_i$  can be blocked for at most the duration of  $\alpha_i = \min(l_i, s_i)$  critical sections, where  $l_i$  is the number of lower-priority tasks that can block  $\tau_i$  and  $s_i$  is the number of distinct semaphores that can block  $\tau_i$ .*

# Computing Blocking Time $B_i$ [1/2]

- A precise evaluation of the blocking factor  $B_i$  is quite complex because each critical section of the lower-priority tasks may interfere with  $\tau_i$  via **direct blocking, push-through blocking, or transitive inheritance**
- Simplified algorithm
  - Assumes **no nested critical sections**, hence **no transitive inheritance**

# Computing Blocking Time $B_i$ [2/2]

- Semaphores that can **directly** block  $\tau_i$  and that are shared by the lower-priority task  $\tau_j$  are  $\sigma_{i,j}^{dir} = \sigma_i \cap \sigma_j$
- Semaphores that can block  $\tau_i$  by **push-through** and that are shared by the lower-priority task  $\tau_j$  are  $\sigma_{i,j}^{pt} = \bigcup_{h:P_h > P_i} \sigma_h \cap \sigma_j$
- Semaphores that can block  $\tau_i$  either **directly or by push-through** and that are shared by the lower-priority task  $\tau_j$ 
  - $\sigma_{i,j} = \sigma_{i,j}^{dir} \cup \sigma_{i,j}^{pt} = \bigcup_{h:P_h \geq P_i} \sigma_h \cap \sigma_j$
- Longest critical sections used by lower-priority task  $\tau_j$  that can block  $\tau_i$  either directly or by push-through is
  - $\gamma_{i,j} = \{Z_{j,k} \mid R_k \in \sigma_{i,j}\}$
- LII critical sections that can block  $\tau_i$  either directly or by push-through is  $\gamma_i = \bigcup_{j:P_j < P_i} \gamma_{i,j}$
- $B_i$  is given by the largest sum of the lengths of the  $\alpha_i$  critical sections in  $\gamma_i$ 
  - The sum should contain only terms  $\delta_{i,k}$  referring to different tasks and different semaphore

# The Priority Ceiling Protocol (PCP)

# PCP vs PIP

- The PIP is a **reactive** locking protocol
  - It only kicks in when resource contention already exists
- **Key PCP insight**
  - Better to **prevent** problematic scenarios rather **than resolve** them
- The PCP is an **anticipatory** locking protocol
  - Exploits the knowledge of resource needs at **design time** to avoids excessive blocking at runtime

# PCP Key Concepts

- **Priority ceilings**

- Each semaphore  $S_k$  is **statically** assigned a priority ceiling  $C_{static}(S_k)$ 
  - $C_{static}(S_k)$  = priority of the highest-priority task that **ever** accesses  $S_k$





- **Current system ceiling**

- At any time  $t$ , a global system ceiling  $C_{global}(t)$  is dynamically computed
  - $C_{global}(t)$  = highest priority ceiling among all semaphores locked at time  $t$  OR  
(if no semaphores are locked) sentinel value  $P_0$  that is **smaller** than all task priorities

- **Protocol**


- Task  $\tau_i$  can acquire semaphore  $S_k$  at time  $t$  only if
  - Its effective priority  $p_i > C_{global}(t)$  OR  $p_i = C_{global}(t)$  and  $\tau_i$  “owns” the ceiling resource
  - OTHERWISE, it transmits its priority to the task  $\tau_j$  that holds semaphore  $S_k$

- **Example**

Task	Priority	Execution Times	Arrival time
$\tau_1$	$P_1$	 Sequential CS	5
$\tau_2$	$P_2$		2
$\tau_3$	$P_3$	 Nested CS	0
			

# Properties of PCP [1/4]

- PCP prevents transitive blocking





# Properties of PCP [2/4]

- PCP prevents deadlocks

# Properties of PCP [3/4]

- A task  $\tau_i$  can be blocked for **at most** the duration of **one** critical section

# Properties of PCP [4/4]

- A critical section  $z_{i,k}$  belonging to task  $\tau_j$  and guarded by semaphore  $S_k$  can block a task  $\tau_i$  only if  $P_j < P_i$  and  $C_{global}(S_k) \geq P_i$

# Computing Blocking Time $B_i$