




Resource Sharing

CPEN 432 Real-Time System Design

Arpan Gujarati
University of British Columbia

Terminology

- Task set $\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$ consists of n periodic tasks
- Each task is characterized by a period T_i and worst-case completion time C_i
- The tasks cooperate through m shared resources R_1, R_2, \dots, R_m
- Each resource R_k is guarded by a distinct **binary semaphore** S_k
 - All critical sections using R_k start and end with operations $wait(S_k)$ and $signal(S_k)$
- Each task is assigned a fixed **base priority** P_i (e.g., using RM)
 - Assumption: priorities are unique and $P_1 > P_2 > \dots > P_n$
- Each task also has an **effective priority** p_i ($\geq P_i$)
 - It is initially set to P_i and can be **dynamically updated**
- B_i denotes the maximum blocking time task τ_i can experience
 - B_i goes into the fixed-priority response-time analysis (recall from previous lectures)
- $z_{i,k}$ denotes any arbitrary critical section of τ_i guarded by semaphore S_k
 - $Z_{i,k}$ denotes the longest among all these critical sections
 - $\delta_{i,k}$ denotes the length of this longest critical section $Z_{i,k}$






What's next?

Protocol Definition


- Unlike NPP, resource holding jobs remain **fully preemptive**
- Tasks are scheduled based on their effective priorities
 - For scheduling purposes, τ_i 's priority is considered to be p_i and not P_i
- Suppose task τ_i tries to enter a critical section by acquiring resource R_k
 - Case 1: R_k is already held by a lower-priority task $\tau_j \implies \tau_i$ is **blocked** by τ_j
 - Case 2: R_k is already held by a higher-priority task $\tau_j \implies \tau_i$ is **interfered** by τ_k
 - Case 3: R_k is not held by any task $\implies \tau_i$ **enters** the critical section
- For Case 1, τ_j **inherits** τ_i 's effective priority
 - τ_j 's dynamic priority is updated as $p_j = p_i$
- In general, τ_j inherits the **highest priority of among all tasks that it blocks**
 - At any point of time, $p_j(R_k) = \max \{P_j, \max_{\forall h} \{p_h \mid \tau_h \text{ is blocked on } R_k\}\}$

Example




Properties of PIP [1/5]

A semaphore S_k can cause push-through blocking to task T_i , only if S_k is accessed both by a task with priority lower than P_i and by a task with priority higher than P_i .



Properties of PIP [2/5]

Transitive priority inheritance can occur only in the presence of nested critical sections.



Properties of PIP [3/5]

If there are l_i lower-priority tasks that can block a task τ_i , then τ_i can be blocked for at most the duration of l_i critical sections, one for each of the l_i lower-priority tasks, regardless of the number of semaphores used by τ_i .

Properties of PIP [4/5]

If there are s_i distinct semaphores that can block a task τ_i , then τ_i can be blocked for at most the duration of s_i critical sections, one for each of the s_i semaphores, regardless of the number of critical sections used by τ_i .

Properties of PIP [5/5]

Under the Priority Inheritance Protocol, a task τ_i can be blocked for at most the duration of $\alpha_i = \min(l_i, s_i)$ critical sections, where l_i is the number of lower-priority tasks that can block τ_i and s_i is the number of distinct semaphores that can block τ_i .

Computing Blocking Time B_i [1/2]

- A precise evaluation of the blocking factor B_i is quite complex because each critical section of the lower-priority tasks may interfere with τ_i via **direct blocking, push-through blocking, or transitive inheritance**
- Simplified algorithm
 - Assumes **no nested critical sections**, hence **no transitive inheritance**

Computing Blocking Time B_i [2/2]

- Semaphores that can **directly** block τ_i and that are shared by the lower-priority task τ_j are $\sigma_{i,j}^{dir} = \sigma_i \cap \sigma_j$
- Semaphores that can block τ_i by **push-through** and that are shared by the lower-priority task τ_j are $\sigma_{i,j}^{pt} = \bigcup_{h:P_h > P_i} \sigma_h \cap \sigma_j$
- Semaphores that can block τ_i either **directly or by push-through** and that are shared by the lower-priority task τ_j
 - $\sigma_{i,j} = \sigma_{i,j}^{dir} \cup \sigma_{i,j}^{pt} = \bigcup_{h:P_h \geq P_i} \sigma_h \cap \sigma_j$
- Longest critical sections used by lower-priority task τ_j that can block τ_i either directly or by push-through is
 - $\gamma_{i,j} = \{Z_{j,k} \mid R_k \in \sigma_{i,j}\}$
- LII critical sections that can block τ_i either directly or by push-through is $\gamma_i = \bigcup_{j:P_j < P_i} \gamma_{i,j}$
- B_i is given by the largest sum of the lengths of the α_i critical sections in γ_i
 - The sum should contain only terms $\delta_{i,k}$ referring to different tasks and different semaphore

The Priority Ceiling Protocol (PCP)

PCP vs PIP

- The PIP is a **reactive** locking protocol
 - It only kicks in when resource contention already exists
- **Key PCP insight**
 - Better to **prevent** problematic scenarios rather **than resolve** them
- The PCP is an **anticipatory** locking protocol
 - Exploits the knowledge of resource needs at **design time** to avoids excessive blocking at runtime

Key Concepts

- **Priority ceilings**

- Each semaphore S_k is **statically** assigned a priority ceiling $C_{static}(S_k)$
 - $C_{static}(S_k)$ = priority of the highest-priority task that **ever** accesses S_k

- **Current system ceiling**

- At any time t , a global system ceiling $C_{global}(t)$ is dynamically computed
 - $C_{global}(t)$ = highest priority ceiling among all semaphores locked at time t OR
(if no semaphores are locked) sentinel value P_0 that is **smaller** than all task priorities

- **Protocol**

- Task τ_i can acquire semaphore S_k at time t only if
 - Its effective priority $p_i > C_{global}(t)$ OR $p_i = C_{global}(t)$ and τ_i “owns” the ceiling resource
 - OTHERWISE, it transmits its priority to the task τ_j that holds semaphore S_k

Example