

Resource Sharing

CPEN 432 Real-Time System Design

Arpan Gujarati
University of British Columbia

Terminology

- Task set $\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$ consists of n periodic tasks
- Each task is characterized by a period T_i and worst-case completion time C_i
- The tasks cooperate through m shared resources R_1, R_2, \dots, R_m
- Each resource R_k is guarded by a distinct **binary semaphore** S_k
 - All critical sections using R_k start and end with operations $wait(S_k)$ and $signal(S_k)$
- Each task is assigned a fixed **base priority** P_i (e.g., using RM)
 - Assumption: priorities are unique and $P_1 > P_2 > \dots > P_n$
- Each task also has an **effective priority** p_i ($\geq P_i$)
 - It is initially set to P_i and can be **dynamically updated**
- B_i denotes the maximum blocking time task τ_i can experience
 - B_i goes into the fixed-priority response-time analysis (recall from previous lectures)
- $z_{i,k}$ denotes any arbitrary critical section of τ_i guarded by semaphore S_k
 - $Z_{i,k}$ denotes the longest among all these critical sections
 - $\delta_{i,k}$ denotes the length of this longest critical section $Z_{i,k}$

The Priority Ceiling Protocol (PCP)

PCP Key Concepts

- **Priority ceilings**

- Each semaphore S_k is **statically** assigned a priority ceiling $C_{static}(S_k)$
 - $C_{static}(S_k)$ = priority of the highest-priority task that **ever** accesses S_k

- **Current system ceiling**

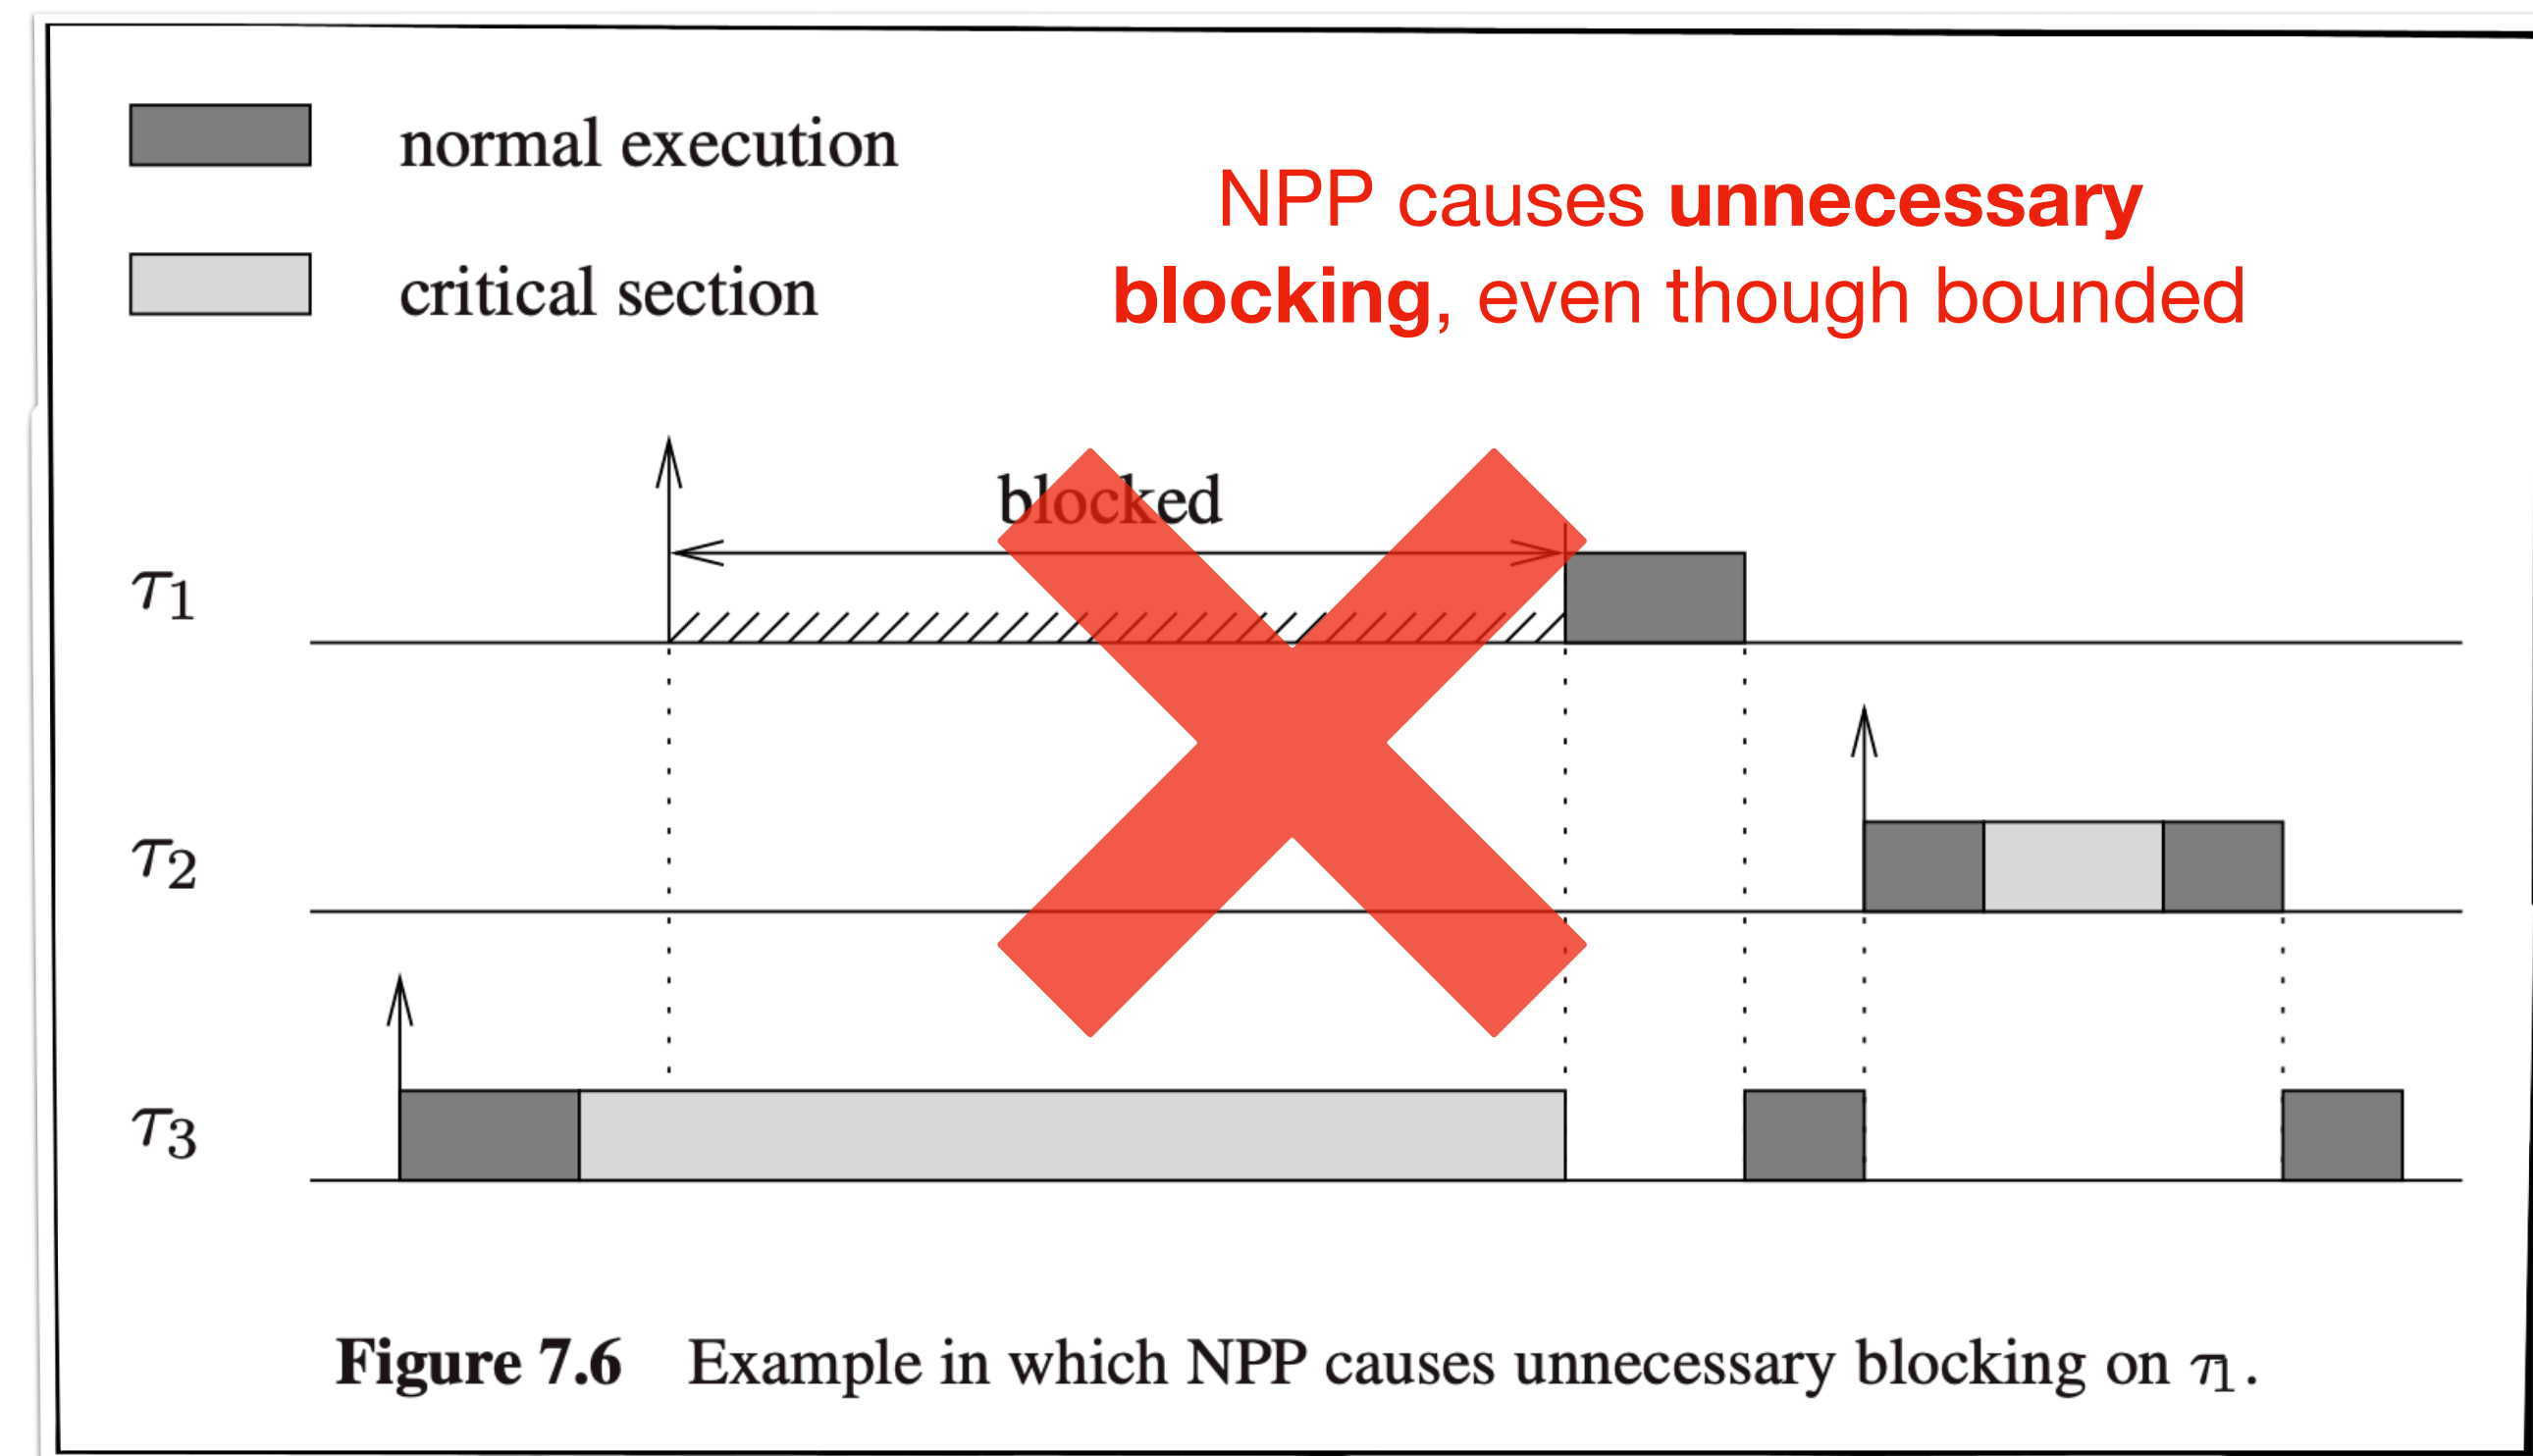

- At any time t , a global system ceiling $C_{global}(t)$ is dynamically computed
 - $C_{global}(t)$ = highest priority ceiling among all semaphores locked at time t OR
(if no semaphores are locked) sentinel value P_0 that is **smaller** than all task priorities

- **Protocol**

- Task τ_i can acquire semaphore S_k at time t only if
 - Its effective priority $p_i > C_{global}(t)$ OR $p_i = C_{global}(t)$ and τ_i “owns” the ceiling resource
 - OTHERWISE, it transmits its priority to the task τ_j that holds semaphore S_k

Analytically, PCP is better than PIP




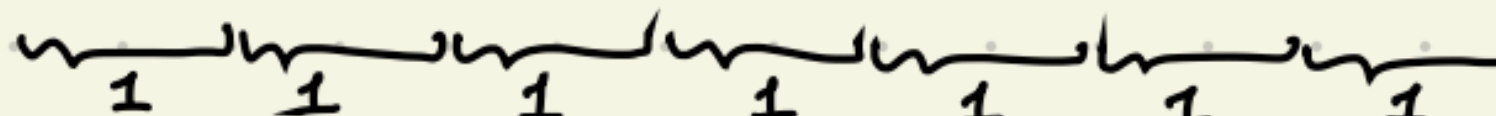
- Like PIP ...



Analytically, PCP is better than PIP

- In addition, unlike PIP
 - PCP prevents transitive blocking
 - PCP prevents deadlocks
 - A task τ_i can be blocked for **at most** the duration of **one** critical section

PCP Example

Task	Priority	Execution Times	Arrival time
τ_1	P_1	 Sequential CS	5
τ_2	P_2		2
τ_3	P_3	 Nested CS 	0

$$P_1 > P_2 > P_3 > P_4$$


$$C(S_A) = 1,$$

$$C(S_B) = 1,$$

$$C(S_C) = 2$$

τ_2 cannot acquire S_C because its priority $P_2 = P_2 \neq C_{global}(t) = P_1$

τ_1 tries to acquire S_A .
But, since $P_1 = P_1 = C_{global}$
& τ_1 does not "own" it
leaving - it cannot



What's Wrong with Context Switches?

- Each contended critical section causes *two additional context switches*.
 - Regular preemption: LO-HI-LO
 - With critical section: LO-HI-LO-HI-LO
 - Context switches can be a main source of overhead: OS context switching code, possibly TLB flushes, **loss of cache affinity**.
 - Switching *back* to a preempted job implies the need for **separate function call stacks** for each task → *memory overhead. Why?*
- Avoiding LO-HI-LO-HI-LO context switches avoids runtime overheads and enables single-stack designs.

Stack Resource Policy (SRP)

The Stack Resource Policy (SRP)

Observation: if a preempting job requires a locked resource, then a LO-HI-LO-HI-LO context switch sequence becomes **inevitable** *only if the preempting job is allowed to start executing.*

Solution: do not allow jobs to commence execution until all (possibly) required resources are available.

→ No more LO-HI-LO-HI-LO context switch sequences...

SRP Definition^{Ba91}

1. Define priority ceilings and system ceilings as under the PCP.
2. When a job is released, it may not commence execution until its (base) priority exceeds the system ceiling (or *preemption threshold*).
3. Whenever a job requires a resource, it gains access immediately.

^{Ba91} T. Baker (1991). Stack-based scheduling for realtime processes. Real-Time Systems, 3(1):67–99.

~~PCP~~ SRP Key Concepts

- **Priority ceilings**

- Each semaphore S_k is **statically** assigned a priority ceiling $C_{static}(S_k)$
 - $C_{static}(S_k)$ = priority of the highest-priority task that **ever** accesses S_k




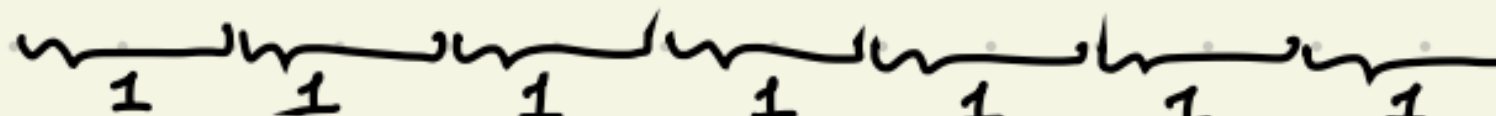
- **Current system ceiling**

- At any time t , a global system ceiling $C_{global}(t)$ is dynamically computed
 - $C_{global}(t)$ = highest priority ceiling among all semaphores locked at time t OR
(if no semaphores are locked) sentinel value P_0 that is **smaller** than all task priorities

- **Protocol**

- Task τ_i can acquire semaphore S_k ~~at time t only if~~ **immediately**
- **Task τ_i may commence its execution only if**
 - Its effective priority $p_i > C_{global}(t)$ OR $p_i = C_{global}(t)$ and τ_i “owns” the ceiling resource
 - OTHERWISE, it transmits its priority to the task τ_j that holds semaphore S_k

PCP Example

Task	Priority	Execution Times	Arrival time
τ_1	P_1	 Sequential CS	5
τ_2	P_2		2
τ_3	P_3	 Nested CS 	0

$$P_1 > P_2 > P_3 > P_4$$


$$C(S_A) = 1,$$

$$C(S_B) = 1,$$

$$C(S_C) = 2$$

τ_2 cannot acquire S_C because its priority $P_2 = P_2 \neq C_{global}(t) = P_1$

τ_1 tries to acquire S_A .
But, since $P_1 = P_1 = C_{global}$
& τ_1 does not "own" it
leaving - it cannot



SRP Blocking Analysis

The bound on *worst-case* pi-blocking under the SRP is *identical* to the PCP's bound.

$$B_i = \max\{Z_{j,k} \mid P_j < P_i \text{ and } C_{global}(S_k) \geq P_i\}$$

- The *actual* pi-blocking differs under the SRP and the PCP.
 - The SRP moves blocking to an earlier point in time.
 - On average, the PCP may yield slightly less blocking. (Why?)

Sharing Runtime Stacks

Example: $prio(\tau_4) > prio(\tau_3) = prio(\tau_2) > prio(\tau_1)$




Figure 7.21 Possible evolution with one stack per task.




Figure 7.22 Possible evolution with a single stack for all tasks.